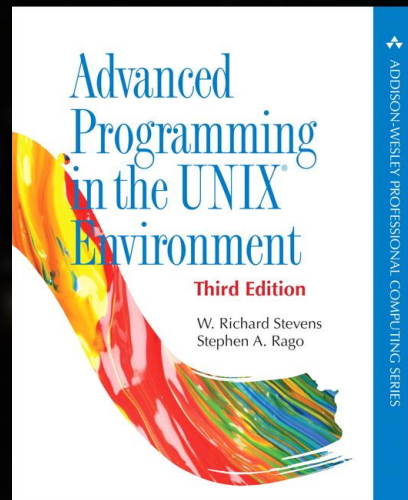






4

The Only Way In Or Out Is Through Phone Lines
- The Matrix (1999), Lana & Lilly Wachowski



Chapter 03: File I/O



Computer

Memory

Kernel: Device Manager

Kernel: Memory Manager


Kernel: File Manager


Kernel: Network Manager

Kernel: Process Manager

Bus

Processor







Device is a Single 1-D Array (String) of Bytes

Even Memory and Processor!

File System: High-Level

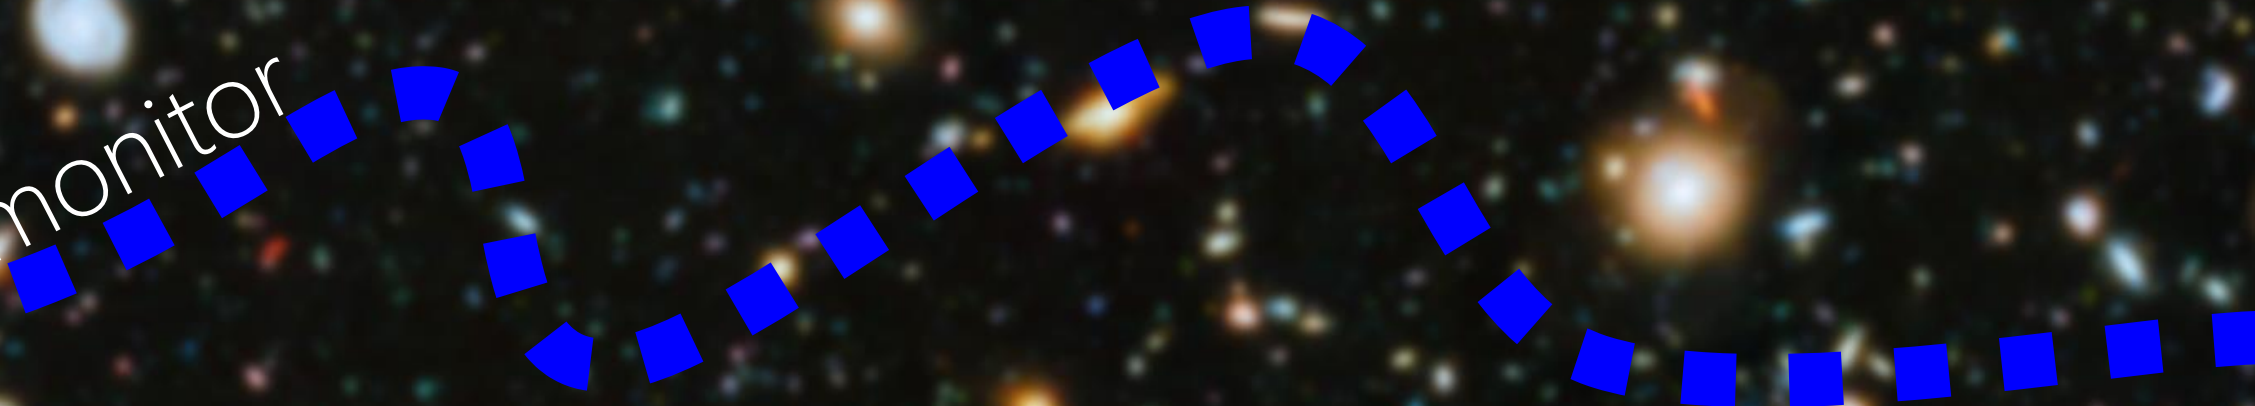
myphoto.jpg



keyboard



monitor





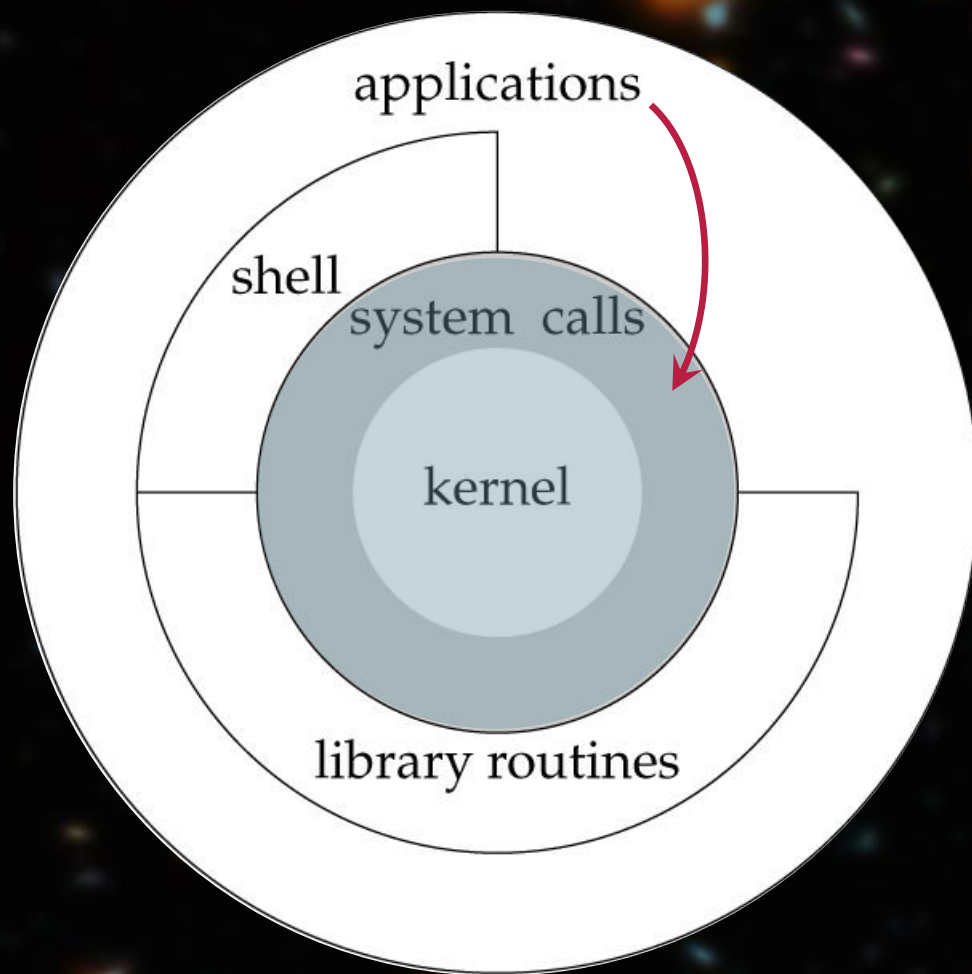
Operator - The Matrix (1999), Lana & Lilly Wachowski



Operation

What do you expect from a kernel about string of bytes | device | file?

File System: High-Level



Header	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
< aio.h>	•	•	•	•	asynchronous I/O
< cpio.h>	•	•	•	•	cpio archive values
< dirent.h>	•	•	•	•	directory entries (Section 4.22)
< dlfcn.h>	•	•	•	•	dynamic linking
< fcntl.h>	•	•	•	•	file control (Section 3.14)
< fnmatch.h>	•	•	•	•	filename-matching types
< glob.h>	•	•	•	•	pathname pattern-matching and generation
< grp.h>	•	•	•	•	group file (Section 6.4)
< iconv.h>	•	•	•	•	codeset conversion utility
< langinfo.h>	•	•	•	•	language information constants
< monetary.h>	•	•	•	•	monetary types and functions
< netdb.h>	•	•	•	•	network database operations
< nl_types.h>	•	•	•	•	message catalogs
< poll.h>	•	•	•	•	poll function (Section 14.4.2)
< pthread.h>	•	•	•	•	threads (Chapters 11 and 12)
< pwd.h>	•	•	•	•	password file (Section 6.2)
< regex.h>	•	•	•	•	regular expressions
< sched.h>	•	•	•	•	execution scheduling
< semaphore.h>	•	•	•	•	semaphores
< strings.h>	•	•	•	•	string operations
< tar.h>	•	•	•	•	tar archive values
< termios.h>	•	•	•	•	terminal I/O (Chapter 18)
< unistd.h>	•	•	•	•	symbolic constants
< wordexp.h>	•	•	•	•	word-expansion definitions
< arpa/inet.h>	•	•	•	•	Internet definitions (Chapter 16)
< net/if.h>	•	•	•	•	socket local interfaces (Chapter 16)
< netinet/in.h>	•	•	•	•	Internet address family (Section 16.3)
< netinet/tcp.h>	•	•	•	•	Transmission Control Protocol definitions
< sys/mman.h>	•	•	•	•	memory management declarations
< sys/select.h>	•	•	•	•	select function (Section 14.4.1)
< sys/socket.h>	•	•	•	•	sockets interface (Chapter 16)
< sys/stat.h>	•	•	•	•	file status (Chapter 4)
< sys/statvfs.h>	•	•	•	•	file system information
< sys/times.h>	•	•	•	•	process times (Section 8.17)
< sys/types.h>	•	•	•	•	primitive system data types (Section 2.8)
< sys/un.h>	•	•	•	•	UNIX domain socket definitions (Section 17.2)
< sys/utsname.h>	•	•	•	•	system name (Section 6.9)
< sys/wait.h>	•	•	•	•	process control (Section 8.6)

creat

POSIX

/dev/cpu/cpuid
/cpu/

/sound/ ..

→ DataType

```
#include <fcntl.h>  
int creat(const char *path, mode_t mode);
```

non-negative number for write-only if OK
-1 on error


File Descriptor (fs)

Number does not Matter, Connection Matters

The Only Way In Or Out Is Through Phone Lines

The number does not matter, the connection is important!

Imagine a dynamic phone#, dynamic postal code, dynamic ip (DHCP)

A cosmic background image featuring a dense field of galaxies in various colors (yellow, orange, blue, white) against a black space. Two horizontal blue lines are positioned above and below the central text.

File Descriptor (fs) != File Identifier

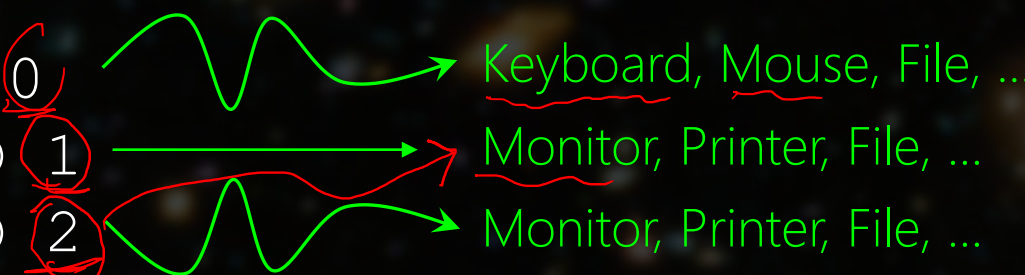
Because kernel reuse them for other files and devices, when available!

File Descriptor (fs)

STDIN_FILENO, STDOUT_FILENO, STDERR_FILENO

unistd.h

```
#define STDIN_FILENO 0
#define STDOUT_FILENO 1
#define STDERR_FILENO 2
```



The diagram illustrates the mapping of standard file descriptors to their respective devices. It features three lines of code from `unistd.h` with the values 0, 1, and 2 circled in red. Green arrows point from these circled values to their corresponding device descriptions: 0 points to "Keyboard, Mouse, File, ...", 1 points to "Monitor, Printer, File, ...", and 2 points to "Monitor, Printer, File, ...". Additionally, a green wavy line connects the circled 0 and 1, and another green wavy line connects the circled 1 and 2, indicating a relationship or grouping between these descriptors.

Keyboard, Mouse, File, ...

Monitor, Printer, File, ...

Monitor, Printer, File, ...

Size of (float)

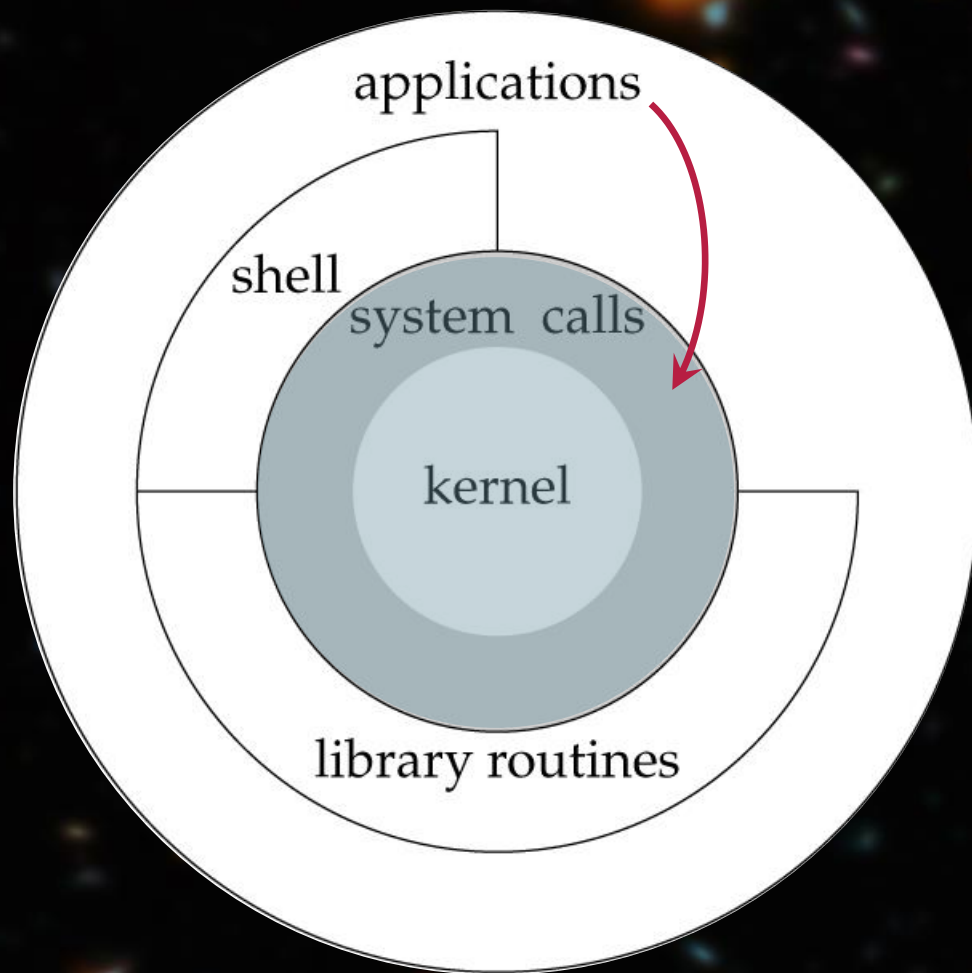
write
POSIX

8
x 4 = 8
2 (float) 2

int nbytes = 4


#include <unistd.h>
ssize_t write(int fd, const void *buf, size_t nbytes);
number of bytes written if OK, -1 on error

8 x



Header	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
< aio.h>	•	•	•	•	asynchronous I/O
< cpio.h>	•	•	•	•	cpio archive values
< dirent.h>	•	•	•	•	directory entries (Section 4.22)
< dlfcn.h>	•	•	•	•	dynamic linking
< fcntl.h>	•	•	•	•	file control (Section 3.14)
< fnmatch.h>	•	•	•	•	filename-matching types
< glob.h>	•	•	•	•	pathname pattern-matching and generation
< grp.h>	•	•	•	•	group file (Section 6.4)
< iconv.h>	•	•	•	•	codeset conversion utility
< langinfo.h>	•	•	•	•	language information constants
< monetary.h>	•	•	•	•	monetary types and functions
< netdb.h>	•	•	•	•	network database operations
< nl_types.h>	•	•	•	•	message catalogs
< poll.h>	•	•	•	•	poll function (Section 14.4.2)
< pthread.h>	•	•	•	•	threads (Chapters 11 and 12)
< pwd.h>	•	•	•	•	password file (Section 6.2)
< regex.h>	•	•	•	•	regular expressions
< sched.h>	•	•	•	•	execution scheduling
< semaphore.h>	•	•	•	•	semaphores
< strings.h>	•	•	•	•	string operations
< tar.h>	•	•	•	•	tar archive values
< termios.h>	•	•	•	•	terminal I/O (Chapter 18)
< unistd.h>	•	•	•	•	symbolic constants
< wordexp.h>	•	•	•	•	word-expansion definitions
< arpa/inet.h>	•	•	•	•	Internet definitions (Chapter 16)
< net/if.h>	•	•	•	•	socket local interfaces (Chapter 16)
< netinet/in.h>	•	•	•	•	Internet address family (Section 16.3)
< netinet/tcp.h>	•	•	•	•	Transmission Control Protocol definitions
< sys/mman.h>	•	•	•	•	memory management declarations
< sys/select.h>	•	•	•	•	select function (Section 14.4.1)
< sys/socket.h>	•	•	•	•	sockets interface (Chapter 16)
< sys/stat.h>	•	•	•	•	file status (Chapter 4)
< sys/statvfs.h>	•	•	•	•	file system information
< sys/times.h>	•	•	•	•	process times (Section 8.17)
< sys/types.h>	•	•	•	•	primitive system data types (Section 2.8)
< sys/un.h>	•	•	•	•	UNIX domain socket definitions (Section 17.2)
< sys/utsname.h>	•	•	•	•	system name (Section 6.9)
< sys/wait.h>	•	•	•	•	process control (Section 8.6)

typedef ~~int~~ PersonId int



typedef


ssize_t, size_t, ..., and many other data types

5 + - Point {
int x,
int y

} Point p;

#include <sys/types.h>

<https://pubs.opengroup.org/onlinepubs/009604599/basedefs/sys/types.h.html>



typedef

int bufsize
(size_t) buf

(ssize_t) size_t, ..., and many other data types

```
#include <sys/types.h>
typedef size_t unsigned long
typedef ssize_t signed long
```

<https://www.ibm.com/docs/en/zos/2.2.0?topic=files-systypesh>



close
POSIX

```
#include <unistd.h>
int close(int fd);
0 if OK, -1 on error
```


close

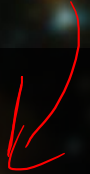
fd: Releases the File Descriptor (Available for Reuse by Kernel)
No Further Access to the File (Device)

```
#include <unistd.h>
int close(int fd);
0 if OK, -1 on error
```



close

Sometimes Optional, but only Sometimes!

When a process terminates, all of its open files are closed automatically by the kernel.
That is all the File Descriptors (fs) are released. 
You can take advantage of this fact and don't explicitly close open files in your programs (not recommended!)


```
hfani@alpha:~$ vi create_file_system_call.c
```

```
#include <fcntl.h>
#include <unistd.h>

#include <sys/types.h>
#include <string.h>

#include <stdio.h>
void main(void) {
    int fd; //file descriptor
    mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH; //for permisison settings
    char *filename = "./my_new_file.txt";

    fd = creat(filename, mode);
    printf("The file descriptor is: %d \n", fd);

    if(fd == -1) {
        printf("Error in creating file!\n");
        return;
    }

    char buf[20];
    size_t nbytes;
    ssize_t bytes_written;

    strcpy(buf, "Hello File!\n");
    nbytes = strlen(buf);

    bytes_written = write(fd, buf, nbytes);
    if(bytes_written != nbytes) {
        printf("Error in writing to the file!")
    }

    int result = close(fd);
    if(result == -1) {
        printf("Error in closing the file!");
    }
}
```



```
#include <fcntl.h>
#include <unistd.h>

#include <sys/types.h>
#include <string.h>

#include <stdio.h>

void main(void) {
    int fd; //file descriptor
    mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH; //for permisison settings
    char *filename = "./my_new_file.txt";

    fd = creat(filename, mode);
    printf("The file descriptor is: %d \n", fd);

    if (fd == -1) {
        printf("Error in creating file!\n");
        return;
    }

    char buf[20];
    size_t nbytes;
    ssize_t bytes_written;

    strcpy(buf, "Hello File!\n");
    nbytes = strlen(buf);

    bytes_written = write(fd, buf, nbytes);
    if (bytes_written != nbytes) {
        printf("Error in writing to the file!");
    }

    int result = close(fd);
    if (result == -1) {
```

```
#include <fcntl.h>
#include <unistd.h>

#include <sys/types.h>
#include <string.h>

#include <stdio.h>
void main(void) {
    int fd; //file descriptor
    mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH; //for permission settings
    char *filename = "./my_new_file.txt";

    fd = creat(filename, mode);
    printf("The file descriptor is: %d \n", fd);

    if(fd == -1){
        printf("Error in creating file!\n");
        return;
    }

    char buf[20];
    size_t nbytes;
    ssize_t bytes_written;

    strcpy(buf, "Hello File!\n");
    nbytes = strlen(buf);

    bytes_written = write(fd, buf, nbytes);
    if(bytes_written != nbytes){
        printf("Error in writing to the file!");
    }

    int result = close(fd);
```



```
#include <fcntl.h>
#include <unistd.h>

#include <sys/types.h>
#include <string.h>

#include <stdio.h>
void main(void){
    int fd;//file descriptor
    mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH;//for permission settings
    char *filename = "../my_new_file.txt";

    fd = creat(filename, mode);
    printf("The file descriptor is: %d \n", fd);

    if(fd == -1){
        printf("Error in creating file!\n");
        return;
    }

    char buf[20];
    size_t nbytes;
    ssize_t bytes_written;

    strcpy(buf, "Hello File!\n");
    nbytes = strlen(buf);

    bytes_written = write(fd, buf, nbytes);
    if(bytes_written != nbytes){
        printf("Error in writing to the file!")
    }

    int result = close(fd);
    if(result == -1){
        printf("Error in closing the file!")
    }
}
```



```
hfani@alpha:~$ cc create_file_system_call.c -o create_file_system_call
hfani@alpha:~$ ./create_file_system_call
The file descriptor is: 3
hfani@alpha:~$
```

```
hfani@alpha:~$ vi my_new_file.txt
```

```
Hello File!
```

```
~
```

```
~
```



```
fd = creat()
```



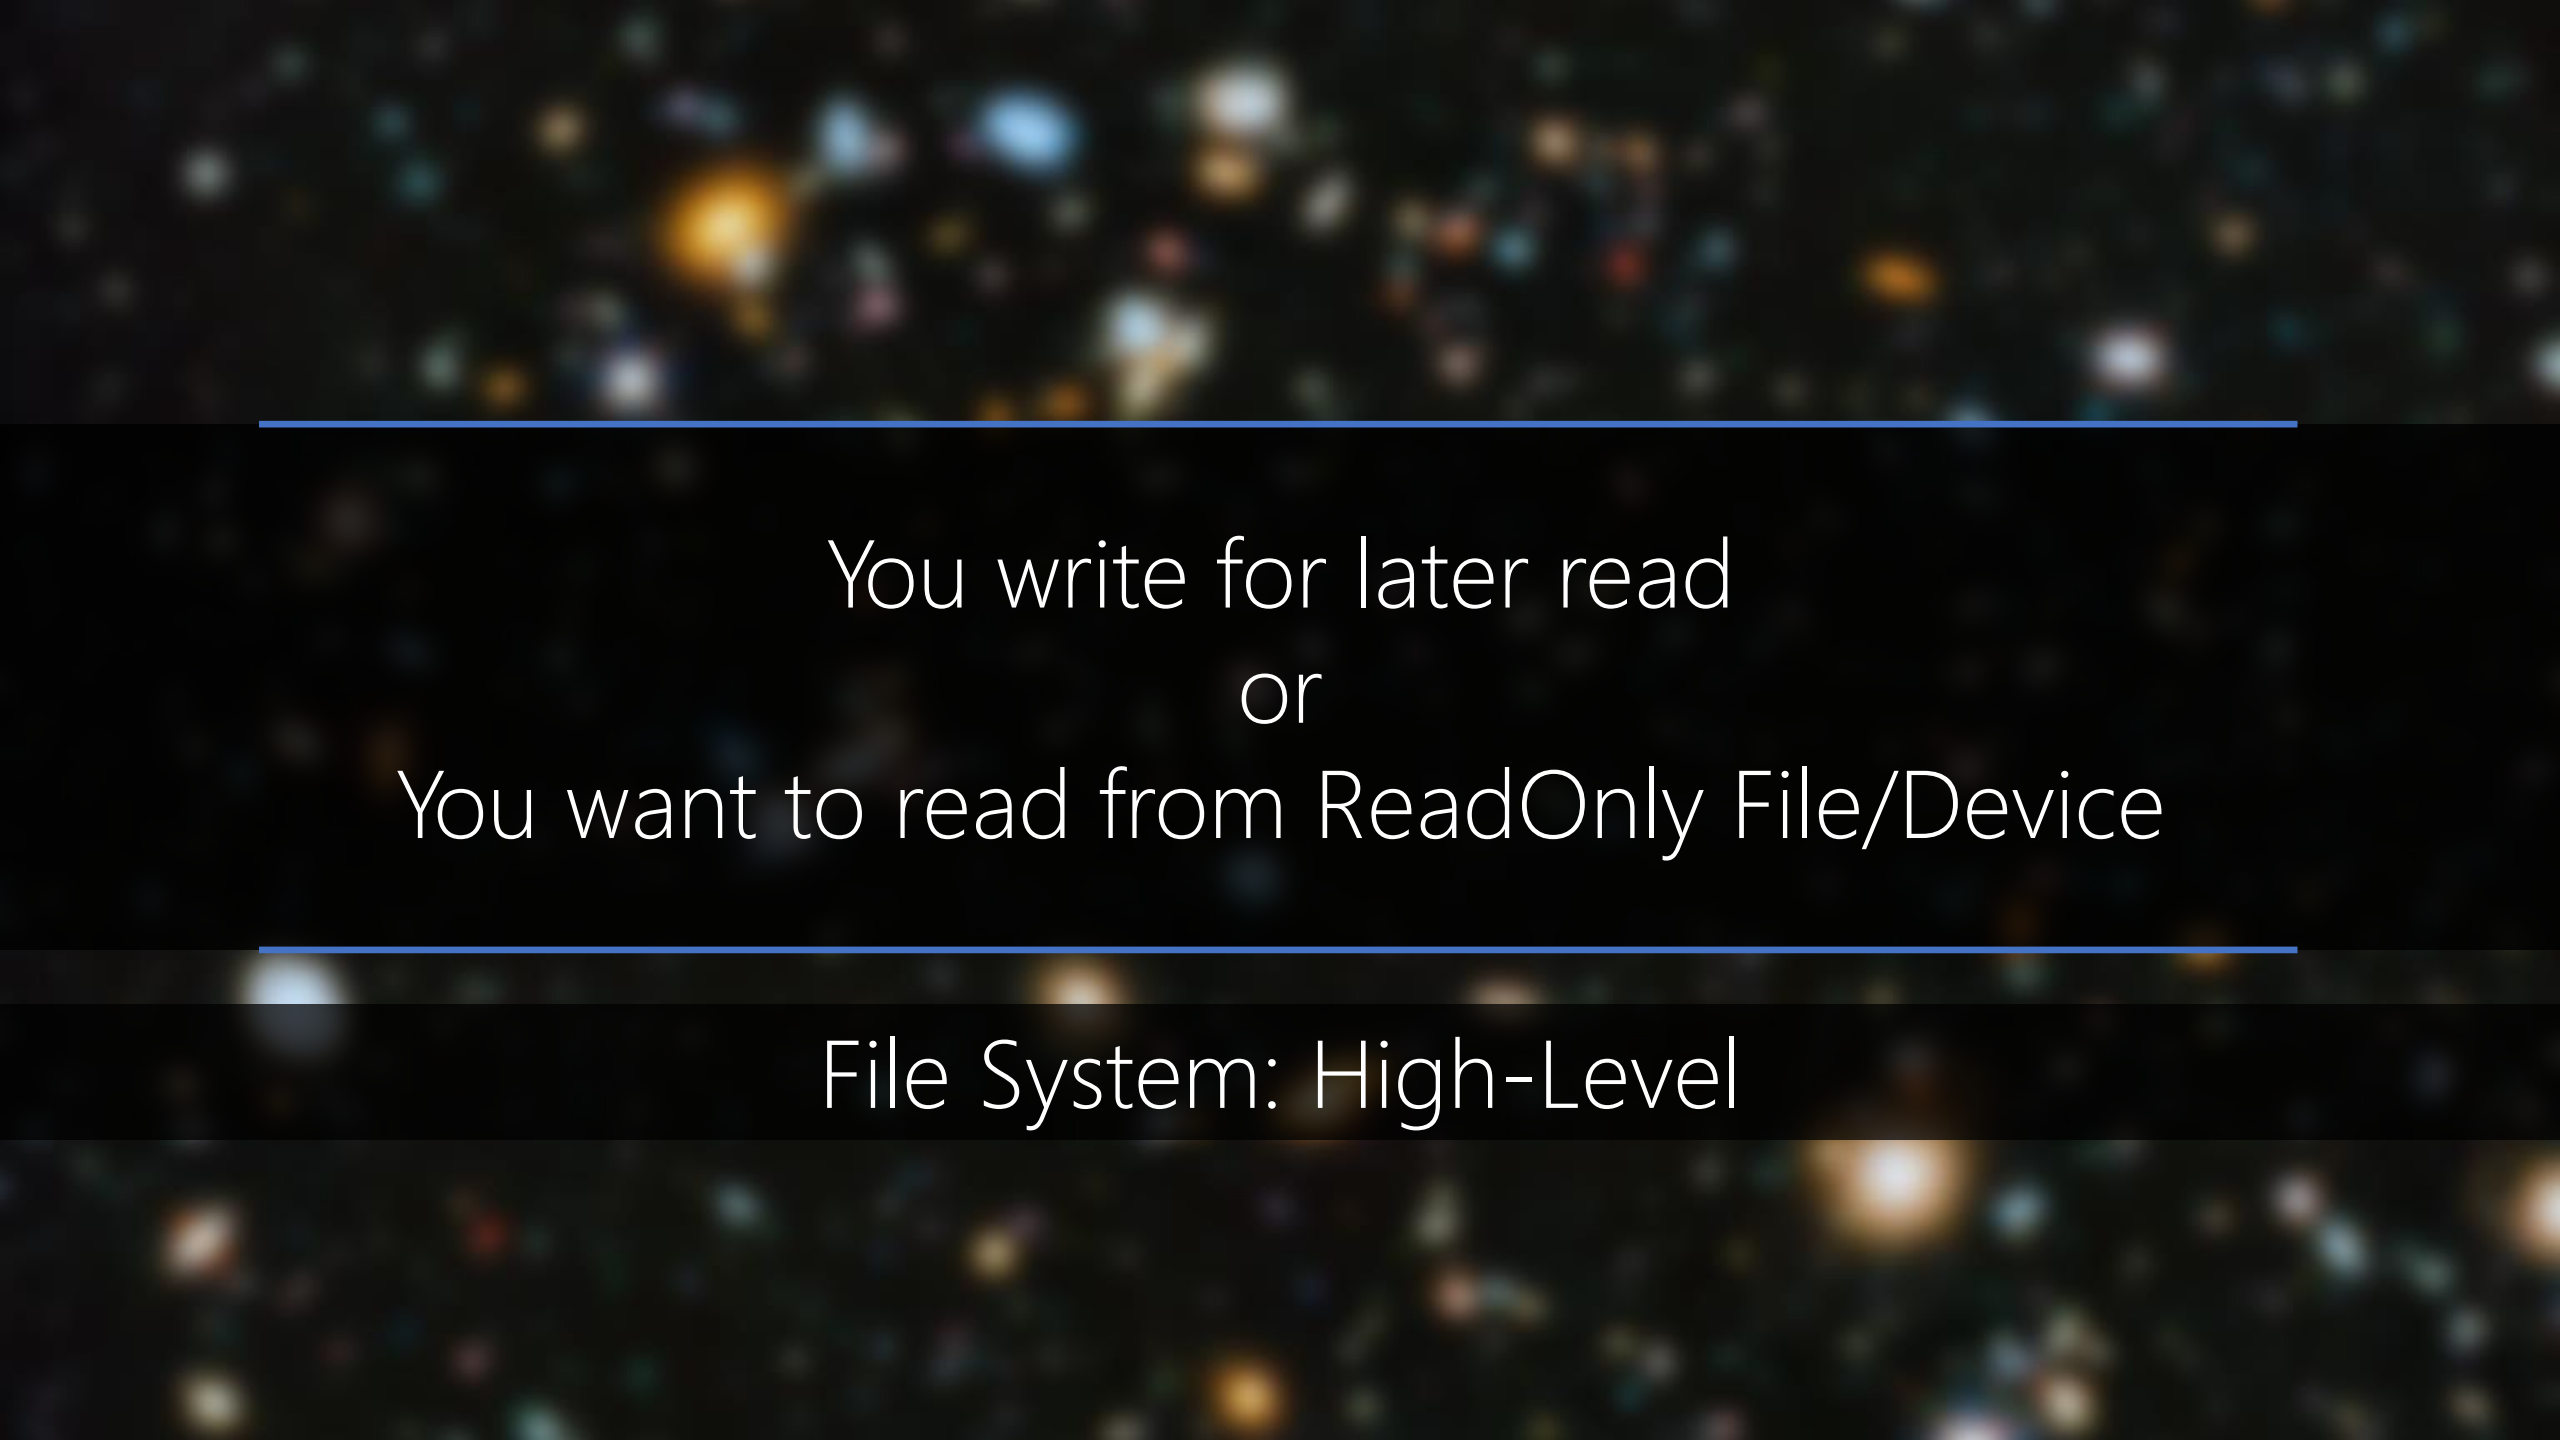
```
write(fd, sth)
```



```
close(fd)
```



You Can't Handle The Truth, A Few Good Men (1992) - Aaron Sorkin



You write for later read
or
You want to read from ReadOnly File/Device

File System: High-Level

```
fd = creat()
```



```
graph TD; A[fd = creat()] --> B[write(fd, sth)]; B --> C[read(fd, sth)]; C --> D[close(fd)];
```

A vertical flowchart on the left side of the image. It consists of four rectangular boxes connected by downward-pointing yellow arrows. The boxes are: 1. Green, containing 'fd = creat()'. 2. Green, containing 'write(fd, sth)'. 3. Blue, containing 'read(fd, sth)'. 4. Green, containing 'close(fd)'.

```
write(fd, sth)
```

```
read(fd, sth)
```

```
close(fd)
```

```
fd = creat()
```



```
graph TD; A[fd = creat()] --> C[read(fd, sth)]; C --> D[close(fd)];
```

A vertical flowchart on the right side of the image. It consists of four rectangular boxes connected by downward-pointing yellow arrows. The boxes are: 1. Green, containing 'fd = creat()'. 2. Blue, containing 'read(fd, sth)'. 3. Green, containing 'close(fd)'. The second box is missing, as indicated by the jump in the arrow from the first to the third box.

```
read(fd, sth)
```

```
close(fd)
```

```
fd = creat()
```

```
write(fd, sth)
```

Not Possible

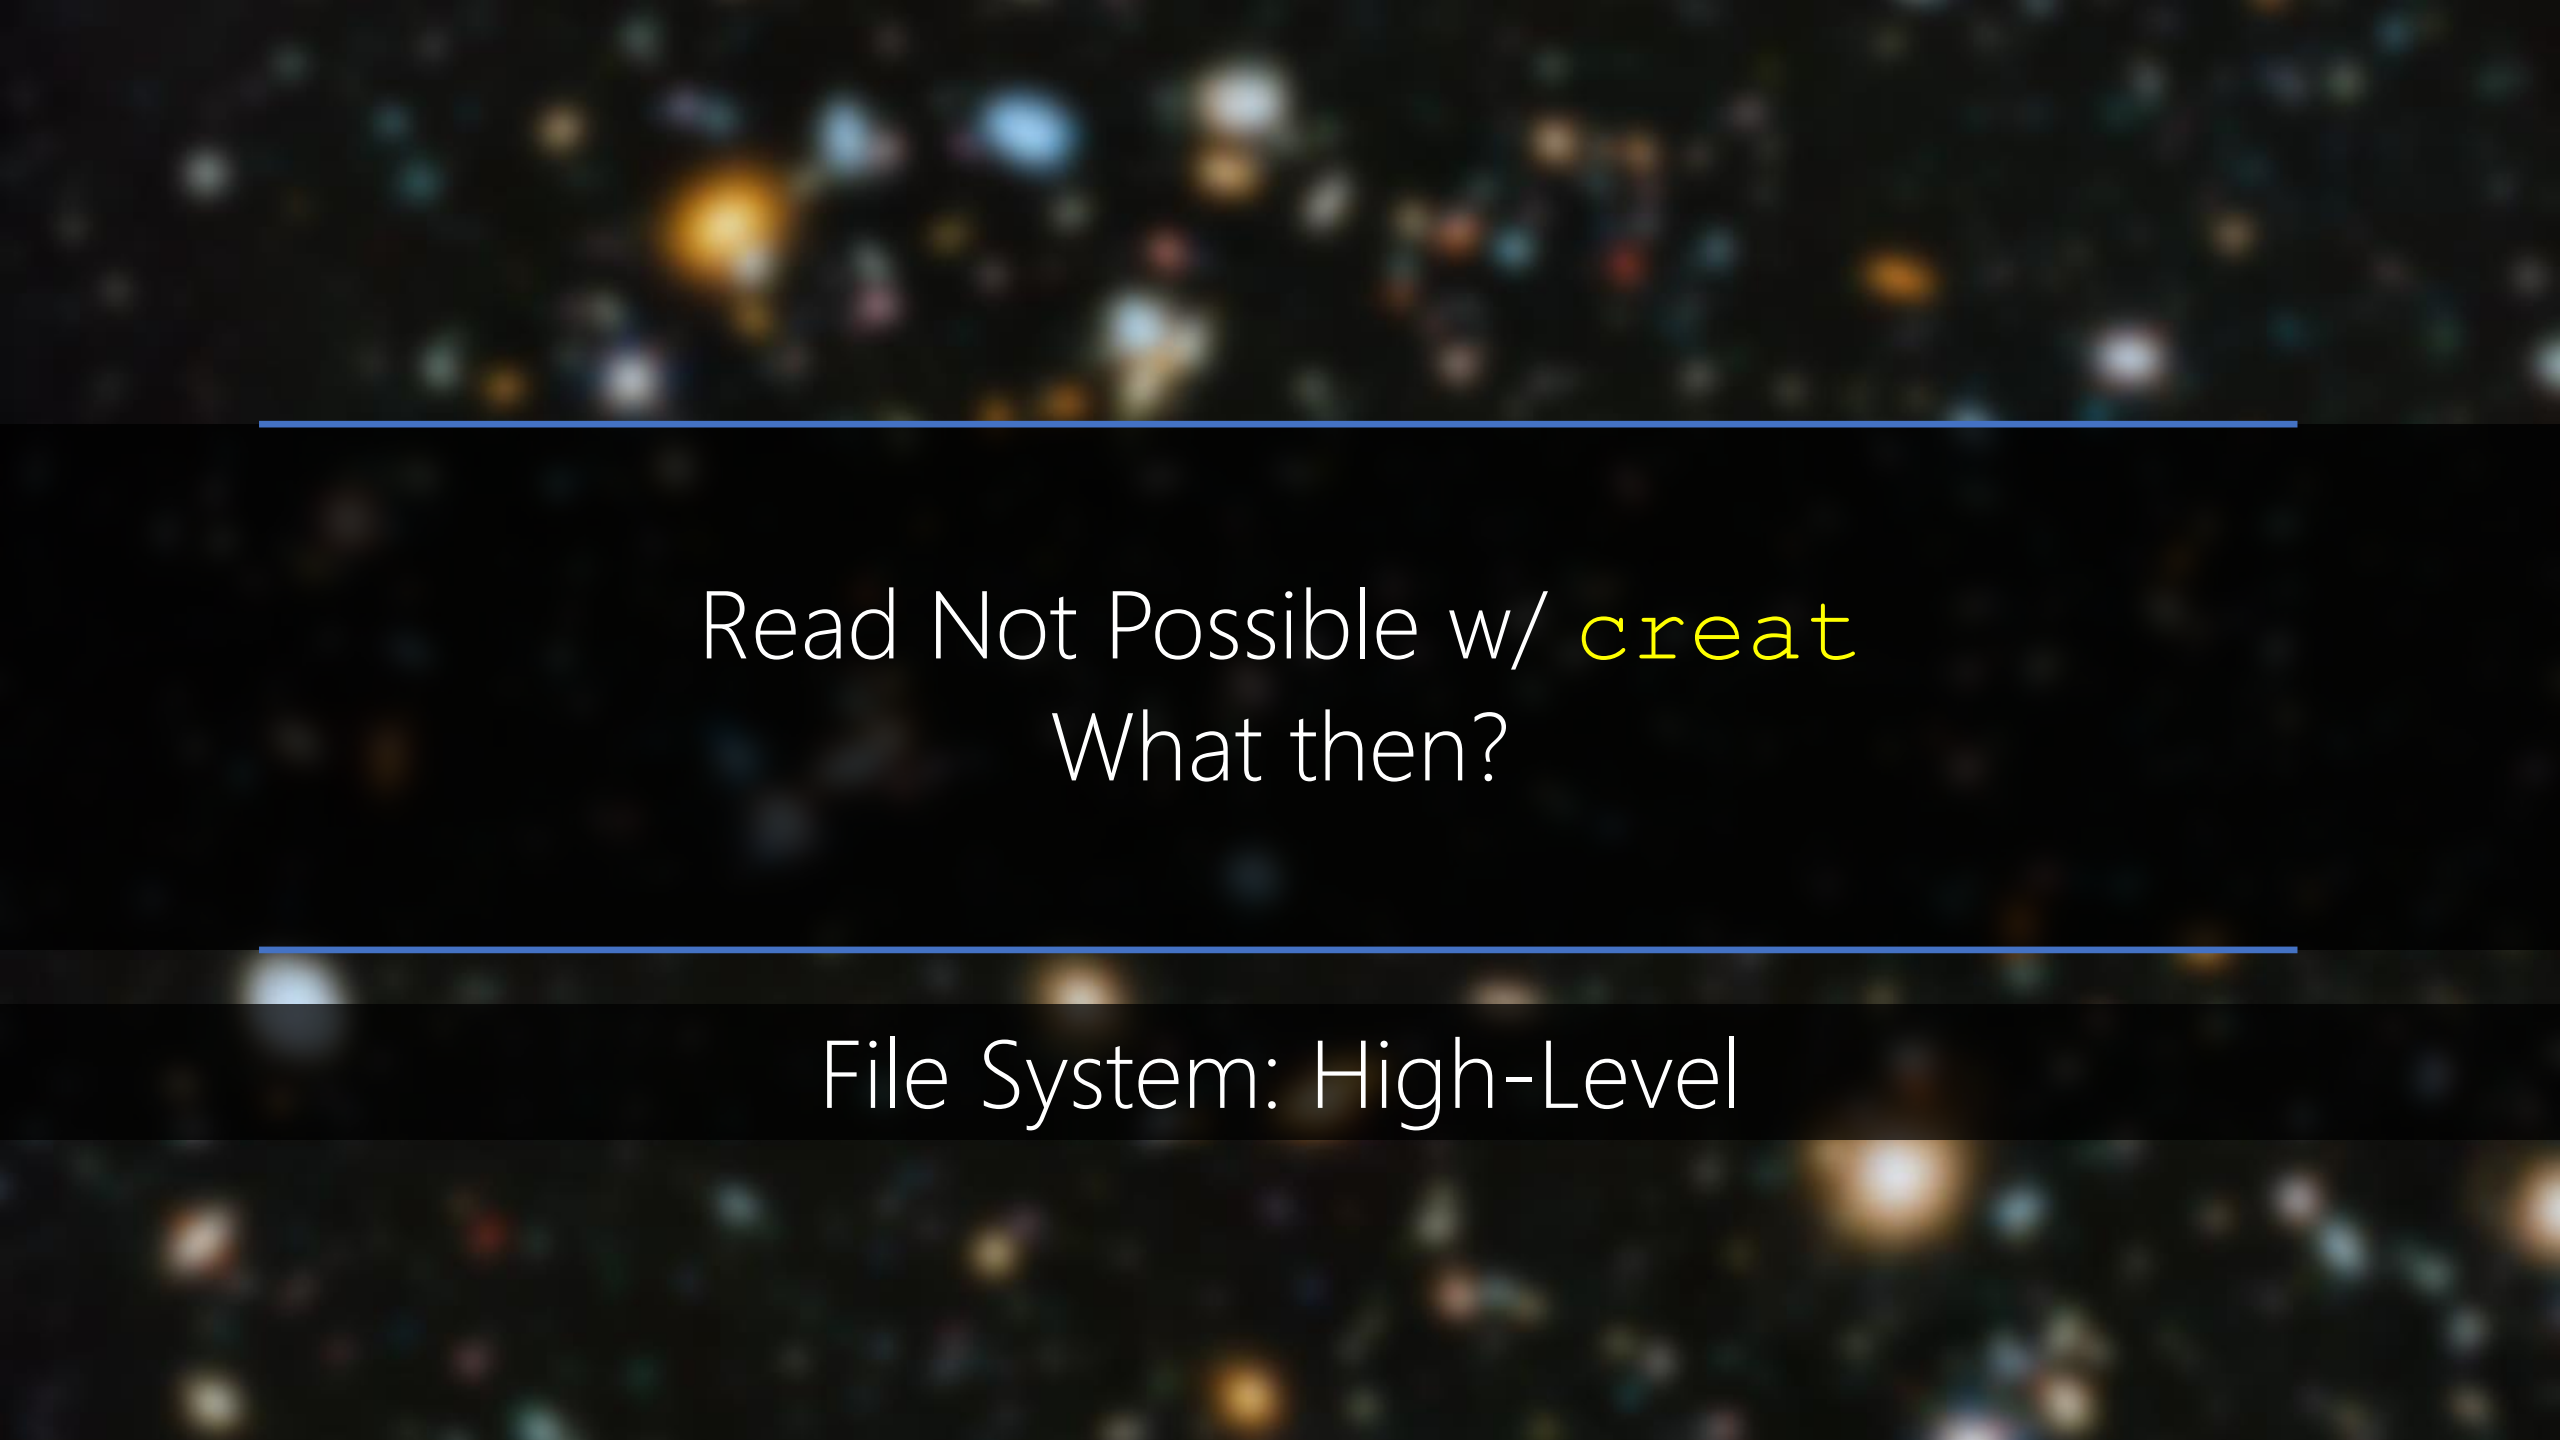
```
read(fd, sth)
```

```
close(fd)
```

```
fd = creat()
```

```
read(fd, sth)
```

```
close(fd)
```

Read Not Possible w/ `creat`
What then?

File System: High-Level

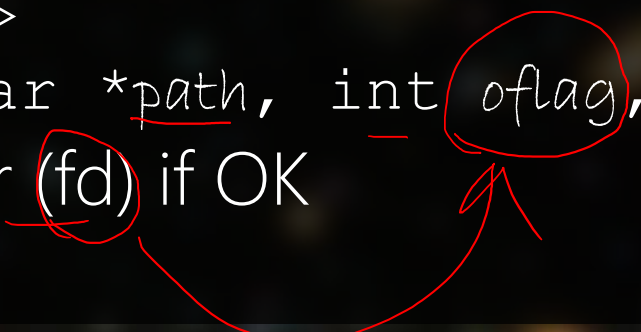
open

POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/open.html>

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error





A quick reminder, why POSIX version?

open

POSIX

```
#include <fcntl.h>  
int open(const char *path, int oflag, ...);  
non-negative number (fd) if OK  
-1 on error
```

open

POSIX

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

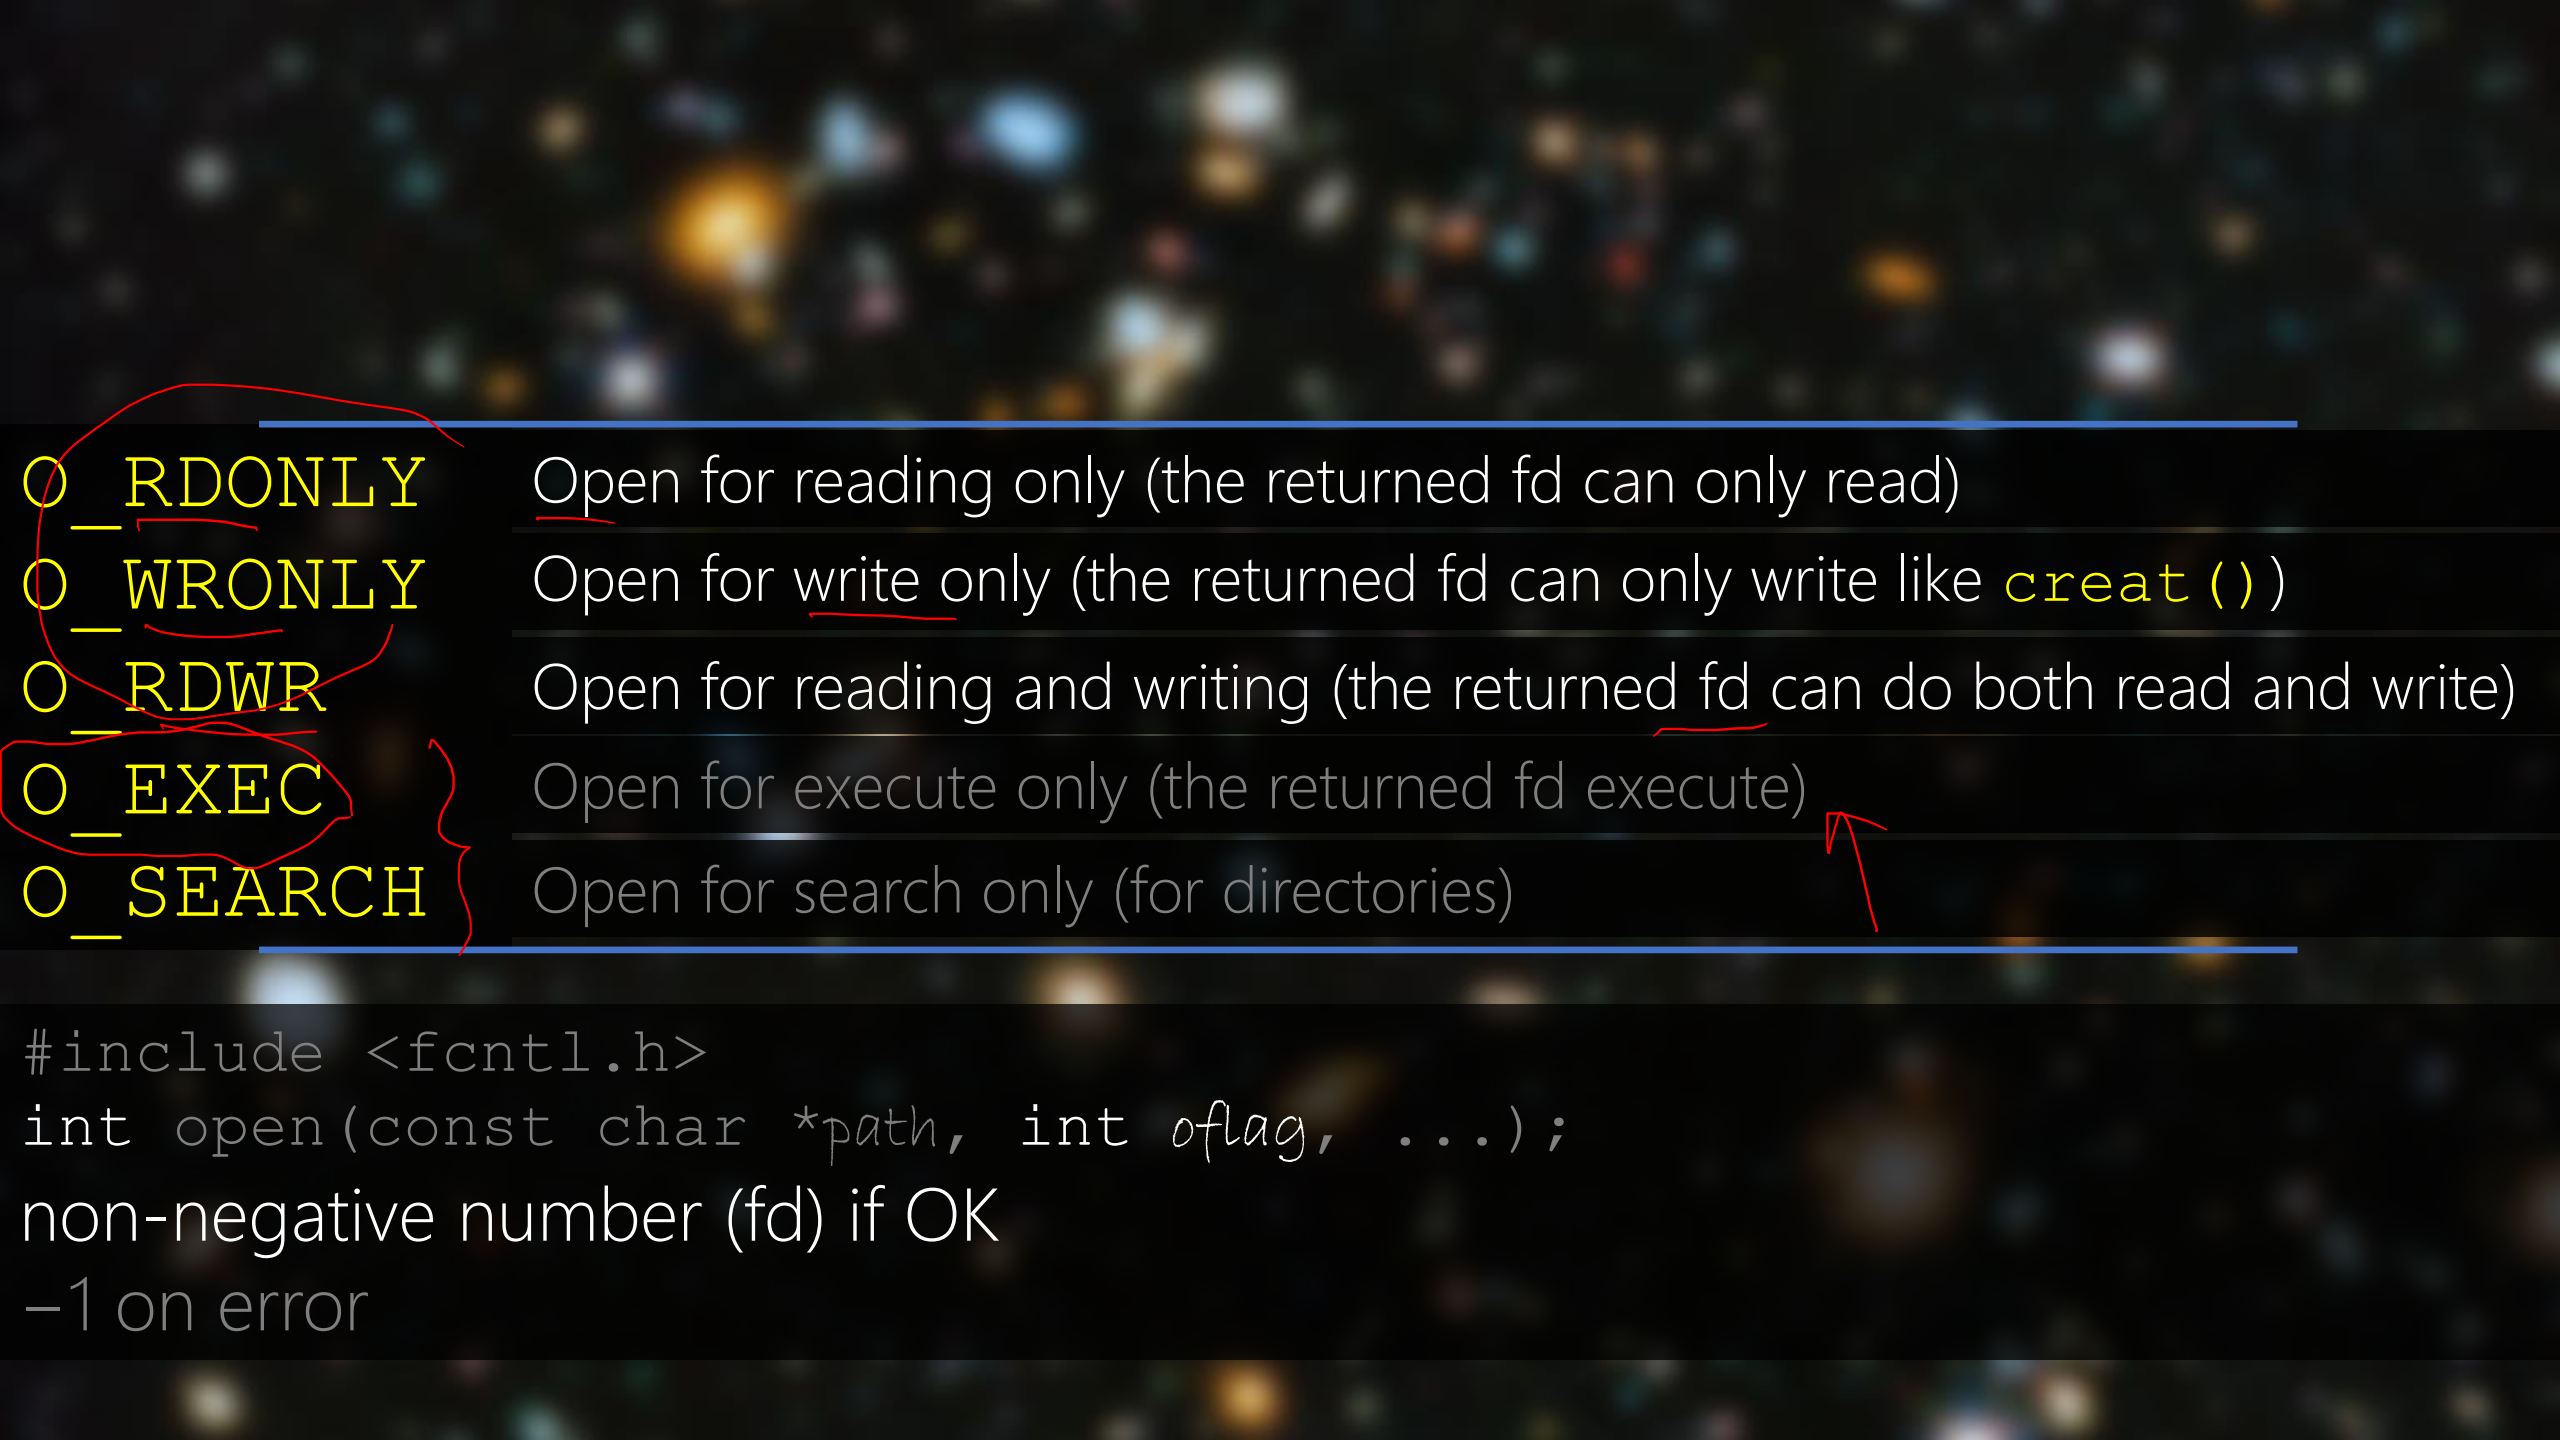
non-negative number (fd) if OK
-1 on error

open

POSIX

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error



<code>O_RDONLY</code>	Open for <u>reading</u> only (the returned fd can only read)
<code>O_WRONLY</code>	Open for <u>write</u> only (the returned fd can only write like <code>creat()</code>)
<code>O_RDWR</code>	Open for reading and writing (the returned <u>fd</u> can do both read and write)
<code>O_EXEC</code>	Open for execute only (the returned fd execute)
<code>O_SEARCH</code>	Open for search only (for directories)

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error

O_RDONLY
O_WRONLY
O_RDWR
O_EXEC
O_SEARCH


Only One of Them

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error

```
hfani@alpha:~$ vi open_rw.c
```

```
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";
    int fd = open(filename, O_RDWR);
    if(fd==-1){
        printf("error happend!\n");
    }
    else if (fd >= 0){
        printf("file successfully opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```




```
hfani@alpha:~$ vi open_rw.c
```

```
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";
    int fd = open(filename, O_RDWR);
    if(fd==-1){
        printf("error happend!\n");
    }
    else if (fd >= 0){
        printf("file successfully opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```

```
hfani@alpha:~$ cc open_rw.c -o open_rw
```

```
hfani@alpha:~$ ./open_rw
```

```
error happend!
```

What's the problem?

creat → close → open
POSIX

```
hfani@alpha:~$ vi creat_close_open_rw.c
```

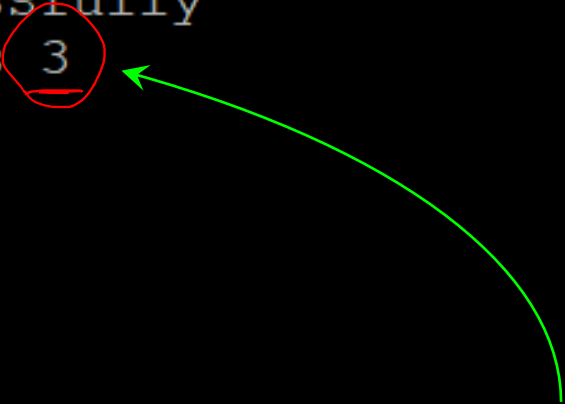
```
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";
    int fd = creat(filename, S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
    if (fd == -1){
        printf("error happened in creating file %s\n", filename);
        return;
    }
    else if (fd > 0){
        printf("file %s has been created successfully and the fd is %d\n", filename, fd);
    }

    int result = close(fd);
    if (result == -1){
        printf("error happened in closing file %s\n", filename);
        return;
    }
    else if (result == 0){
        printf("file %s has been closed successfully\n", filename);
    }

    fd = open(filename, O_RDWR);
    if(fd== -1){
        printf("error happend in opening file %s!\n", filename);
    }
    else if (fd > 0){
        printf("file successfully opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```



```
hfani@alpha:~$ cc creat_close_open_rw.c -o creat_close_open_rw
creat_close_open_rw.c: In function 'main':
creat_close_open_rw.c:14:15: warning: implicit declaration of function 'close'
   14 |     int result = close(fd);
      |                  ^~~~~
      |                  pclose
hfani@alpha:~$ ./creat_close_open_rw
file open_rw_test.txt has been created successfully and the fd is 3
file open_rw_test.txt has been closed successfully
file successfully opened for 2 and the fd is 3
```



O_CREAT

O_EXCL

Create the file if it does not exist (you have to specify the mode_t)

Raise error (fd == -1) if create and the file already exists

<https://stackoverflow.com/questions/48388212/what-is-the-written-out-word-for-o-excl>

```
#include <fcntl.h>
```

```
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK

-1 on error

O_CREAT
O_EXCL

In Combination w/ the Main Five Options

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error


```
hfani@alpha:~$ vi creat_open_rw.c
```

```
#include <fcntl.h>
```

```
#include <stdio.h>
```

```
void main(void){
```

```
    char filename[20]="open_rw_test.txt";
```

```
    int fd = open(filename, O_RDWR | O_CREAT, S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
```

```
    if(fd==-1){
```

```
        printf("error happend in creating the file %s (if not exists) or opening file!\n.", filename);
```

```
    }
```

```
    else if (fd > 0){
```

```
        printf("file successfully created (if not exist) and opened for %d and the fd is %d\n", O_RDWR, fd);
```

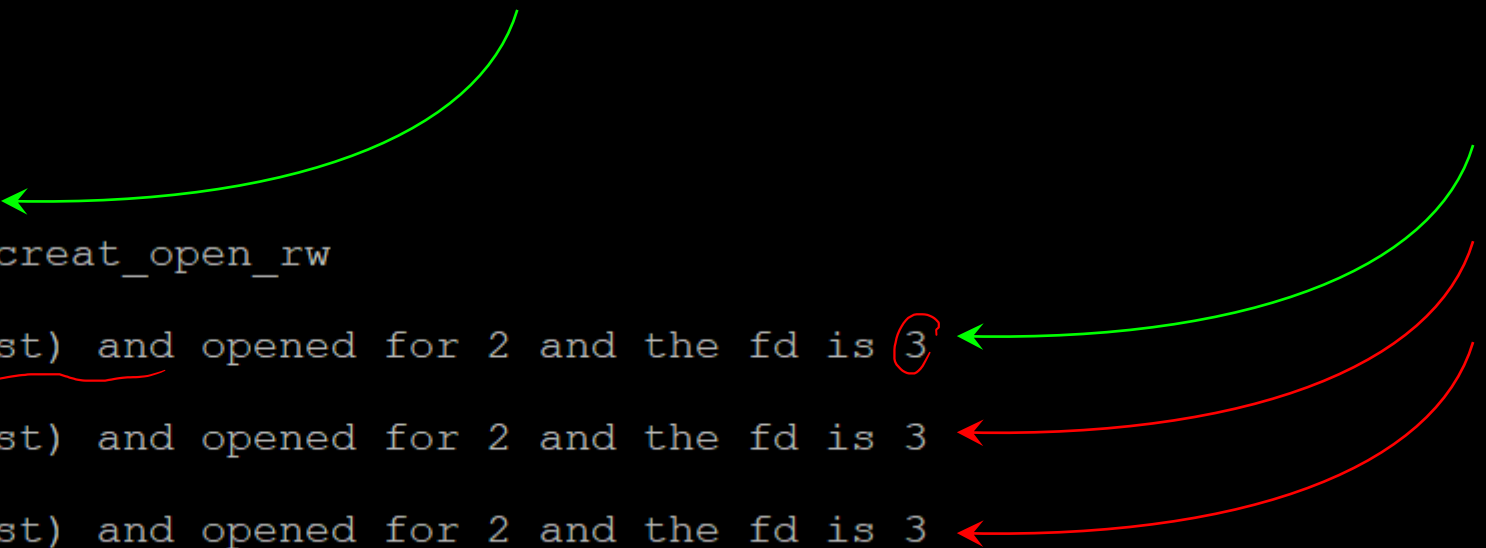
```
    }
```

```
}
```

```
hfani@alpha:~$ vi creat_open_rw.c
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";

    int fd = open(filename, O_RDWR | O_CREAT, S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
    if(fd==-1){
        printf("error happend in creating the file %s (if not exists) or opening file!\n.", filename);
    }
    else if (fd > 0){
        printf("file successfully created (if not exist) and opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```

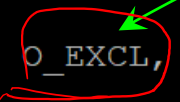
```
hfani@alpha:~$ rm ./open_rw_test.txt
hfani@alpha:~$ cc creat_open_rw.c -o creat_open_rw
hfani@alpha:~$ ./creat_open_rw
file successfully created (if not exist) and opened for 2 and the fd is 3
hfani@alpha:~$ ./creat_open_rw
file successfully created (if not exist) and opened for 2 and the fd is 3
hfani@alpha:~$ ./creat_open_rw
file successfully created (if not exist) and opened for 2 and the fd is 3
hfani@alpha:~$
```



hfani@alpha:~\$ vi creat_excl_open_rw.c

```
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";

    int fd = open(filename, O_RDWR | O_CREAT | O_EXCL, S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
    if(fd==-1){
        printf("error happend in creating the file %s (may be it exists) or opening file!\n.", filename);
    }
    else if (fd > 0){
        printf("file successfully created (it didn't exist) and opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```



.hfani@alpha:~\$ rm open_rw_test.txt

hfani@alpha:~\$./creat_excl_open_rw

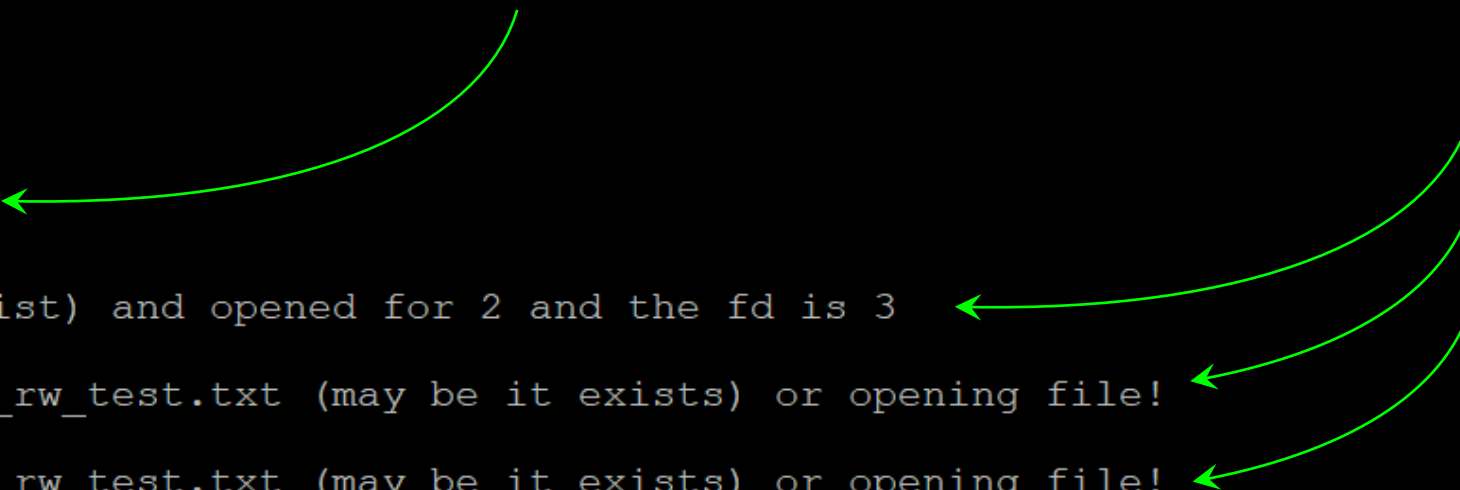
file successfully created (it didn't exist) and opened for 2 and the fd is 3

hfani@alpha:~\$./creat_excl_open_rw

error happend in creating the file open_rw_test.txt (may be it exists) or opening file!

.hfani@alpha:~\$./creat_excl_open_rw

error happend in creating the file open_rw_test.txt (may be it exists) or opening file!



O_APPEND Append to the end of file on each write

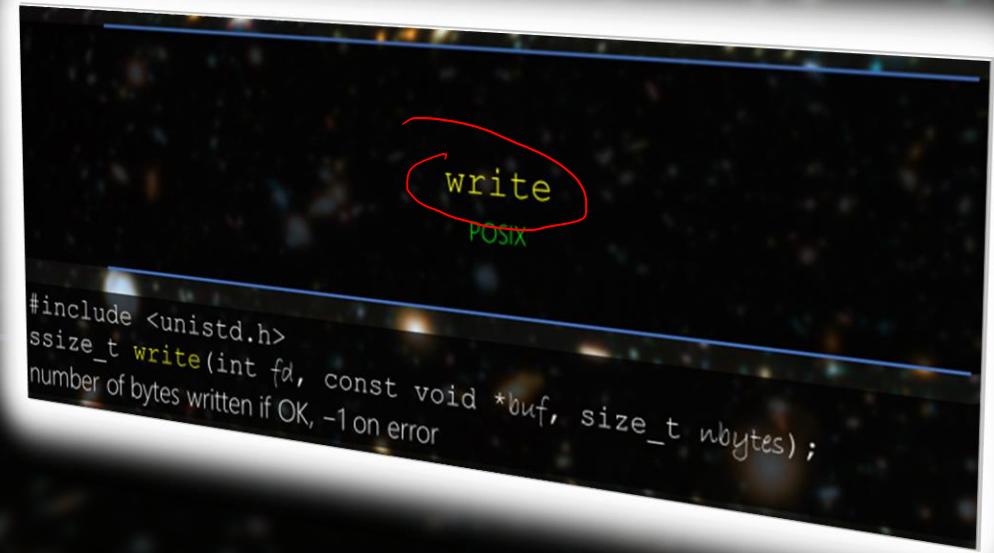
```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error

→ read
POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/read.html>

```
#include <unistd.h>
ssize_t read(int fd, const void *buf, size_t nbytes);
number of bytes written if OK, -1 on error
```

read

POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/read.html>

```
#include <unistd.h>
ssize_t read(int fd, const void *buf, size_t nbytes);
number of bytes written if OK, -1 on error
```



```
fd = creat()
```

```
write(fd, sth)
```

```
close(fd)
```

```
fd = open(creat if not exists)
```

```
read(fd, sth)
```

```
write(fd, sth)
```

```
close(fd)
```



You Want me on that Wall, You Need me on that Wall

<https://www.youtube.com/watch?v=xd1fs5nWol>

A Few Good Men - Aaron Sorkin

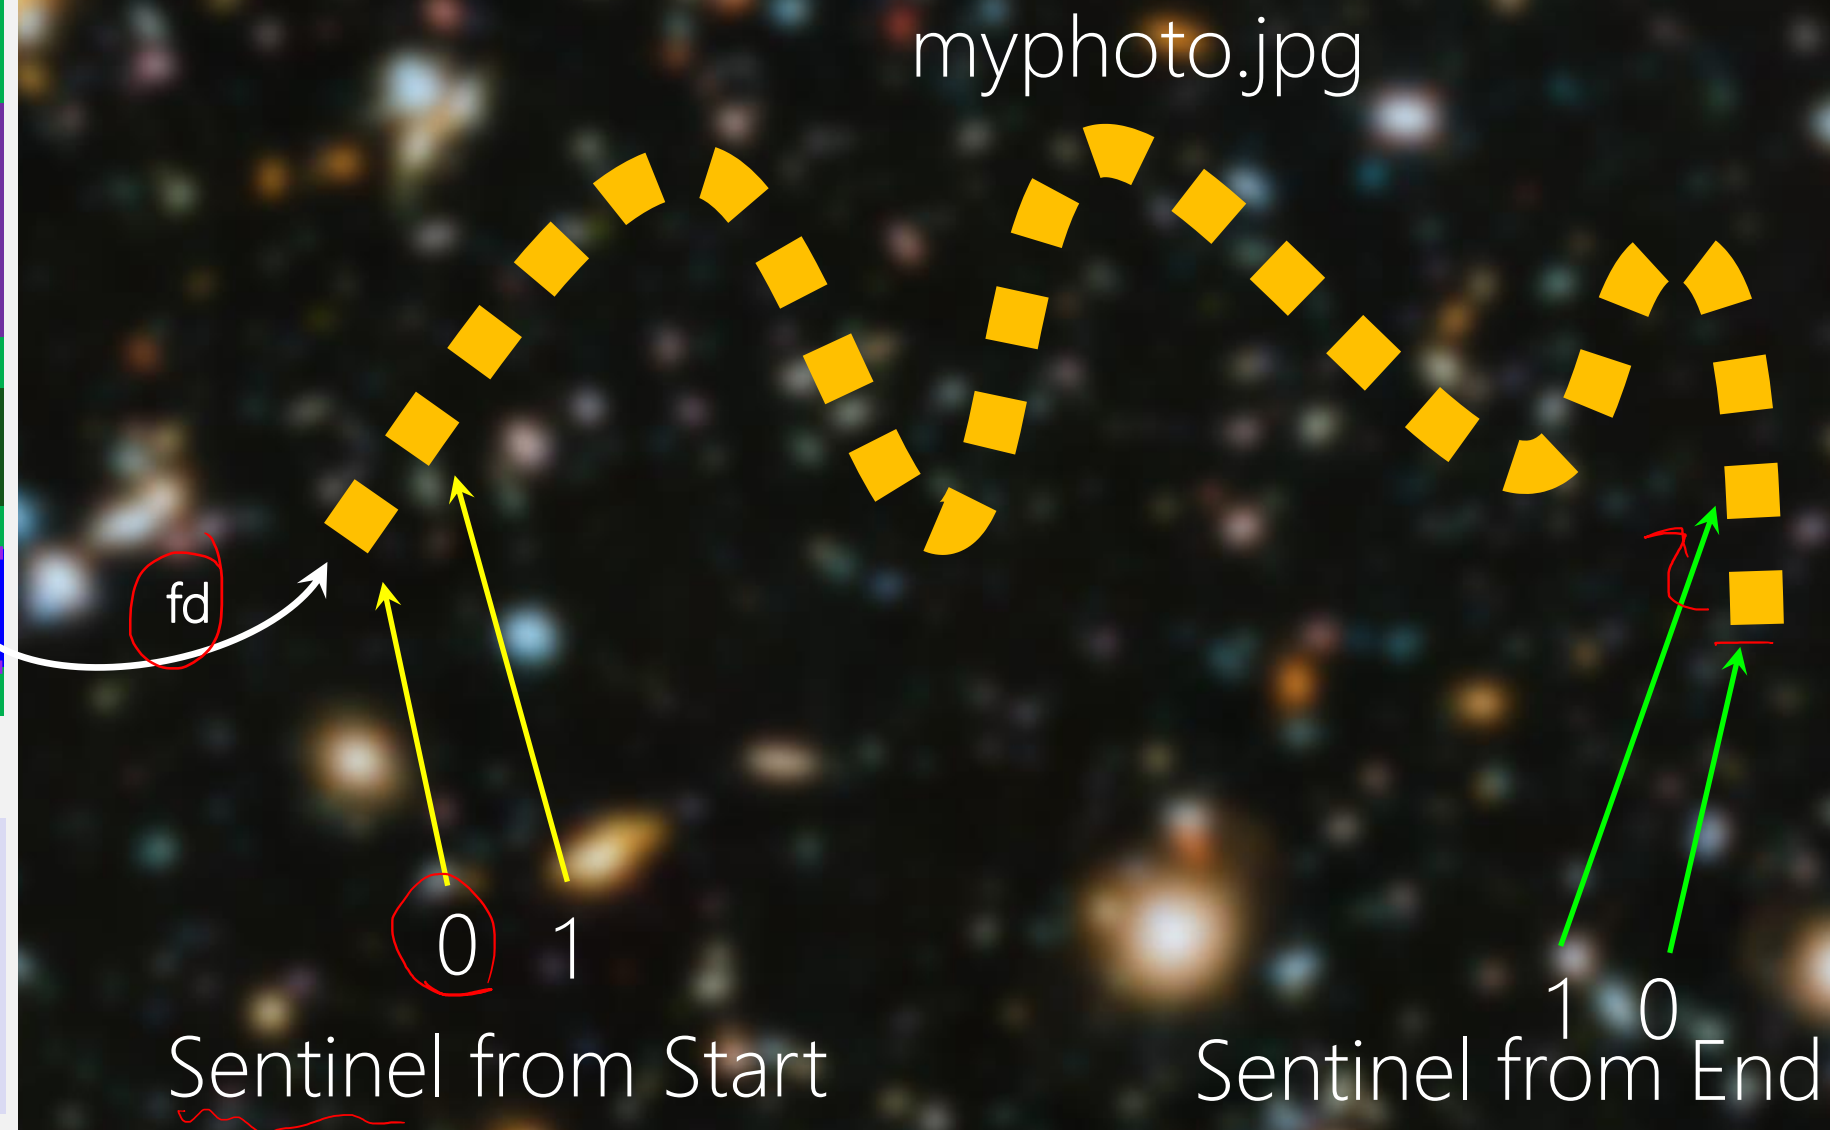
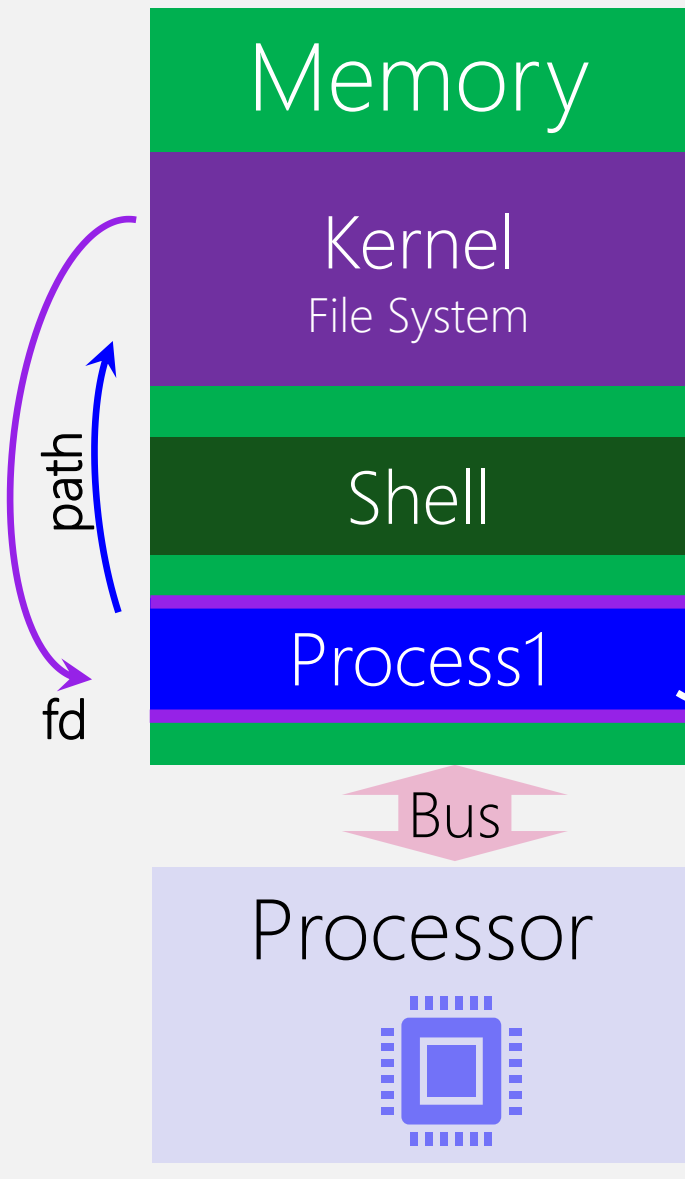
lseek

POSIX

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
```

file's new offset if OK (can be negative)
-1 on error

Computer



Every opened file has one sentinel: current offset
Measures the number of bytes from the beginning of the file.

Every opened file has one sentinel: `current offset`
Measures the number of bytes from the beginning of the file.

`creat()` or `open()`



Every opened file has one sentinel: `current offset`
Measures the number of bytes from the beginning of the file.

`creat()` or `open()` → 0

Every opened file has one sentinel: `current offset`
Measures the number of bytes from the beginning of the file.

`read()` or `write()` → ++ actual number of bytes read or written

`read()` or `write()` → always move forward

myphoto.jpg

fd

current offset: 11

lseek

POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/lseek.html>

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```

lseek

POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/lseek.html>

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
```

file's new offset if OK (can be negative)
-1 on error

lseek

How many bytes move the current offset

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```

```
#include <sys/types.h>
typedef off_t signed long
```


lseek


How many bytes move the current offset from what place or origin?

SEEK_SET, SEEK_CUR, SEEK_END

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
```

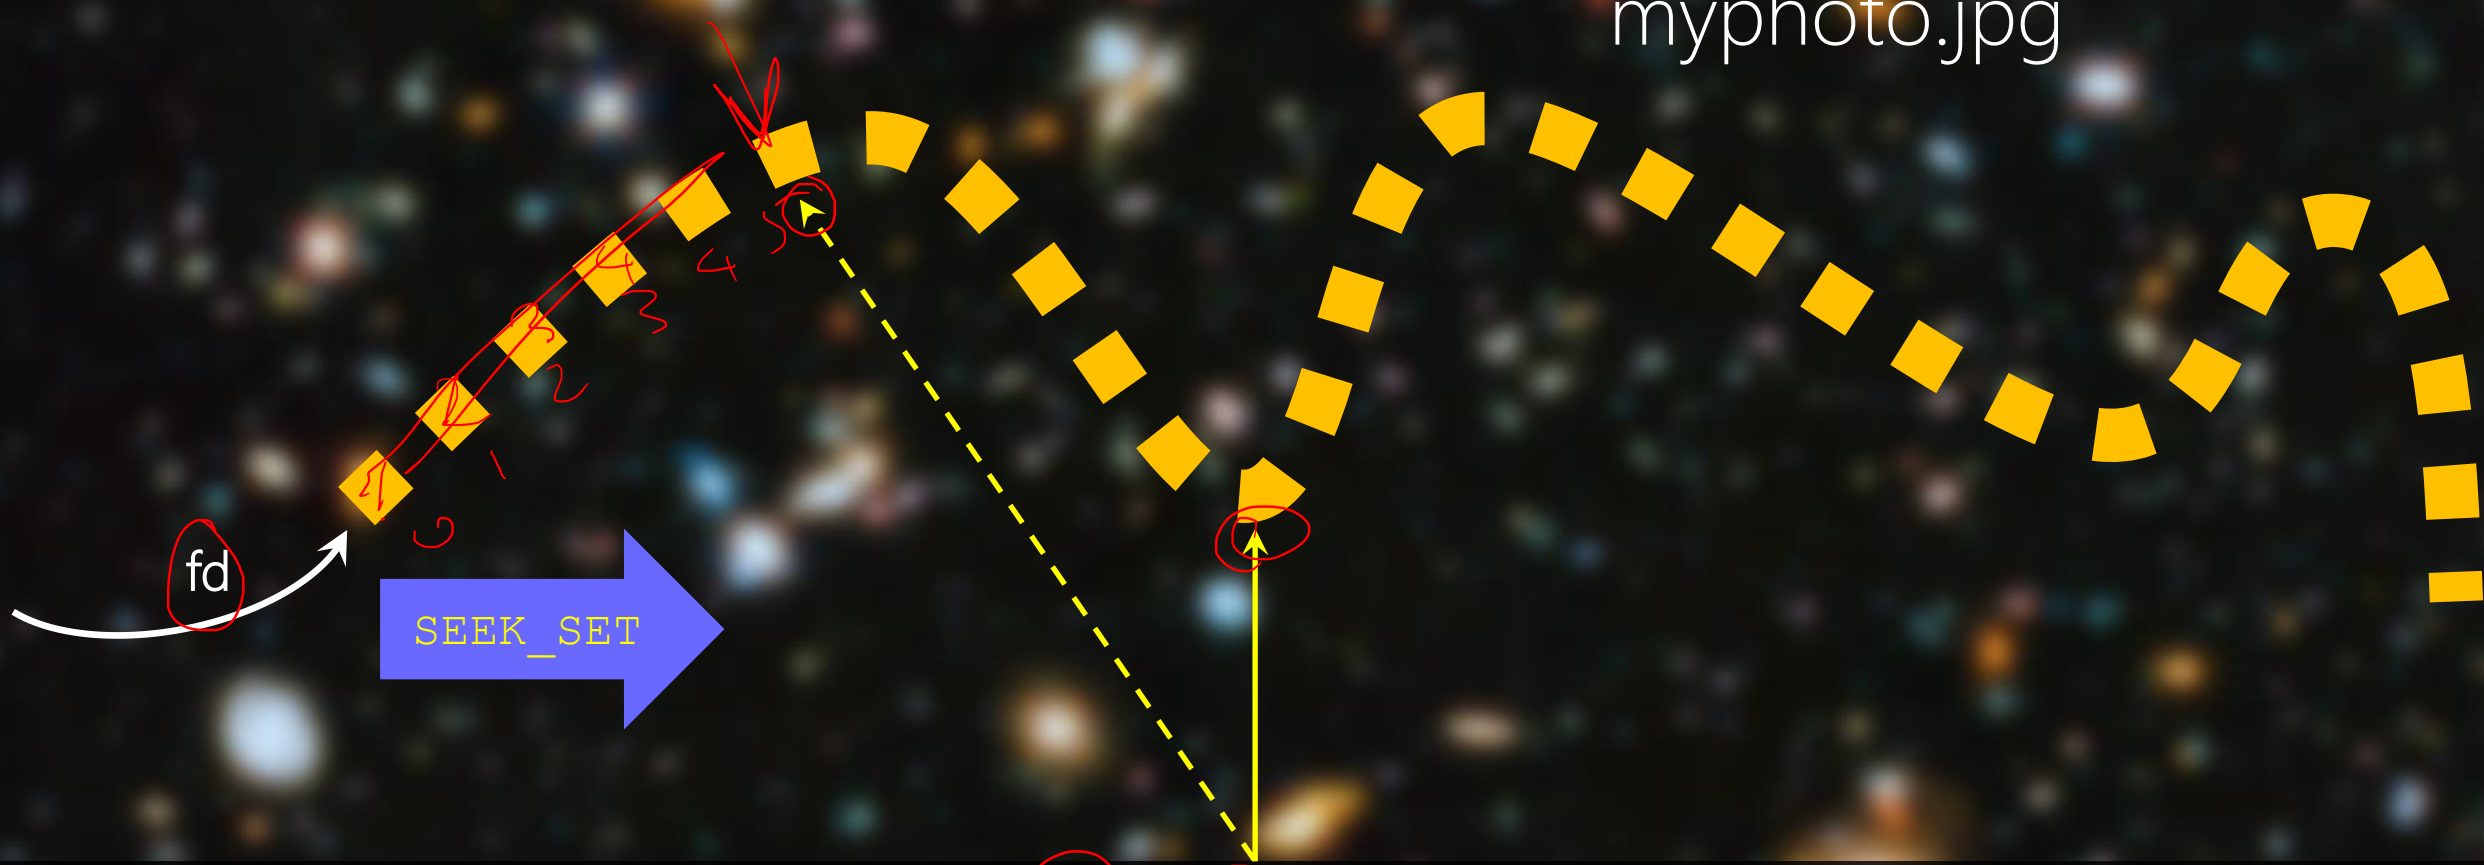
file's new offset if OK (can be negative)
-1 on error

myphoto.jpg



current offset: 11


myphoto.jpg



lseek(fd, 5, SEEK_SET)

current offset: 11 → 5


myphoto.jpg



current offset: 11 → 25

`lseek(fd, 5, SEEK_END)`


myphoto.jpg



current offset: $11 \rightarrow +5 = 16$

`lseek(fd, 5, SEEK_CUR)`

myphoto.jpg



current offset: 11 \rightarrow -5 = 6

`lseek(fd, -5, SEEK_CUR)`

lseek

New current offset

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
```

file's new offset if OK (can be negative)
-1 on error



lseek

How to know the value of current offset?



lseek

How to know the value of current offset?

```
long cur_offset;  
cur_offset = lseek(fd, 0, SEEK_CUR);
```

lseek

How to know the file can be seekable?

```
long cur_offset;  
cur_offset = lseek(fd, 0, SEEK_CUR);  
if (cur_offset == -1){  
    /* either the fd is on a file which is not seekable  
       or an error occurred */  
}
```