



You Can't Handle The Truth, A Few Good Men (1992) - Aaron Sorkin

Keys to Midterm Exam

Exams



Midterm Exam

[Midterm Exam Instructions.pdf](#) ▼

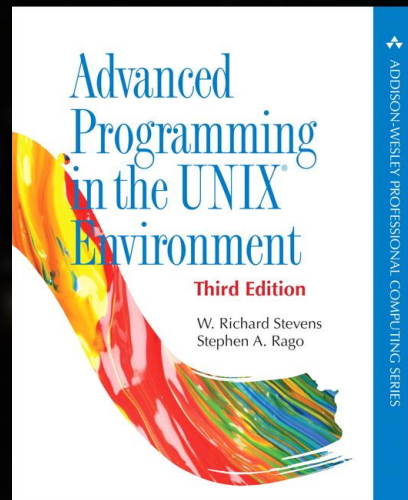
[Midterm Exam I Keys.pdf](#) ▼

[Midterm Exam II Keys.pdf](#) ▼

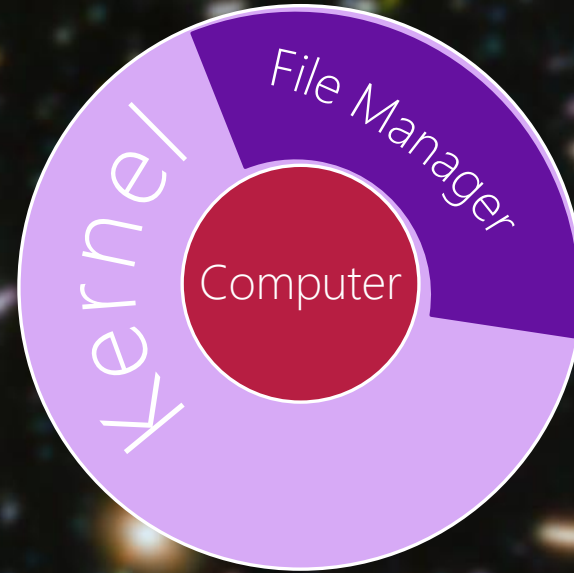
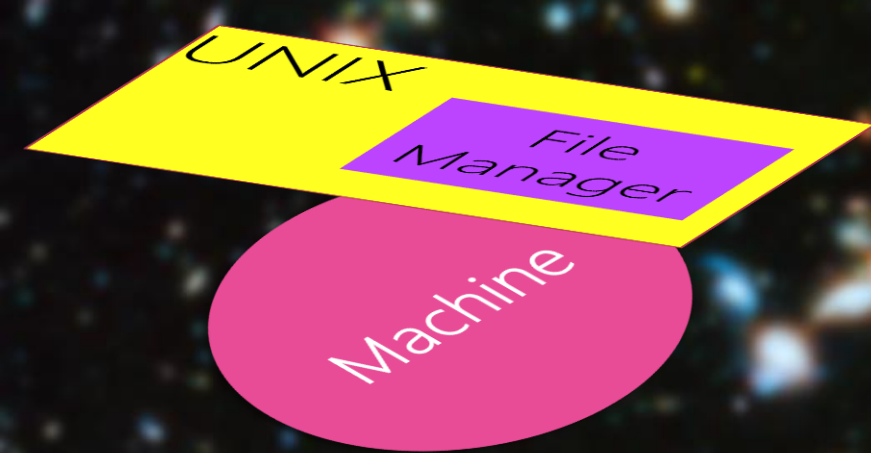
A deep-field astronomical image showing a vast field of galaxies in various colors (blue, orange, white) against a black background. Two horizontal blue lines are positioned above and below the central text.

Survey

<https://forms.gle/RcTadsRcgJExWSd66>



Chapter 03: File I/O



Computer

Memory

Kernel: Device Manager

Kernel: Memory Manager

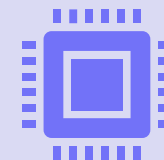
Kernel: File Manager

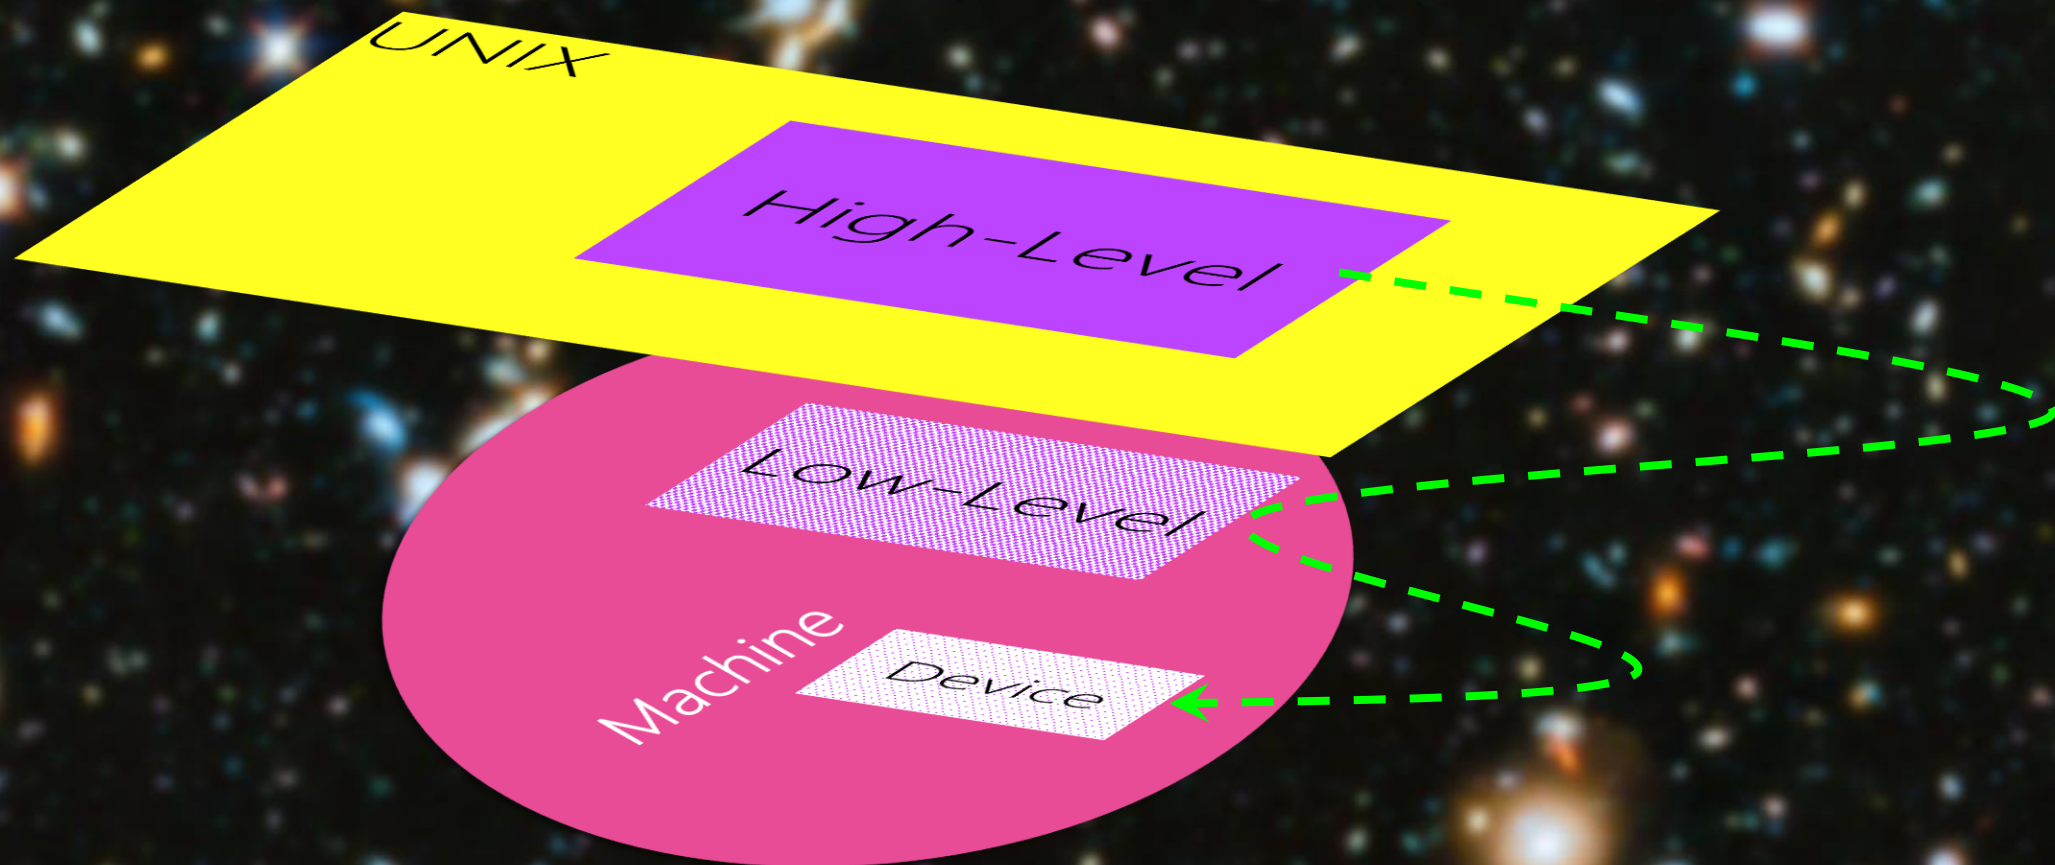
Kernel: Network Manager

Kernel: Process Manager

Bus

Processor







File Manager

widely known as File System

High-Level



Device is a Single 1-D Array (String) of Bytes

Even Memory and Processor!

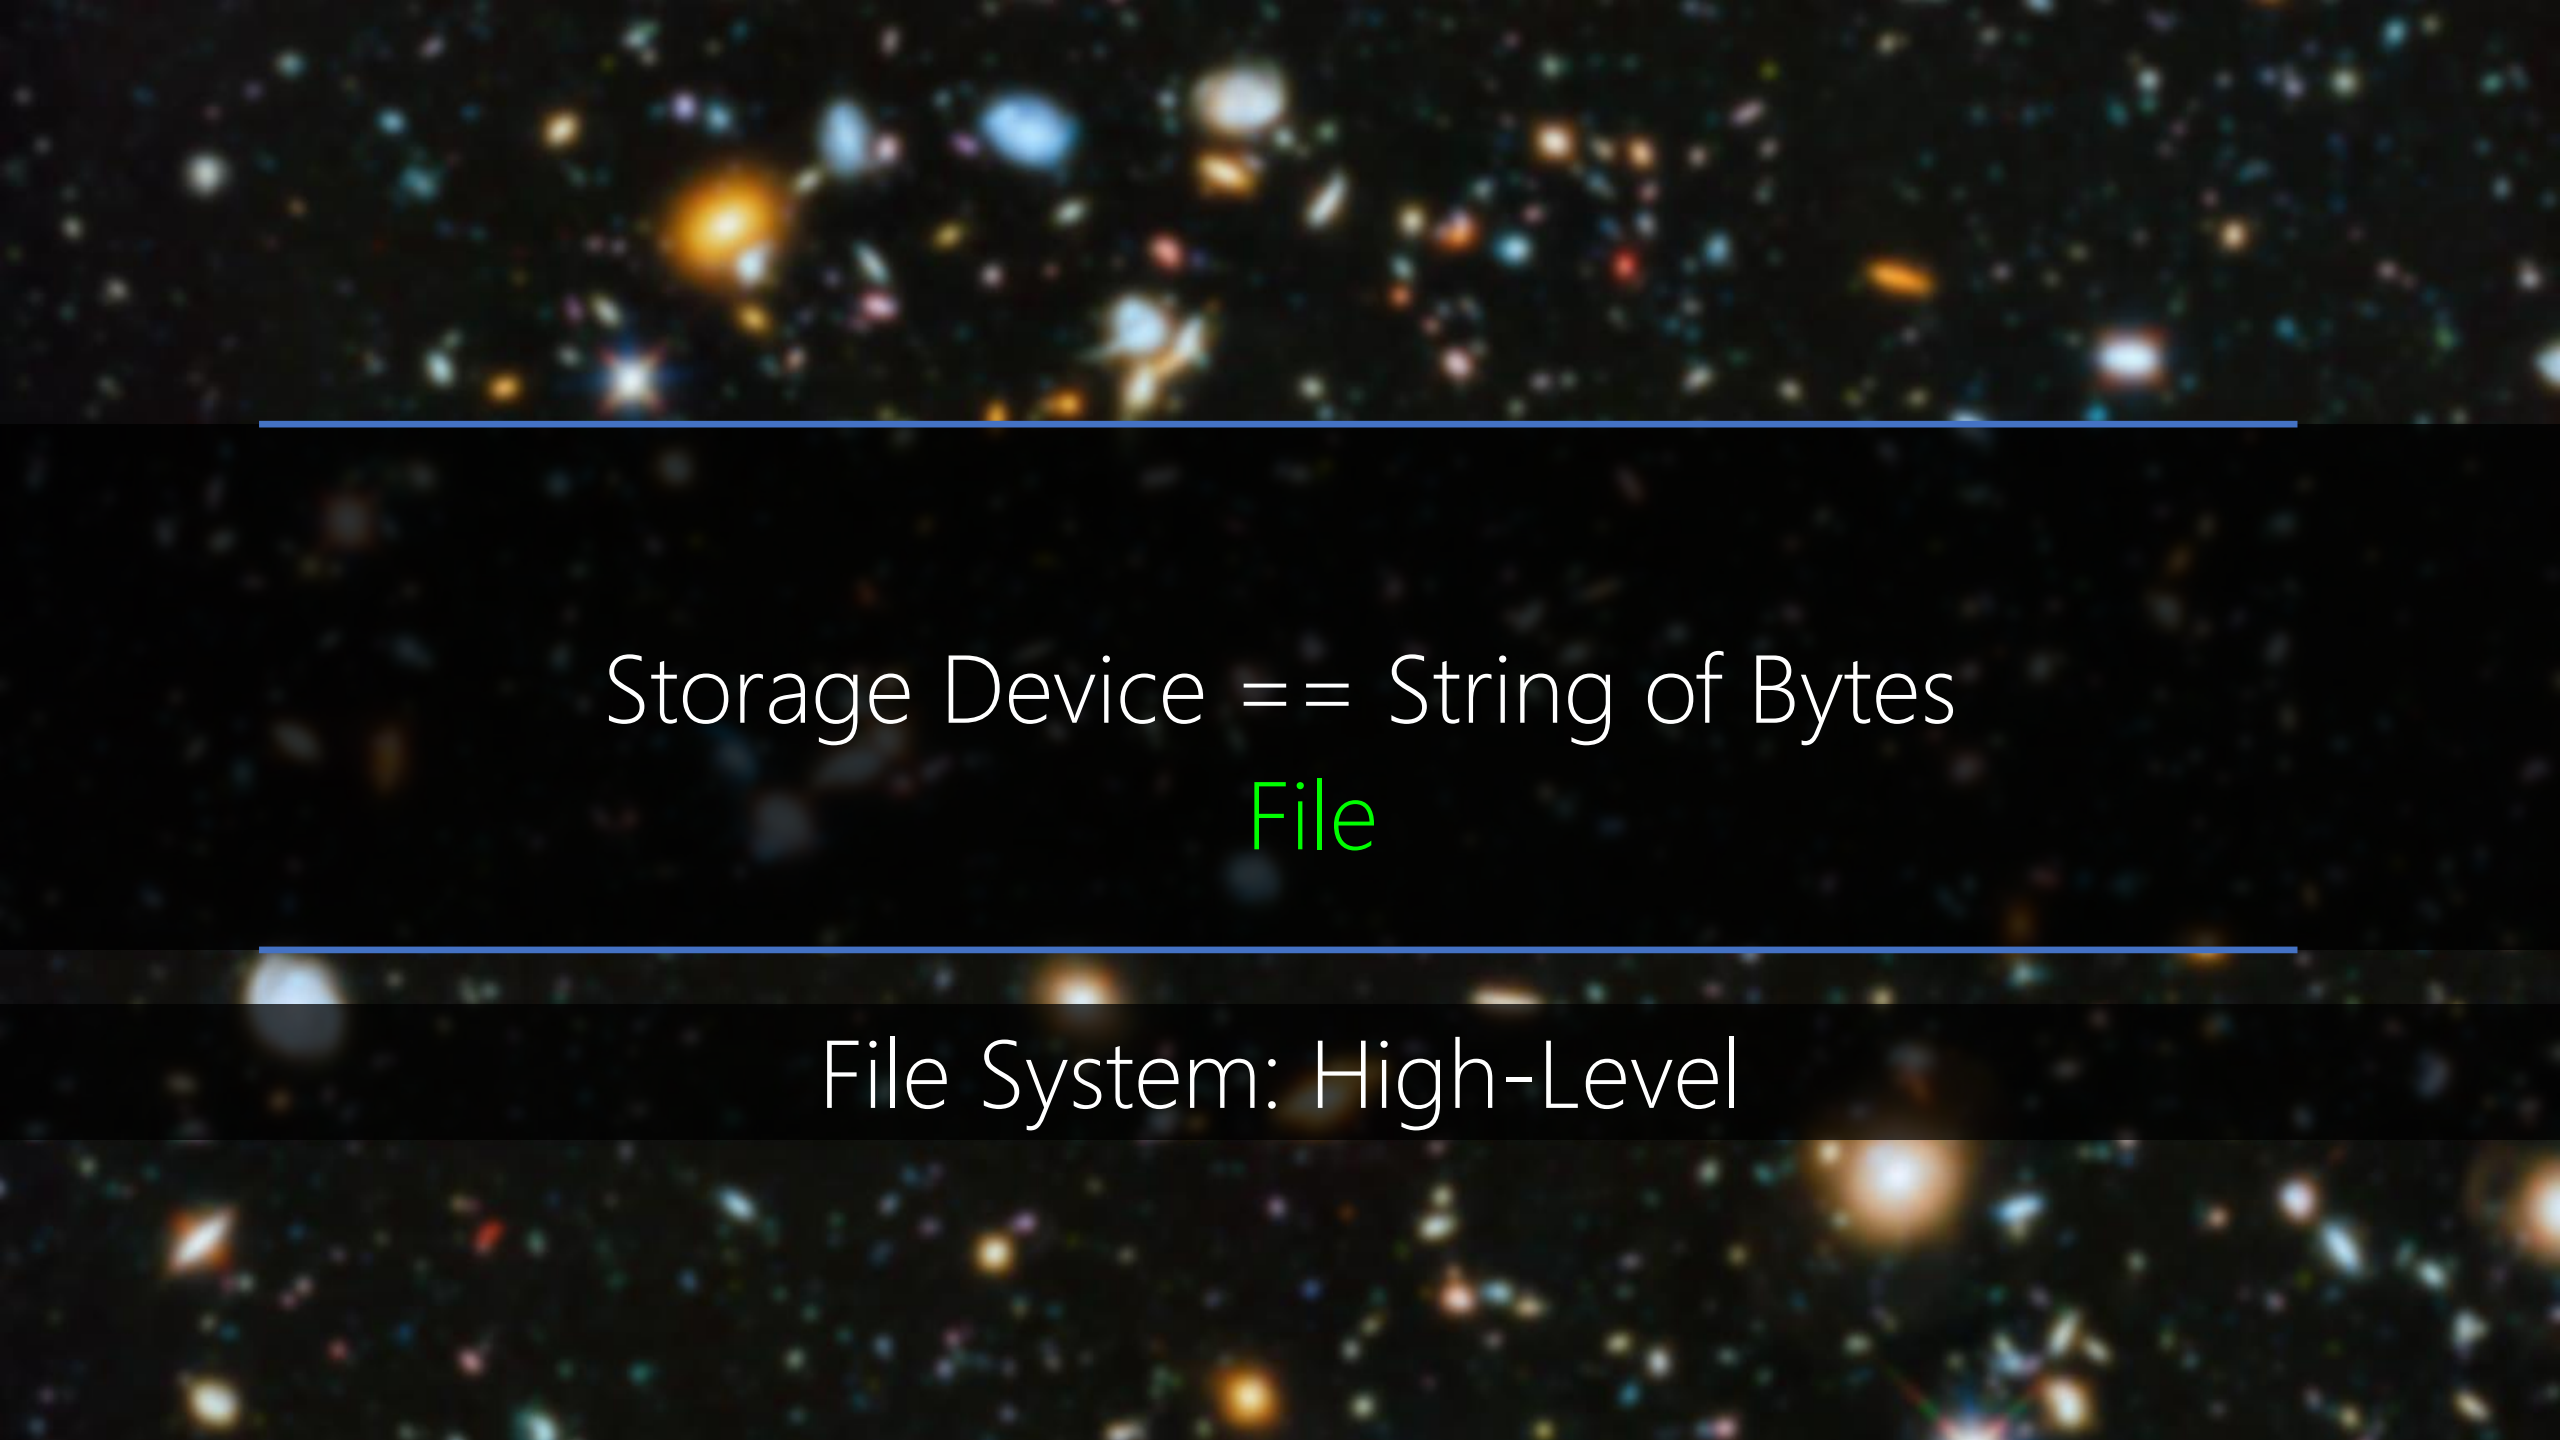
File System: High-Level



Device is a Single 1-D Array (String) of Bytes

Please give up memory & processor.
Leave them for Process Manager!

File System: High-Level

The background of the slide is a deep space image showing numerous galaxies in various colors (blue, orange, white) against a black sky. A solid blue horizontal bar spans the width of the slide, positioned behind the text.

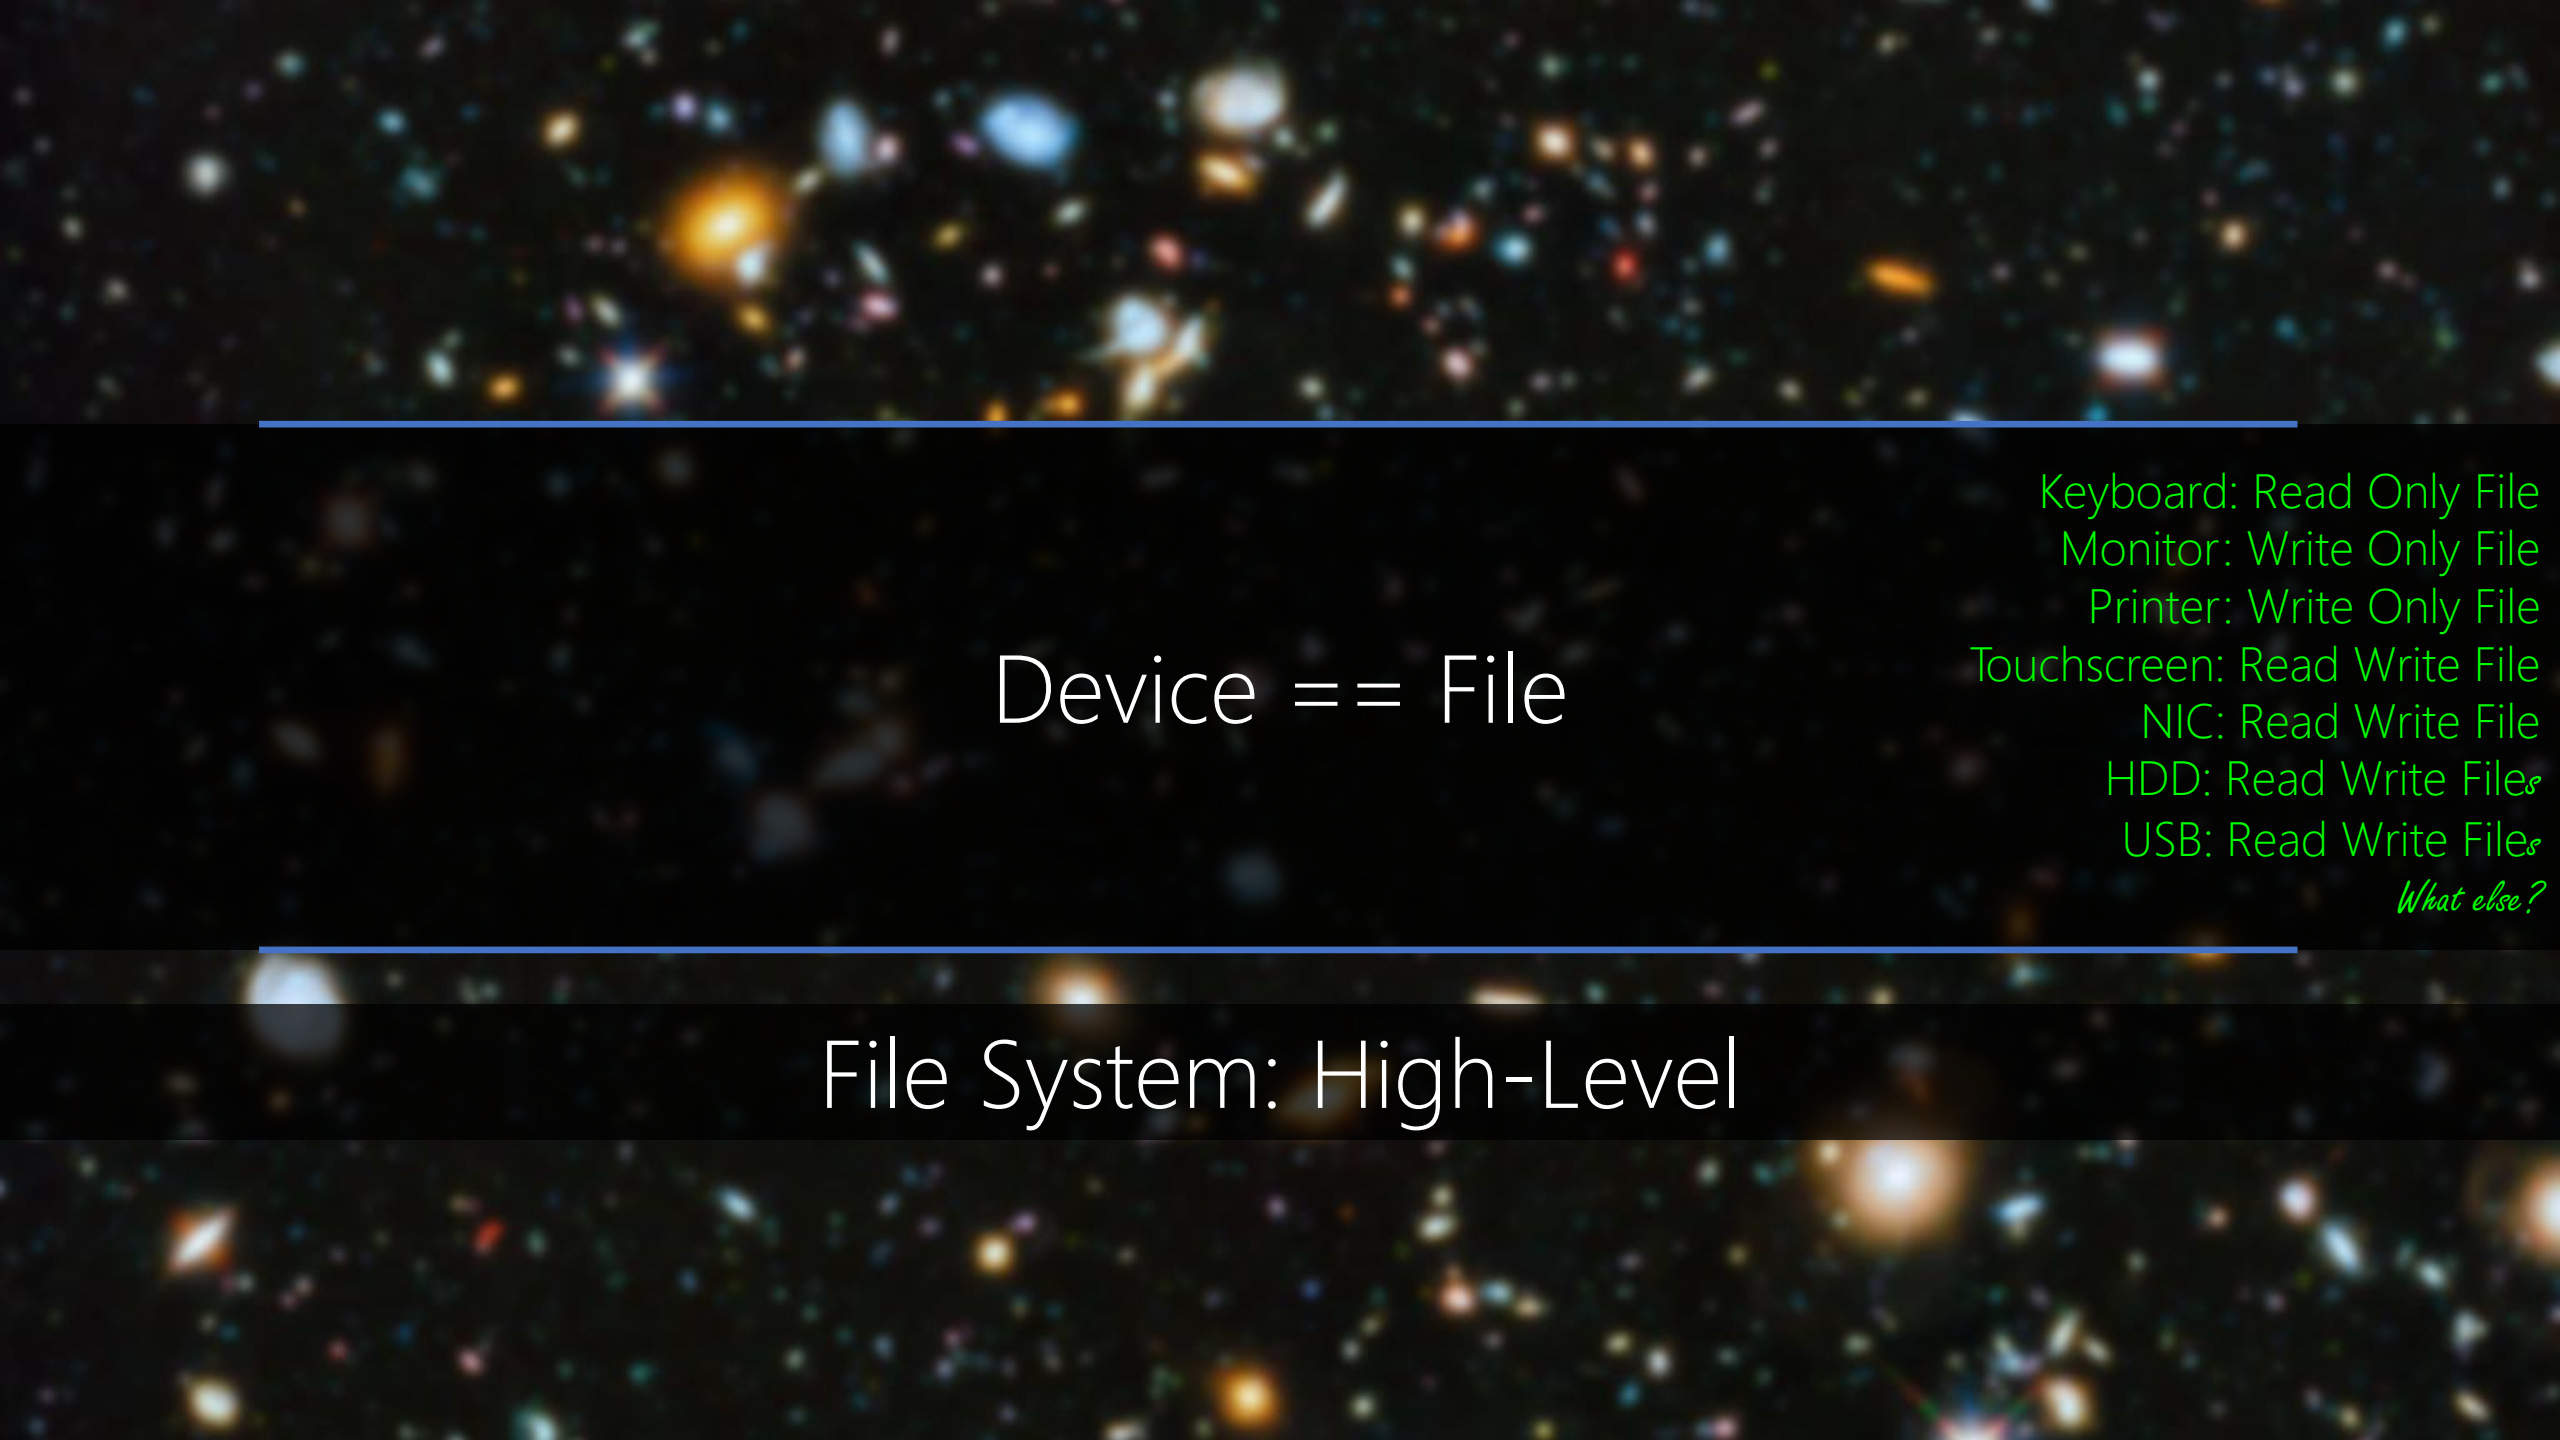
Storage Device == String of Bytes
File

File System: High-Level



Large Storage Device == *Set of Files*
set of sub-devices

File System: High-Level



Device == File

Keyboard: Read Only File
Monitor: Write Only File
Printer: Write Only File
Touchscreen: Read Write File
NIC: Read Write File
HDD: Read Write Files
USB: Read Write Files
What else?

File System: High-Level

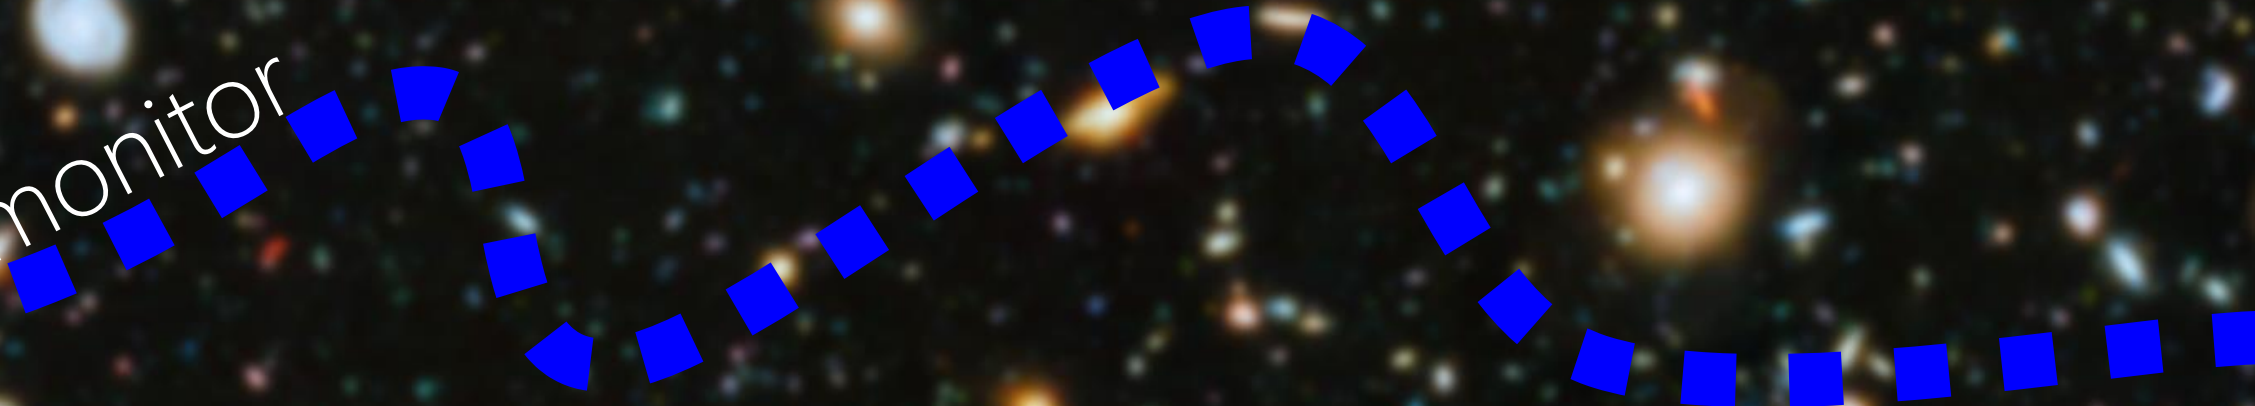
myphoto.jpg



keyboard



monitor





Operator - The Matrix (1999), Lana & Lilly Wachowski



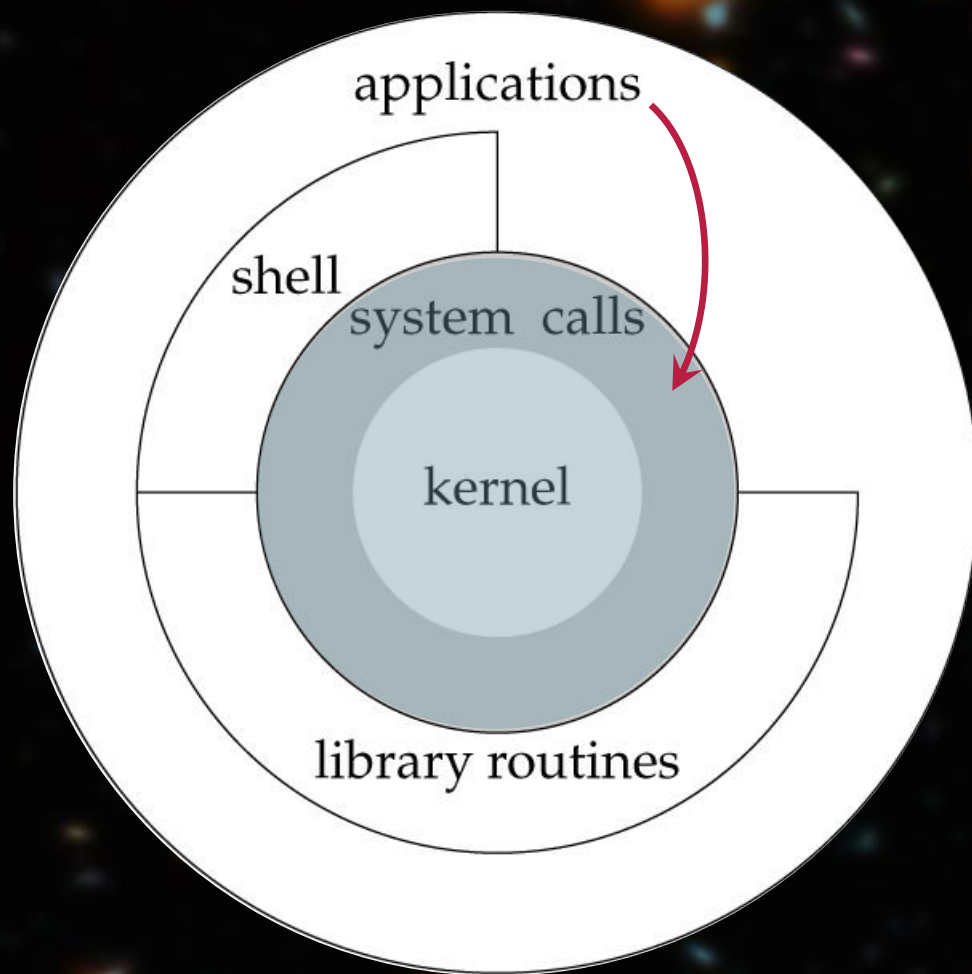
Operation

What do you expect from a kernel about string of bytes | device | file?

File System: High-Level

Create a New One
Open an Existing One
Write to an Opened One
Read from an Opened One
Move Forward/Backward in an Opened One
Delete an Existing One
Check the Existence of One
Hide an Existing One
Prevent Others to Open an Existing One
Prevent Others to Write to an Existing One
What else?

File System: High-Level



Header	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
< aio.h>	•	•	•	•	asynchronous I/O
< cpio.h>	•	•	•	•	cpio archive values
< dirent.h>	•	•	•	•	directory entries (Section 4.22)
< dlfcn.h>	•	•	•	•	dynamic linking
< fcntl.h>	•	•	•	•	file control (Section 3.14)
< fnmatch.h>	•	•	•	•	filename-matching types
< glob.h>	•	•	•	•	pathname pattern-matching and generation
< grp.h>	•	•	•	•	group file (Section 6.4)
< iconv.h>	•	•	•	•	codeset conversion utility
< langinfo.h>	•	•	•	•	language information constants
< monetary.h>	•	•	•	•	monetary types and functions
< netdb.h>	•	•	•	•	network database operations
< nl_types.h>	•	•	•	•	message catalogs
< poll.h>	•	•	•	•	poll function (Section 14.4.2)
< pthread.h>	•	•	•	•	threads (Chapters 11 and 12)
< pwd.h>	•	•	•	•	password file (Section 6.2)
< regex.h>	•	•	•	•	regular expressions
< sched.h>	•	•	•	•	execution scheduling
< semaphore.h>	•	•	•	•	semaphores
< strings.h>	•	•	•	•	string operations
< tar.h>	•	•	•	•	tar archive values
< termios.h>	•	•	•	•	terminal I/O (Chapter 18)
< unistd.h>	•	•	•	•	symbolic constants
< wordexp.h>	•	•	•	•	word-expansion definitions
< arpa/inet.h>	•	•	•	•	Internet definitions (Chapter 16)
< net/if.h>	•	•	•	•	socket local interfaces (Chapter 16)
< netinet/in.h>	•	•	•	•	Internet address family (Section 16.3)
< netinet/tcp.h>	•	•	•	•	Transmission Control Protocol definitions
< sys/mman.h>	•	•	•	•	memory management declarations
< sys/select.h>	•	•	•	•	select function (Section 14.4.1)
< sys/socket.h>	•	•	•	•	sockets interface (Chapter 16)
< sys/stat.h>	•	•	•	•	file status (Chapter 4)
< sys/statvfs.h>	•	•	•	•	file system information
< sys/times.h>	•	•	•	•	process times (Section 8.17)
< sys/types.h>	•	•	•	•	primitive system data types (Section 2.8)
< sys/un.h>	•	•	•	•	UNIX domain socket definitions (Section 17.2)
< sys/utsname.h>	•	•	•	•	system name (Section 6.9)
< sys/wait.h>	•	•	•	•	process control (Section 8.6)

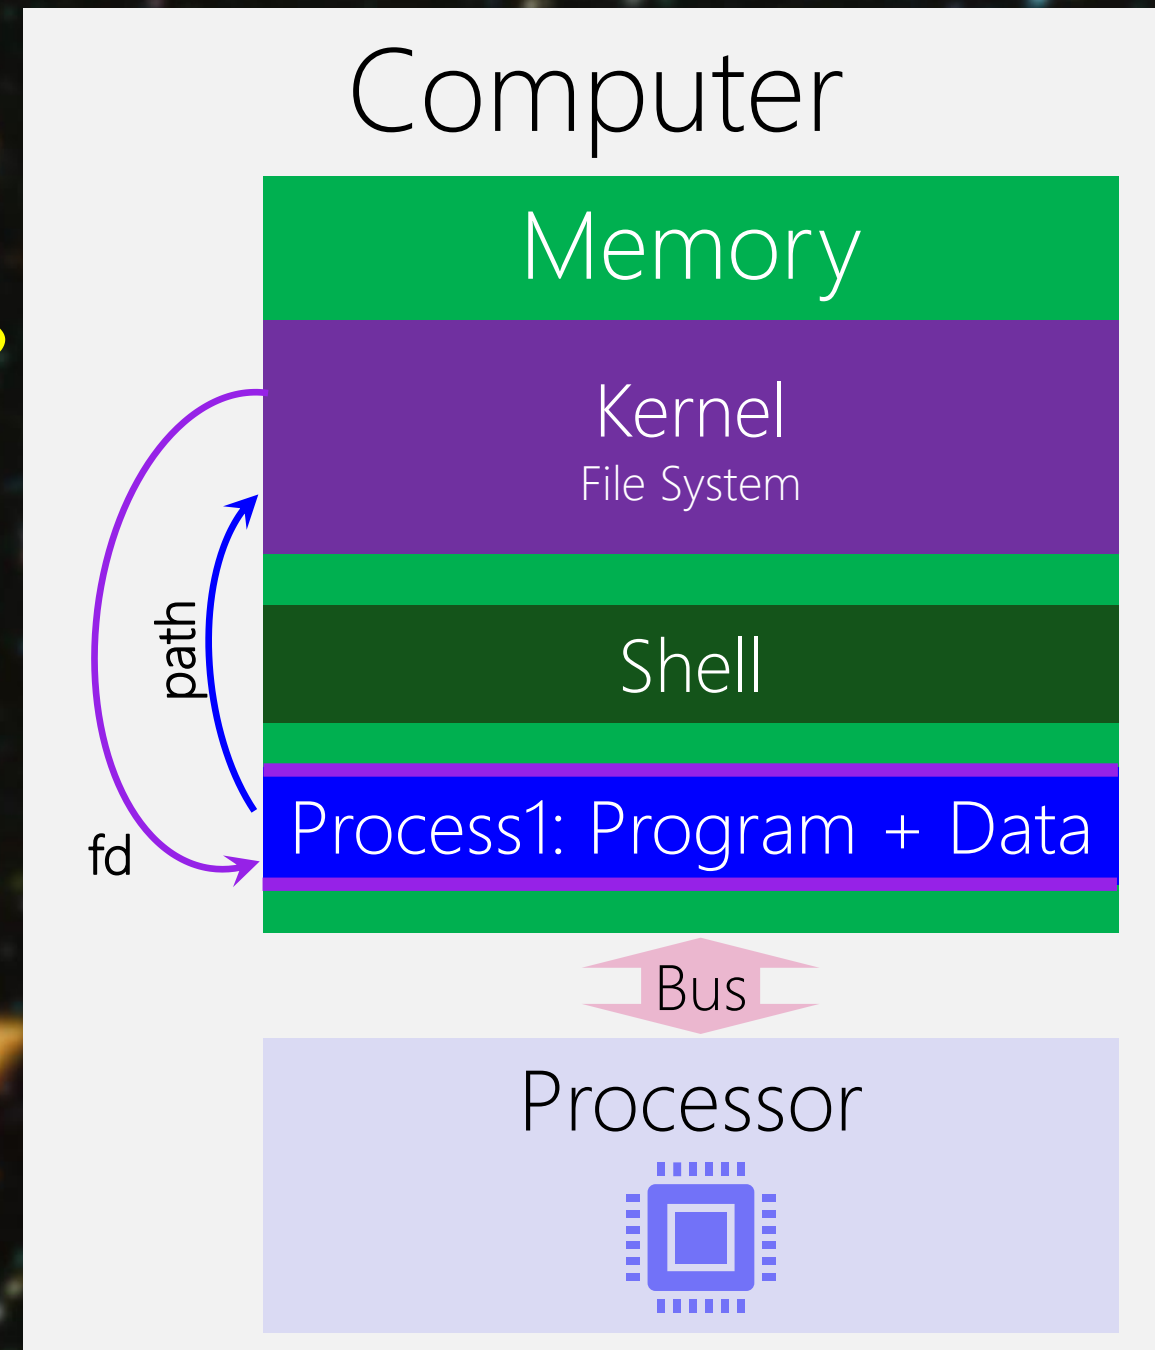
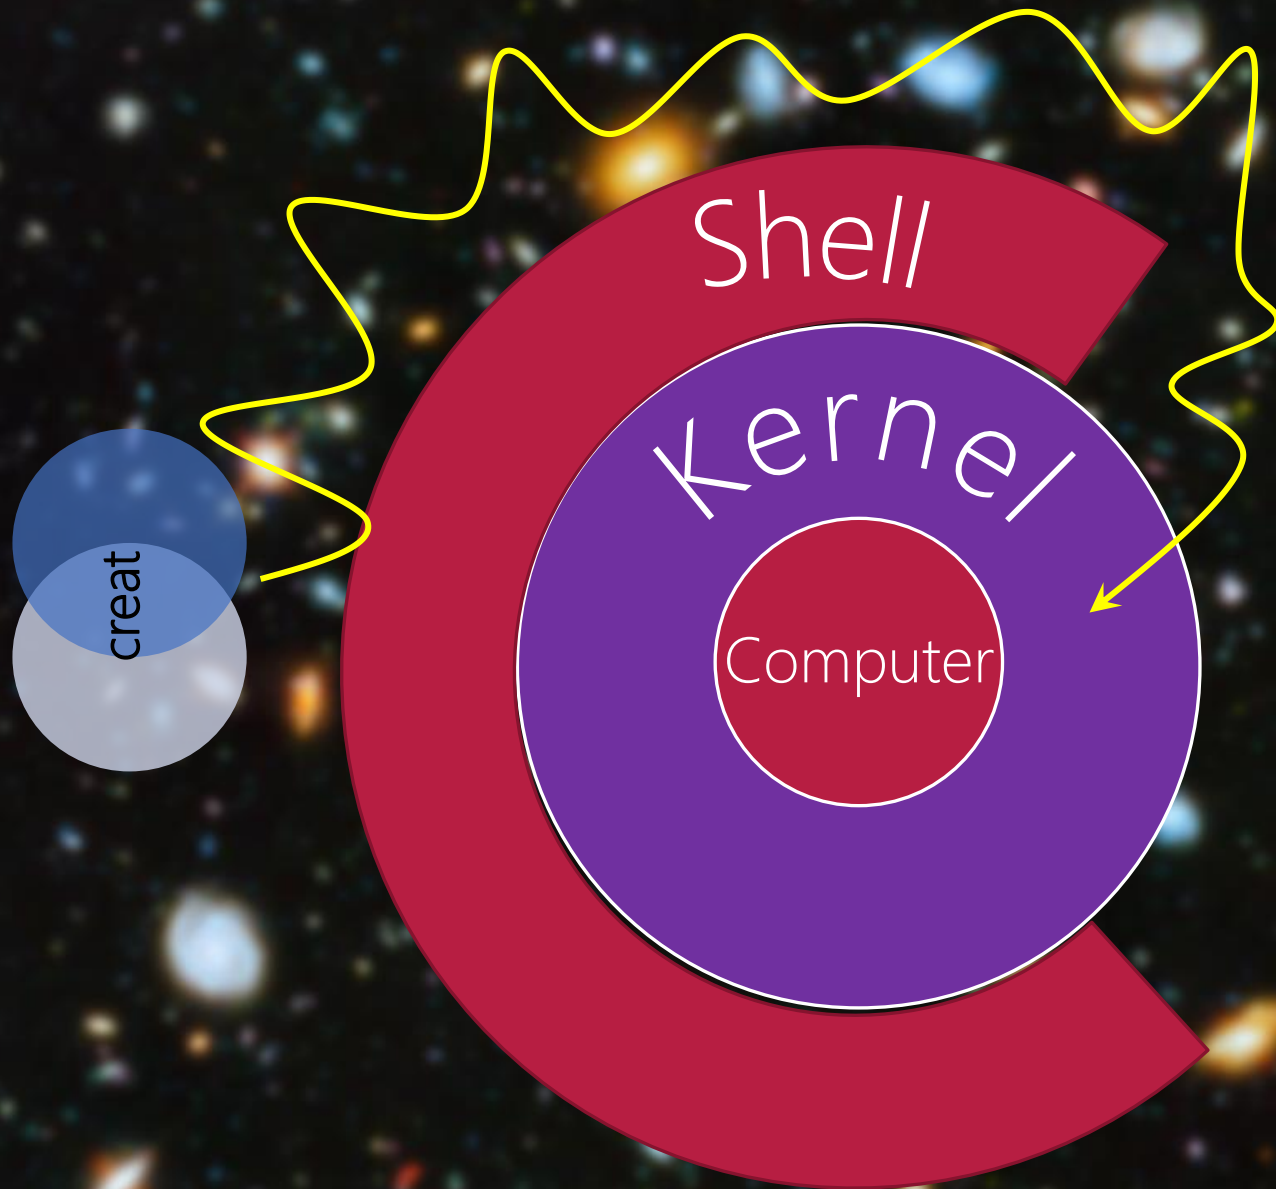


creat

POSIX

```
#include <fcntl.h>
int creat(const char *path, mode_t mode);
```

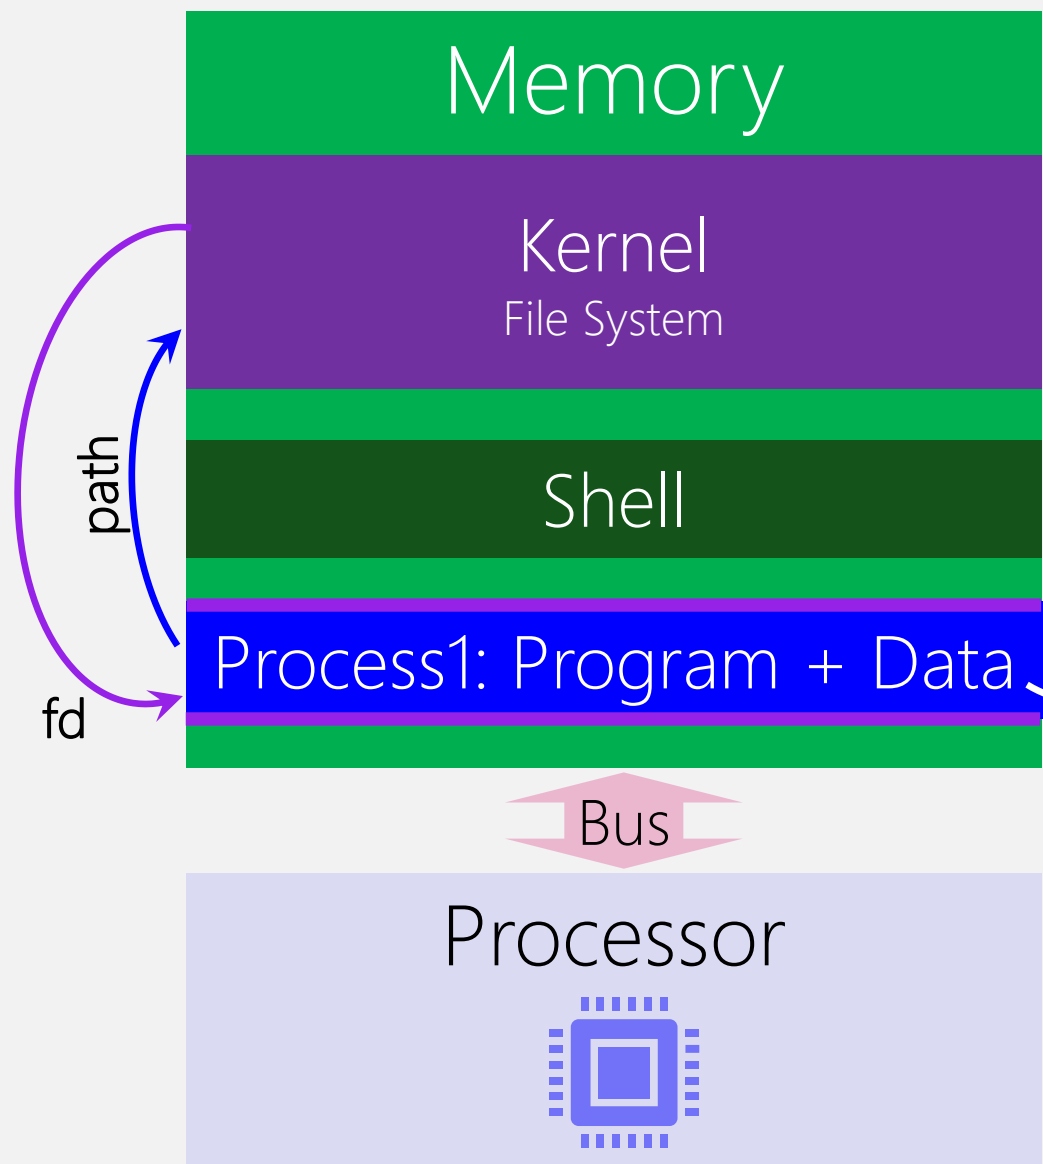
non-negative number for write-only if OK
-1 on error





Trinity escapes from Agents
- The Matrix (1999), Lana & Lilly Wachowski

Computer





File Descriptor (fd)

Number does not Matter, Connection Matters

The Only Way In Or Out Is Through Phone Lines
The number does not matter, the connection is important!
Imagine a dynamic phone#, dynamic postal code, dynamic ip (DHCP)

A cosmic background image featuring a dense field of galaxies in various colors (yellow, orange, blue, red) against a black space. Two horizontal blue lines are positioned above and below the central text.

File Descriptor (fd) != File Identifier

Because kernel reuse them for other files and devices, when available!

File Descriptor (fd)

$fd \in [0 : OPEN_MAX - 1]$

unistd.h

- #define OPEN_MAX 20
- #define OPEN_MAX 63
- No limit, maximum integer number supported by the system

File Descriptor (fd)

STDIN_FILENO, STDOUT_FILENO, STDERR_FILENO

```
unistd.h
```

```
#define STDIN_FILENO 0
```



Keyboard, Mouse, File, ...

```
#define STDOUT_FILENO 1
```



Monitor, Printer, File, ...

```
#define STDERR_FILENO 2
```



Monitor, Printer, File, ...



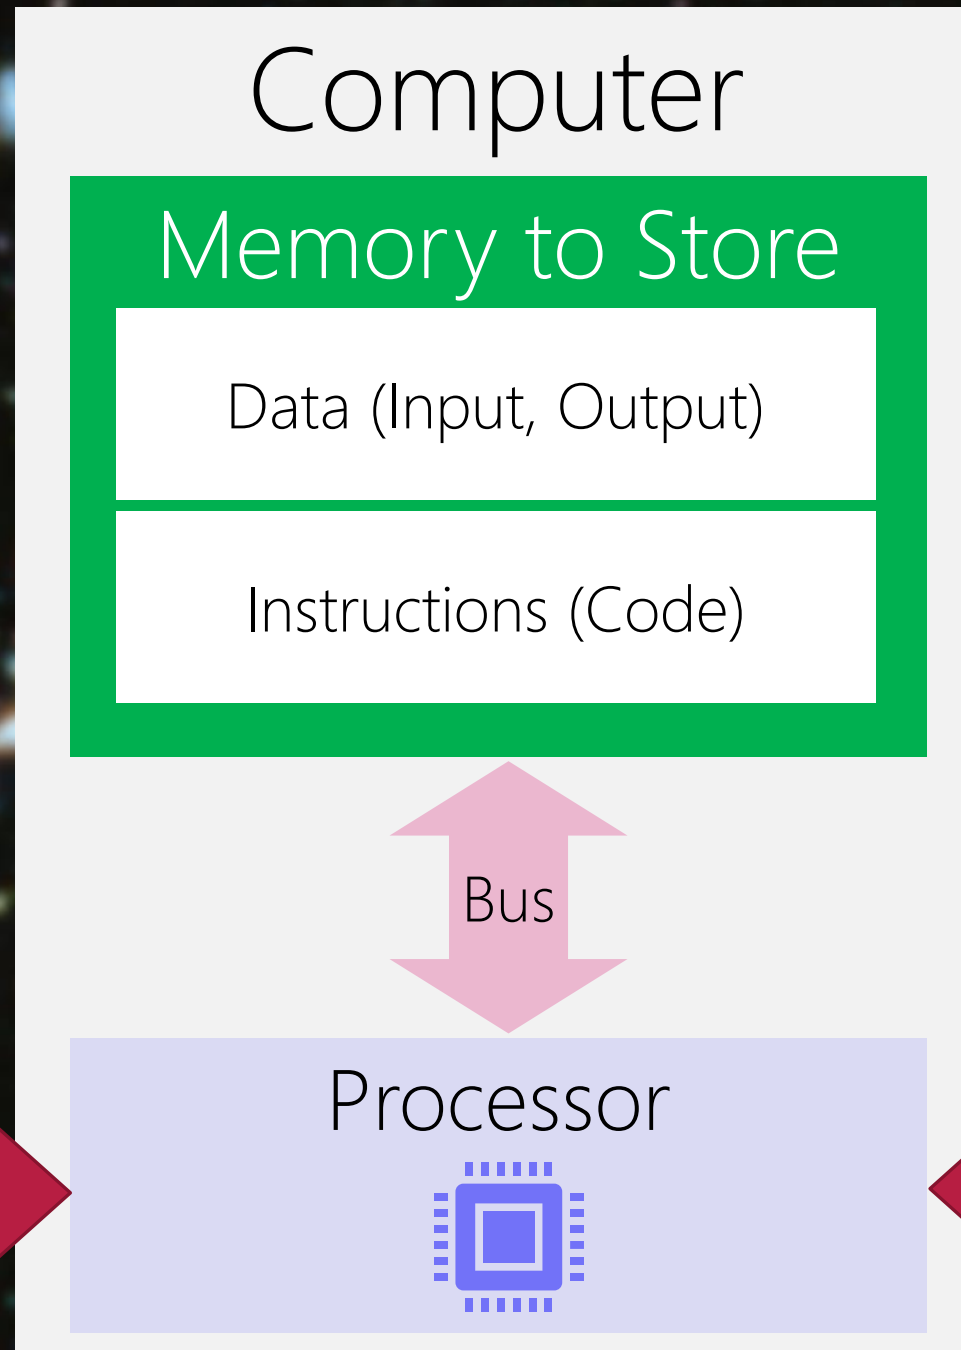
write

POSIX

```
#include <unistd.h>
ssize_t write(int fd, const void *buf, size_t nbytes);
number of bytes written if OK, -1 on error
```


Computer System

Input/Output
Devices



Computer

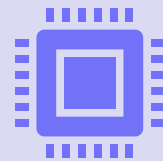
Memory to Store

Data (Input, Output)

Instructions (Code)

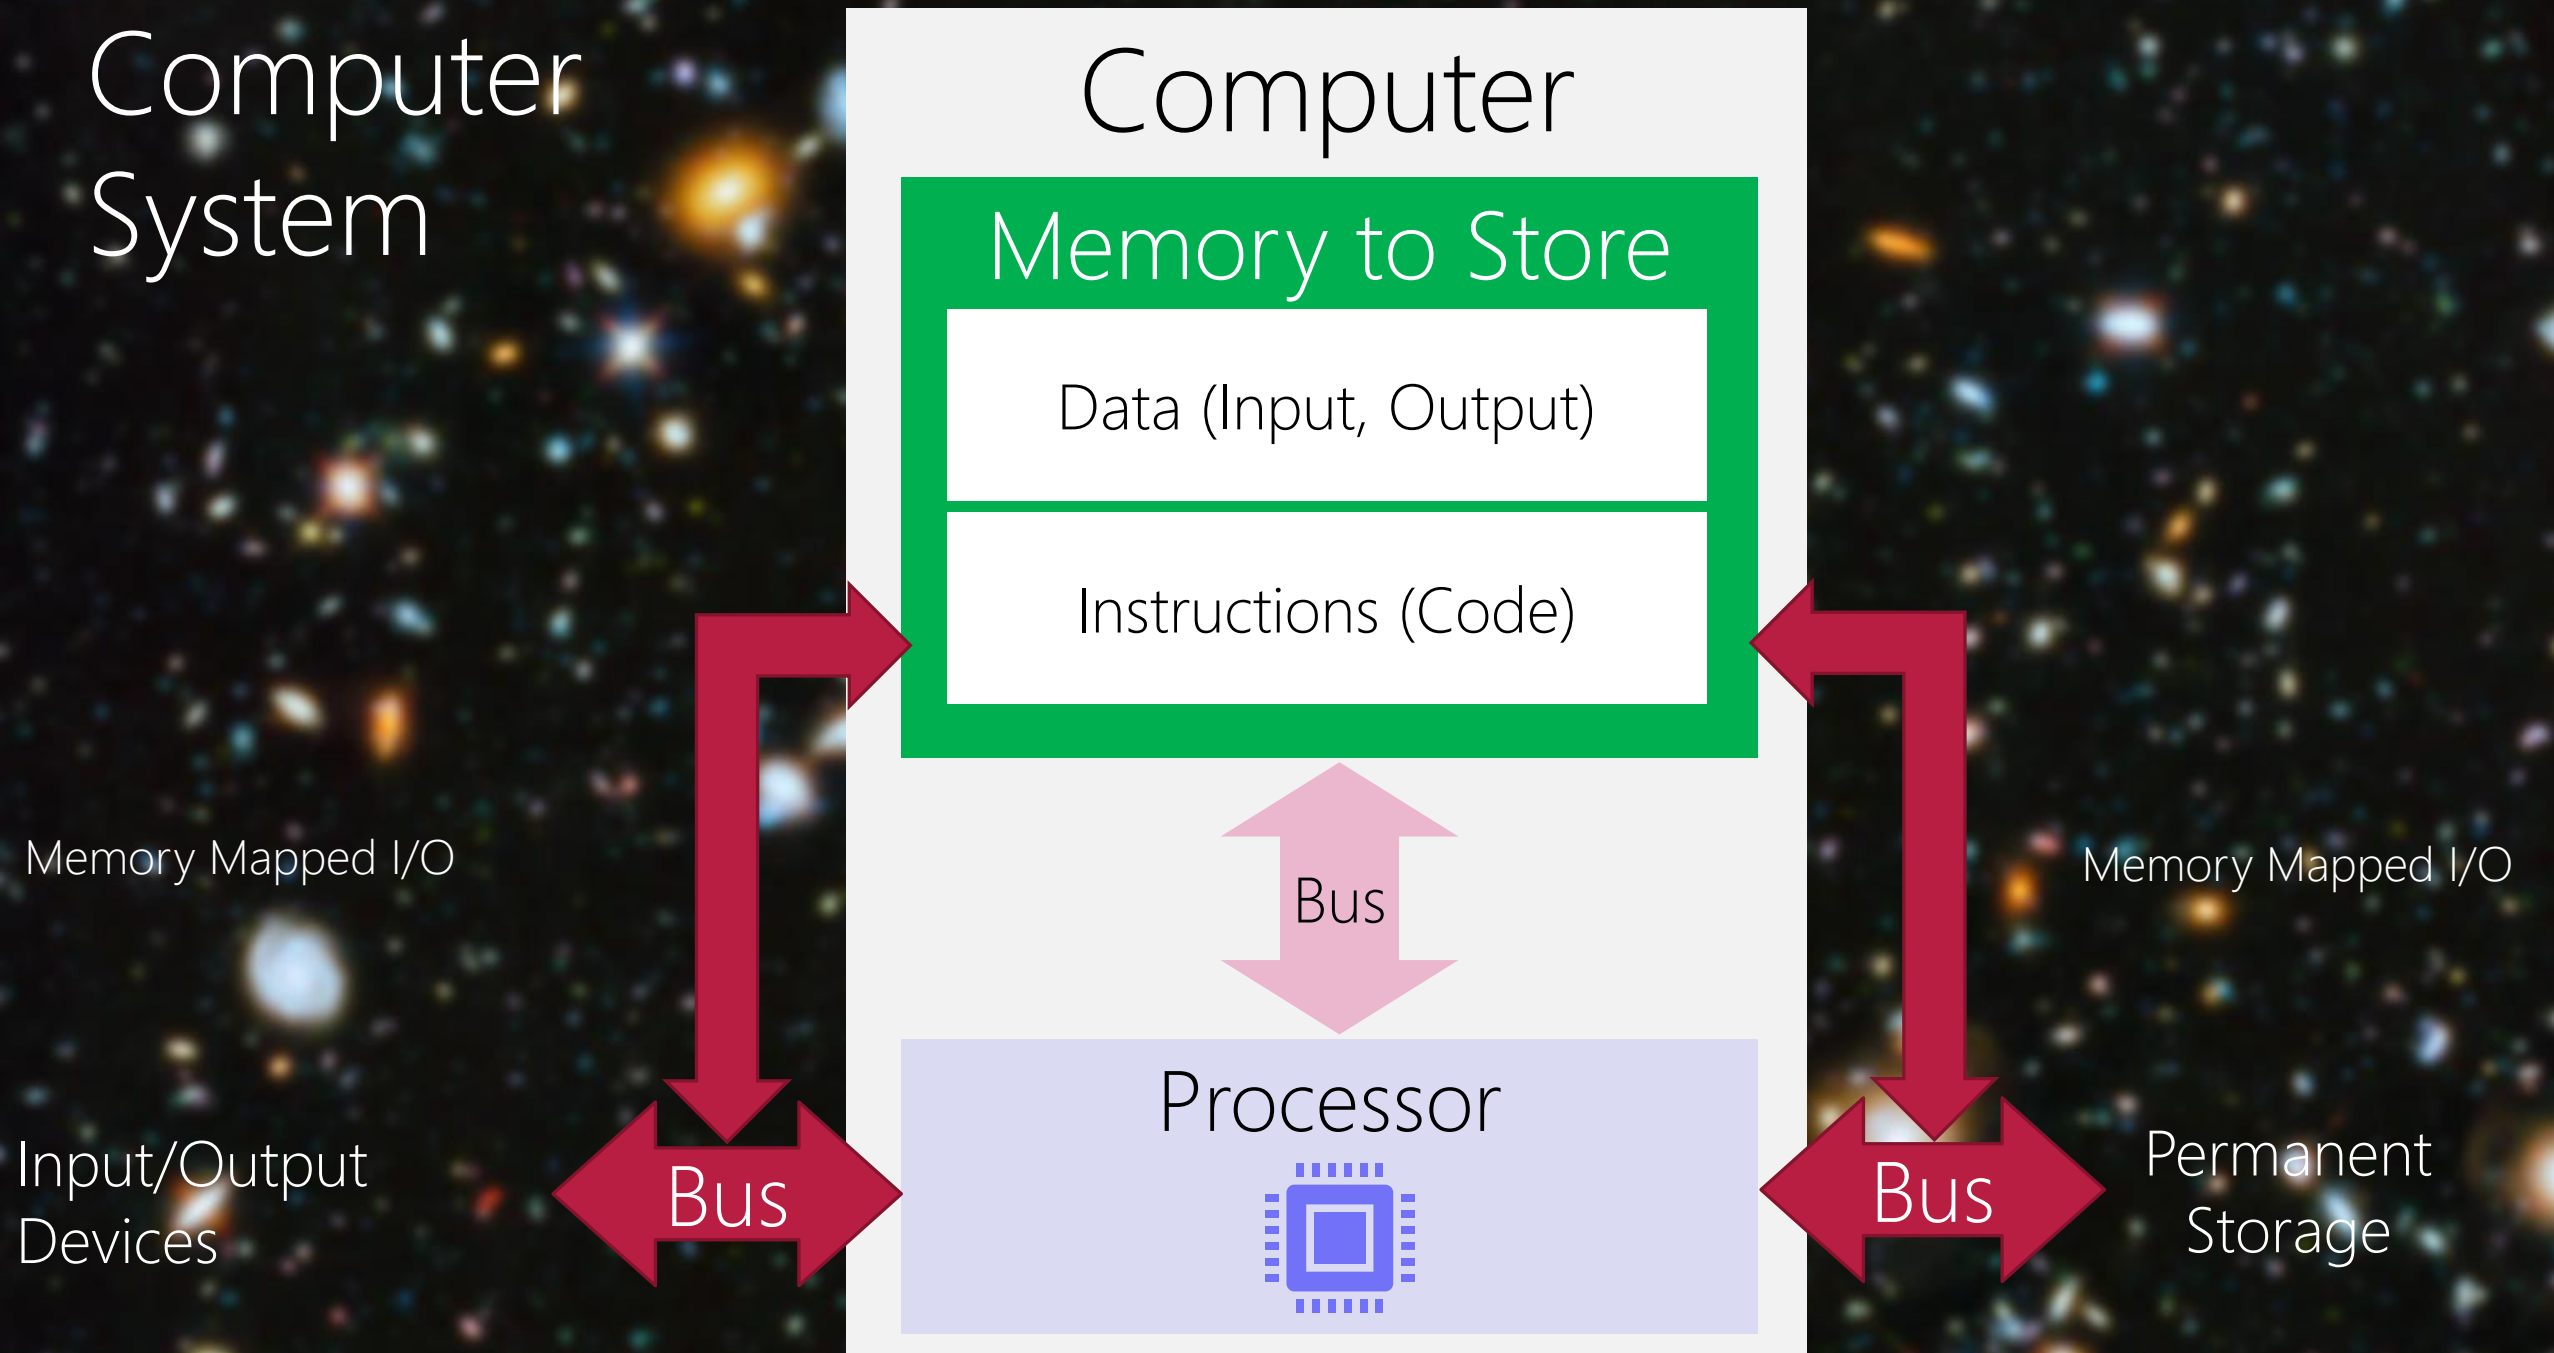
Bus

Processor



Permanent
Storage

Computer System



write

nbytes: Write this Amount of Byte to the File (Device)
Your Responsibility to Provide a Correct Conversion to Number of Bytes!

```
#include <unistd.h>
ssize_t write(int fd, const void *buf, size_t nbytes);
number of bytes written if OK, -1 on error
```

typedef

ssize_t, size_t, ..., and many other data types

```
#include <sys/types.h>
typedef size_t unsigned long
typedef ssize_t signed long
```

<https://www.ibm.com/docs/en/zos/2.2.0?topic=files-systypesh>



close
POSIX

```
#include <unistd.h>
int close(int fd);
0 if OK, -1 on error
```



close

Sometimes Optional, but only Sometimes!

When a process terminates, all of its open files are closed automatically by the kernel.
That is all the File Descriptors (fs) are released.
You can take advantage of this fact and don't explicitly close open files in your programs (not recommended!)


```
#include <fcntl.h>
#include <unistd.h>

#include <sys/types.h>
#include <string.h>

#include <stdio.h>
void main(void) {
    int fd; //file descriptor
    mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH; //for permisison settings
    char *filename = "./my_new_file.txt";

    fd = creat(filename, mode);
    printf("The file descriptor is: %d \n", fd);

    if(fd == -1) {
        printf("Error in creating file!\n");
        return;
    }

    char buf[20];
    size_t nbytes;
    ssize_t bytes_written;

    strcpy(buf, "Hello File!\n");
    nbytes = strlen(buf);

    bytes_written = write(fd, buf, nbytes);
    if(bytes_written != nbytes) {
        printf("Error in writing to the file!")
    }

    int result = close(fd);
    if(result == -1) {
        printf("Error in closing the file!");
    }
}
```

```
#include <fcntl.h>
#include <unistd.h>

#include <sys/types.h>
#include <string.h>

#include <stdio.h>

void main(void) {
    int fd; //file descriptor
    mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH; //for permission settings
    char *filename = "./my_new_file.txt";

    fd = creat(filename, mode);
    printf("The file descriptor is: %d \n", fd);

    if (fd == -1) {
        printf("Error in creating file!\n");
        return;
    }

    char buf[20];
    size_t nbytes;
    ssize_t bytes_written;

    strcpy(buf, "Hello File!\n");
    nbytes = strlen(buf);

    bytes_written = write(fd, buf, nbytes);
    if (bytes_written != nbytes) {
        printf("Error in writing to the file!");
    }

    int result = close(fd);
    if (result == -1) {
```



```
#include <fcntl.h>
#include <unistd.h>

#include <sys/types.h>
#include <string.h>

#include <stdio.h>
void main(void) {
    int fd; //file descriptor
    mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH; //for permission settings
    char *filename = "./my_new_file.txt";

    fd = creat(filename, mode);
    printf("The file descriptor is: %d \n", fd);

    if(fd == -1){
        printf("Error in creating file!\n");
        return;
    }

    char buf[20];
    size_t nbytes;
    ssize_t bytes_written;

    strcpy(buf, "Hello File!\n");
    nbytes = strlen(buf);

    bytes_written = write(fd, buf, nbytes);
    if(bytes_written != nbytes){
        printf("Error in writing to the file!");
    }

    int result = close(fd);
```

```
#include <fcntl.h>
#include <unistd.h>

#include <sys/types.h>
#include <string.h>

#include <stdio.h>
void main(void){
    int fd;//file descriptor
    mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH;//for permission settings
    char *filename = "../my_new_file.txt";

    fd = creat(filename, mode);
    printf("The file descriptor is: %d \n", fd);

    if(fd == -1){
        printf("Error in creating file!\n");
        return;
    }

    char buf[20];
    size_t nbytes;
    ssize_t bytes_written;

    strcpy(buf, "Hello File!\n");
    nbytes = strlen(buf);

    bytes_written = write(fd, buf, nbytes);
    if(bytes_written != nbytes){
        printf("Error in writing to the file!")
    }

    int result = close(fd);
    if(result == -1){
        printf("Error in closing the file!")
    }
}
```



```
hfani@alpha:~$ cc create_file_system_call.c -o create_file_system_call
hfani@alpha:~$ ./create_file_system_call
The file descriptor is: 3
hfani@alpha:~$
```



```
hfani@alpha:~$ vi my_new_file.txt
```

```
Hello File!
```

```
~
```

```
~
```

```
fd = creat()
```



```
write(fd, sth)
```



```
close(fd)
```




You write for later read
or
You want to read from ReadOnly File/Device

File System: High-Level

```
fd = creat()
```



```
graph TD; A[fd = creat()] --> B[write(fd, sth)]; B --> C[read(fd, sth)]; C --> D[close(fd)];
```

A vertical flowchart on the left side of the image. It consists of four rectangular boxes connected by downward-pointing yellow arrows. The boxes are: 1. Green, containing 'fd = creat()'. 2. Green, containing 'write(fd, sth)'. 3. Blue, containing 'read(fd, sth)'. 4. Green, containing 'close(fd)'.

```
write(fd, sth)
```

```
read(fd, sth)
```

```
close(fd)
```

```
fd = creat()
```



```
graph TD; A[fd = creat()] --> C[read(fd, sth)]; C --> D[close(fd)];
```

A vertical flowchart on the right side of the image. It consists of three rectangular boxes connected by downward-pointing yellow arrows. The boxes are: 1. Green, containing 'fd = creat()'. 2. Blue, containing 'read(fd, sth)'. 3. Green, containing 'close(fd)'.

```
read(fd, sth)
```

```
close(fd)
```

Not Possible w/ `creat`
What then?

File System: High-Level

open

POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/open.html>

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error



A quick reminder, why **POSIX** version?

open

POSIX

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error

open

POSIX

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error

open

POSIX

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error

<code>O_RDONLY</code>	Open for reading only (the returned fd can only read)
<code>O_WRONLY</code>	Open for write only (the returned fd can only write like <code>creat()</code>)
<code>O_RDWR</code>	Open for reading and writing (the returned fd can do both read and write)
<code>O_EXEC</code>	Open for execute only (the returned fd execute)
<code>O_SEARCH</code>	Open for search only (for directories)

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error

O_RDONLY
O_WRONLY
O_RDWR
O_EXEC
O_SEARCH

Only One of Them

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error

```
hfani@alpha:~$ vi open_rw.c
```

```
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";
    int fd = open(filename, O_RDWR);
    if(fd==-1){
        printf("error happend!\n");
    }
    else if (fd >= 0){
        printf("file successfully opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```

```
hfani@alpha:~$ vi open_rw.c
```

```
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";
    int fd = open(filename, O_RDWR);
    if(fd==-1){
        printf("error happend!\n");
    }
    else if (fd >= 0){
        printf("file successfully opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```

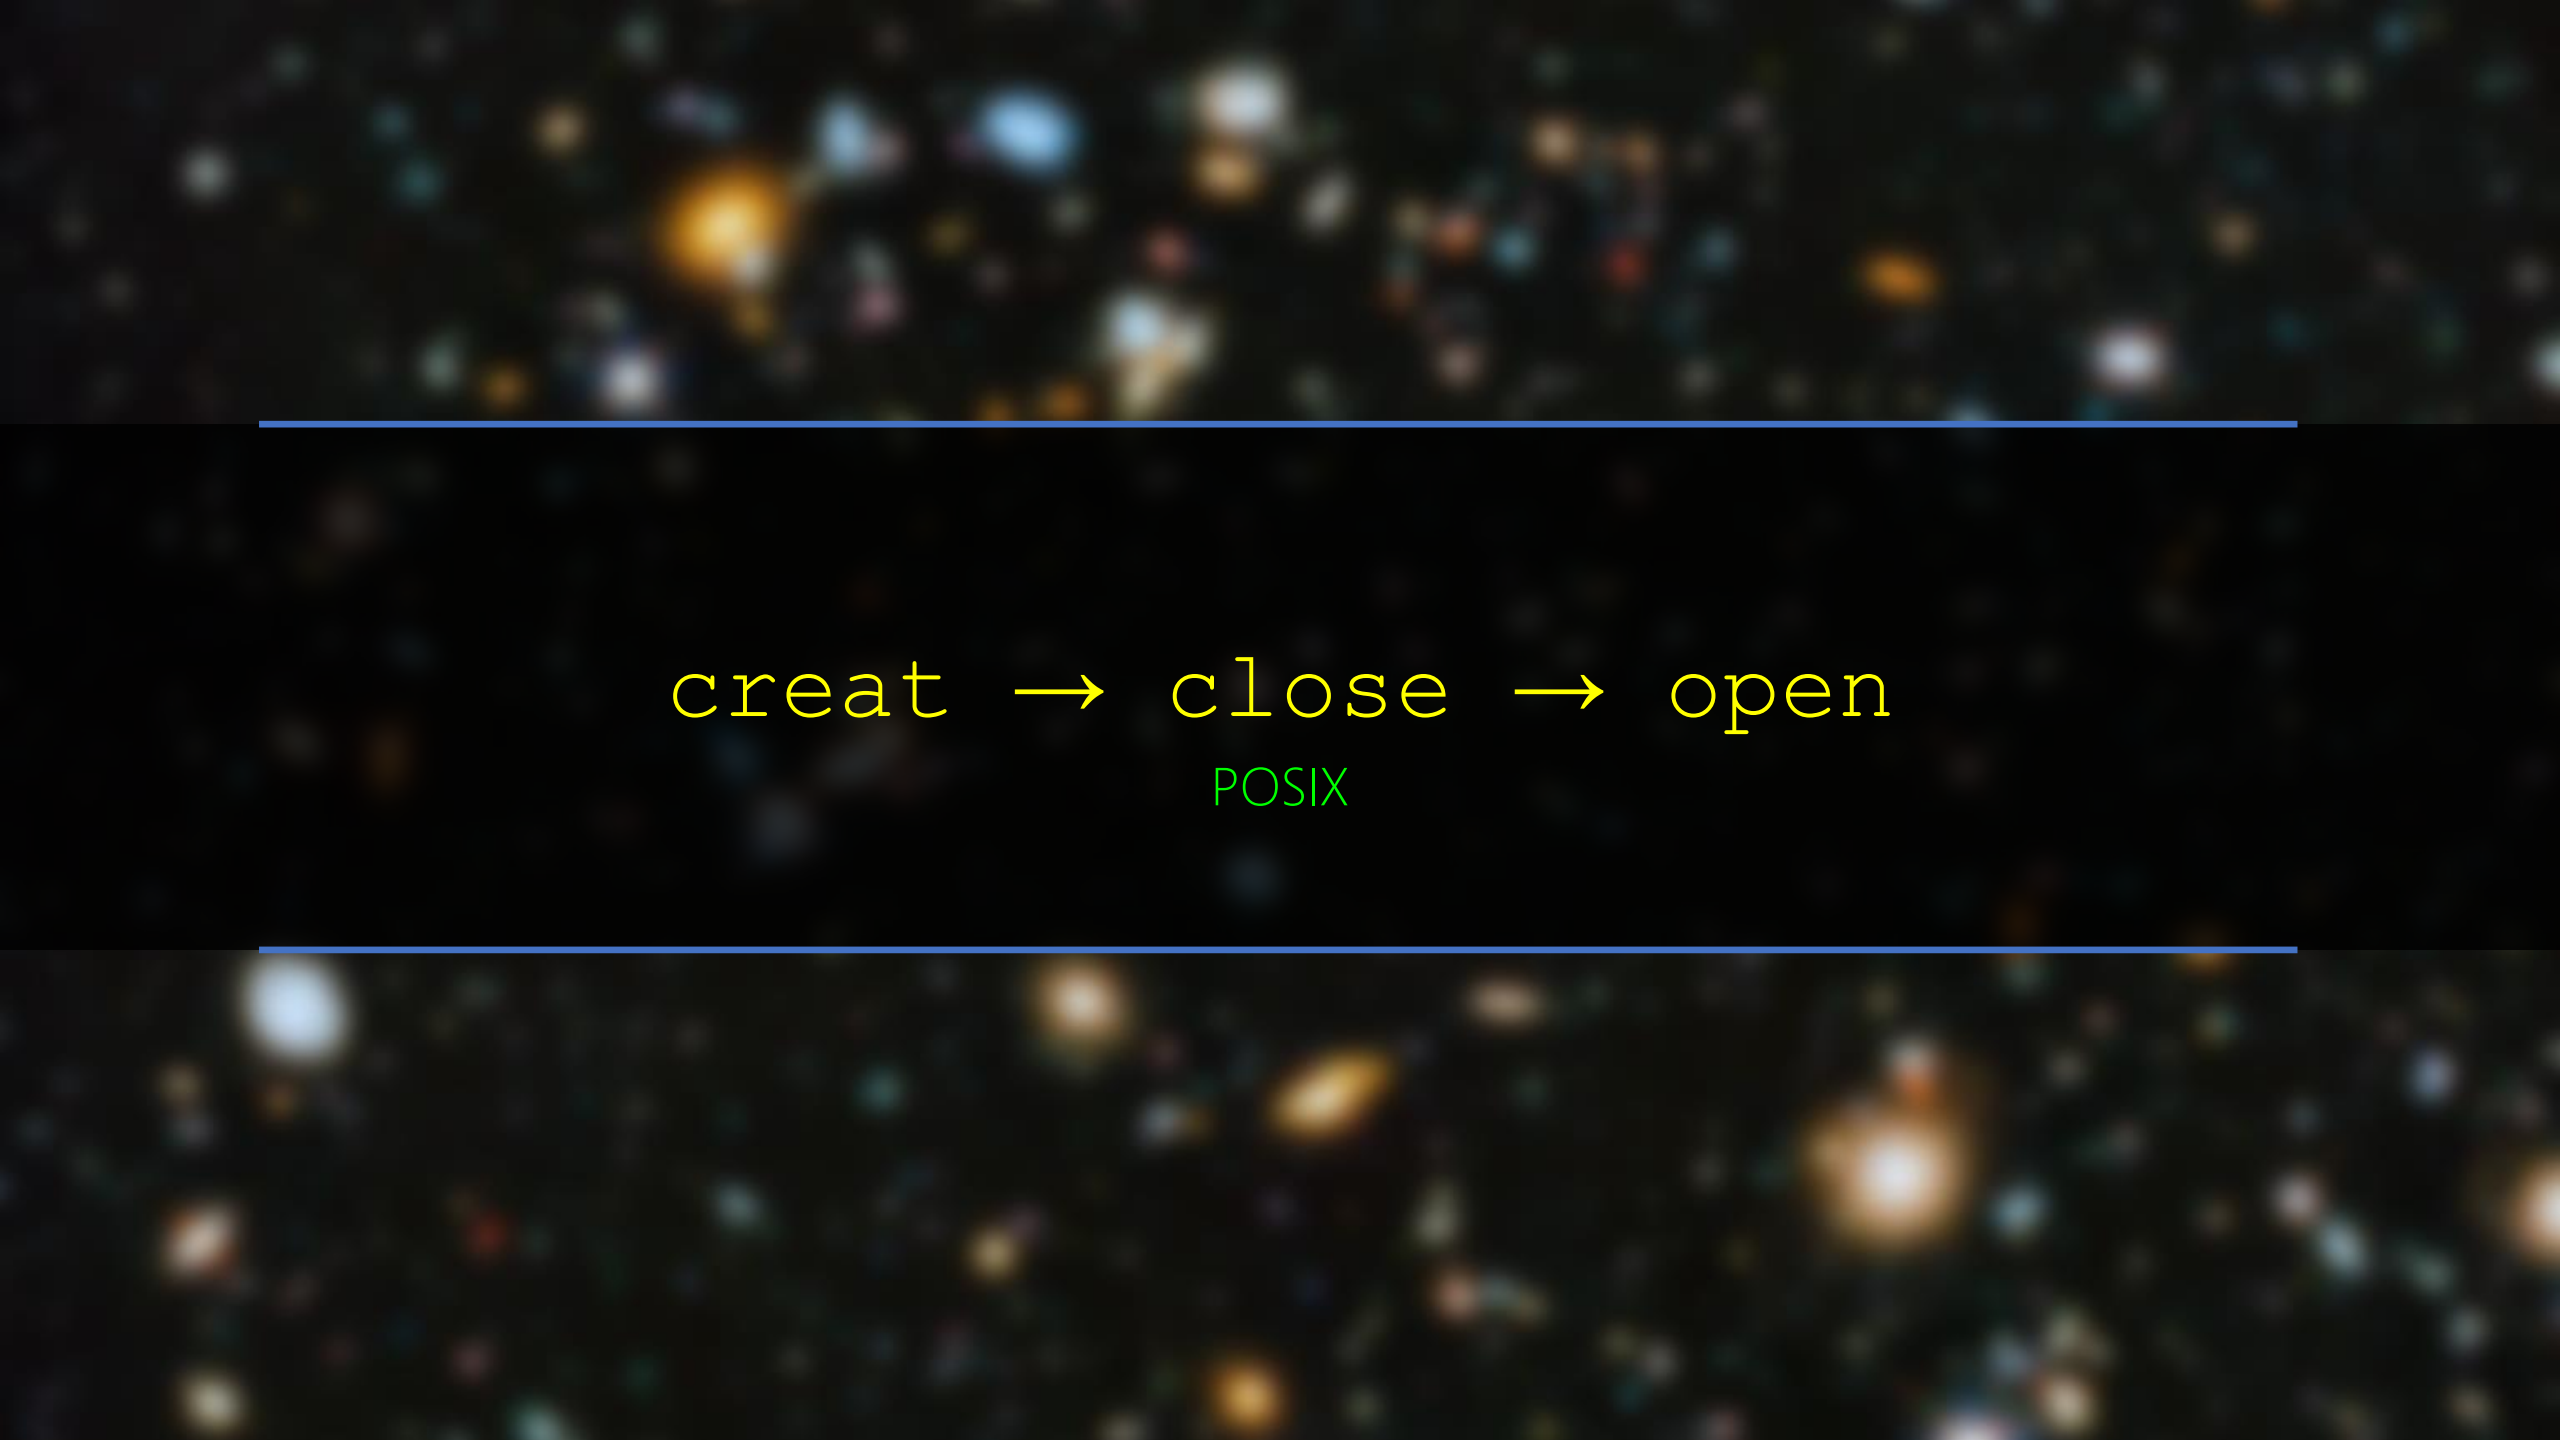
```
hfani@alpha:~$ cc open_rw.c -o open_rw
```

```
hfani@alpha:~$ ./open_rw
```

```
error happend! _
```



What's the problem?



creat → close → open

POSIX

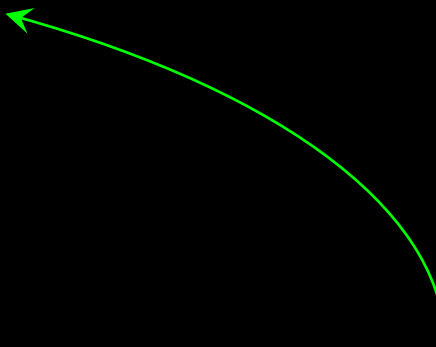
```
hfani@alpha:~$ vi creat_close_open_rw.c
```

```
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";
    int fd = creat(filename, S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
    if (fd == -1){
        printf("error happened in creating file %s\n", filename);
        return;
    }
    else if (fd > 0){
        printf("file %s has been created successfully and the fd is %d\n", filename, fd);
    }

    int result = close(fd);
    if (result == -1){
        printf("error happened in closing file %s\n", filename);
        return;
    }
    else if (result == 0){
        printf("file %s has been closed successfully\n", filename);
    }

    fd = open(filename, O_RDWR);
    if(fd==-1){
        printf("error happend in opening file %s!\n", filename);
    }
    else if (fd > 0){
        printf("file successfully opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```

```
hfani@alpha:~$ cc creat_close_open_rw.c -o creat_close_open_rw
creat_close_open_rw.c: In function 'main':
creat_close_open_rw.c:14:15: warning: implicit declaration of function 'close'
   14 |     int result = close(fd);
      |                   ^~~~~
      |                   pclose
hfani@alpha:~$ ./creat_close_open_rw
file open_rw_test.txt has been created successfully and the fd is 3
file open_rw_test.txt has been closed successfully
file successfully opened for 2 and the fd is 3
```



<code>O_CREAT</code>	Create the file if it does not exist (you have to specify the <code>mode_t</code>)
<code>O_EXCL</code>	Raise error (<code>fd == -1</code>) if the file already exists

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (`fd`) if OK
-1 on error

O_CREAT
O_EXCL

In Combination w/ the Main Five Options

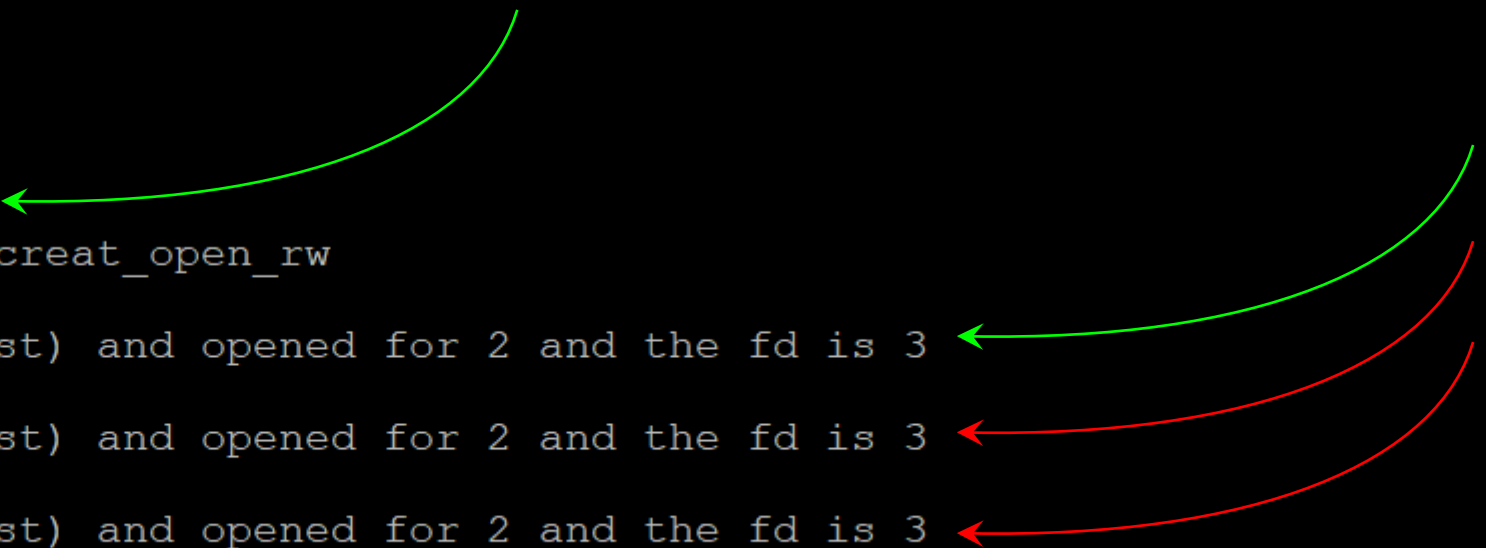
```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error


```
hfani@alpha:~$ vi creat_open_rw.c
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";

    int fd = open(filename, O_RDWR | O_CREAT, S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
    if(fd==-1){
        printf("error happend in creating the file %s (if not exists) or opening file!\n.", filename);
    }
    else if (fd > 0){
        printf("file successfully created (if not exist) and opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```

```
hfani@alpha:~$ rm ./open_rw_test.txt
hfani@alpha:~$ cc creat_open_rw.c -o creat_open_rw
hfani@alpha:~$ ./creat_open_rw
file successfully created (if not exist) and opened for 2 and the fd is 3
hfani@alpha:~$ ./creat_open_rw
file successfully created (if not exist) and opened for 2 and the fd is 3
hfani@alpha:~$ ./creat_open_rw
file successfully created (if not exist) and opened for 2 and the fd is 3
hfani@alpha:~$
```



hfani@alpha:~\$ vi creat_excl_open_rw.c

```
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="open_rw_test.txt";

    int fd = open(filename, O_RDWR | O_CREAT | O_EXCL, S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH);
    if(fd==-1){
        printf("error happend in creating the file %s (may be it exists) or opening file!\n.", filename);
    }
    else if (fd > 0){
        printf("file successfully created (it didn't exist) and opened for %d and the fd is %d\n", O_RDWR, fd);
    }
}
```

)

.hfani@alpha:~\$ rm open_rw_test.txt

hfani@alpha:~\$./creat_excl_open_rw

file successfully created (it didn't exist) and opened for 2 and the fd is 3

hfani@alpha:~\$./creat_excl_open_rw

error happend in creating the file open_rw_test.txt (may be it exists) or opening file!

.hfani@alpha:~\$./creat_excl_open_rw

error happend in creating the file open_rw_test.txt (may be it exists) or opening file!

O_APPEND Append to the end of file on each write

```
#include <fcntl.h>
int open(const char *path, int oflag, ...);
```

non-negative number (fd) if OK
-1 on error



read

POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/read.html>

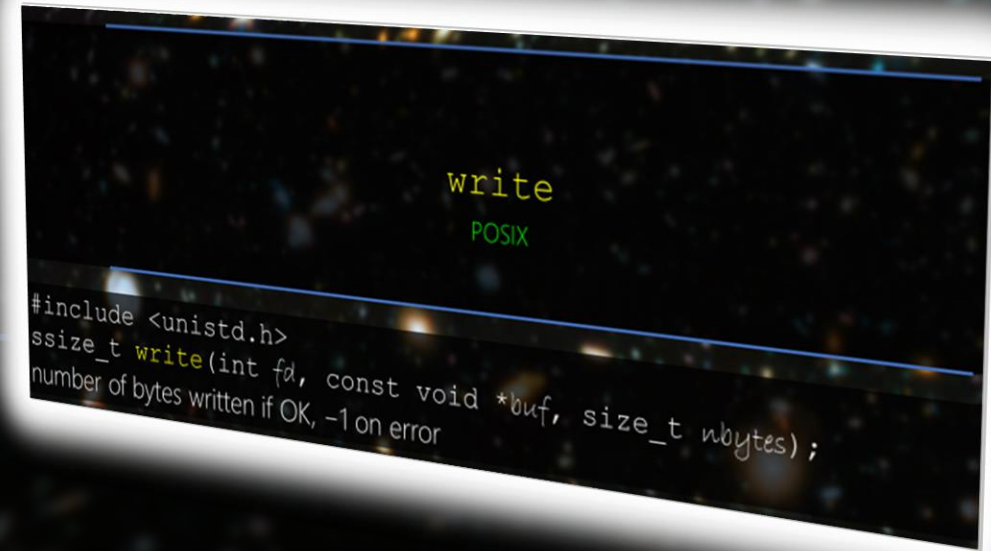
```
#include <unistd.h>
ssize_t read(int fd, const void *buf, size_t nbytes);
number of bytes written if OK, -1 on error
```

read

POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/read.html>

```
#include <unistd.h>
ssize_t read(int fd, const void *buf, size_t nbytes);
number of bytes written if OK, -1 on error
```



```
fd = creat()
```

```
write(fd, sth)
```

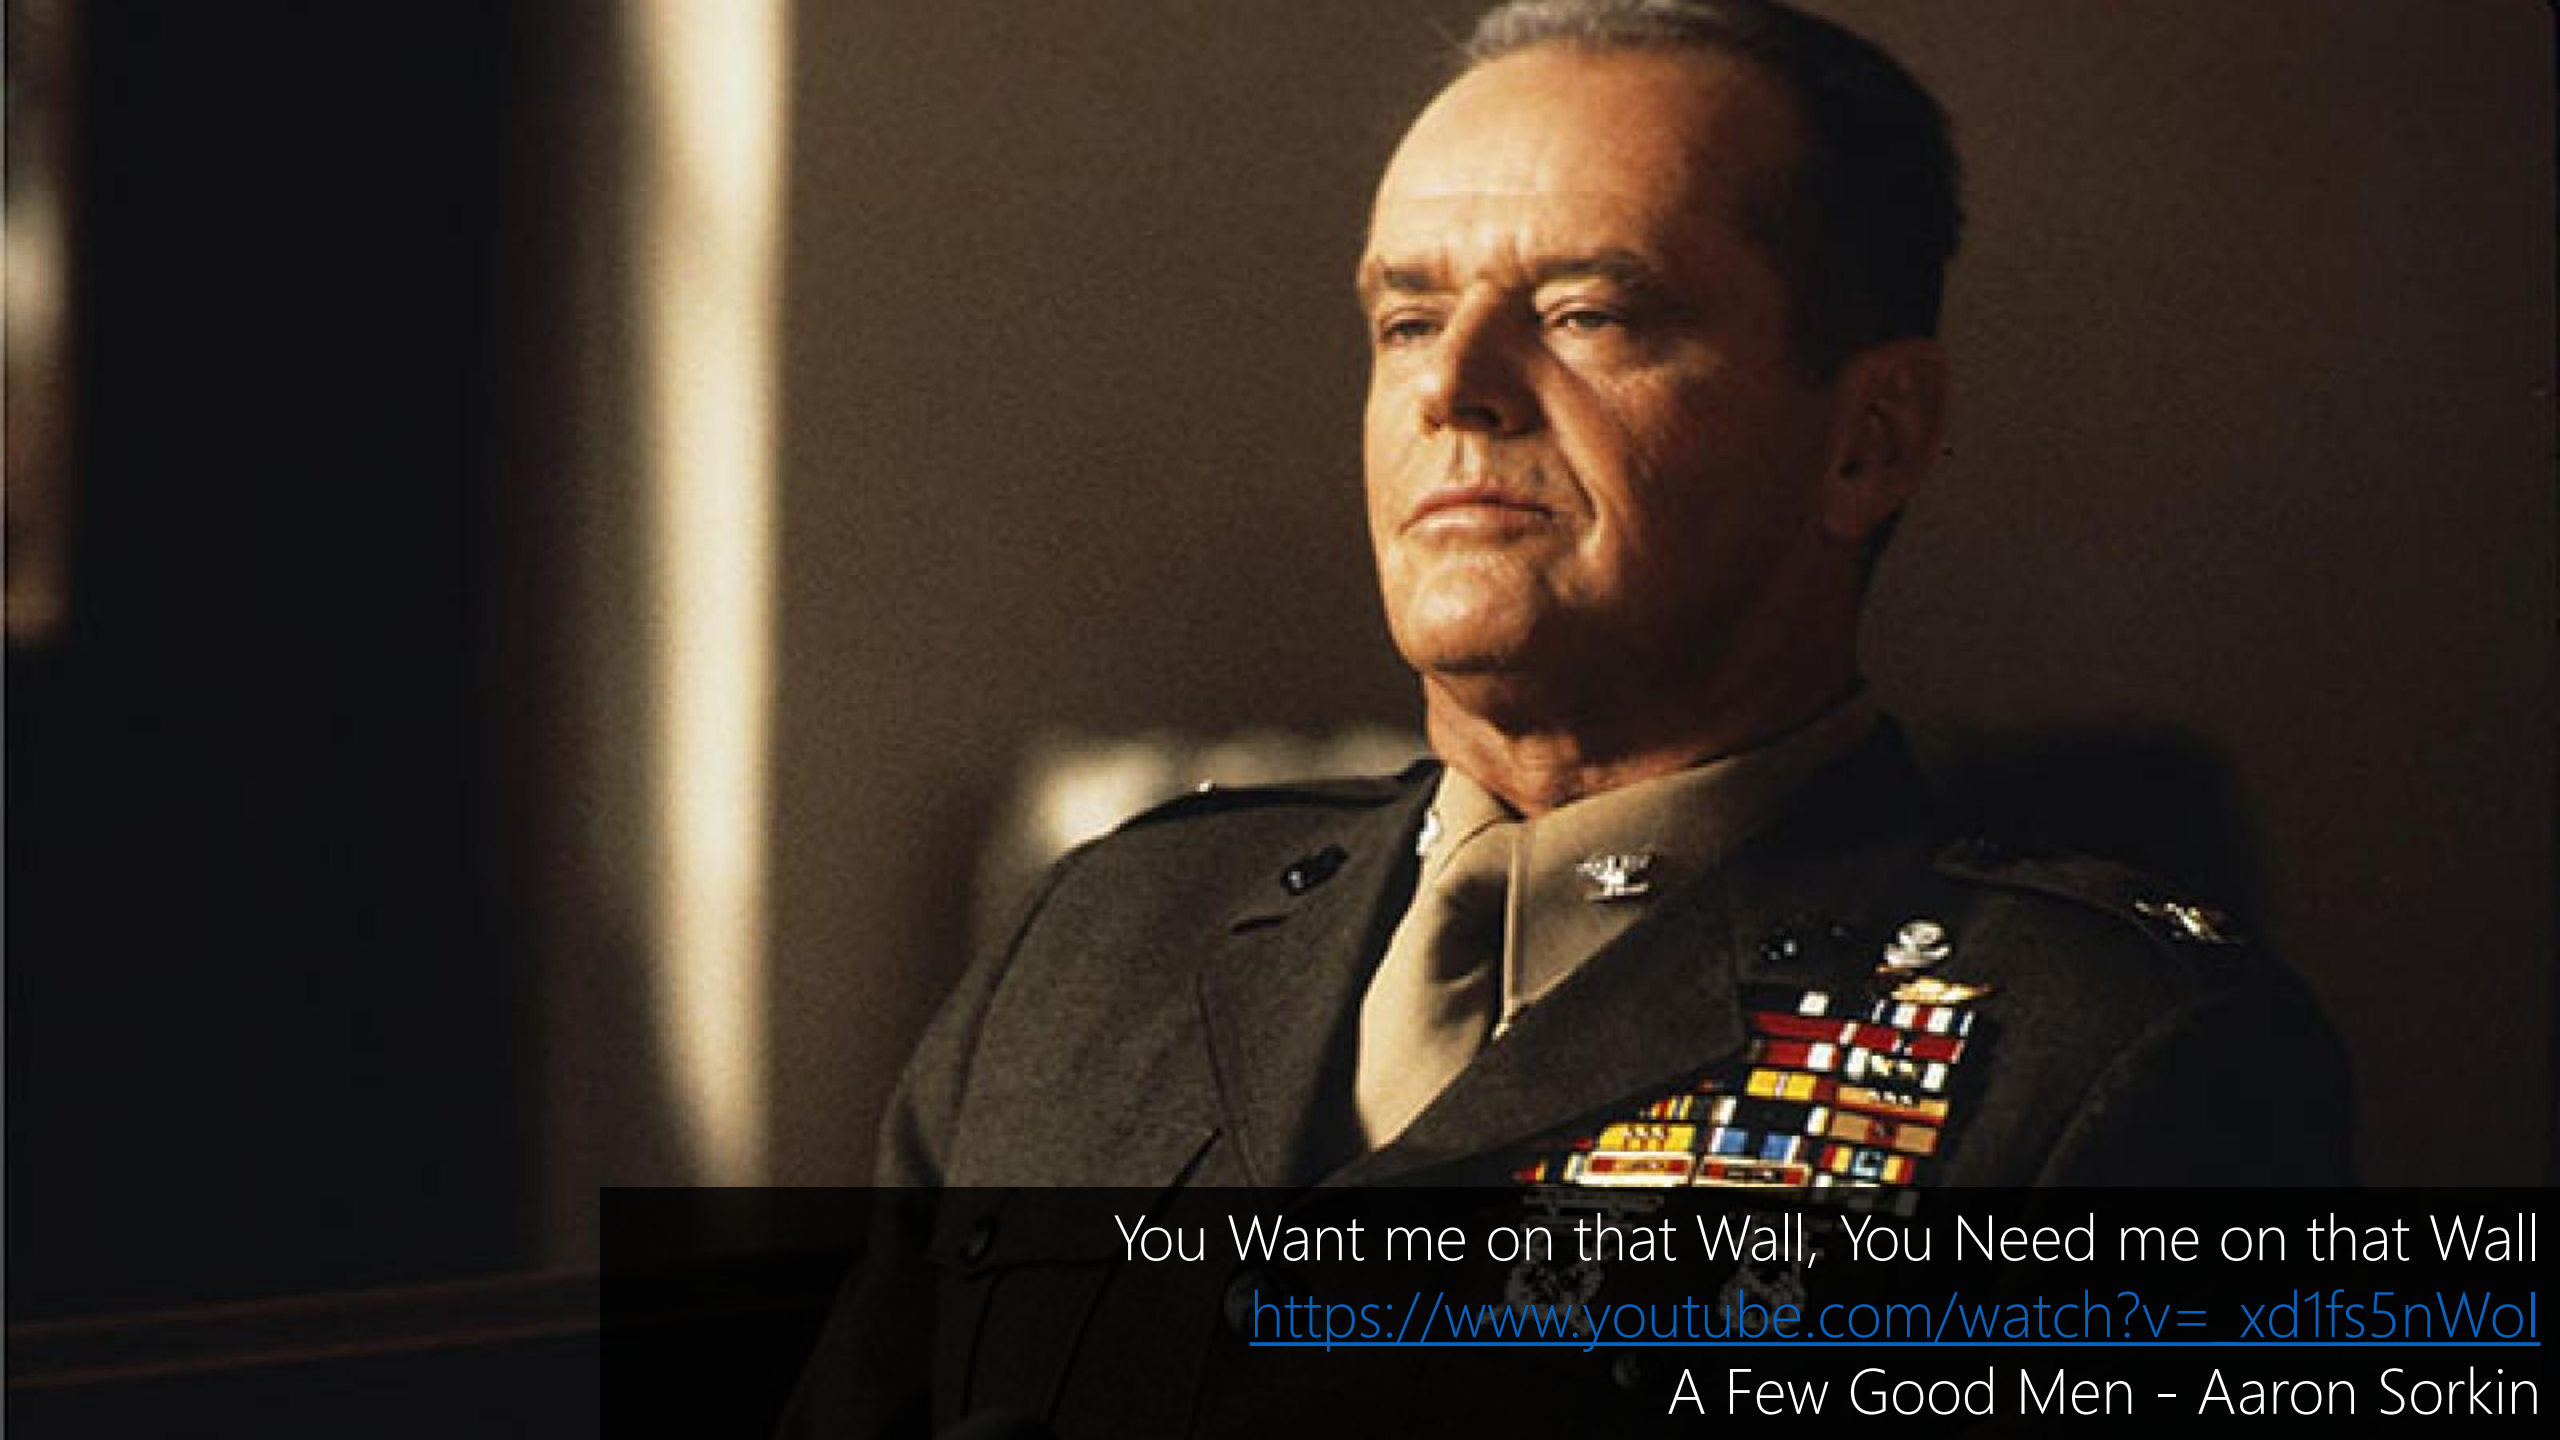
```
close(fd)
```

```
fd = open(creat if not exists)
```

```
read(fd, sth)
```

```
write(fd, sth)
```

```
close(fd)
```

You Want me on that Wall, You Need me on that Wall

<https://www.youtube.com/watch?v=xd1fs5nWol>

A Few Good Men - Aaron Sorkin

lseek

POSIX

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```

Computer

Memory

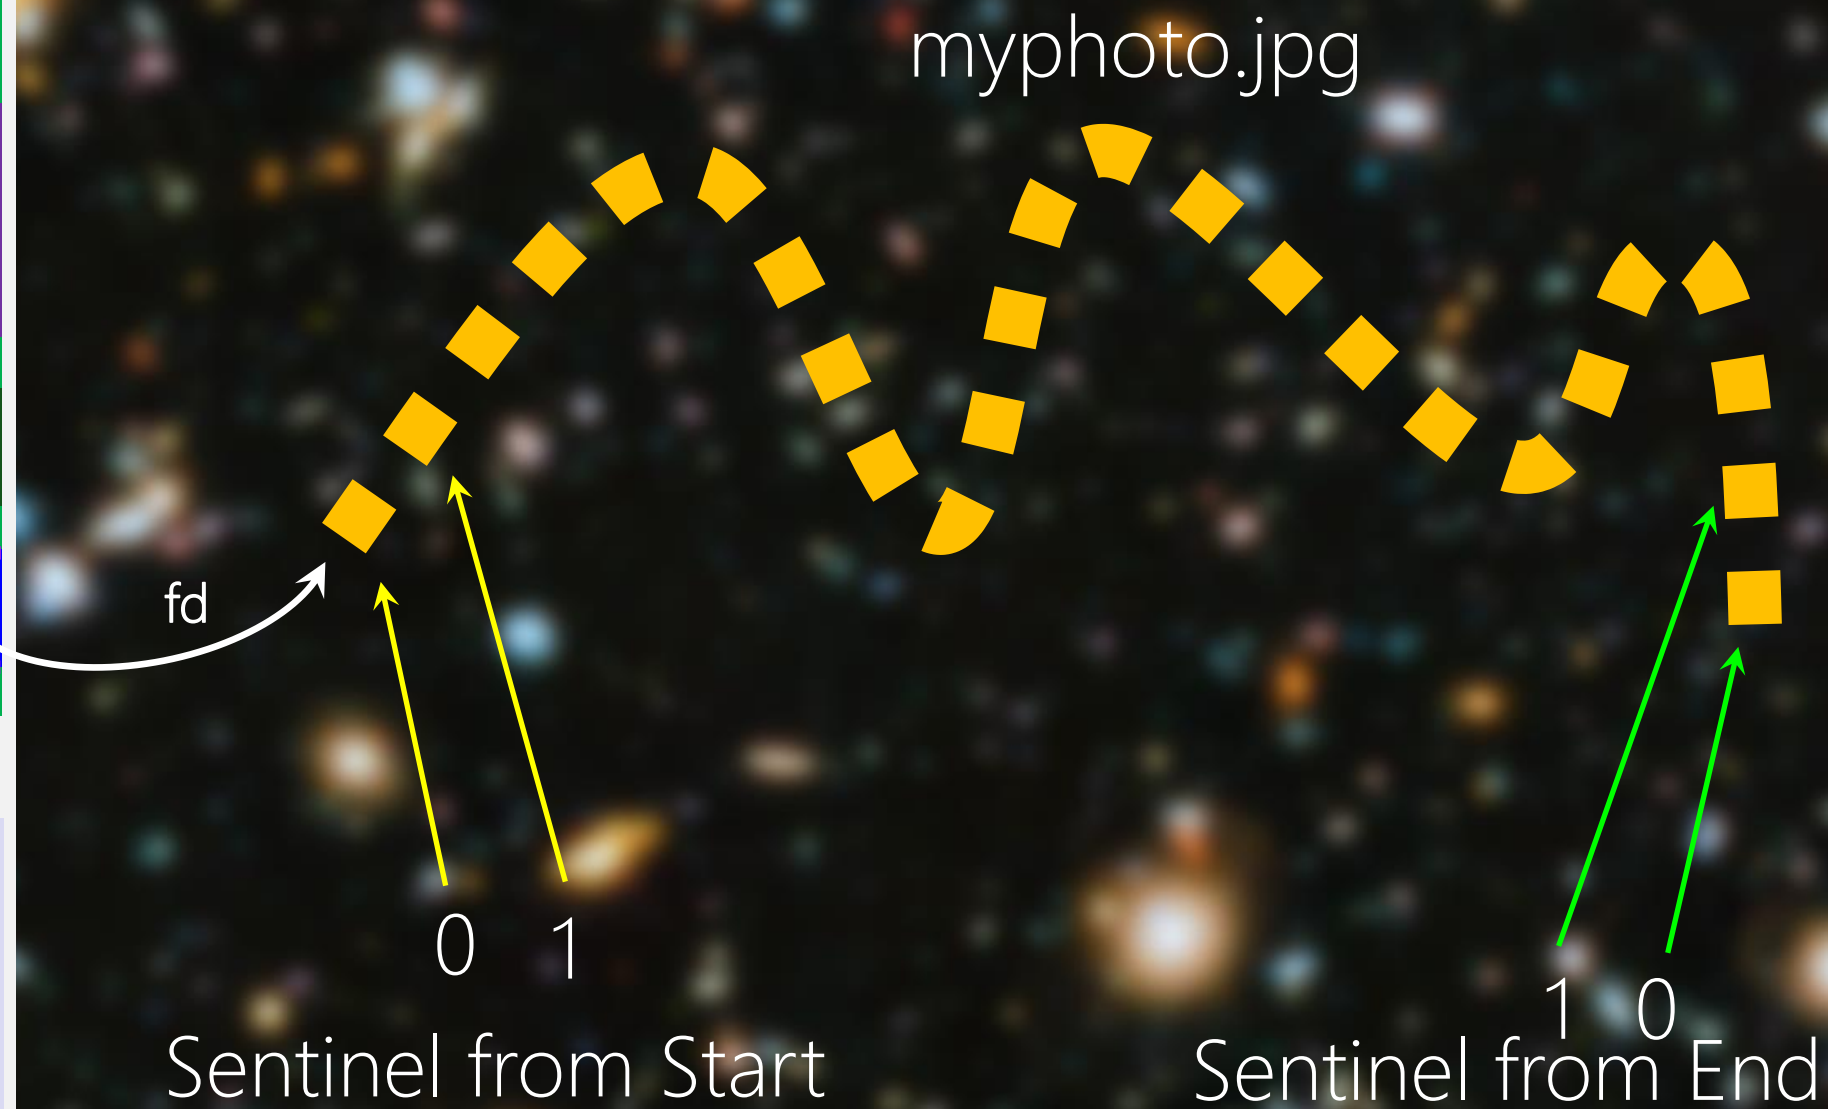
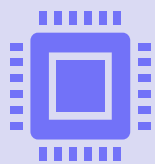
Kernel
File System

Shell

Process1

Bus

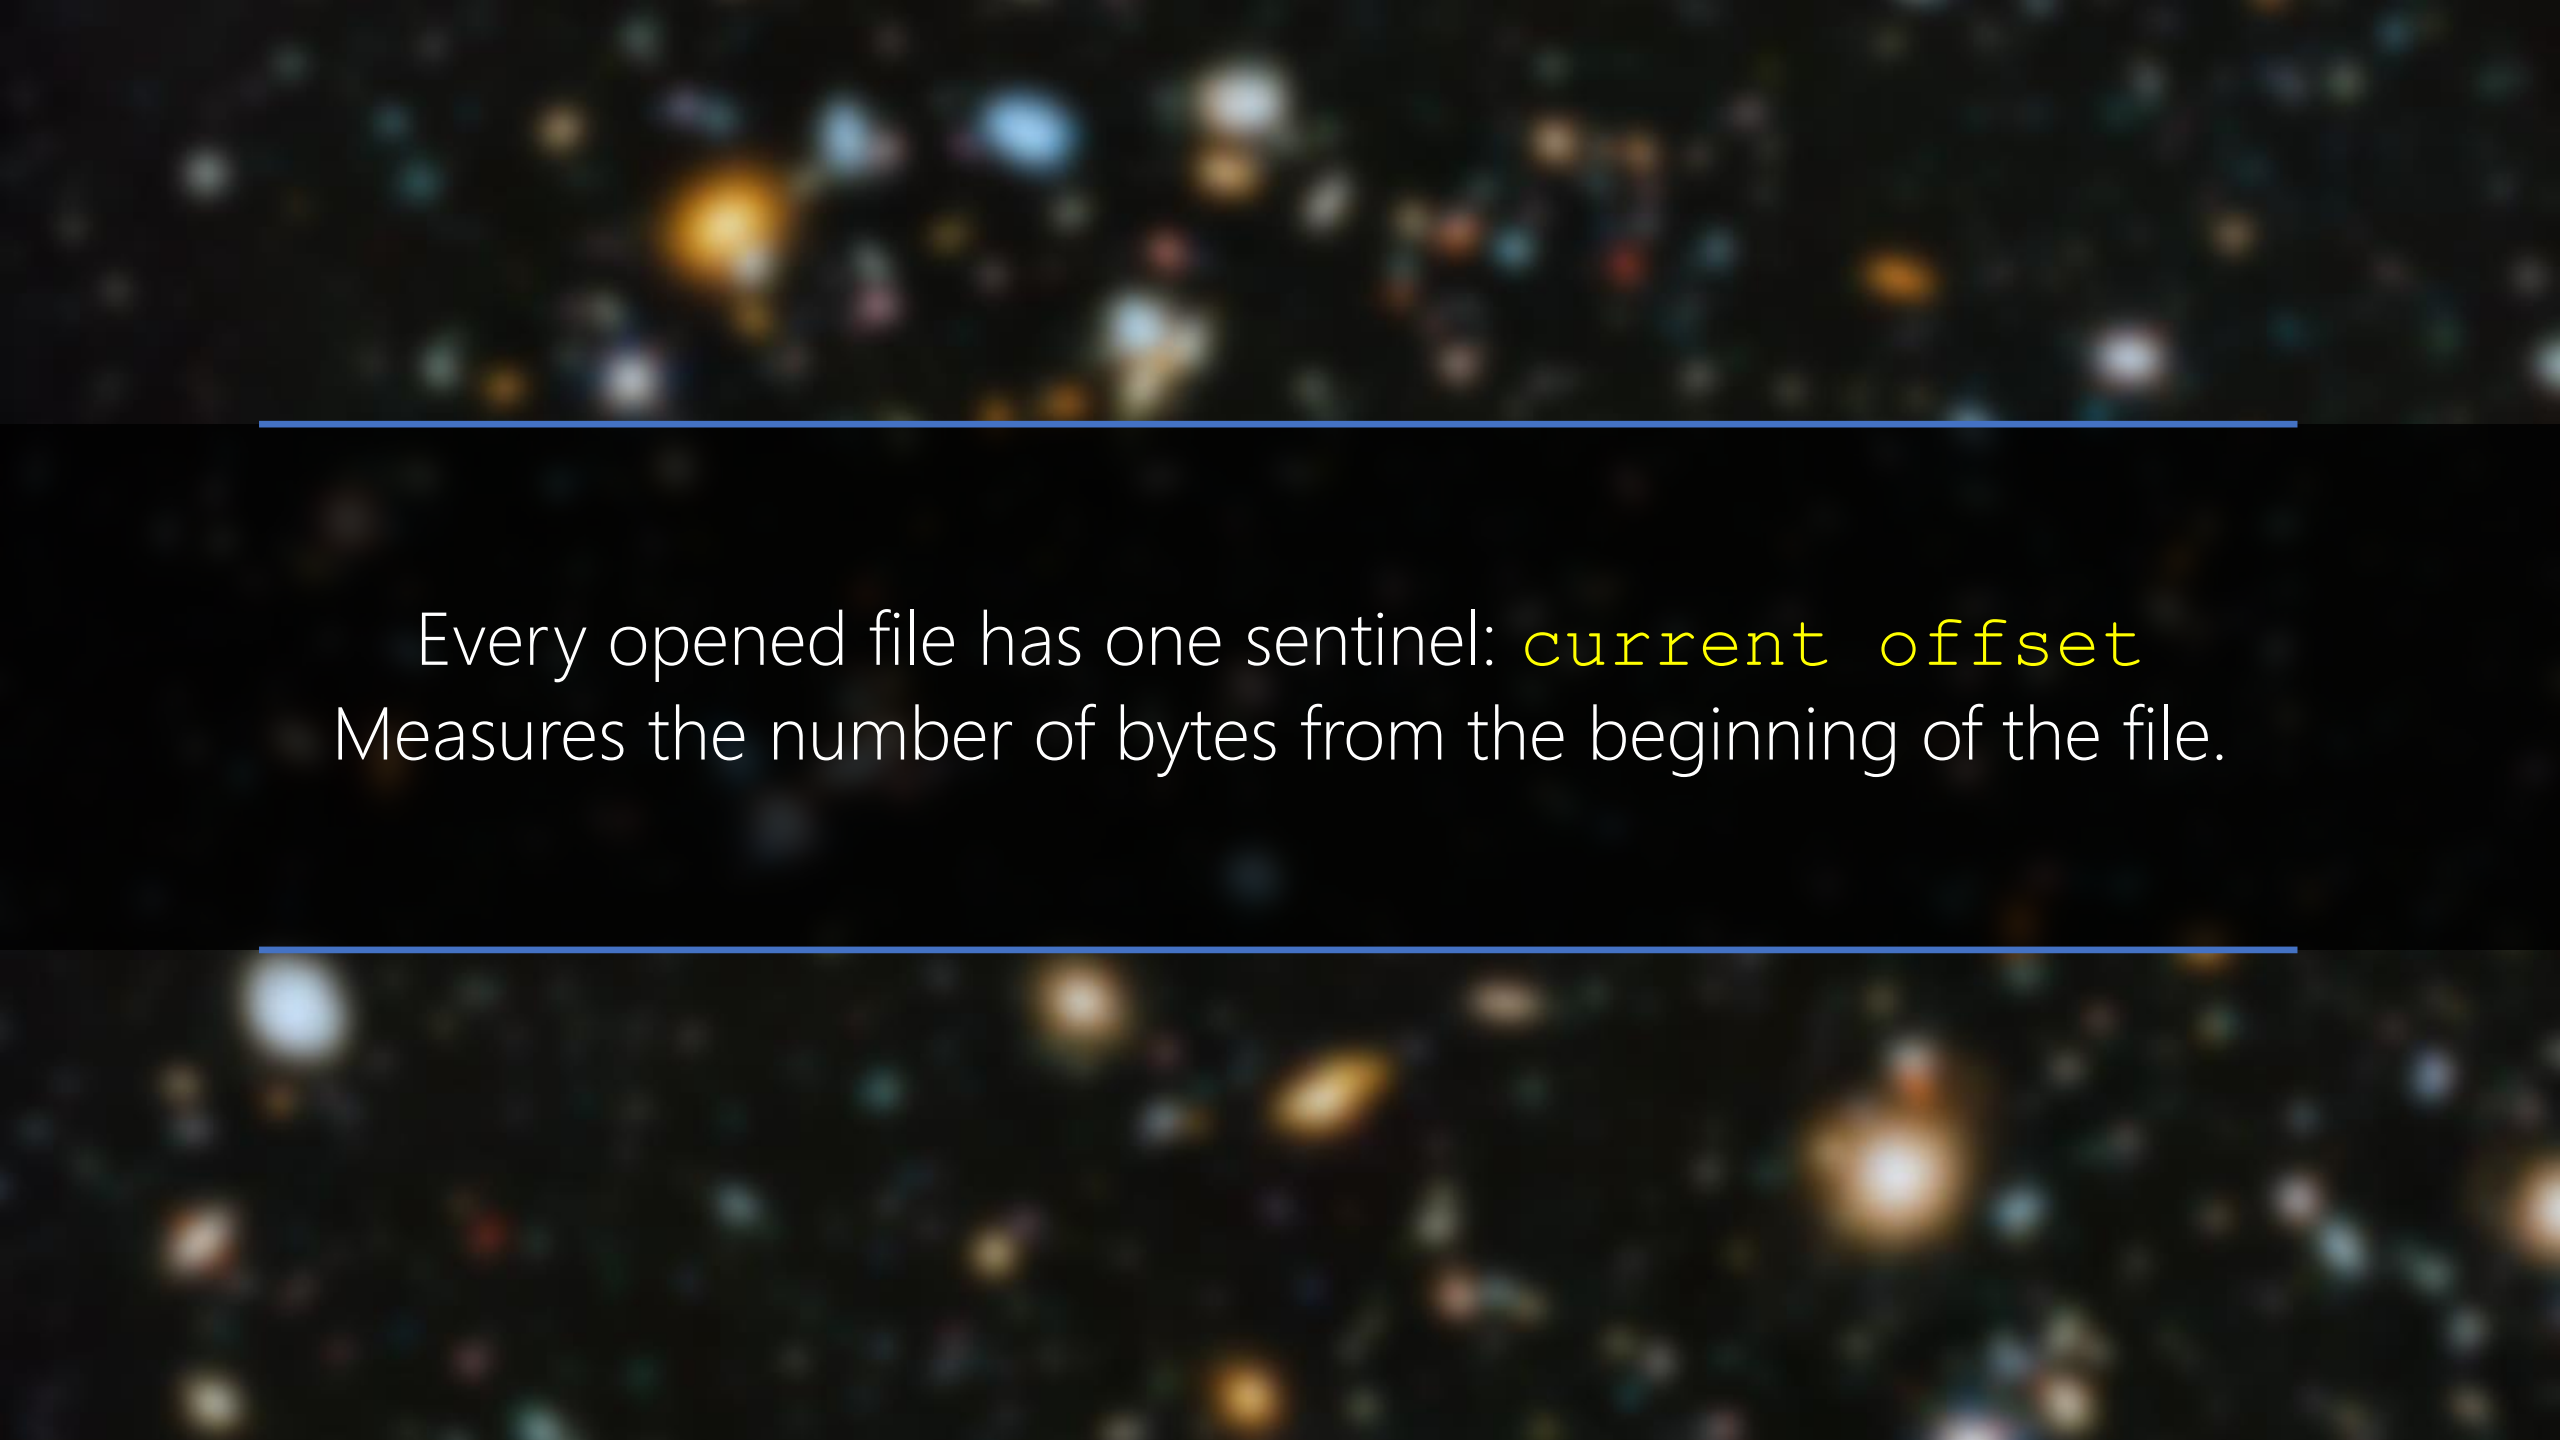
Processor



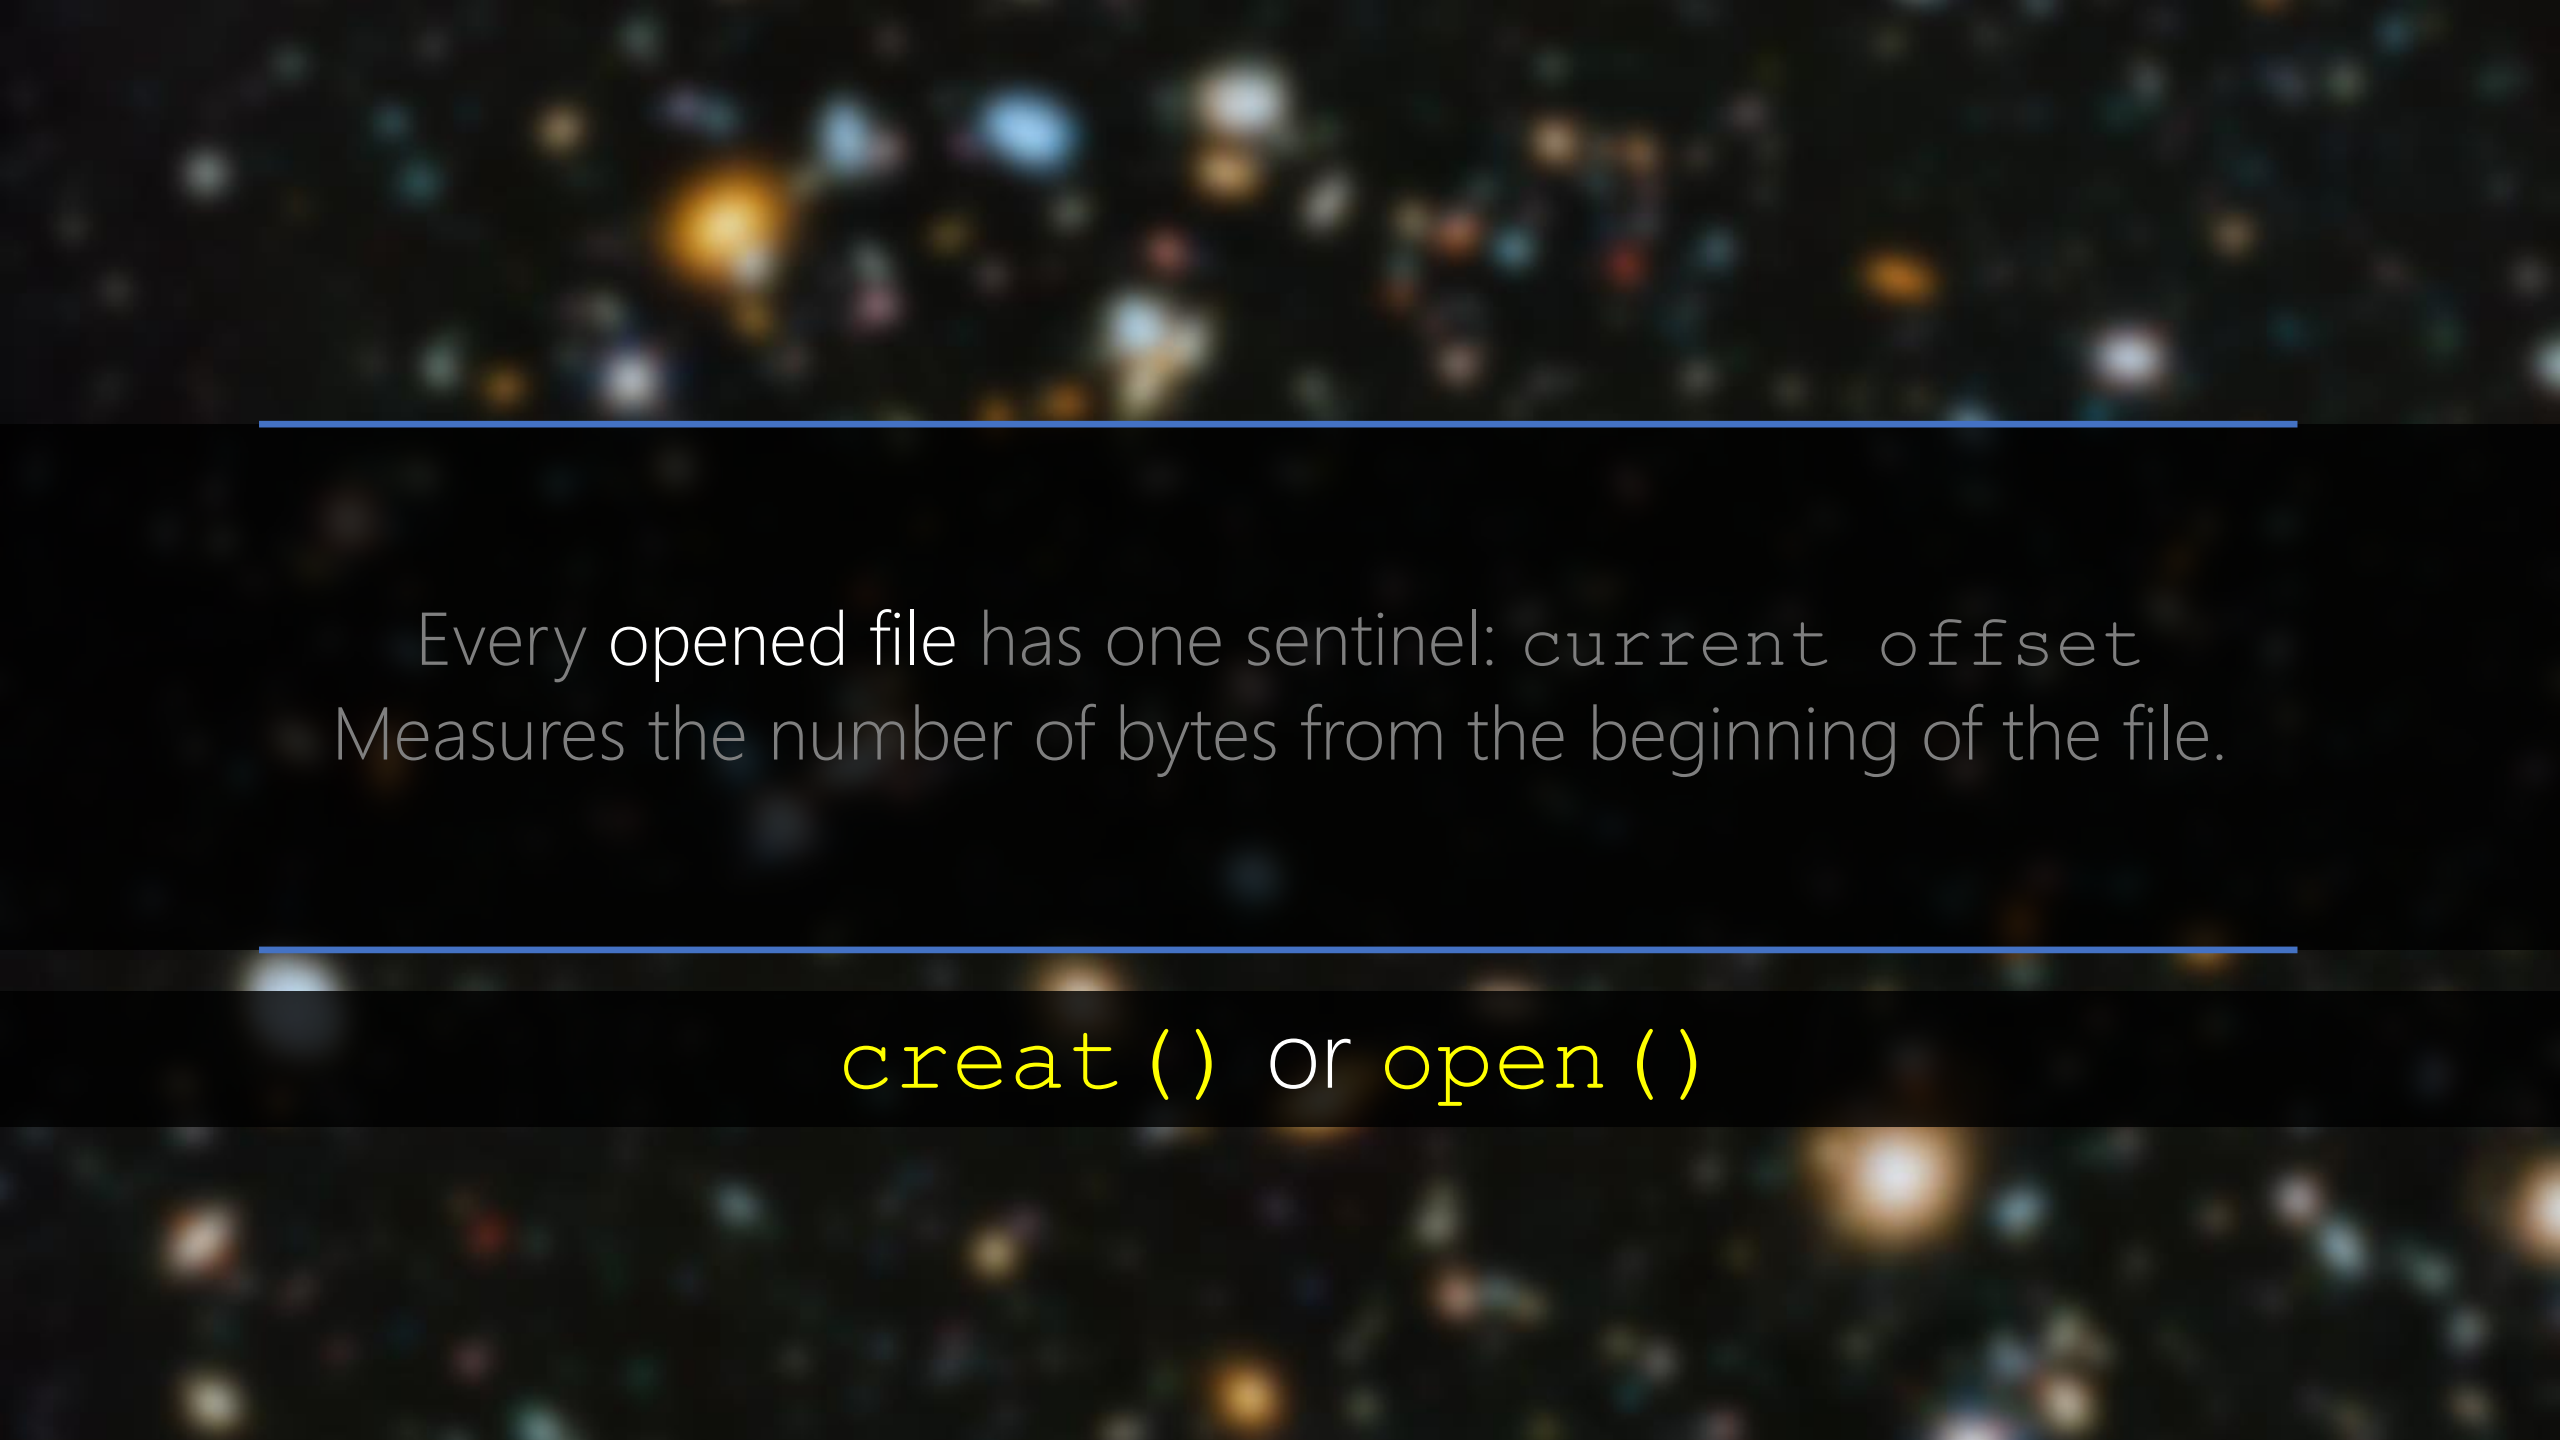
myphoto.jpg

Sentinel from Start

Sentinel from End



Every opened file has one sentinel: `current offset`
Measures the number of bytes from the beginning of the file.



Every opened file has one sentinel: `current_offset`
Measures the number of bytes from the beginning of the file.

`creat()` or `open()`

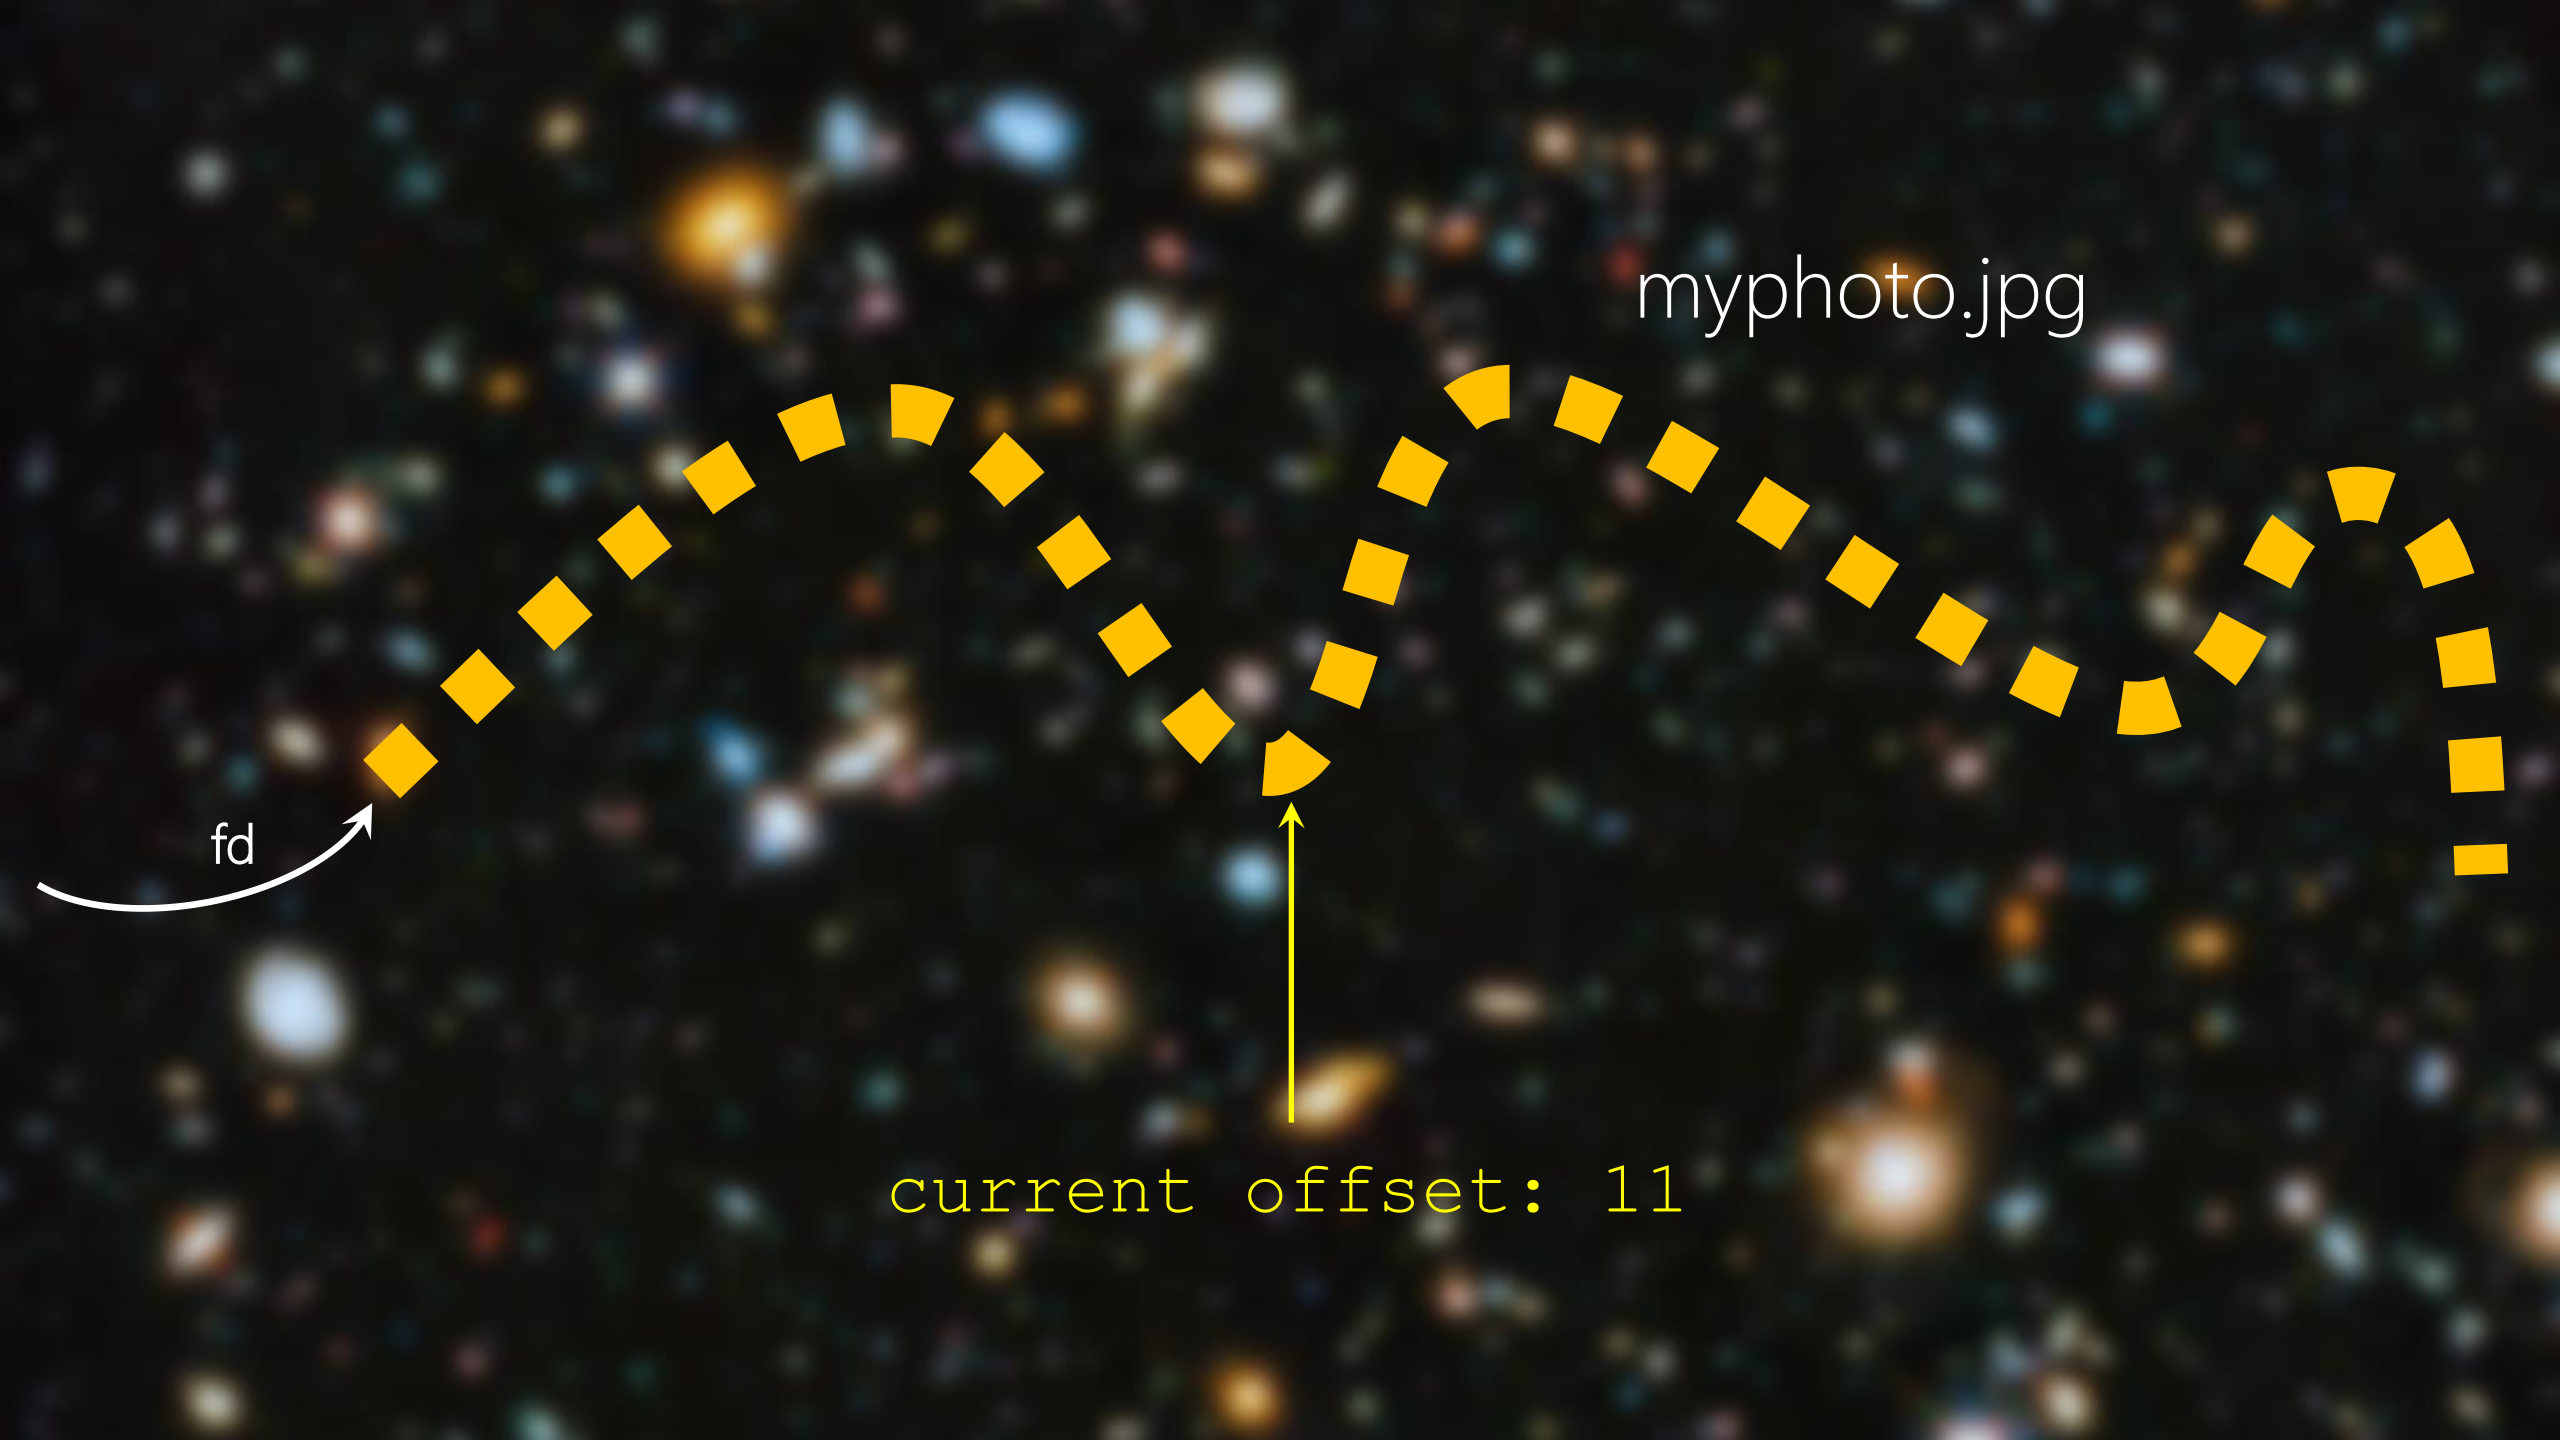
Every opened file has one sentinel: `current offset`
Measures the number of bytes from the beginning of the file.

`creat()` or `open()` \rightarrow 0

Every opened file has one sentinel: `current offset`
Measures the number of bytes from the beginning of the file.

`read()` or `write()` → ++ actual number of bytes read or written

`read()` or `write()` → always move forward



myphoto.jpg

fd

current offset: 11

lseek

POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/lseek.html>

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```


lseek

POSIX

<https://pubs.opengroup.org/onlinepubs/9699919799/functions/lseek.html>

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```

lseek

How many bytes move the current offset

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```

```
#include <sys/types.h>
typedef off_t signed long
```

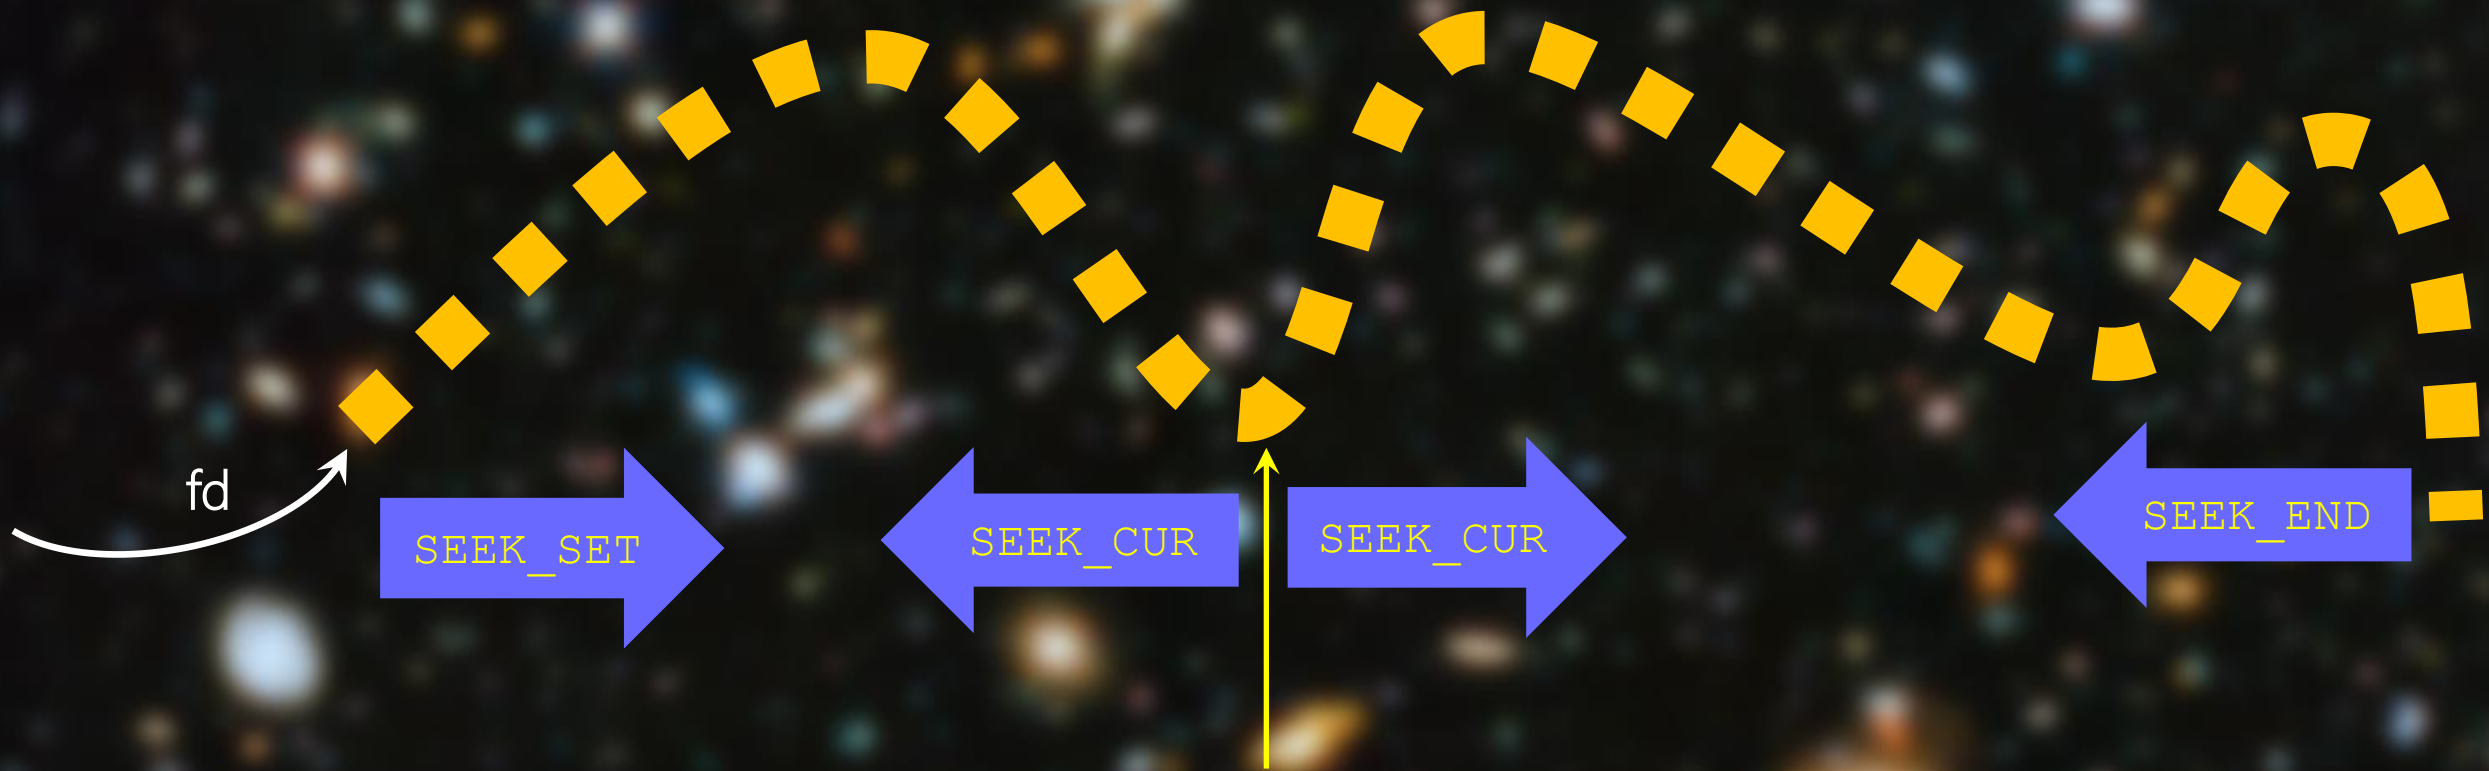
lseek

How many bytes move the current offset from what place or origin?

SEEK_SET, SEEK_CUR, SEEK_END

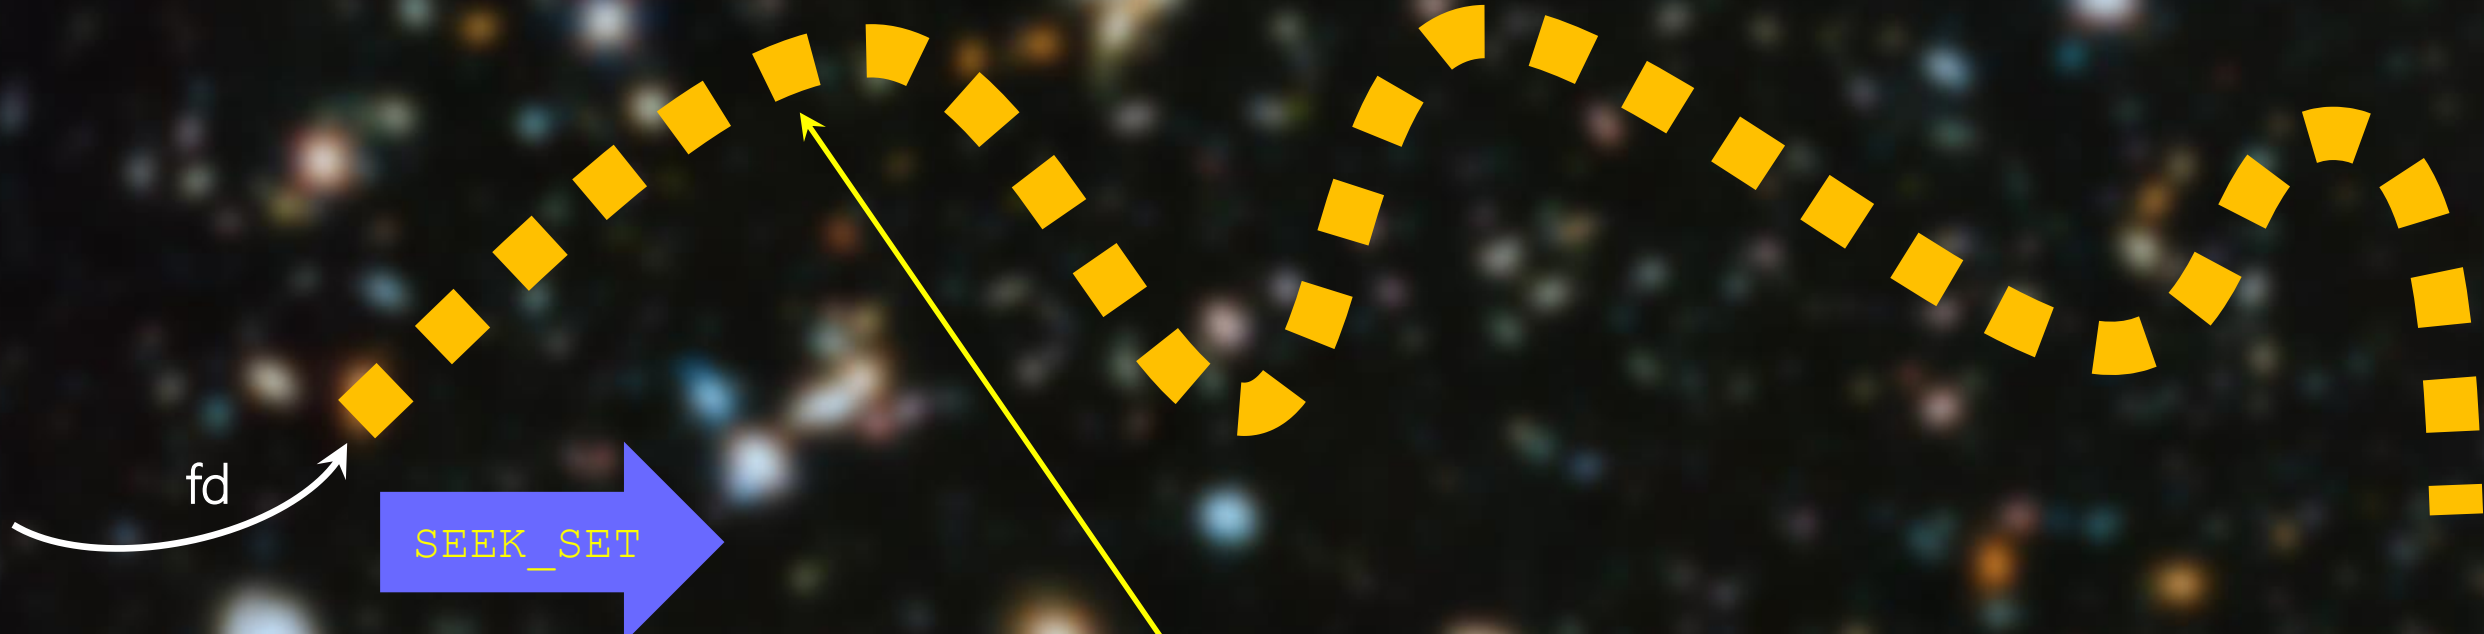
```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```


myphoto.jpg



current offset: 11

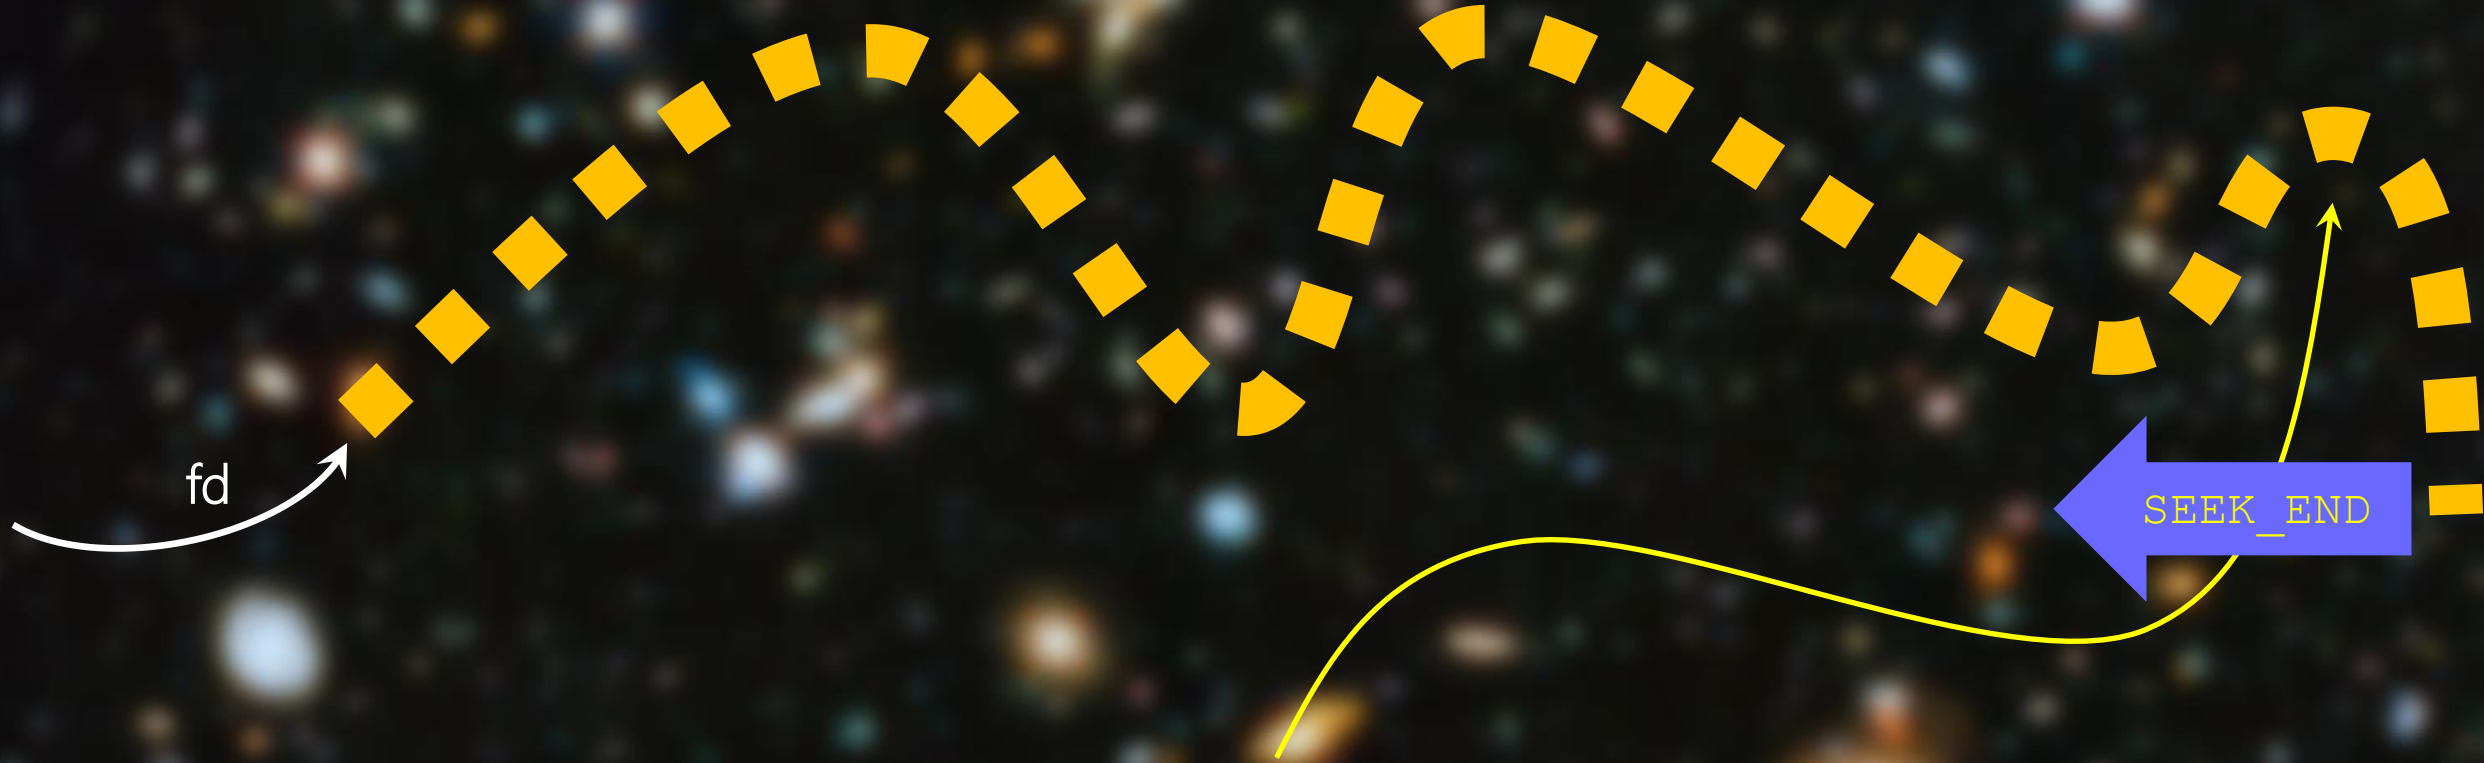
myphoto.jpg



current offset: 11 → 5

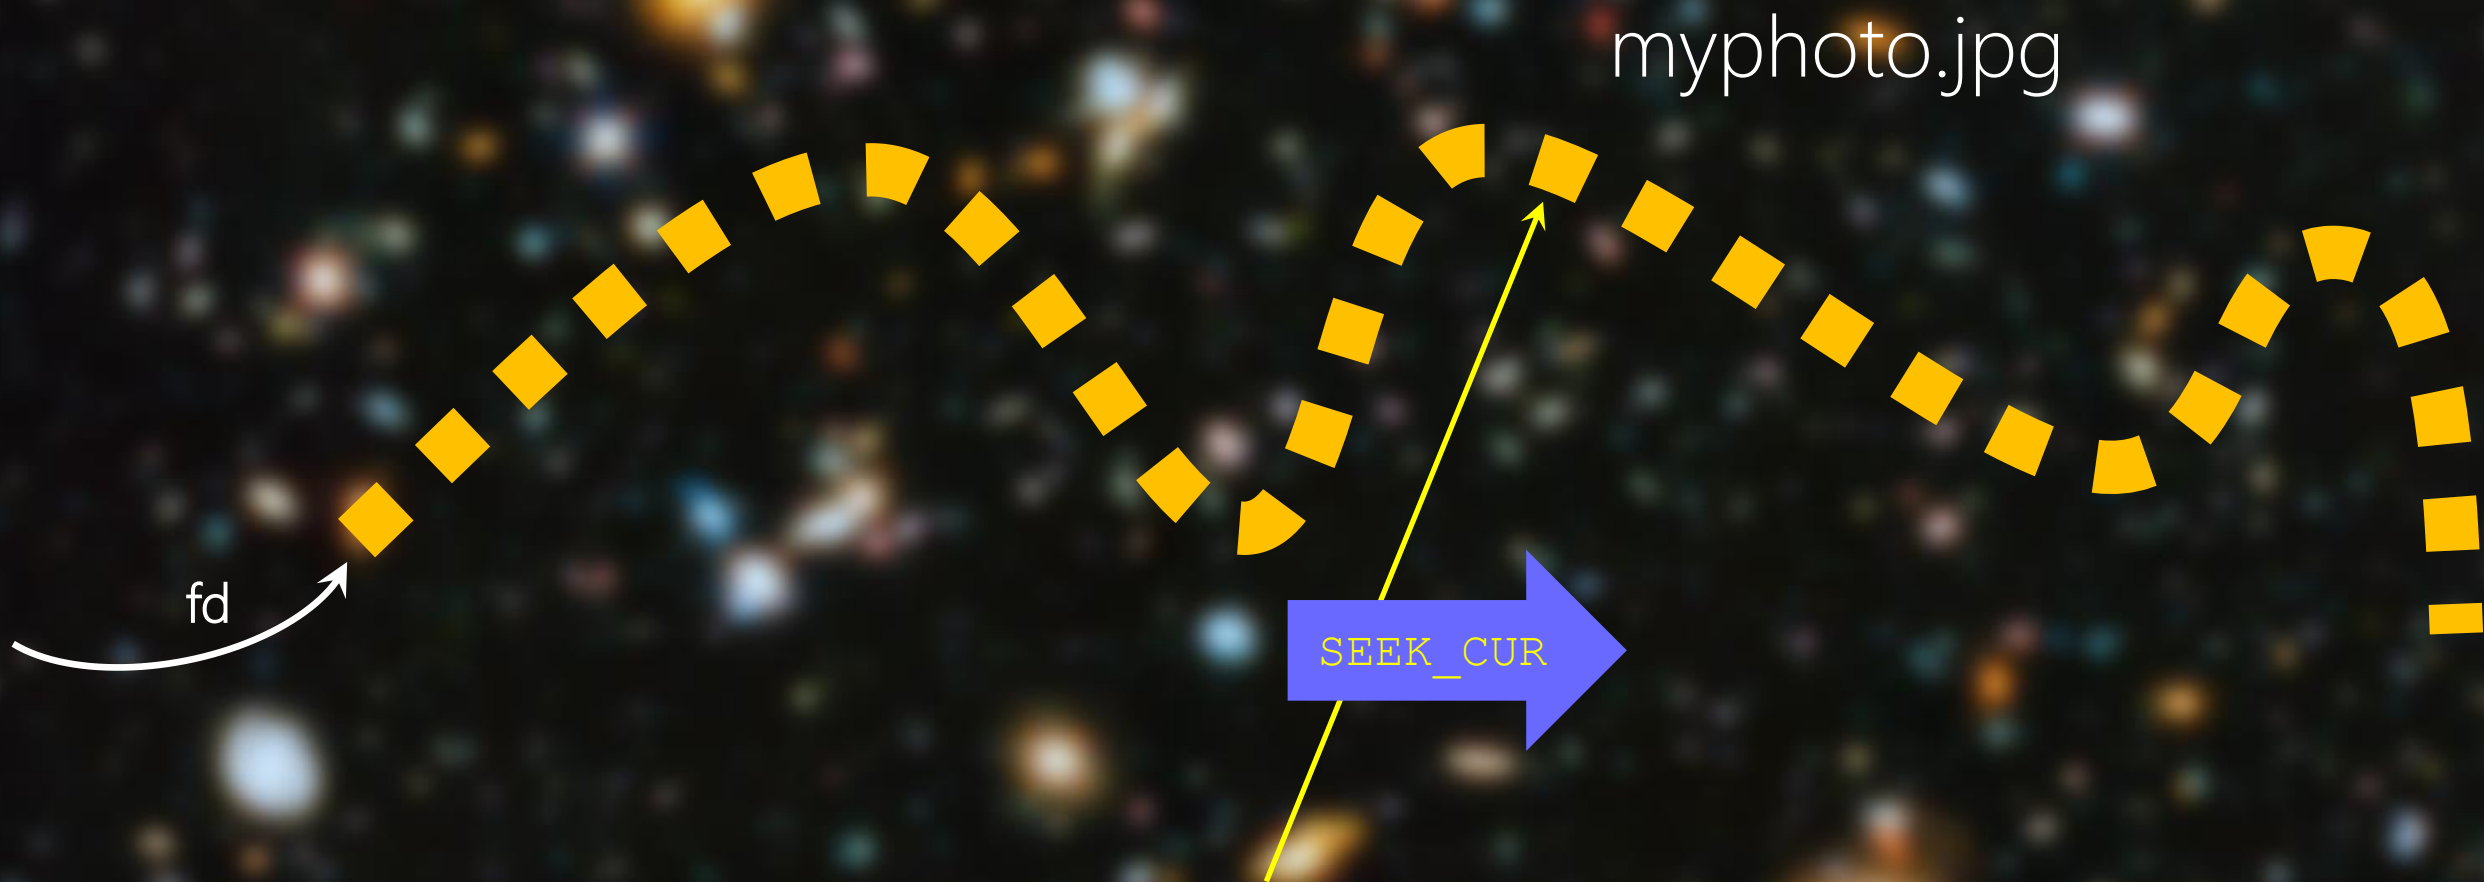
```
lseek(fd, 5, SEEK_SET)
```

myphoto.jpg



current offset: 11 → 25

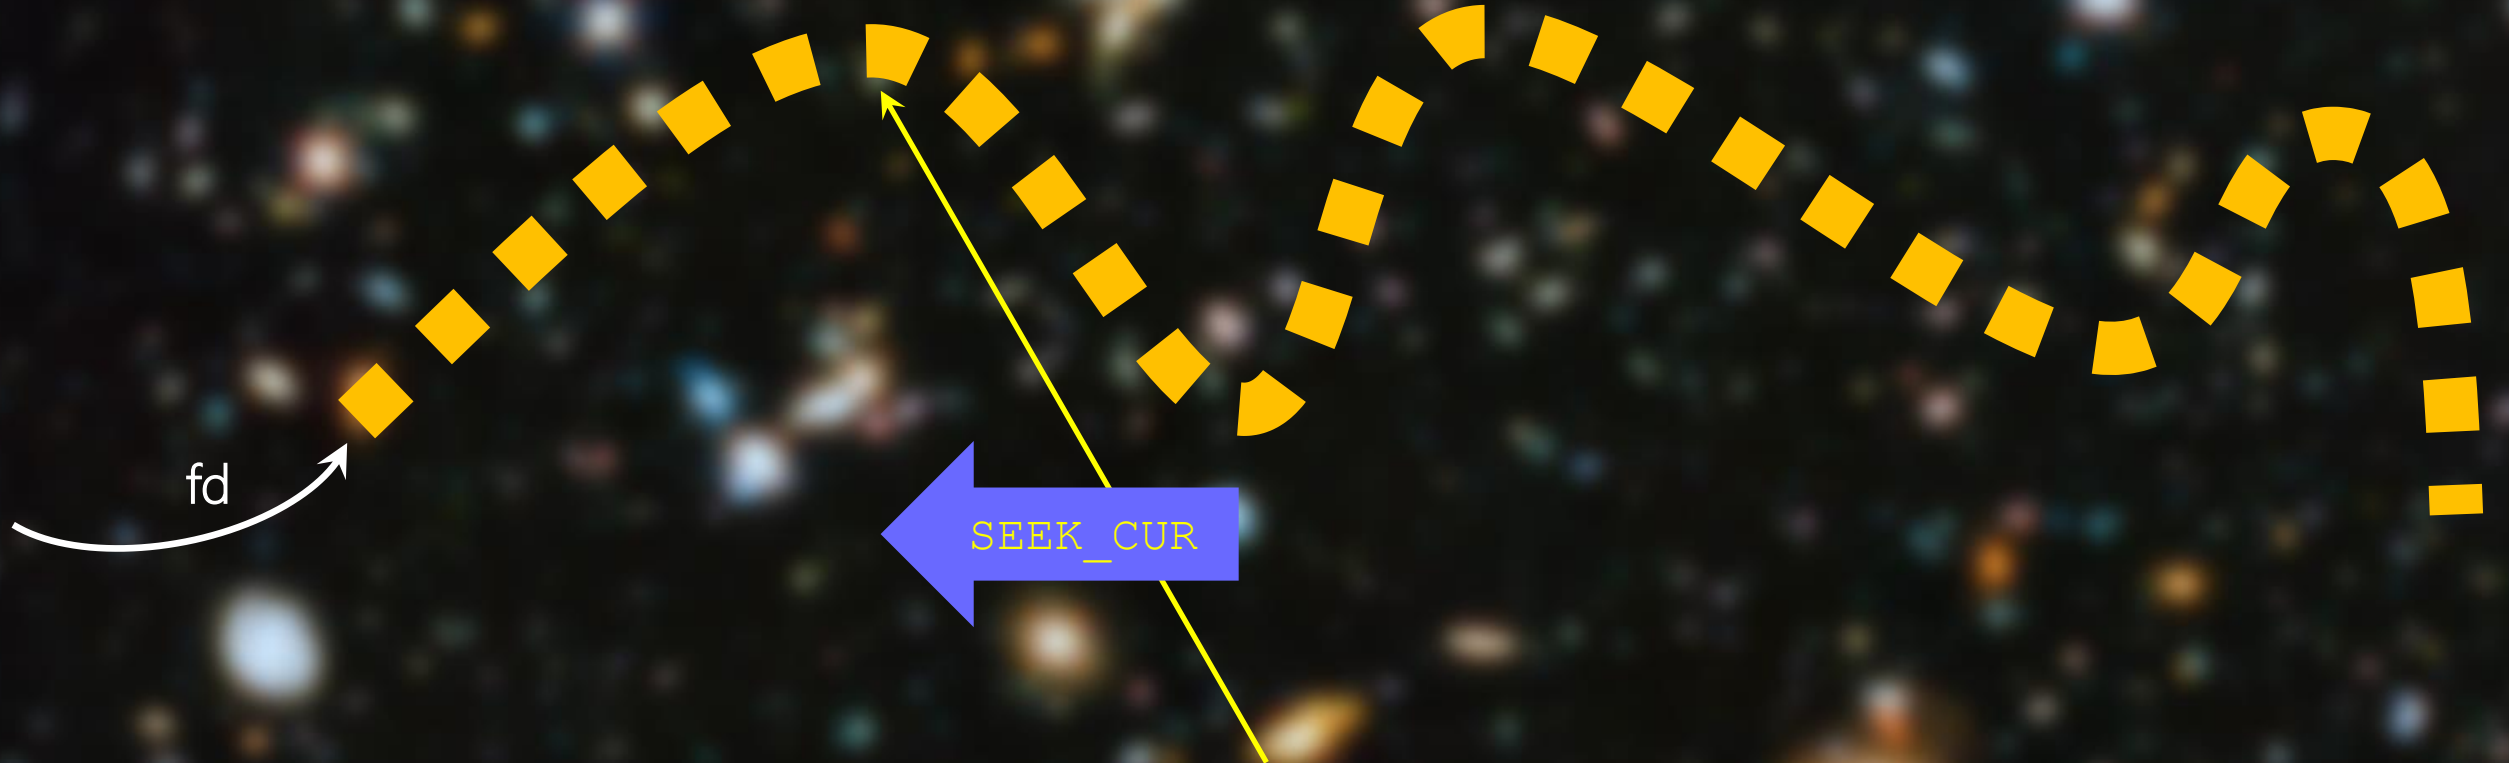
```
lseek(fd, 5, SEEK_END)
```

current offset: 11 \rightarrow +5 = 16

`lseek(fd, 5, SEEK_CUR)`

myphoto.jpg



current offset: $11 \rightarrow -5 = 6$

`lseek(fd, -5, SEEK_CUR)`

lseek

New current offset

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```


The background of the slide is a deep space image showing a dense field of galaxies in various colors (yellow, orange, blue, white) against a black background. A solid black horizontal bar spans the width of the slide, positioned in the middle. Two thin blue horizontal lines are located just above and just below the black bar.

lseek

How to know the value of current offset?