



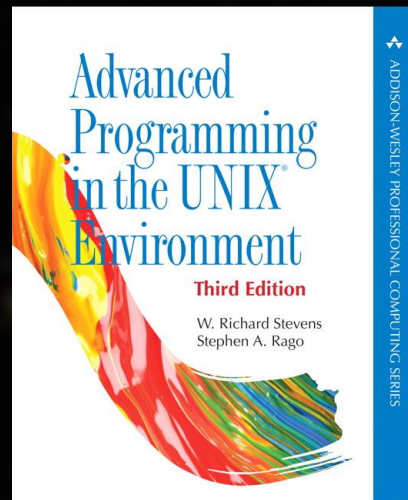
1-Week Extension to Lab06, Lec06

A photograph showing a line of people in graduation gowns and caps. Instead of human faces, the heads are replaced by various types of robot heads, some with long blonde hair, some with dark hair, and some with more complex, metallic-looking features. They are standing in a line against a brick wall.

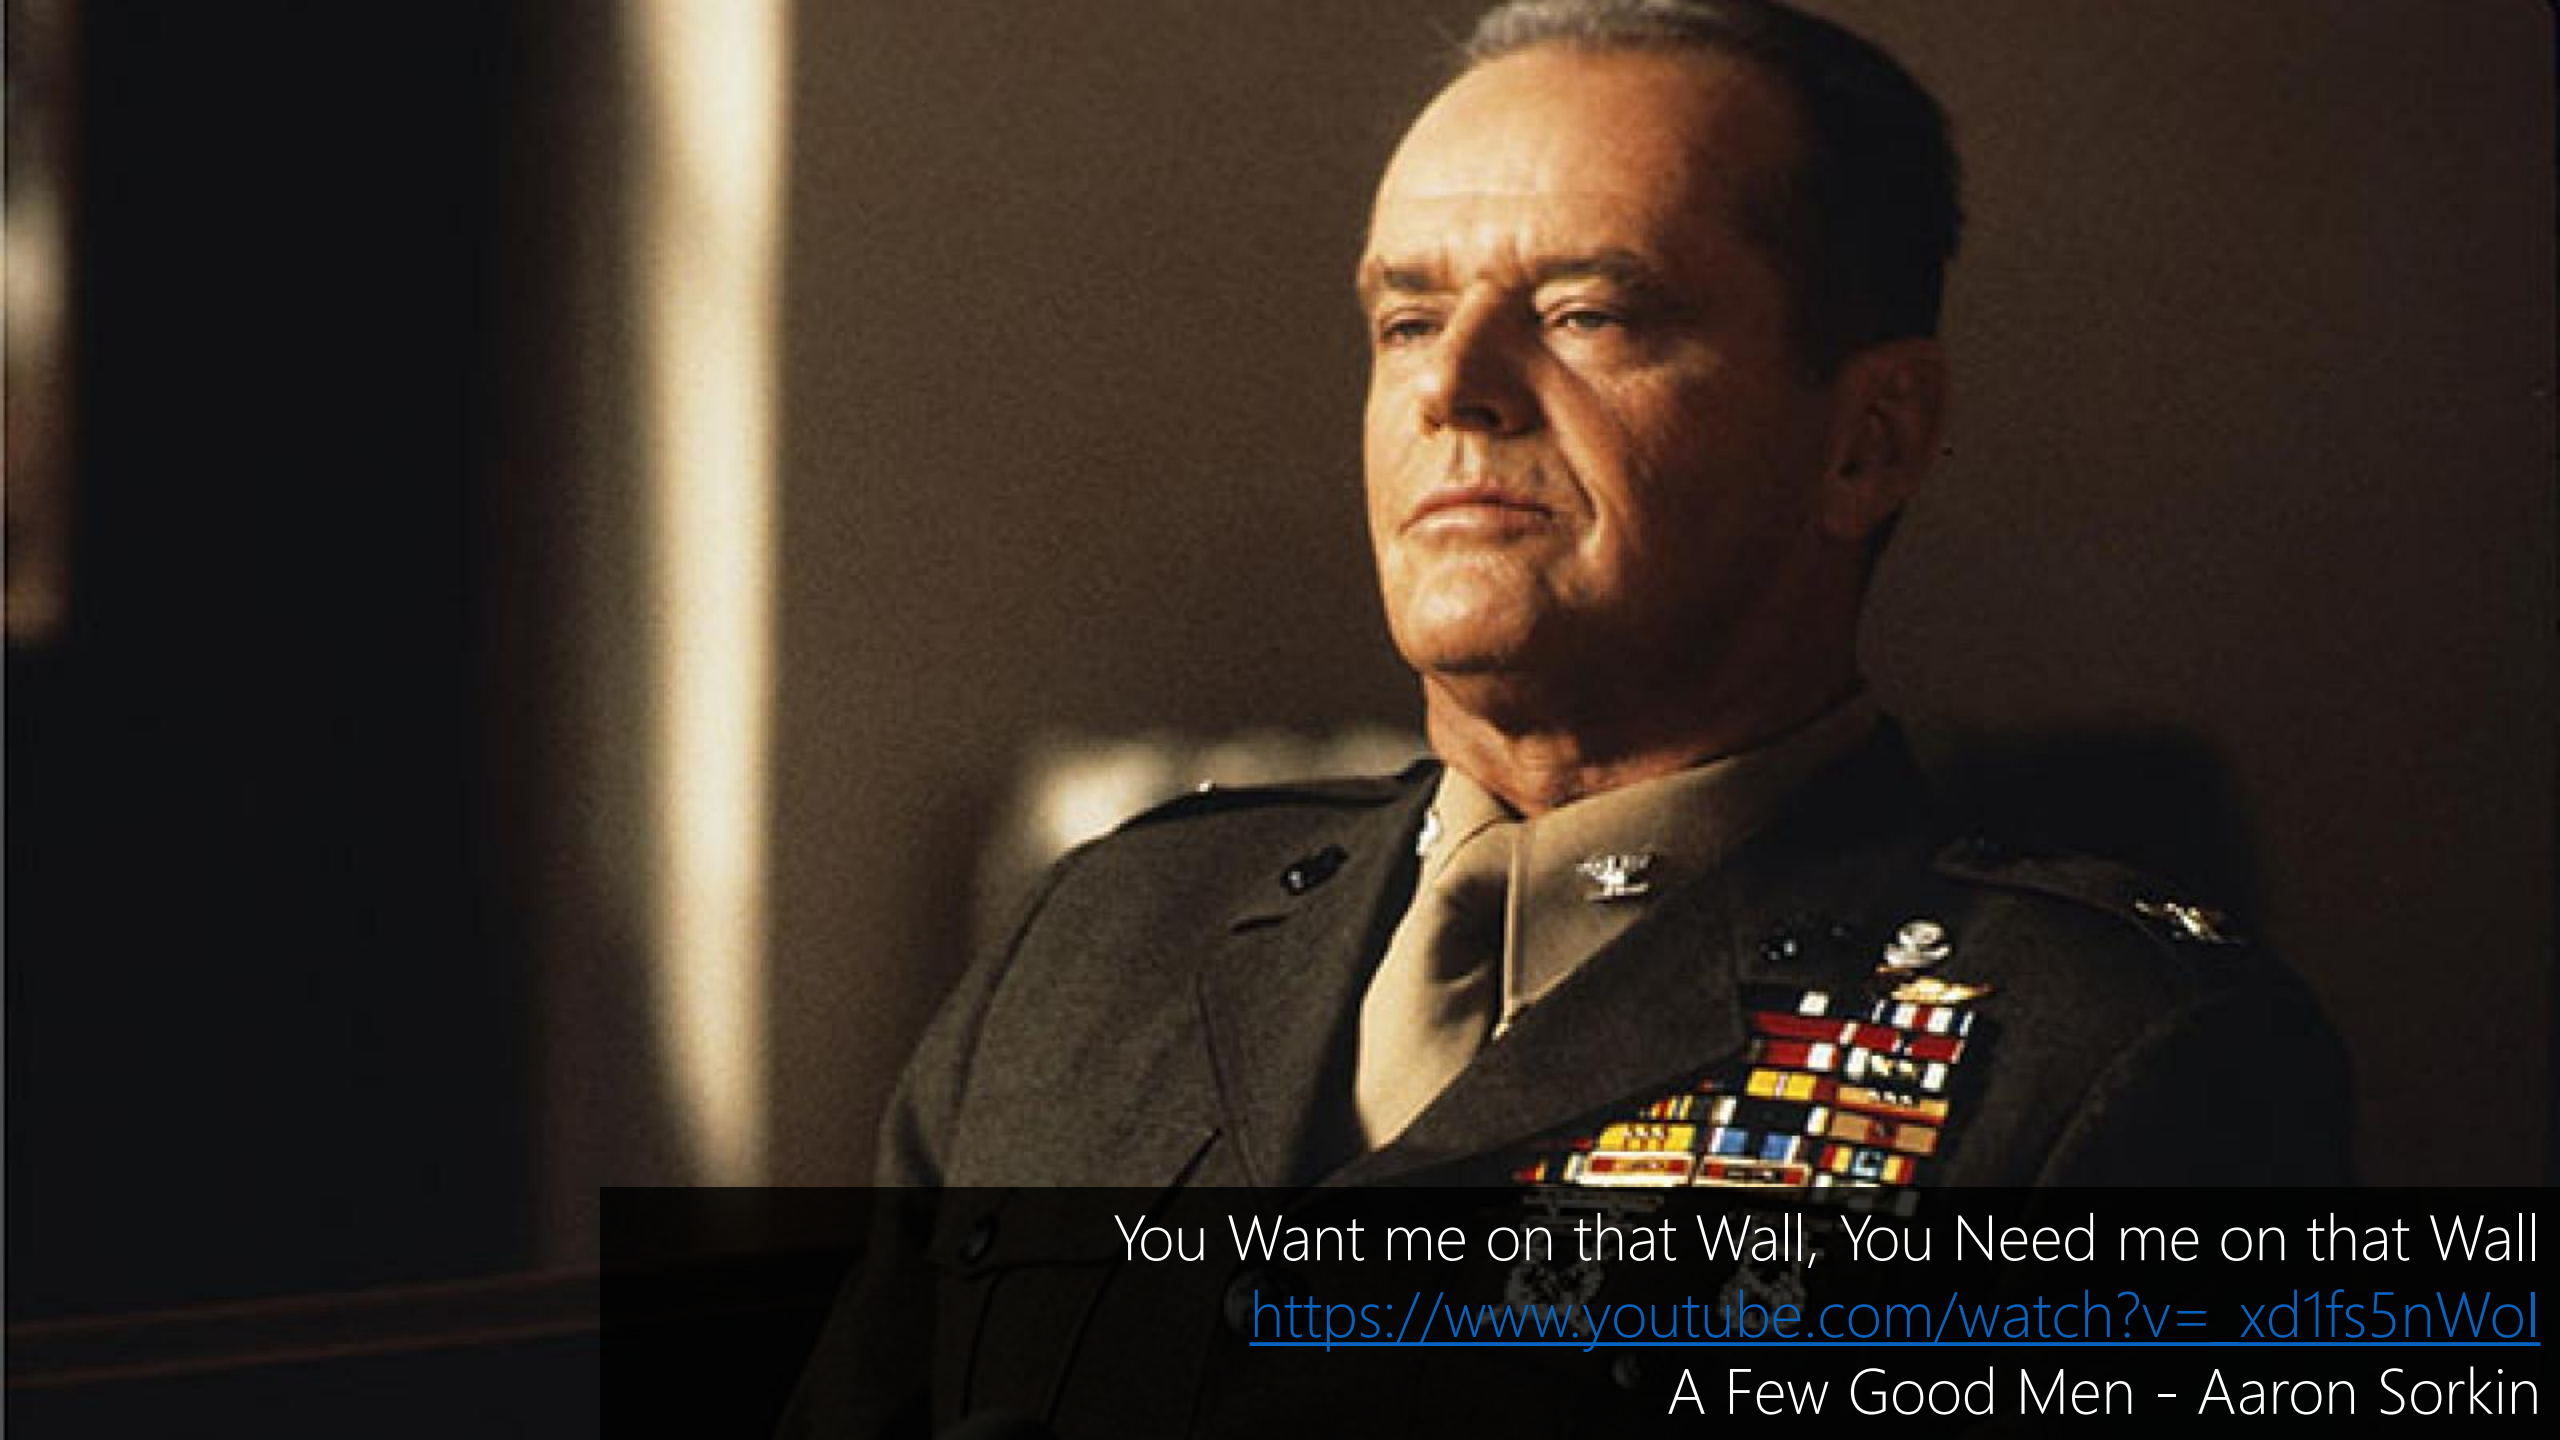
Pink Floyd - Another Brick In The Wall

Are robots becoming more human or humans becoming more robotic?





## Chapter 04: File I/O



You Want me on that Wall, You Need me on that Wall

<https://www.youtube.com/watch?v=xd1fs5nWol>

A Few Good Men - Aaron Sorkin

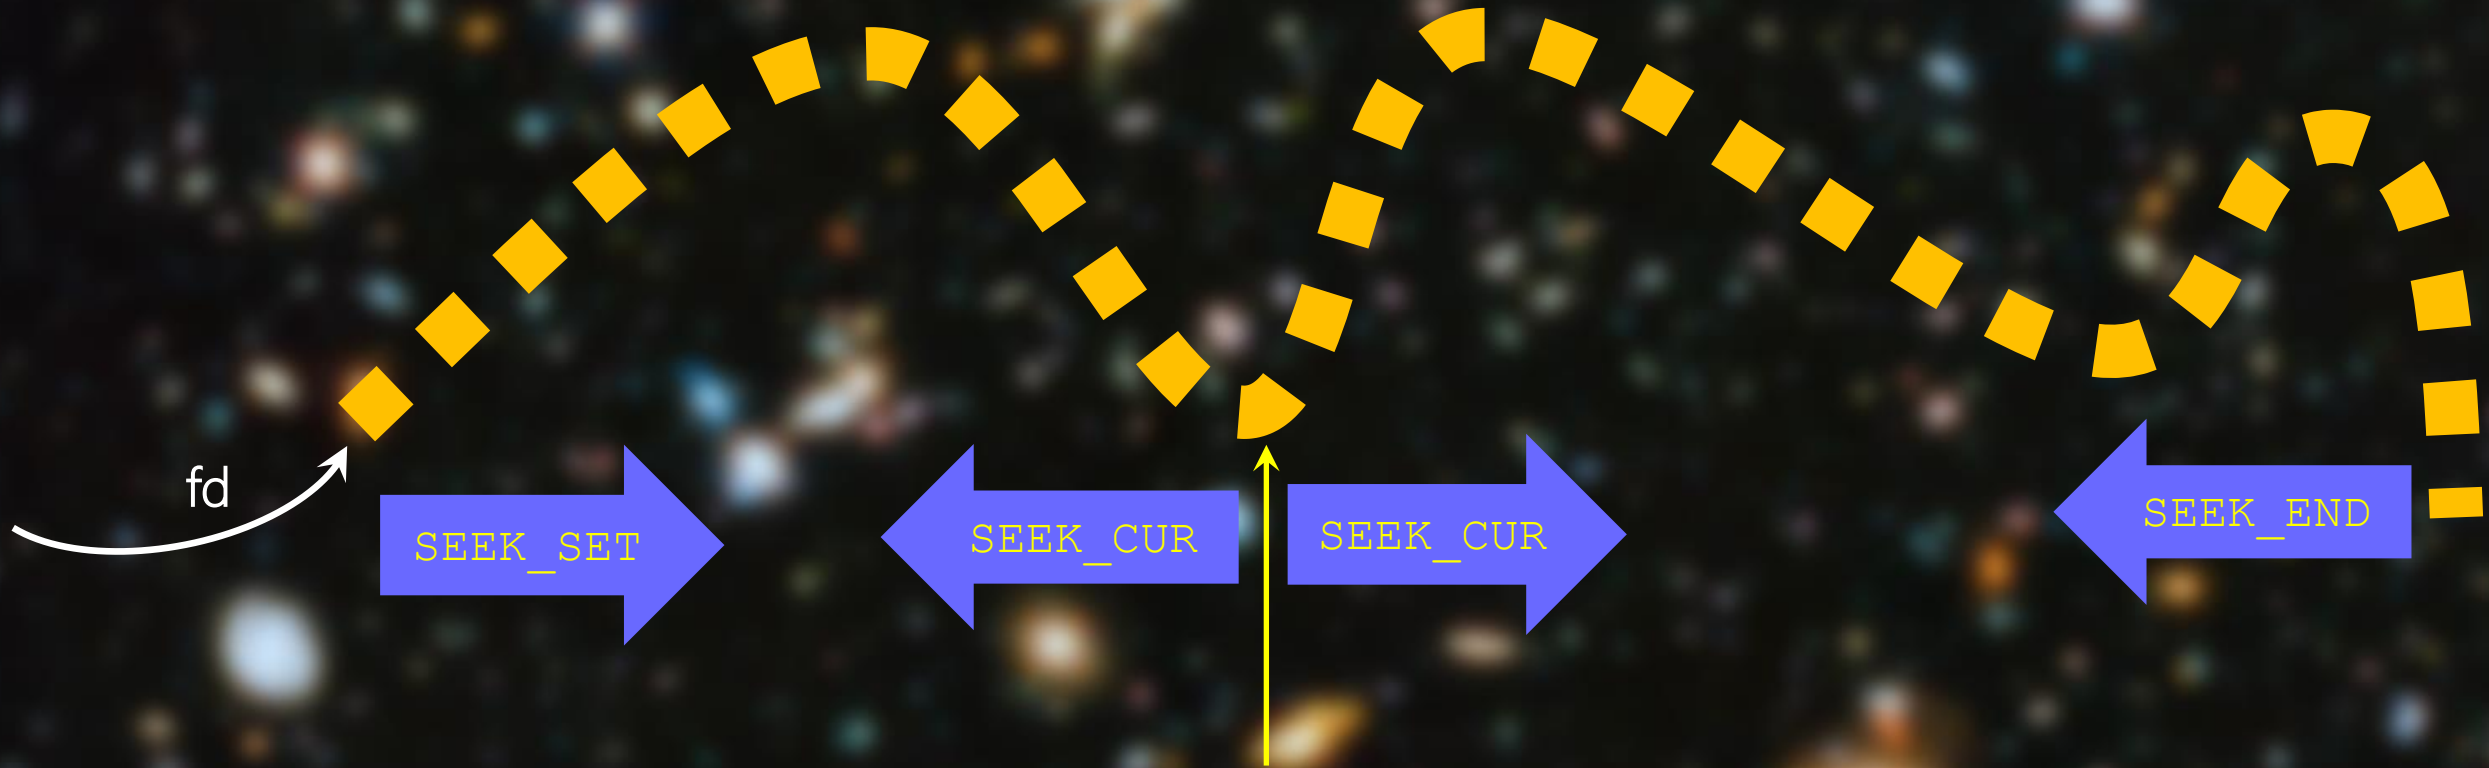
# lseek

POSIX

```
#include <unistd.h>
int lseek(int fd, off_t offset, int whence);
file's new offset if OK (can be negative)
-1 on error
```



myphoto.jpg



current offset: 11

```
hfani@charlie:~$ vi hole.c
```

```
#include <fcntl.h>
#include <unistd.h>
```

```
void main(void){
```

```
    int fd = open("./hole_test.txt", O_RDWR | O_CREAT, S_IRUSR | S_IWUSR);
```

```
    int cur_offset = lseek(fd, 10, SEEK_SET);
```

```
    char buf[20] = "write after the hole.";
```

```
    write(fd, buf, 20);
```

```
}
```

```
~
```

current offset = 0

move it 10 bytes ahead from start

write new bytes

```
hfani@charlie:~$ cc hole.c -o hole
```

```
hole.c: In function 'main':
```

```
hole.c:6:17: warning: initializer-string for array of 'char' is too long
```

```
    6 |   char buf[20] = "write after the hole.";
      |                   ^~~~~~~~~~~~~~~~~~~~~~
```

```
hfani@charlie:~$ ./hole
```

```
hfani@charlie:~$ vi hole_test.txt
```

```
^@^@^@^@^@^@^@^@^@^@^@^@write after the hole
```

```
~
```

```
hfani@charlie:~$ hexdump hole_test.txt
```

```
00000000 0000 0000 0000 0000 0000 7277 7469 2065
```

```
00000010 6661 6574 2072 6874 2065 6f68 656c
```

```
0000001e
```

```
hfani@charlie:~$ od -c hole_test.txt
```

```
00000000 \0 \0 \0 \0 \0 \0 \0 \0 \0 w r i t e
```

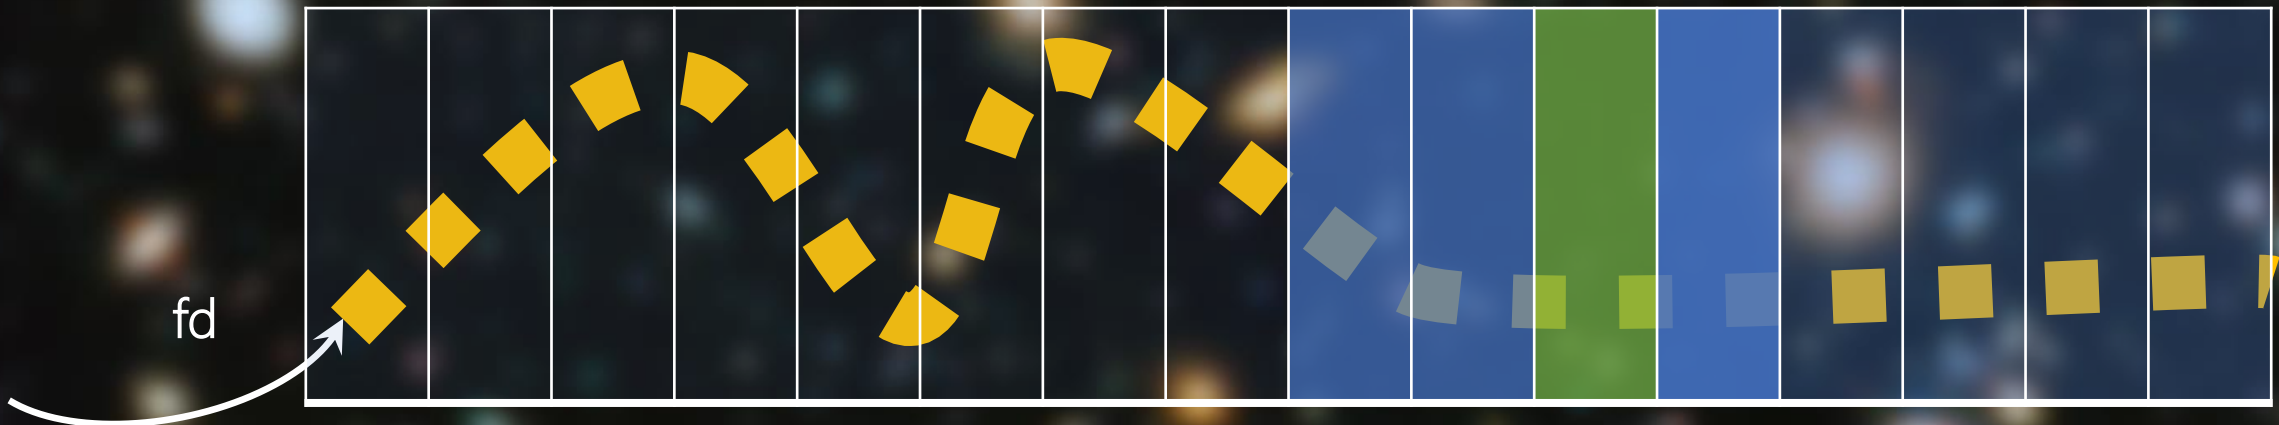
```
00000020 a f t e r t h e h o l e
```

```
00000036
```

---

# lseek

A Better User Case: Binary Search







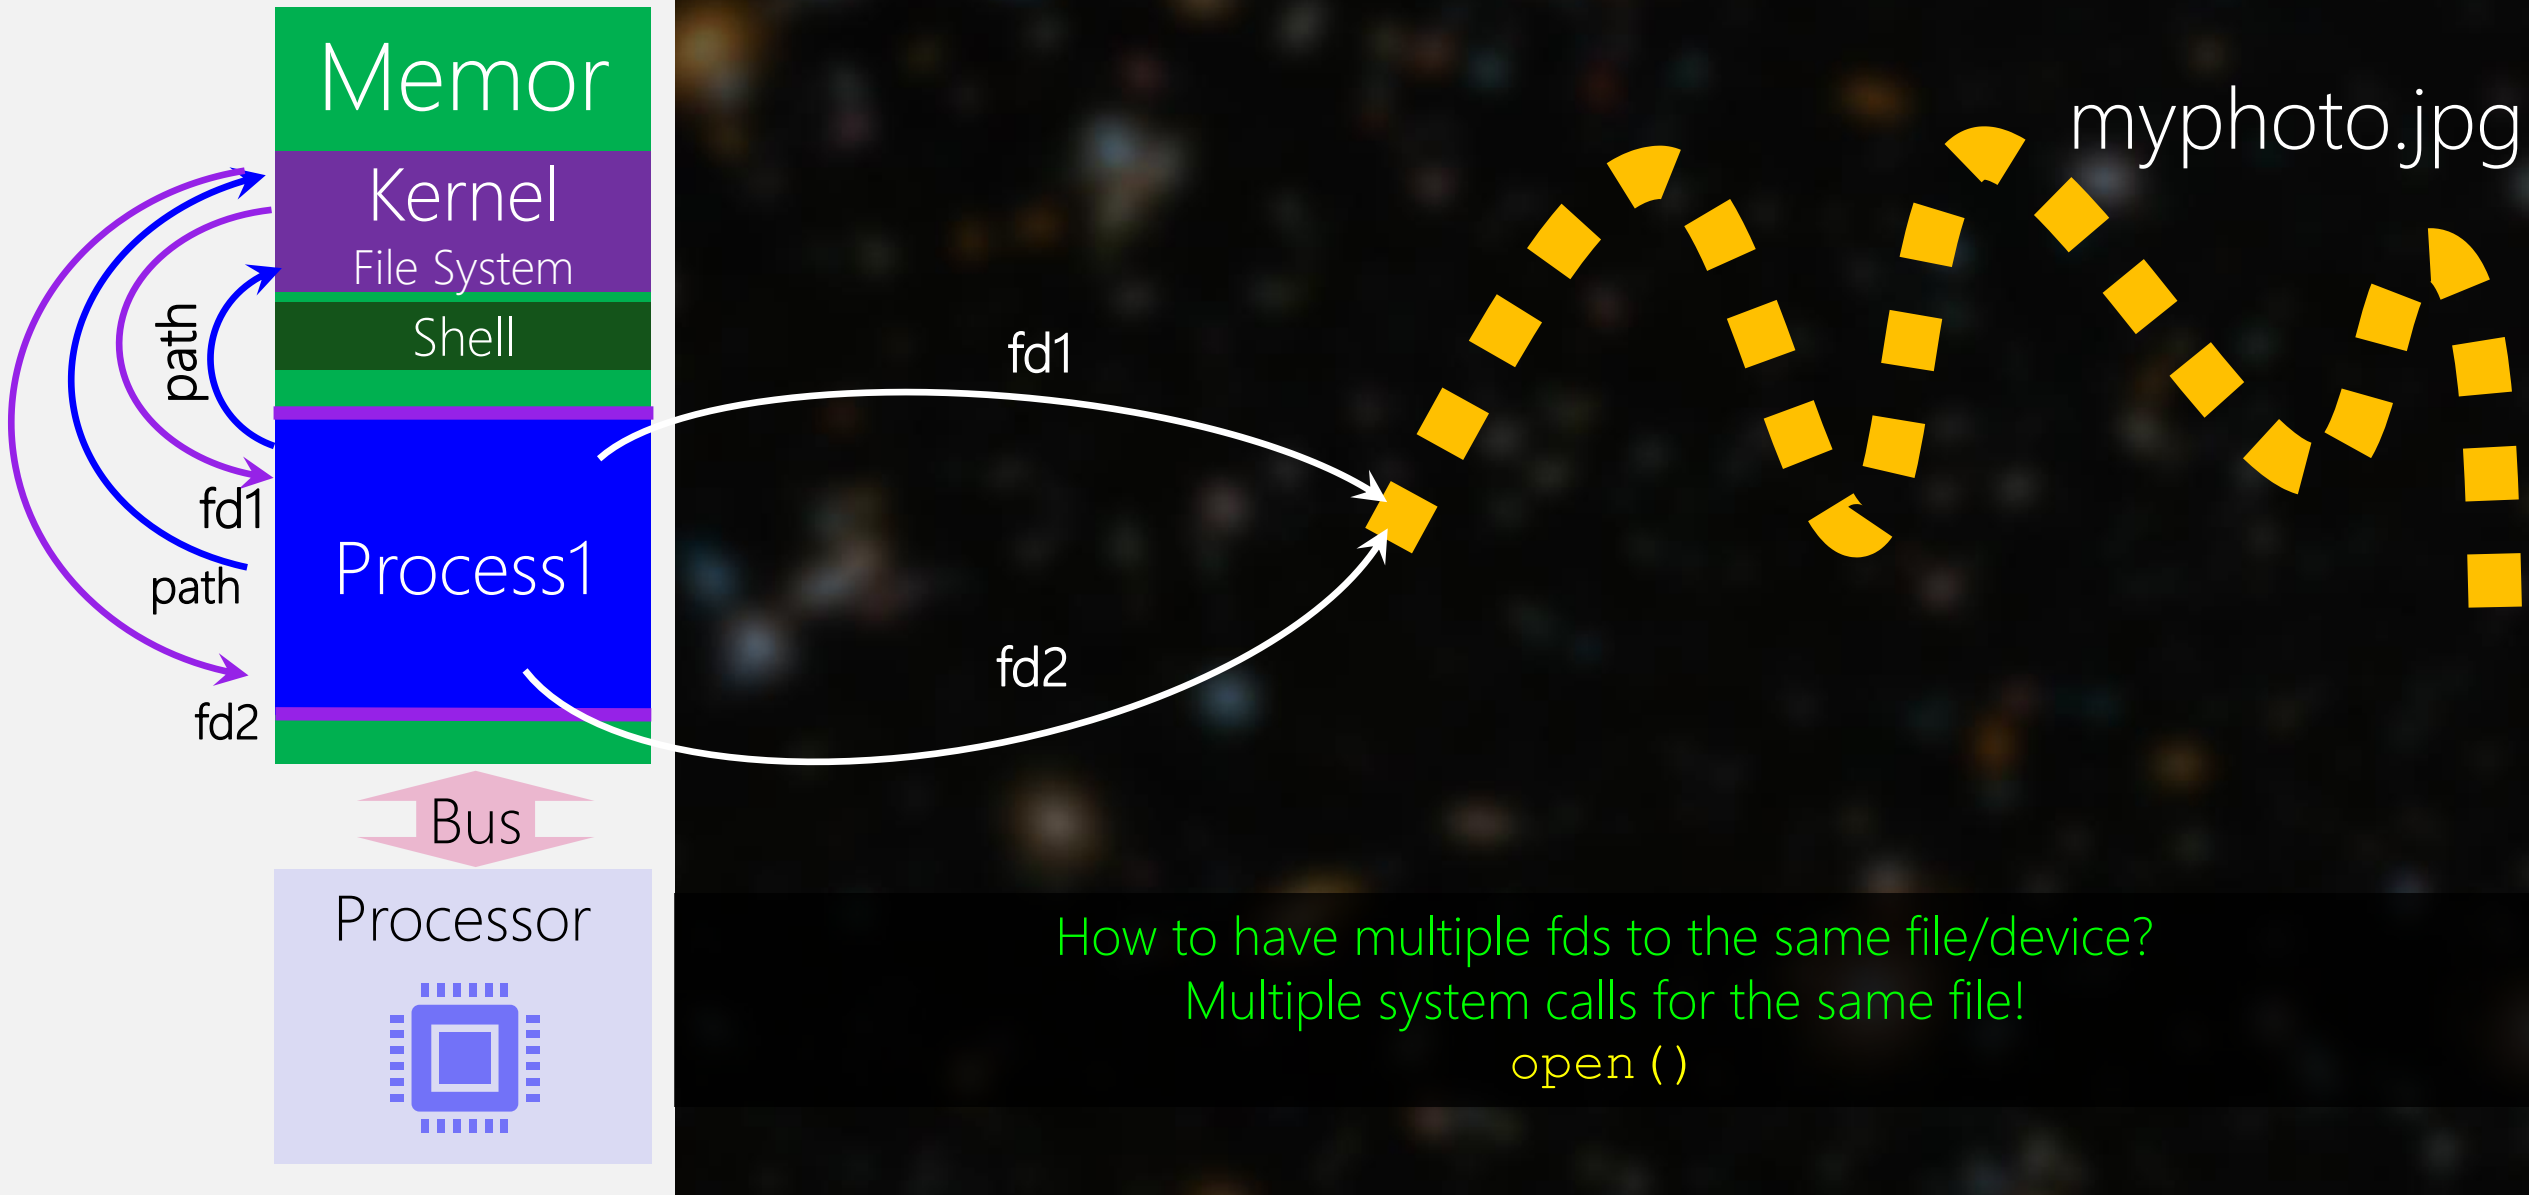
---

dup and dup2

---

File System: High-Level

# Computer



Kernel  
File System

Process1

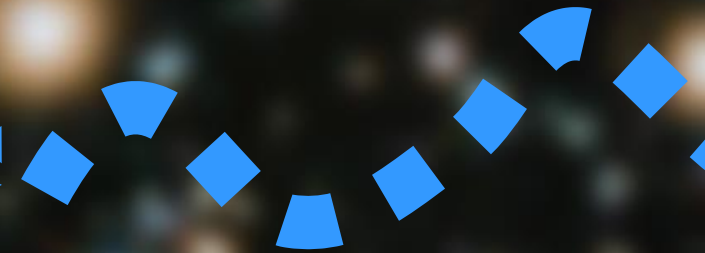
File Descriptors	File Pointer
fd1	
fd2	
fd3	
fd4	
...	

File Pointer
current file offset
other flags
pointer to first byte

File Pointer
current file offset
other flags
pointer to first byte

File Pointer
current file offset
other flags
pointer to first byte

File Pointer
current file offset
other flags
pointer to first byte





---

# dup

---

```
#include <unistd.h>
int dup(int fd);
```

new file descriptor **to the same file** if OK, -1 on error

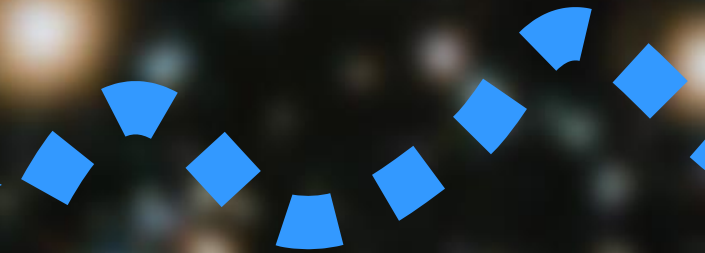
Kernel  
File System

Process1

File Descriptors	File Pointer
fd1	
fd2	
fd3	
fd4	
...	

File Pointer
current file offset
other flags
pointer to first byte

File Pointer
current file offset
other flags
pointer to first byte



---

dup

Why do we duplicate fd?

---

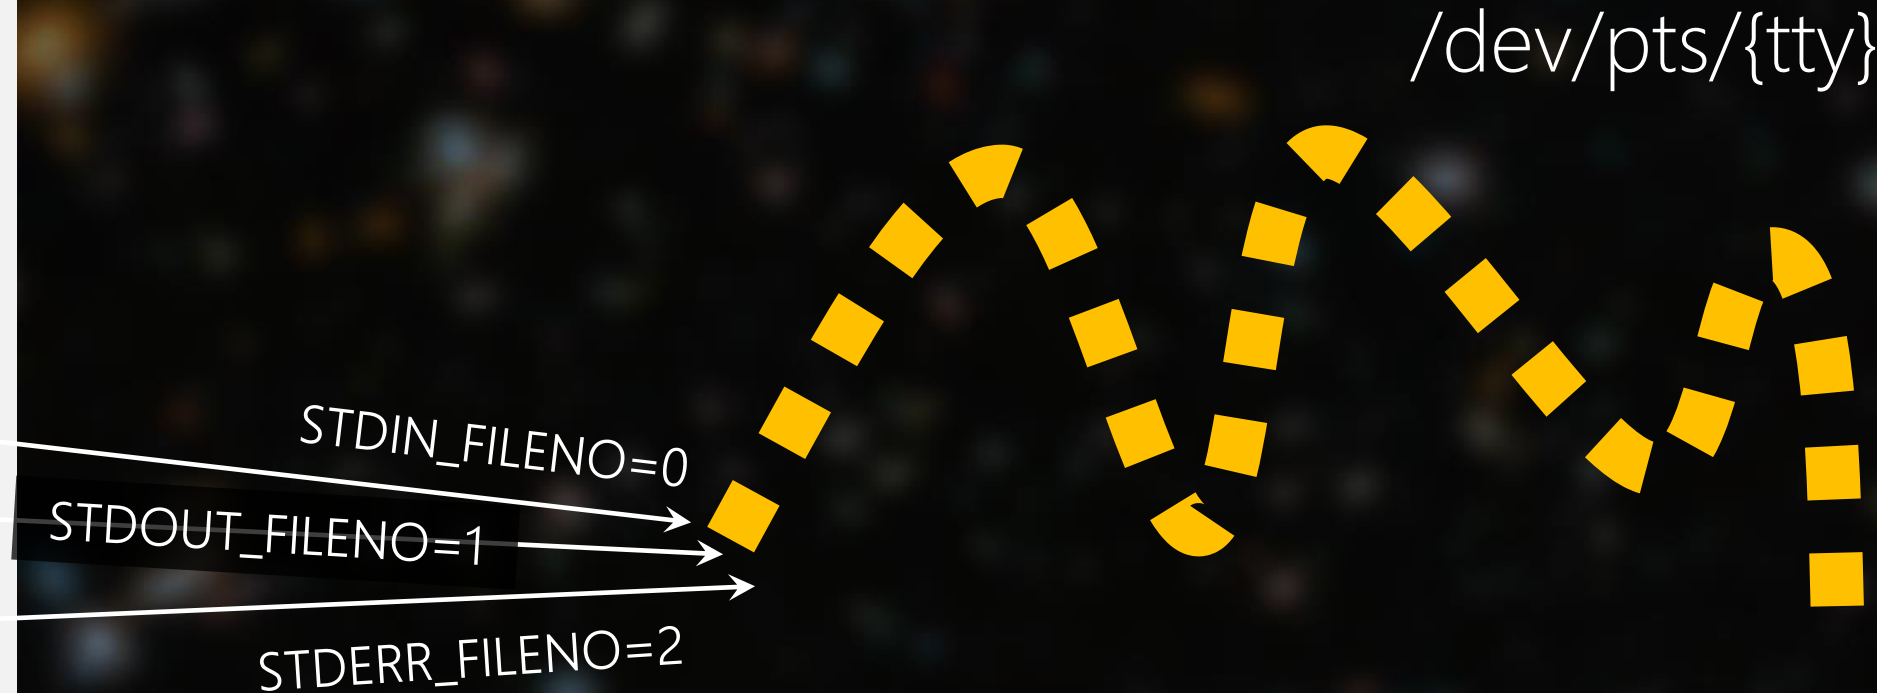
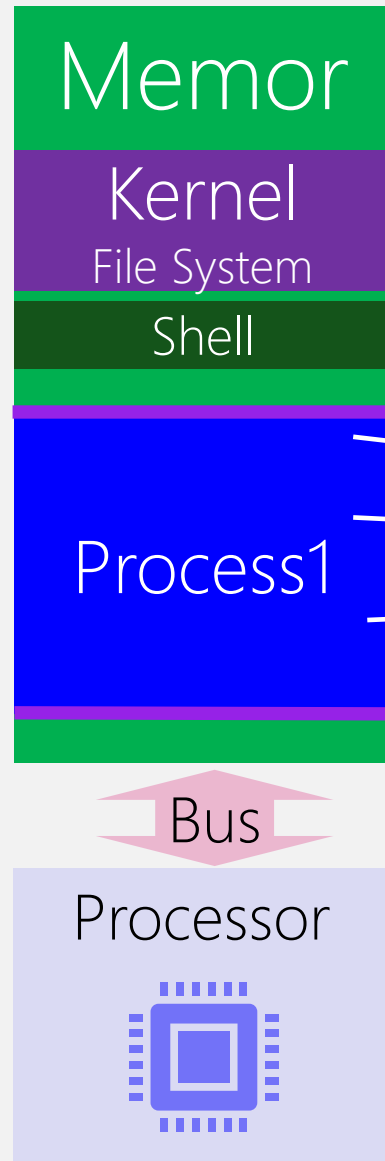
I/O Redirection

STDIO ↔ File

STDERR → STDOUT



# Computer

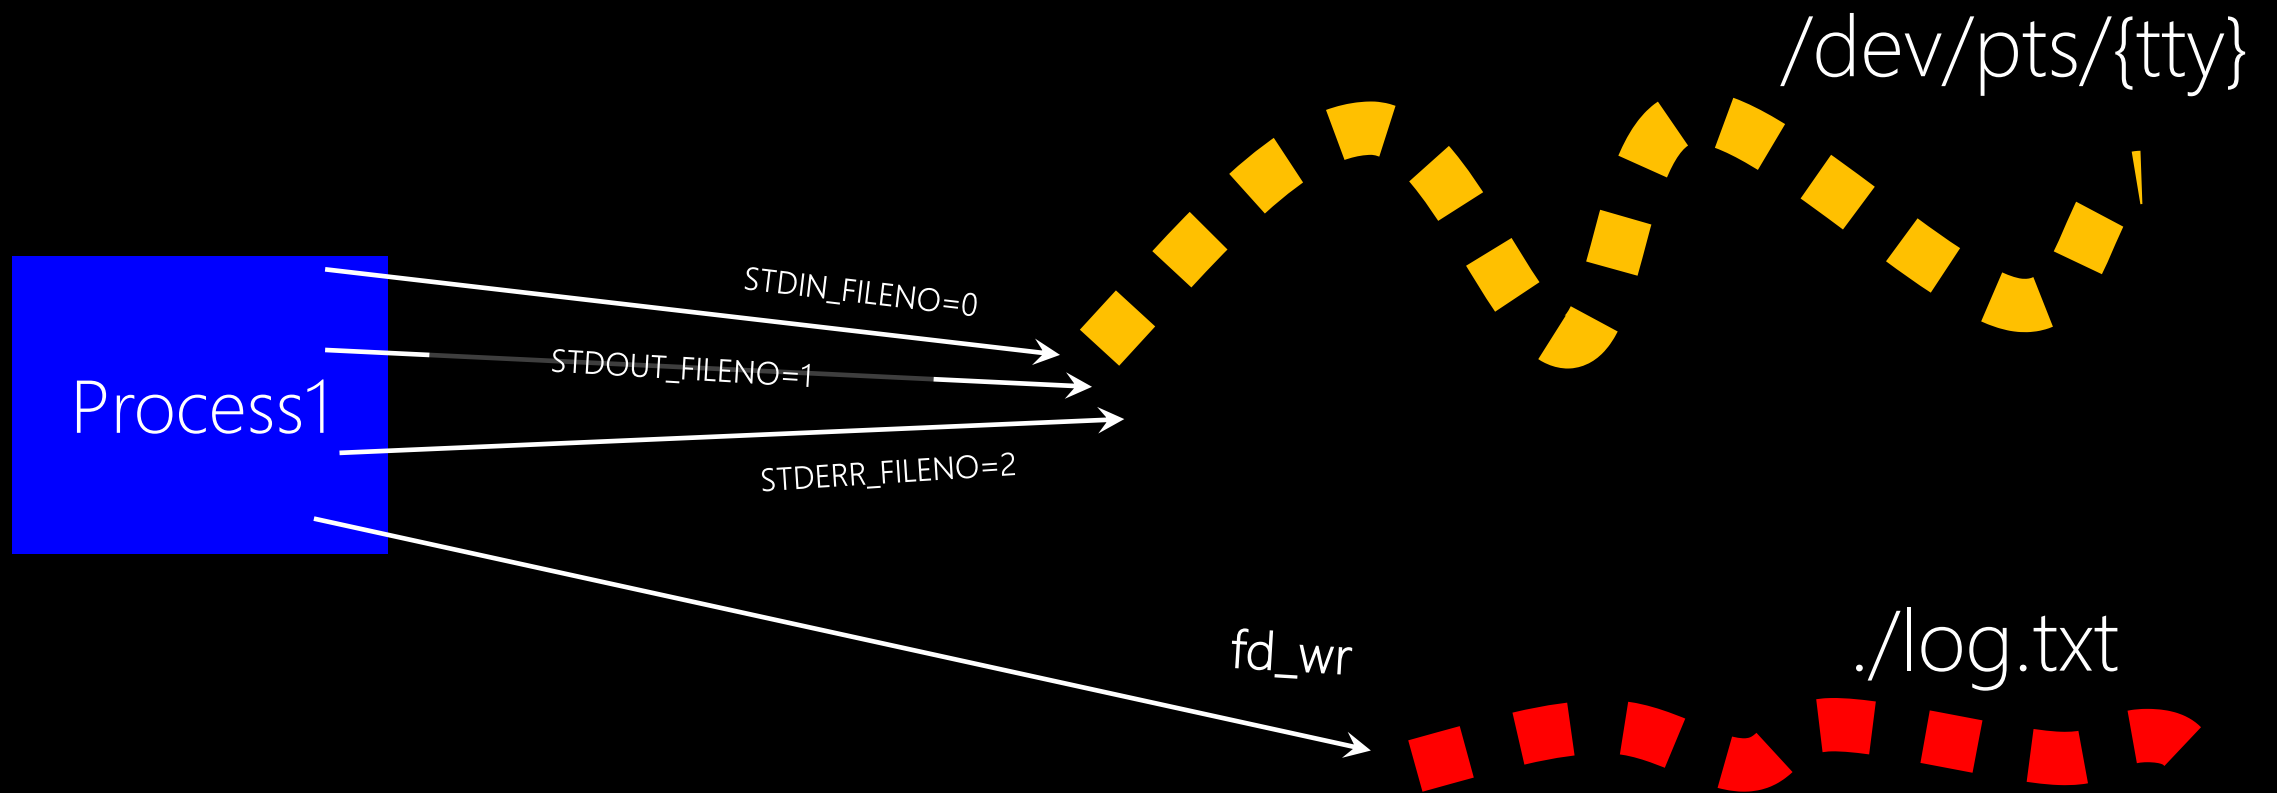


When Shell bootstraps a program, it automatically opens three fds for the program (process):

`STDIN_FILENO = 0 : O_RDONLY`  
`STDOUT_FILENO = 1 : O_WRONLY`  
`STDERR_FILENO = 2 : O_WRONLY`

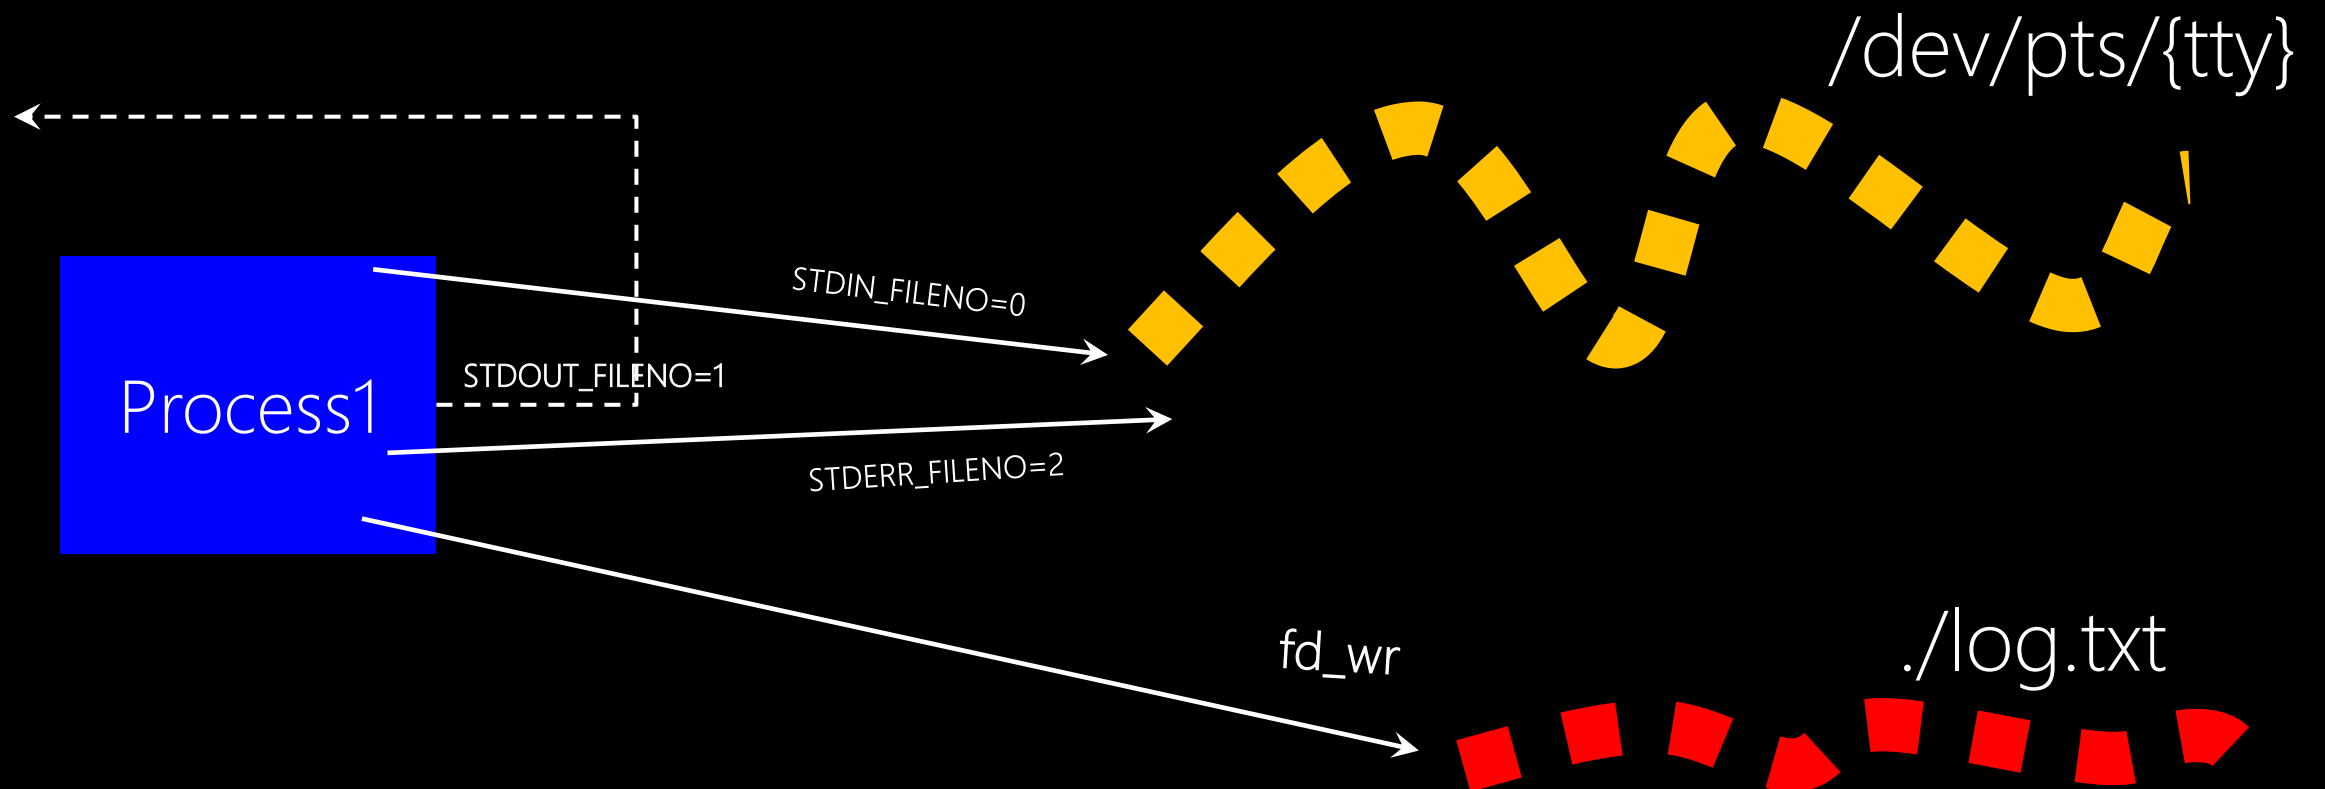
```
#include <fcntl.h>
#include <unistd.h>
void main(void){
    char buf_rd[20];

    int fd_wr = open("./log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
```



```
#include <fcntl.h>
#include <unistd.h>
void main(void){
    char buf_rd[20];

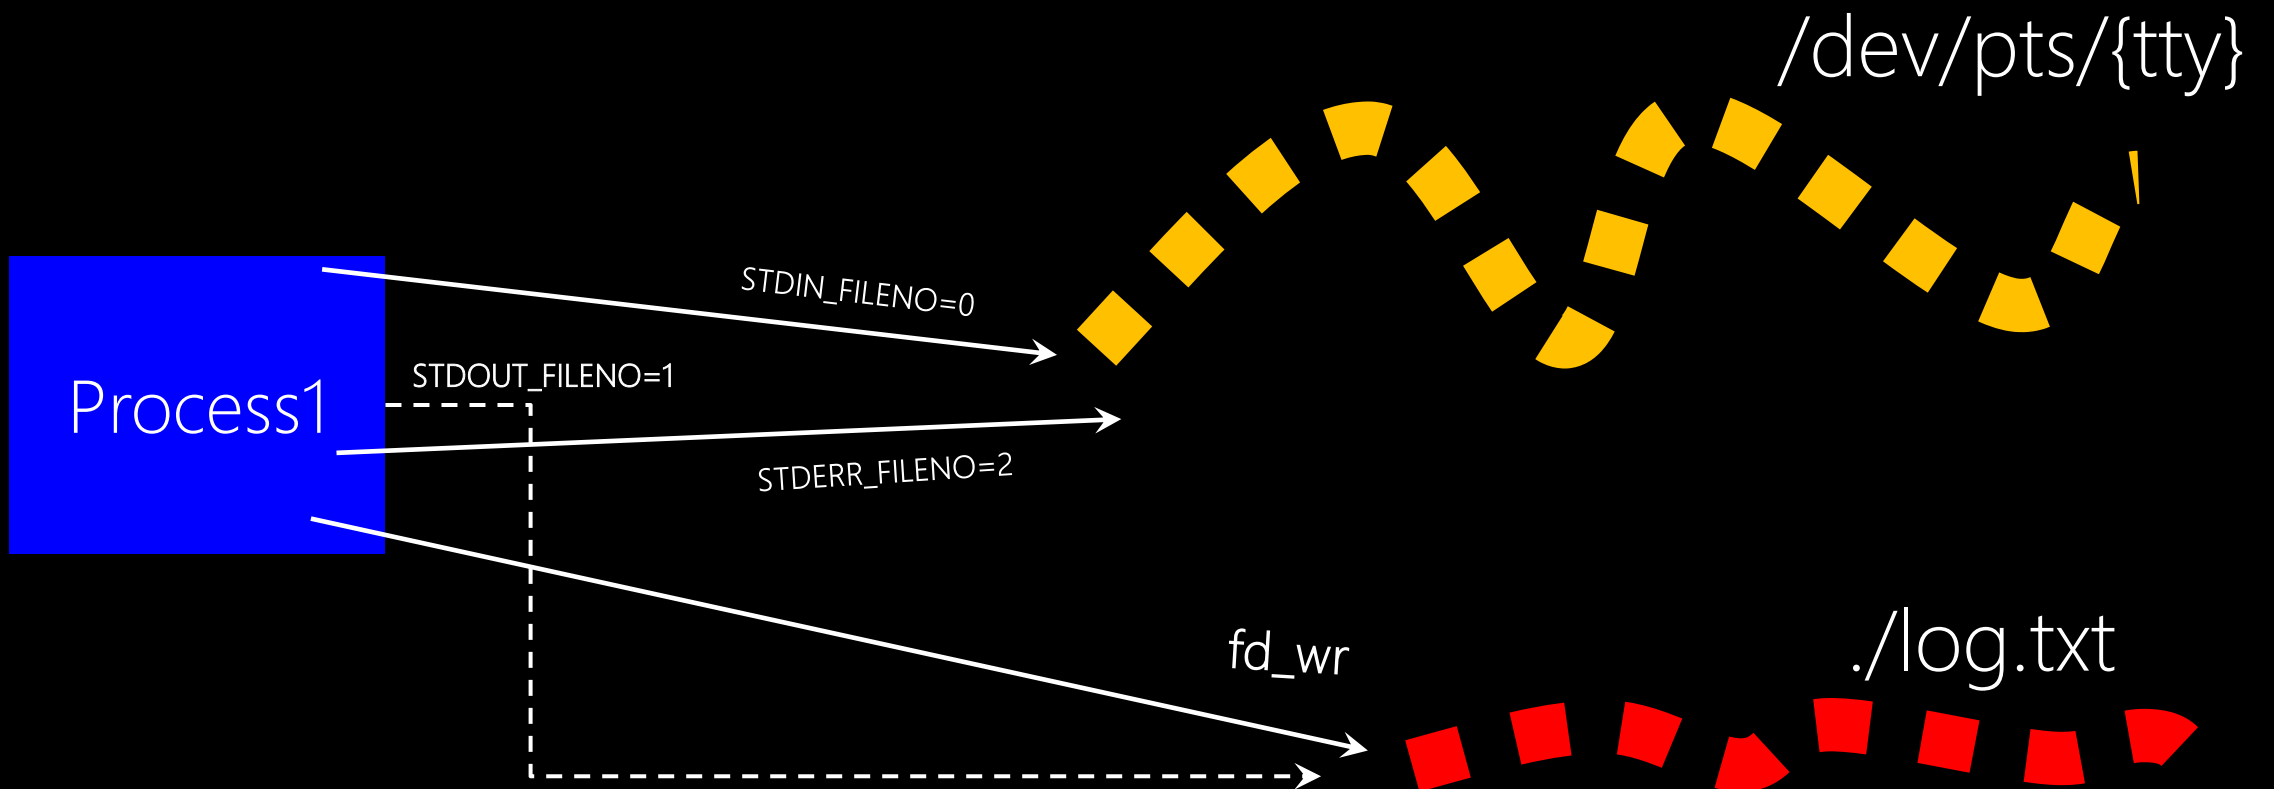
    int fd_wr = open("./log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    //now the fd with value 1 is free
```





```
#include <fcntl.h>
#include <unistd.h>
void main(void) {
    char buf_rd[20];

    int fd_wr = open("./log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    //now the fd with value 1 is free
    //let's get it for our log file
    int new_fd = dup(fd_wr);
    //now the value of new_fd is 1
    //both fd_wr and new_fd are pointing to log.txt
```



```

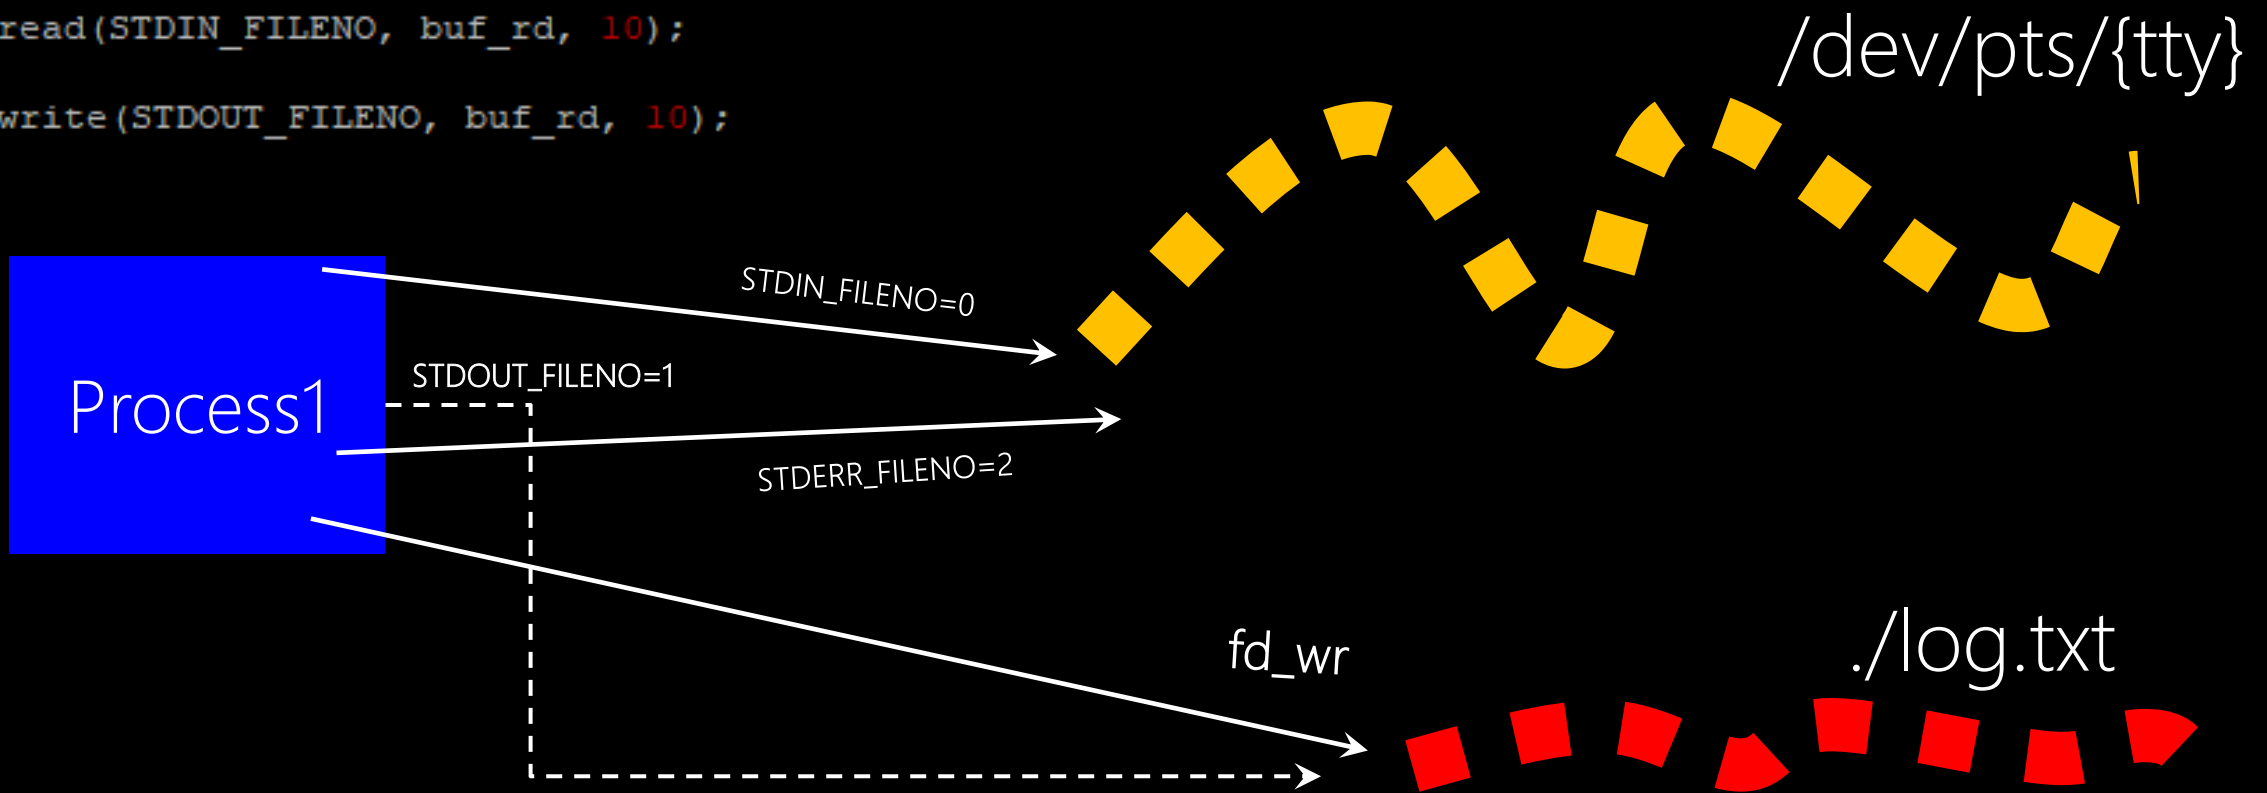
#include <fcntl.h>
#include <unistd.h>
void main(void){
    char buf_rd[20];

    int fd_wr = open("./log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    //now the fd with value 1 is free
    //let's get it for our log file
    int new_fd = dup(fd_wr);
    //now the value of new_fd is 1
    //both fd_wr and new_fd are pointing to log.txt

    read(STDIN_FILENO, buf_rd, 10);

    write(STDOUT_FILENO, buf_rd, 10);

```



# Shell

```
$ ./program > log.txt
```

We can ask the shell to do this redirection for us

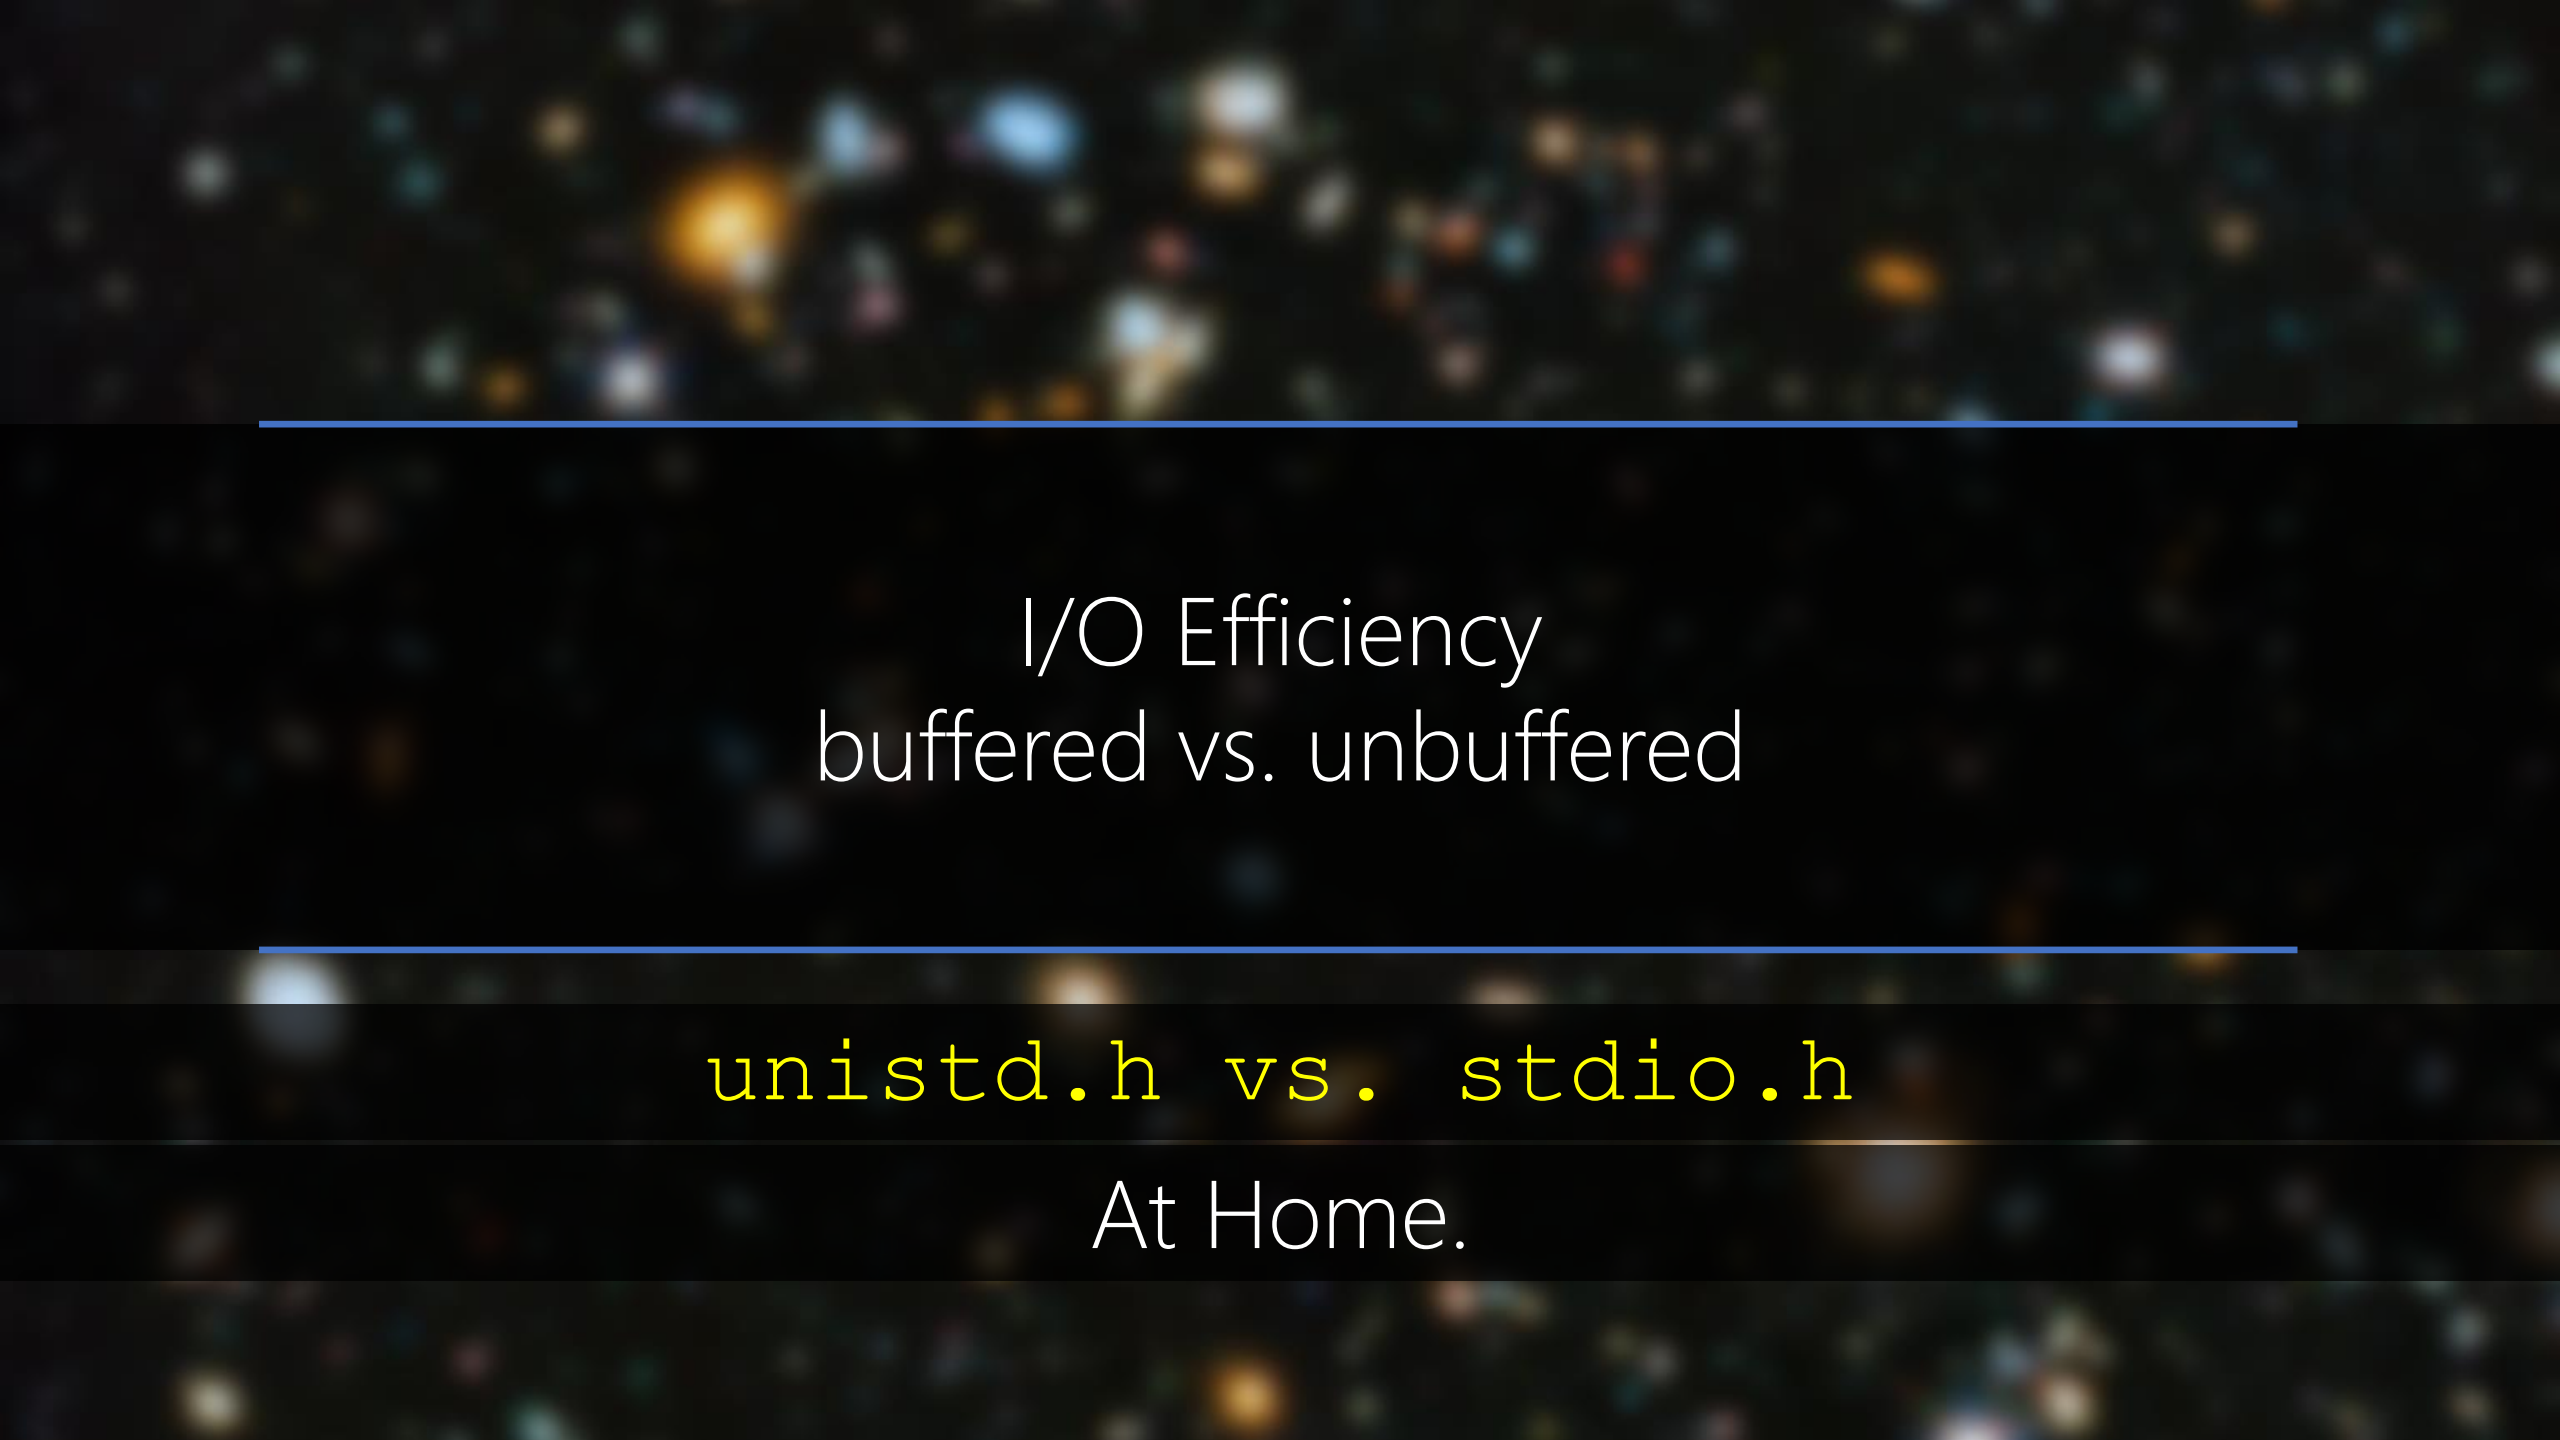
*`{program file} > {new destination for STDOUT_FILENO}`*

The image features a cosmic background of distant galaxies and stars in various colors (yellow, orange, blue, white) against a black space. A horizontal black band runs across the middle of the image. Two thin blue lines are positioned above and below this band, spanning most of the image width.

dup2

At Home





I/O Efficiency  
buffered vs. unbuffered

`unistd.h` vs. `stdio.h`

At Home.



---

File Sharing

---

Advanced! We won't cover it.



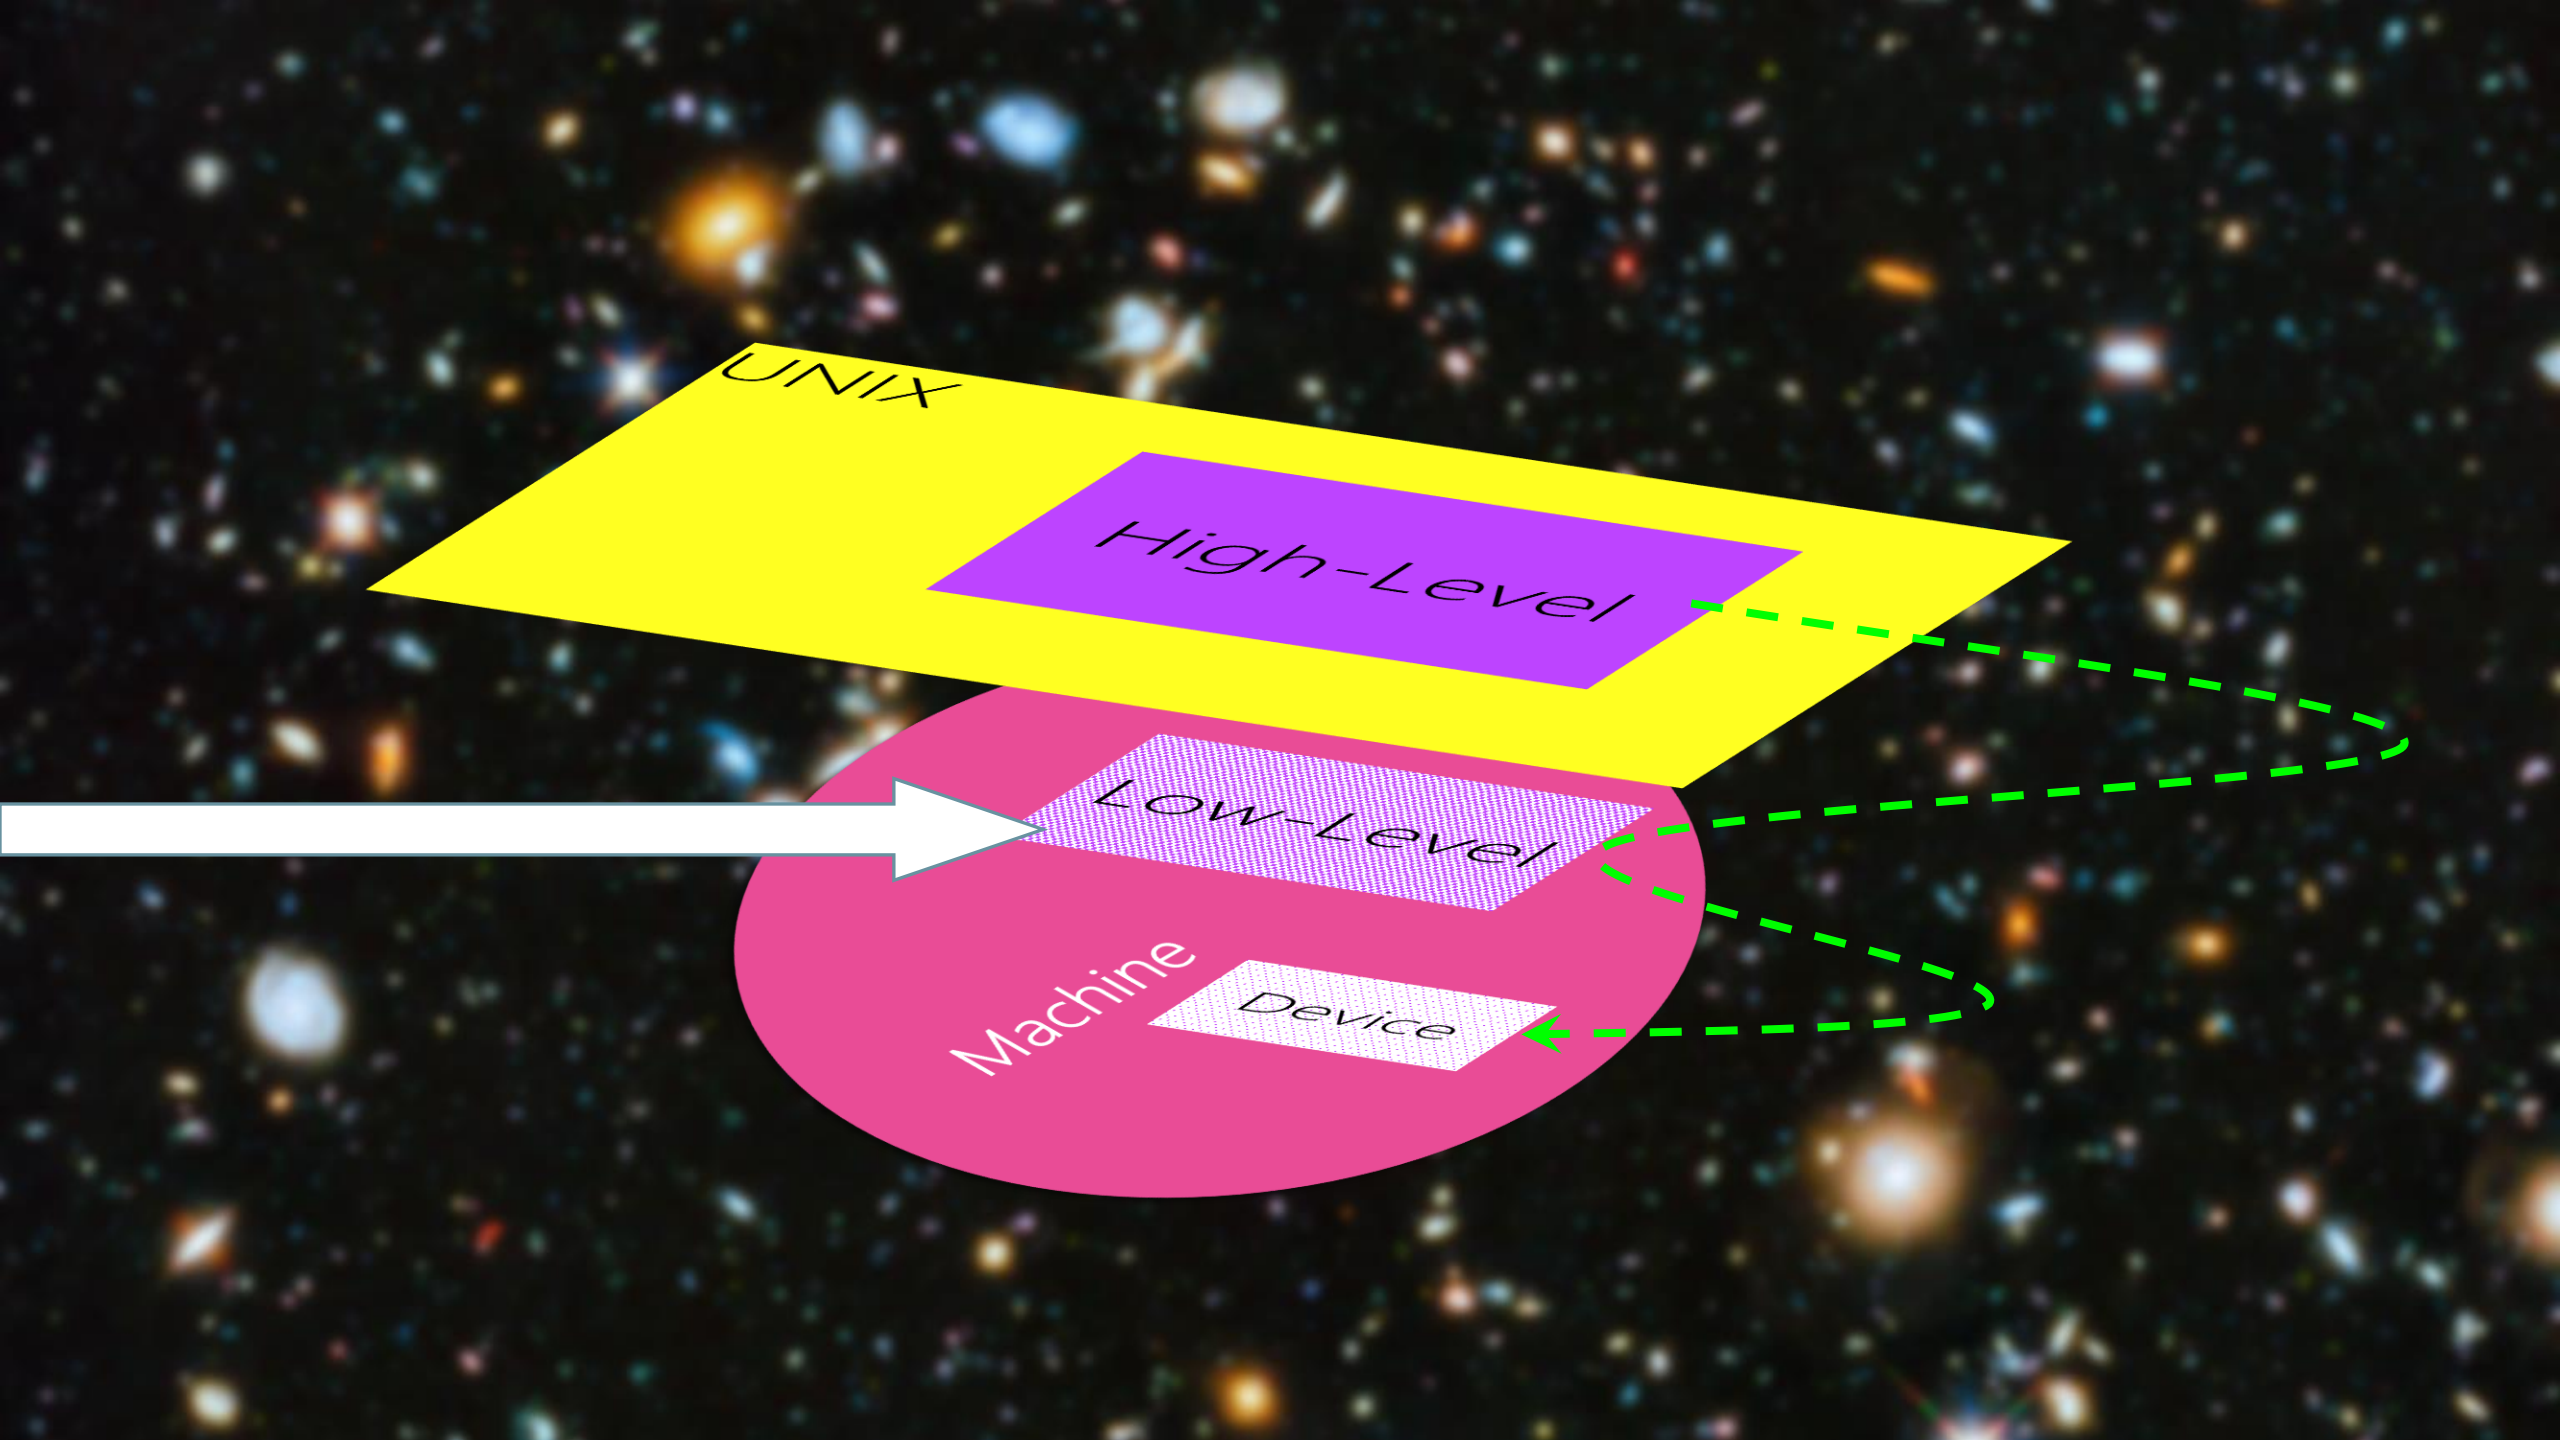
---

Atomic Operation

---

Advanced! We won't cover it.





UNIX

High-Level

Low-Level

Machine

Device



# Storage File System

aka. Disk File System

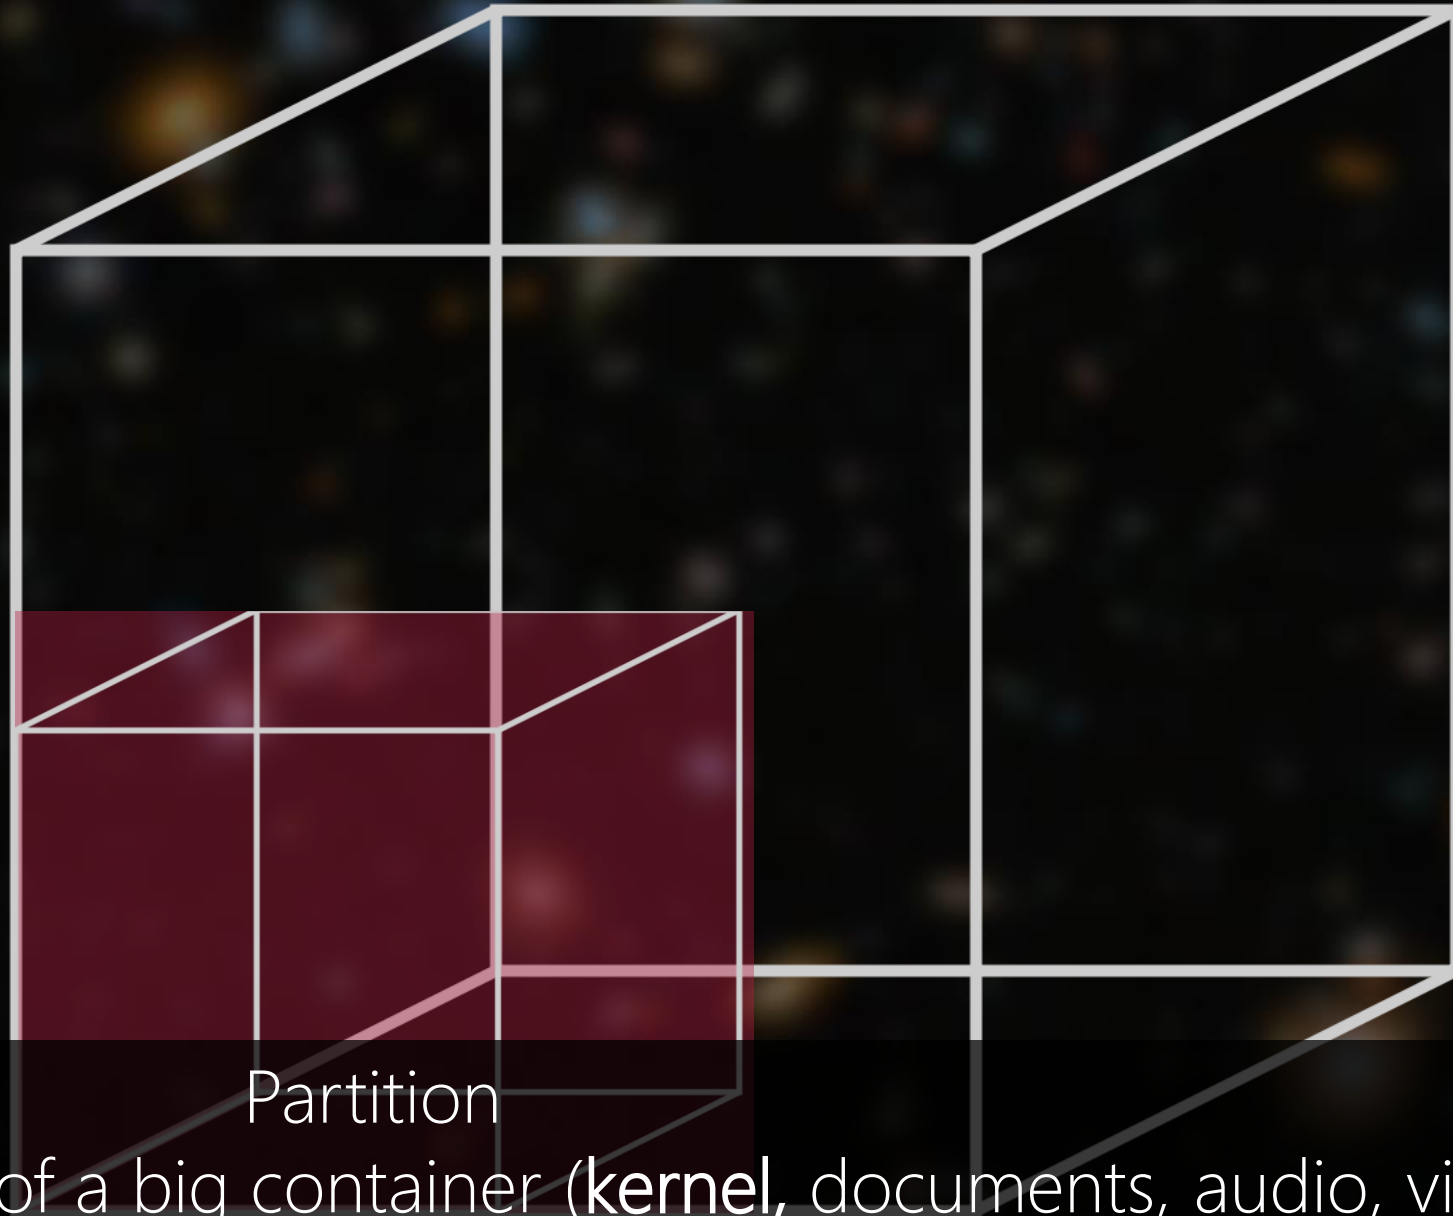
UNIX File System (UFS) vs. NTFS, ext4, ...

File System: Low Level

A white wireframe cube is centered on a dark background filled with numerous small, out-of-focus stars and galaxies, creating a cosmic or space-themed aesthetic. The cube is drawn with simple white lines, showing its three-dimensional structure with a central square face and receding edges.

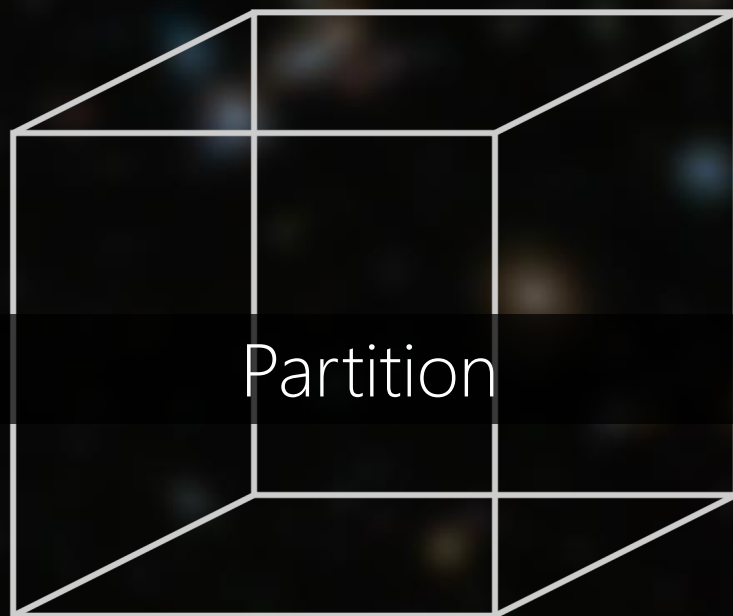
## Storage

Regardless of physical differences, it's a container for information  
HDD-IDE, HDD-SATA, CD-RW, SSD (USB, NVMe), ... are assumed to be the same!

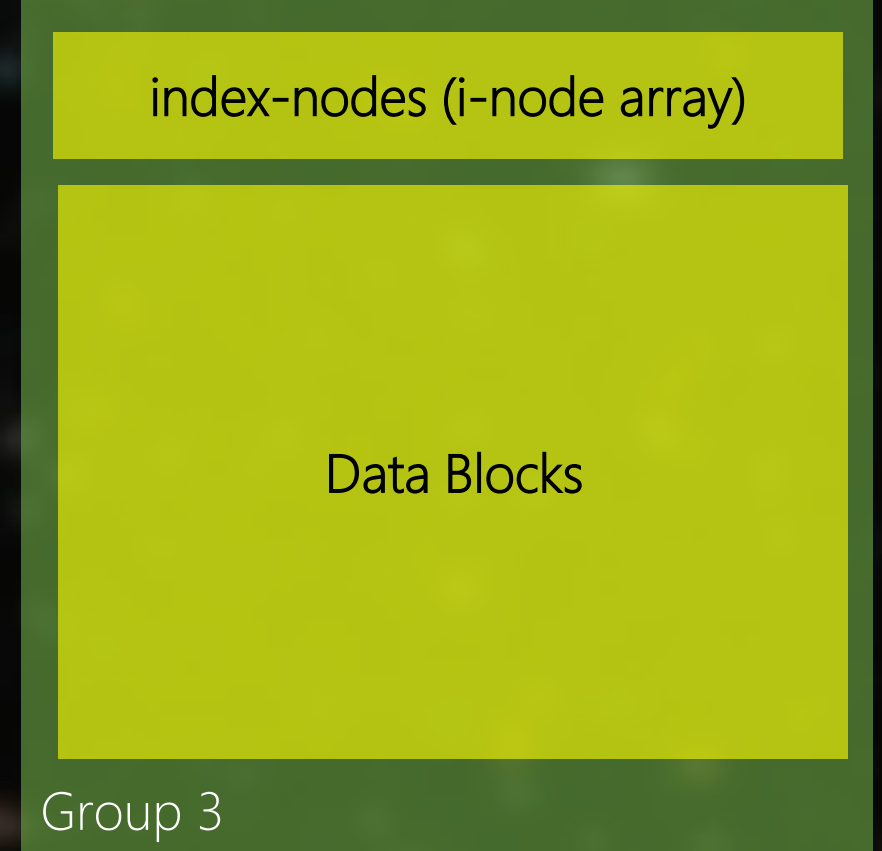
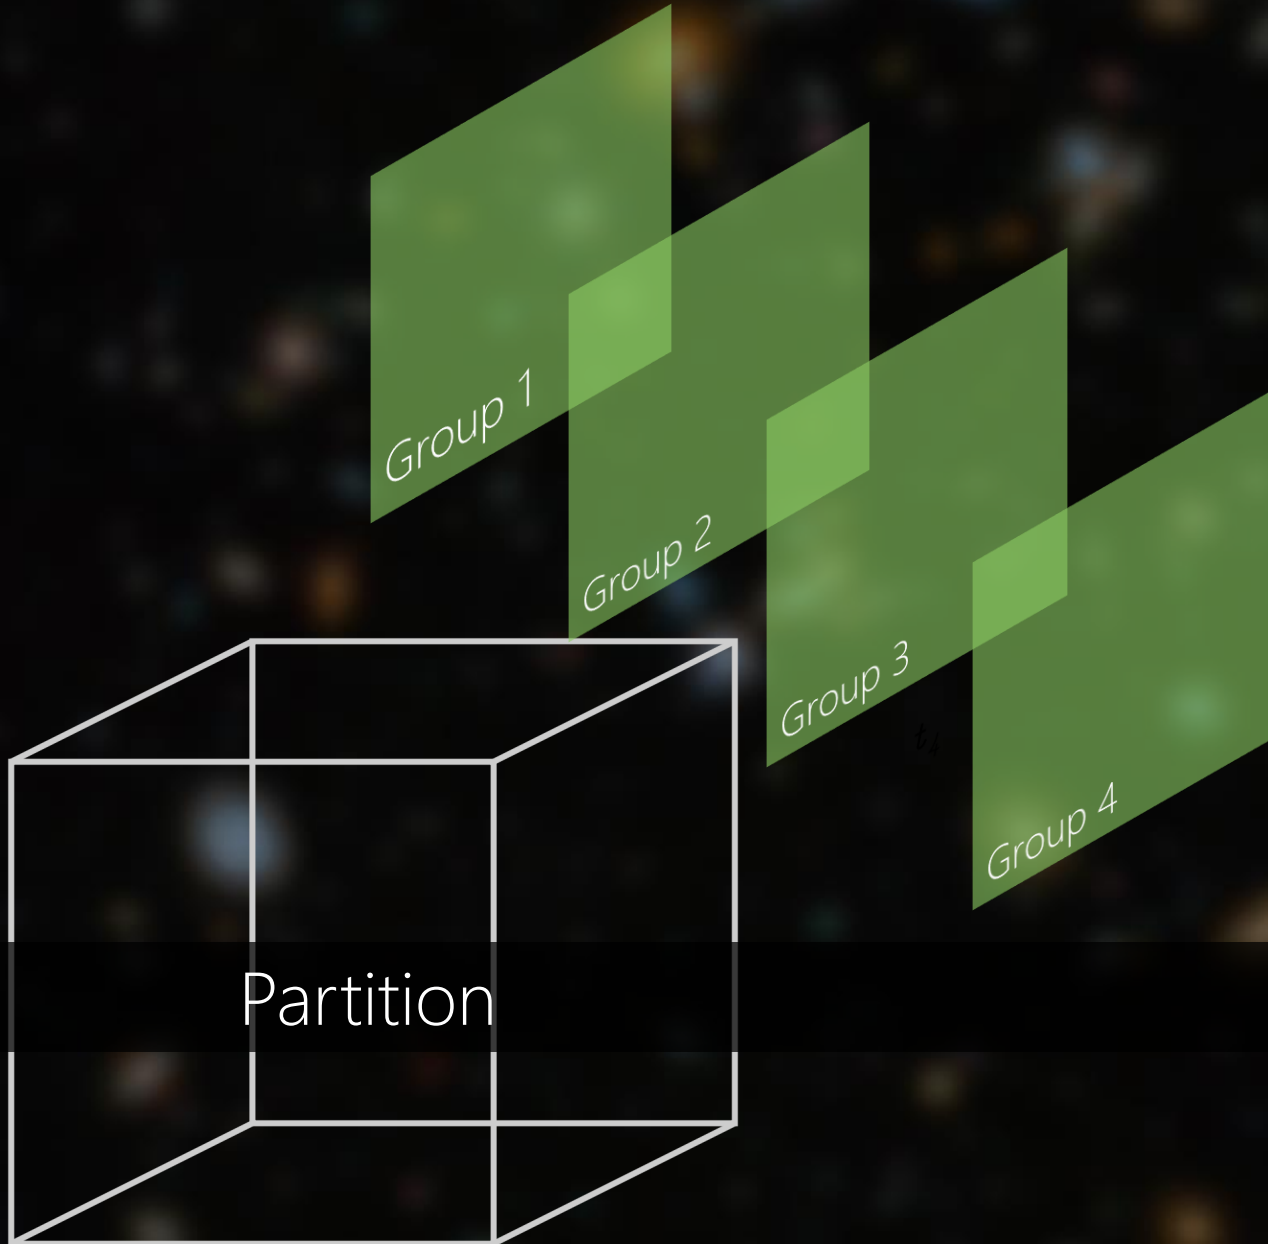


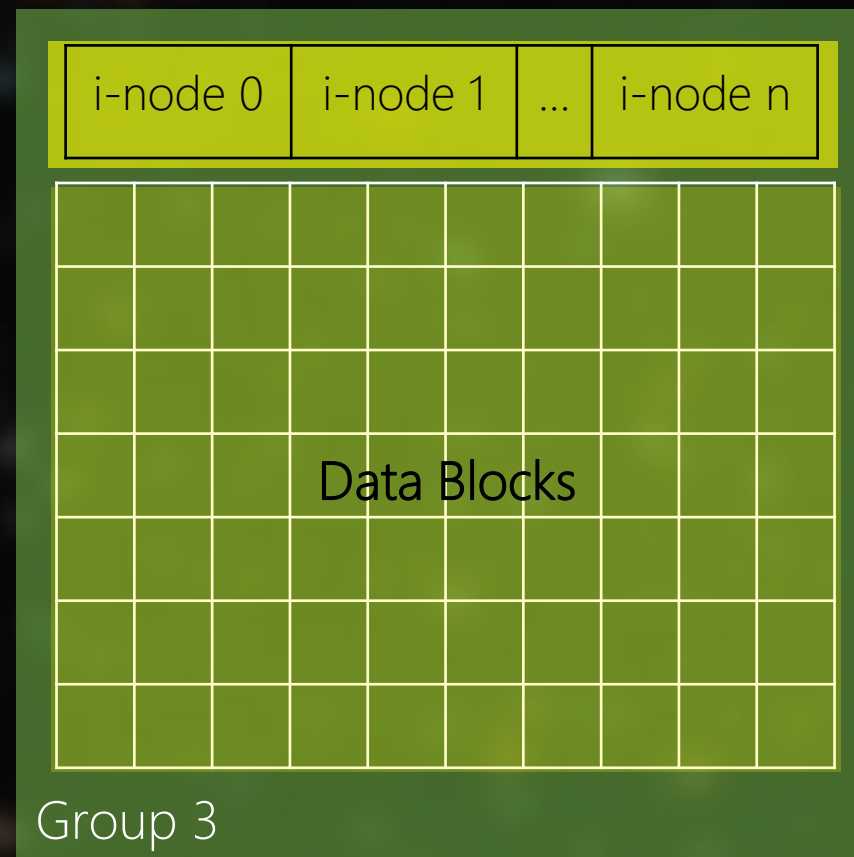
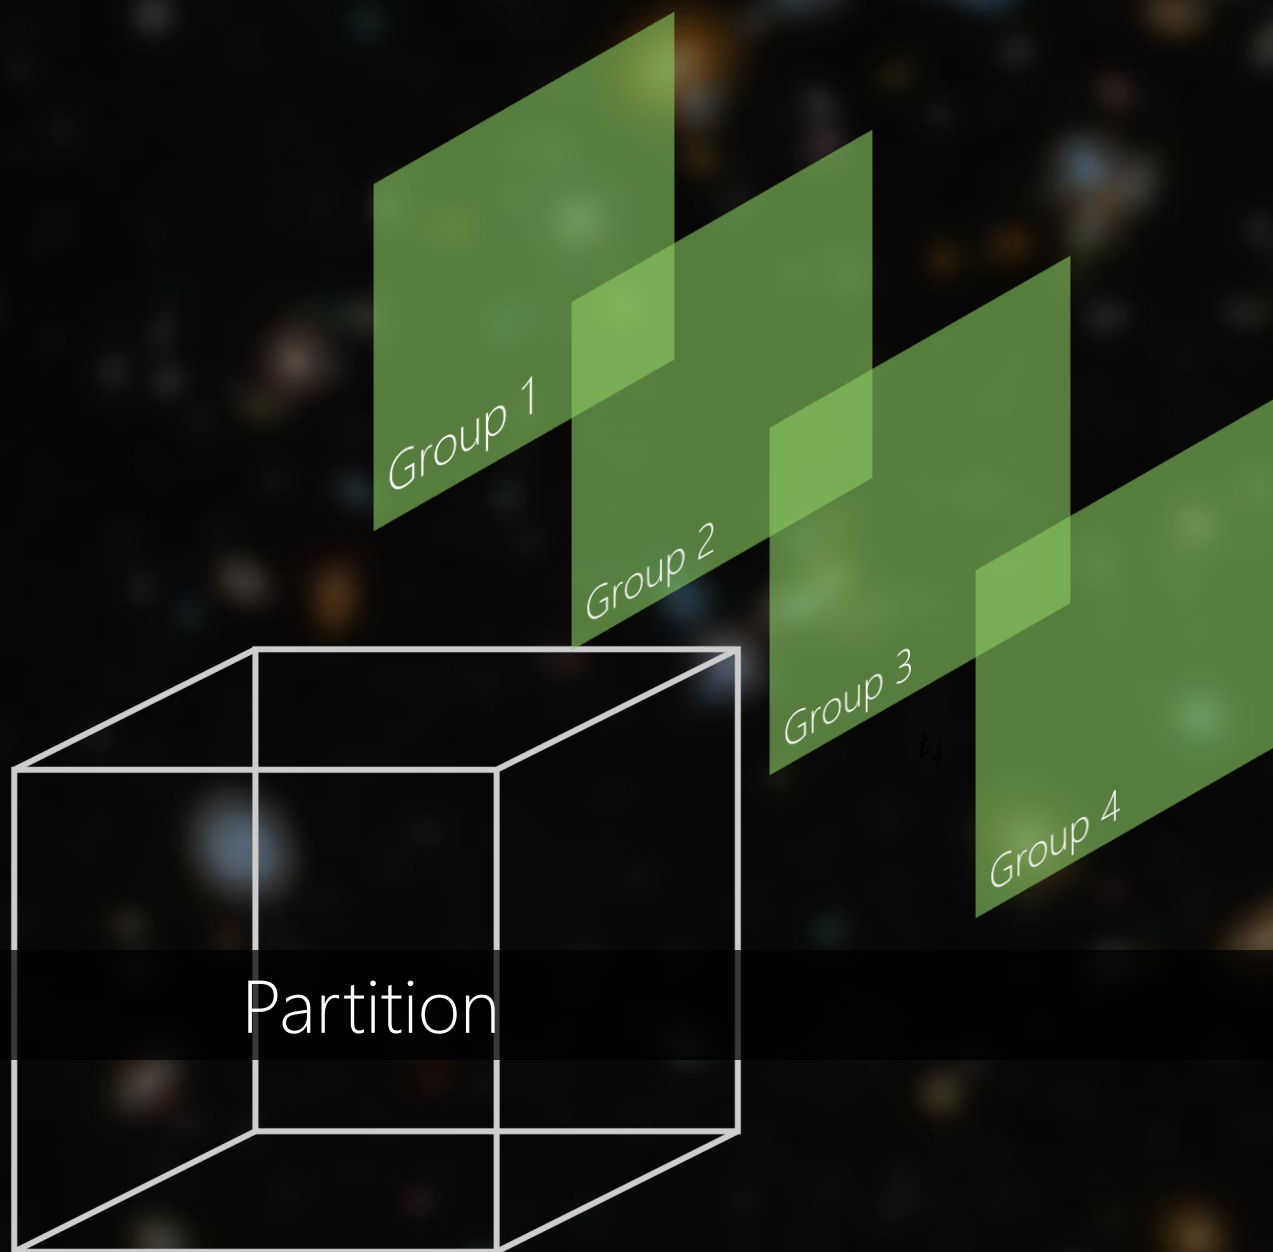
Partition

Physical division of a big container (**kernel**, documents, audio, video, ...)  
Some special treatment (customization) for each partition









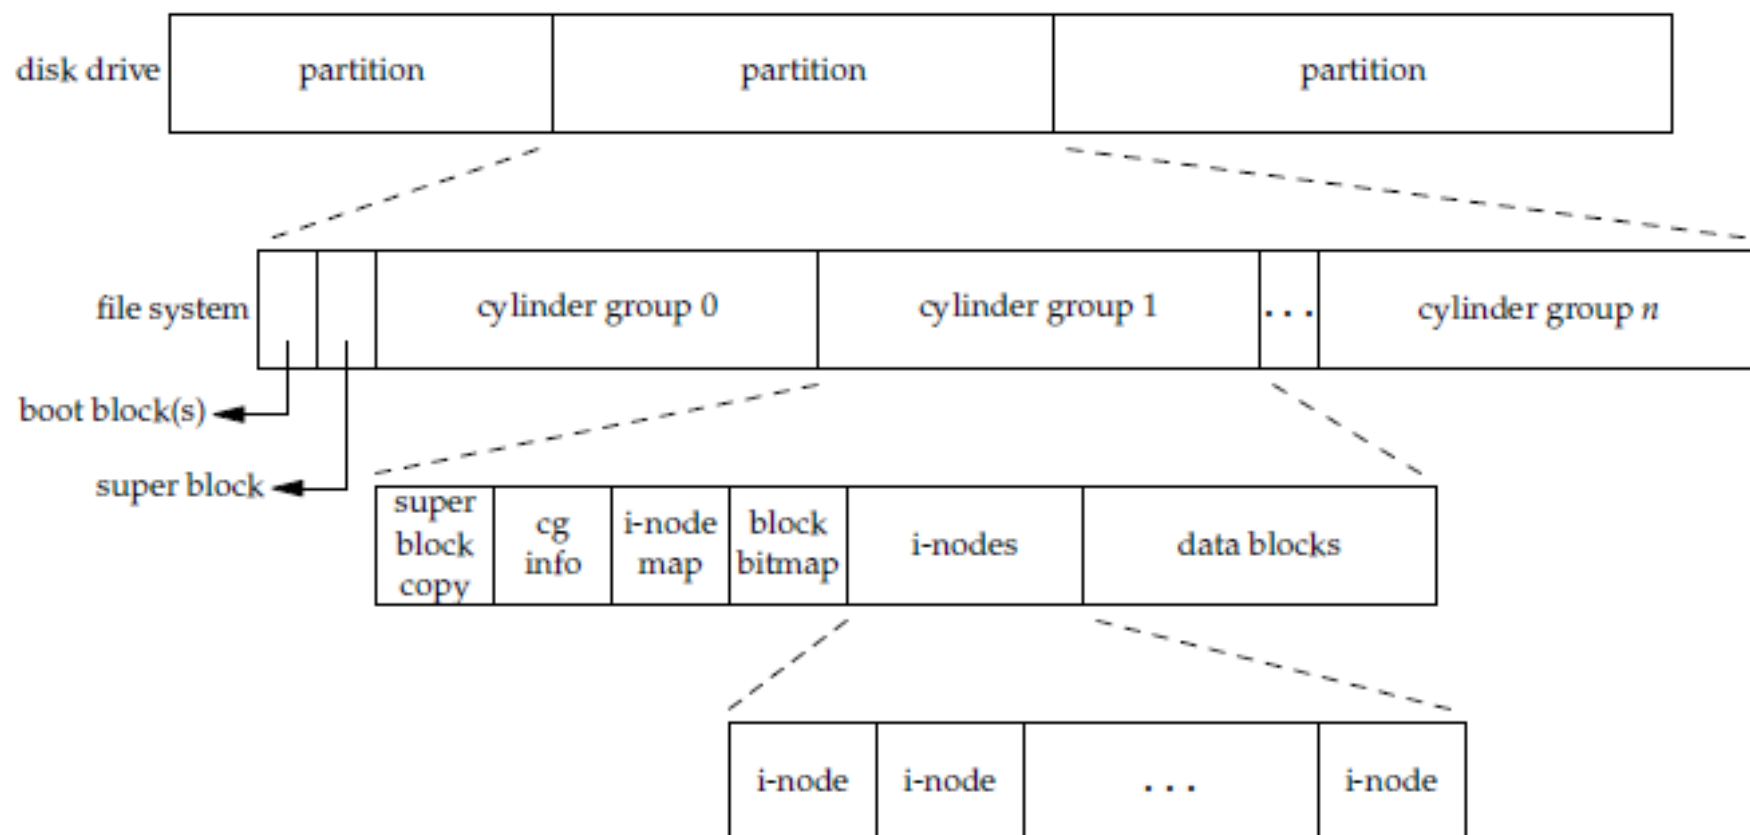
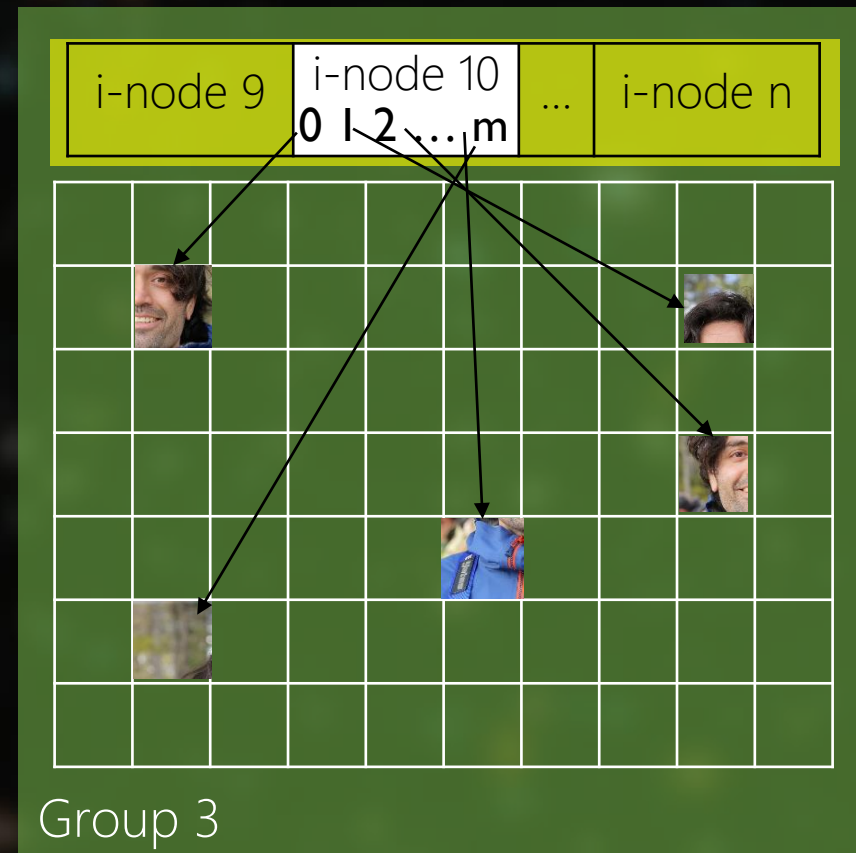
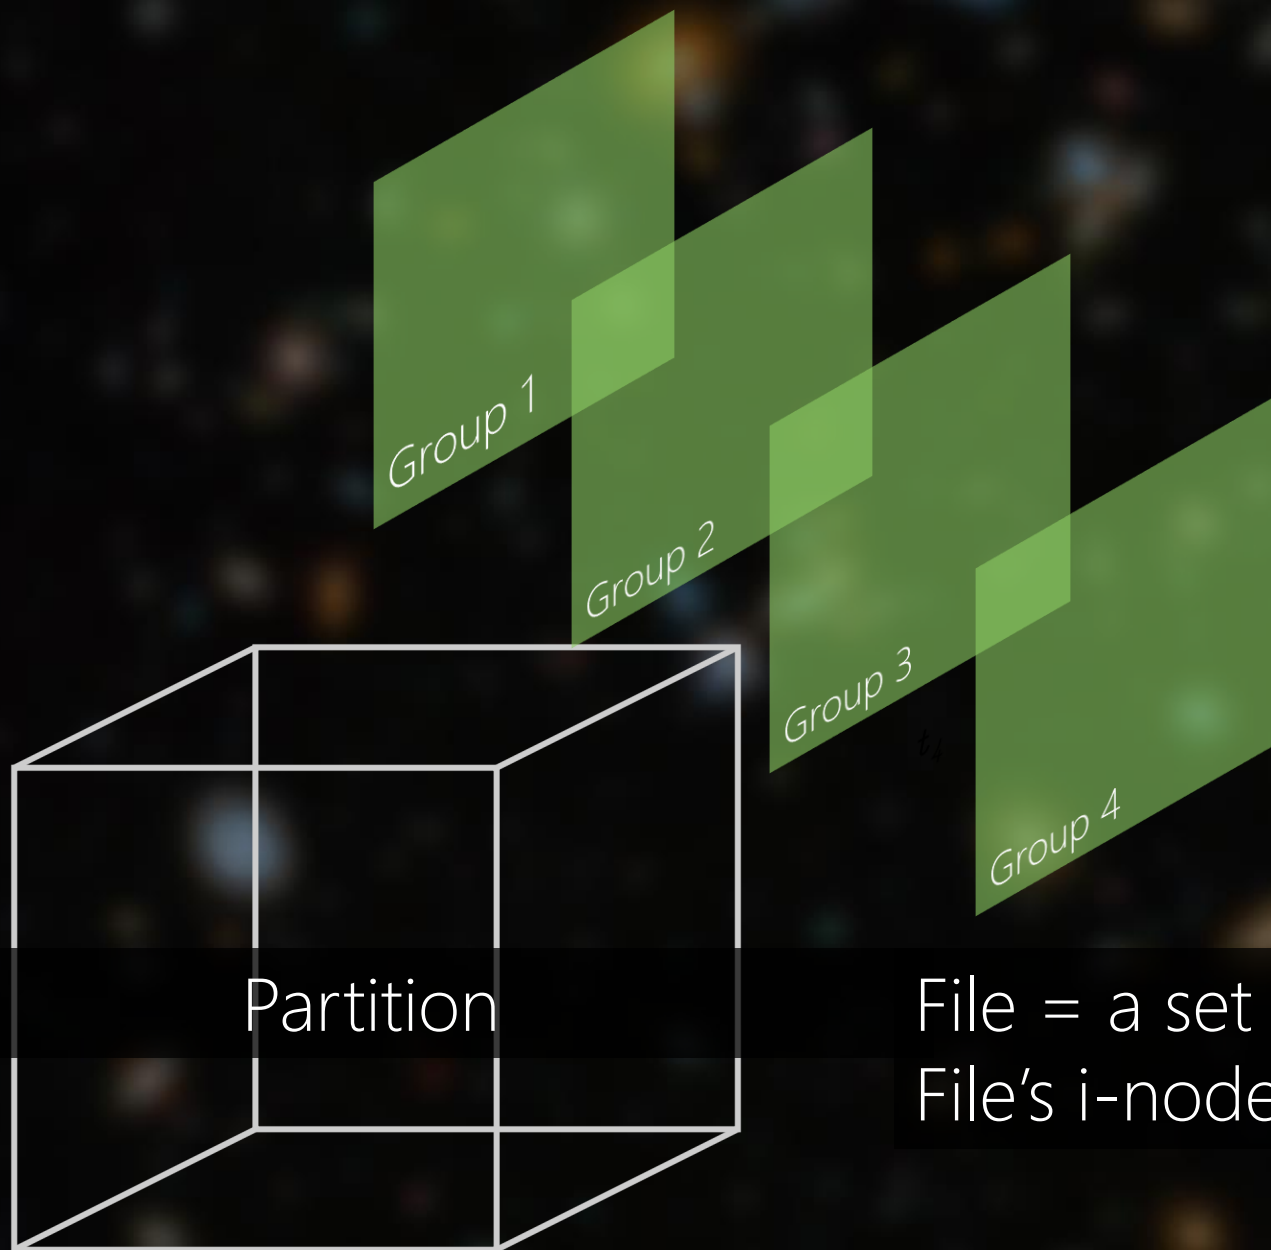


Figure 4.13 Disk drive, partitions, and a file system



File = a set or heap (not list) of data blocks  
File's i-node knows the order!



---

# Each File has one index-node (i-node)

```
$ ls -i {filename}
```

---

```
hfani@charlie:~$ ls -i ./hfani.jpeg  
94552308 ./hfani.jpeg
```

# Each File has many data blocks

Shell → `stat {filename}`

```
hfani@charlie:~$ stat ./hfani.jpeg
  File: ./hfani.jpeg
  Size: 113327          Blocks: 223          IO Block: 1048576 regular file
Device: 29h/41d Inode: 94552308    Links: 1
Access: (0644/-rw-r--r--)  Uid: (239080/  hfani)   Gid: ( 400/   temp)
Access: 2021-07-05 22:48:41.000000000 -0400
Modify: 2021-07-05 22:48:41.000000000 -0400
Change: 2021-10-27 05:15:35.999746375 -0400
 Birth: -
```

---

# Each File has many data blocks

System Call → stat, fstat

---

```
#include <sys/stat.h>
int stat(const char *restrict pathname, struct stat *restrict buf);
int fstat(int fd, struct stat *buf);
Return 0 if OK -1 on error
```

```
struct stat {
    mode_t st_mode;           /* file type & mode (permissions) */
    ino_t st_ino;             /* i-node number (serial number) */
    dev_t st_dev;             /* device number (file system) */
    dev_t st_rdev;            /* device number for special files */
    nlink_t st_nlink;         /* number of links */
    uid_t st_uid;             /* user ID of owner */
    gid_t st_gid;             /* group ID of owner */
    off_t st_size;            /* size in bytes, for regular files */
    struct timespec st_atim;  /* time of last access */
    struct timespec st_mtim;  /* time of last modification */
    struct timespec st_ctim;  /* time of last file status change */
    blksize_t st_blksize;     /* best I/O block size */
    blkcnt_t st_blocks;       /* number of disk blocks allocated */
};
```



```
#include <sys/stat.h>
#include <fcntl.h>
#include <stdio.h>

void main(void)
{
    struct stat sb;

    char *file = "./hfani.jpeg";
    int fd = open(file, O_RDONLY);

    if (fstat(fd, &sb) == -1) {
        printf("error in fetching stat on file %s", file);
        return;
    }

    printf("I-node number: %d\n", sb.st_ino);
    printf("Preferred I/O block size: %d bytes\n", sb.st_blksize);
    printf("File size: %d bytes\n", sb.st_size);
    printf("Blocks allocated: %d\n", sb.st_blocks);
}
```

hfani@charlie:~\$ cc stat.c -o stat

hfani@charlie:~\$ ./stat

I-node number: 94552308

Preferred I/O block size: 1048576 bytes

File size: 113327 bytes

Blocks allocated: 223

The size of data blocks are determined at the time of partition creation (aka partitioning & formatting) and is fixed afterward.

Large (10MB) vs. Small (1MB)

## Fragmentation

```
hfani@charlie:~$ stat ./hfani.jpeg
  File: ./hfani.jpeg
  Size: 113327          Blocks: 223          IO Block: 1048576 regular file
Device: 29h/41d Inode: 94552308    Links: 1
Access: (0644/-rw-r--r--)  Uid: (239080/  hfani)   Gid: ( 400/   temp)
Access: 2021-07-05 22:48:41.000000000 -0400
Modify: 2021-07-05 22:48:41.000000000 -0400
Change: 2021-10-27 05:15:35.999746375 -0400
 Birth: -
```

---

/Path/Filename → i-node → Data Blocks

/home/hfani/hfani.jpeg → 94552308 →



---

/Path/Filename → i-node → Data Blocks

/home/hfani/hfani.jpeg → 94552308 →





---

/Path/Filename → i-node → Data Blocks

/home/hfani/hfani.jpeg → 94552308 →



---

# Directories

They are files. So, each of them has one i-node.  
The content is not an image, audio, ... but mapping between filenames and their i-nodes

Logical (not physical) division of files

i-node 2 for root /  
Data block b1

i-node#	filename
2	.
2	..
i1	home

i-node i1 for root '/home'  
Data block b2

i-node#	filename
i1	.
2	..
i3	hfani

i-node i3 for root '/home/hfani'  
Data block b3

i-node#	filename
i3	.
?????	..
i4	hfani.jpeg

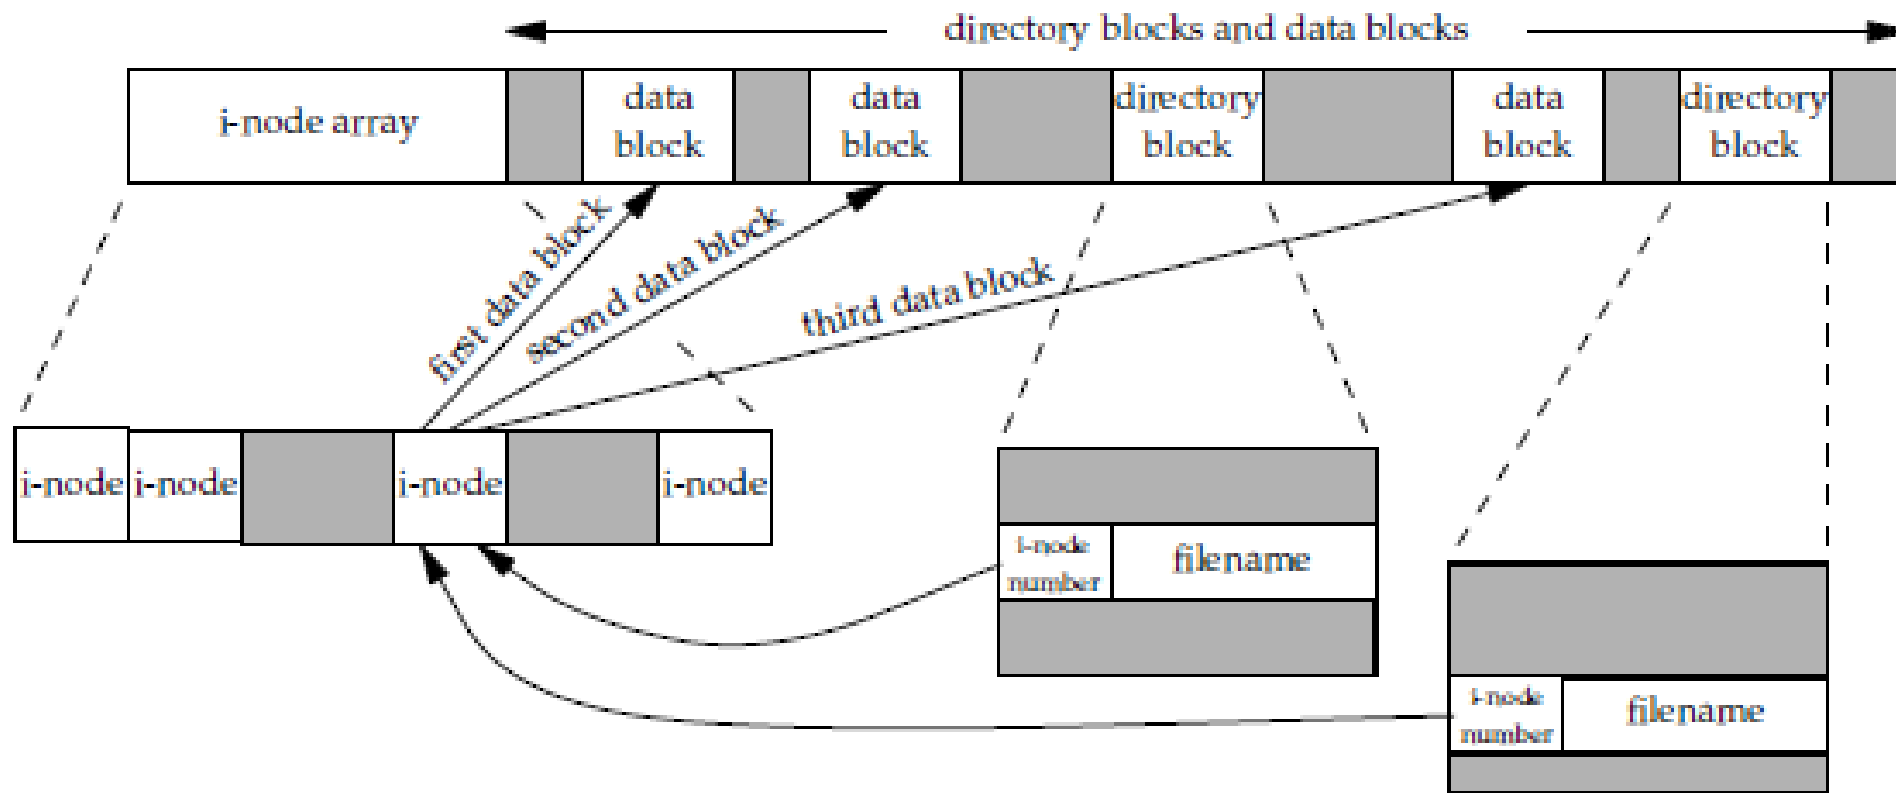


Figure 4.14 Cylinder group's i-nodes and data blocks in more detail



---

# Directories

Is it possible to have multiple links to the same file? Yes.  
How?

---



i-node 2 for root /  
Data block b1

i-node#	filename
2	.
2	..
i4	hosseinfani.png
i1	

i-node i1 for root '/home'  
Data block b2

i-node#	filename
i1	.
2	..
i3	hfani

i-node i3 for root '/home/hfani'  
Data block b3

i-node#	filename
i3	.
?????	..
i4	hfani.jpeg

No Duplication!  
Still one single file.  
Two or more links.  
hardlink, shortcut, ....

---

What happens when you **delete** a file?

---

Zero all data blocks of the file?

Zero all data blocks of the file and Zero its i-node pointers?

Zero its entry (i-node, name) in the parent directory?

i-node 2 for root /  
Data block b1

i-node#	filename
2	.
2	..
i1	home

i-node i1 for root '/home'  
Data block b2

i-node#	filename
i1	.
2	..
i3	hfani

i-node i3 for root '/home/hfani'  
Data block b3

i-node#	filename
i3	.
i1	..
i4	hfani.jpeg



---

What happens when you **delete** a file?

---

Is it able to **recover** a deleted file?



What happens when you **move** a file?





---

What happens when you **copy** a file?

---

# File Types

- Regular Files → text or binary
- Directory Files → (i-node, filename) pairs
- Special Files (Devices)
  - Block → HDD, SSD, CD, ...
  - Character (Stream) → TTY, Keyboard, Mouse, Printer, ...
- Socket (Networking)
- FIFO (Pipes)
- Symbolic Link

---

RD/WR on a Regular File: `open("./myfile.txt")`  
vs.

RD/WR on a Storage (Lab06): `open("/dev/sda1")`

---

- Regular Files → text or binary
- Directory Files → (i-node, filename) pairs
- Special Files (Devices)
  - Block → HDD, SSD, CD, ...
  - Character (Stream) → TTY, Keyboard, Mouse, Printer, ...
- Socket (Networking)
- FIFO (Pipes)
- Symbolic Link



---

# Process Manager

aka. Process Control

---