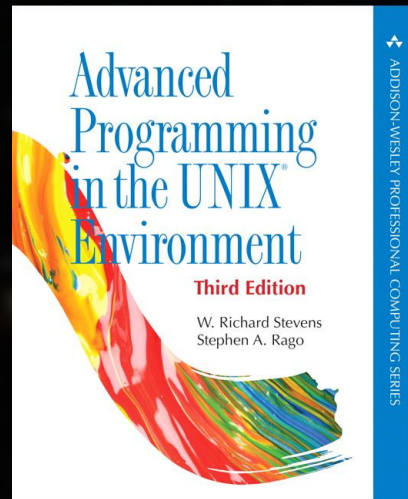


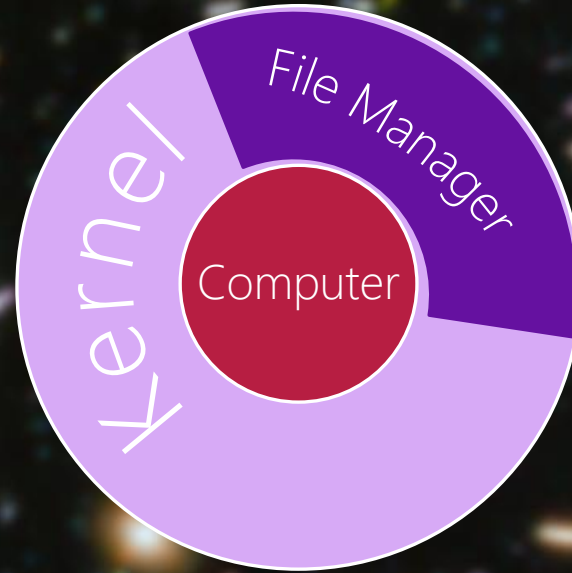
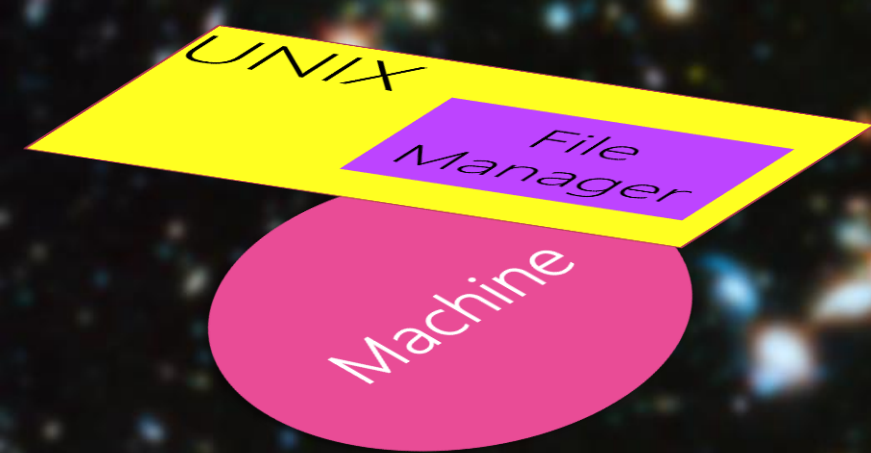
A deep-field astronomical image showing a vast field of galaxies in various colors (blue, orange, white) against a black background. Two horizontal blue lines frame the central text.

Midterm Exam Marks + Lab07 + Lec07

- 
- Review File System lseek, dup, I/O Redirection
 - File System for Storage Devices: i-node, Directories



Chapter 03: File I/O Chapter 04: File I/O



Computer

Memory

Kernel: Device Manager

Kernel: Memory Manager

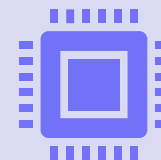
Kernel: File Manager

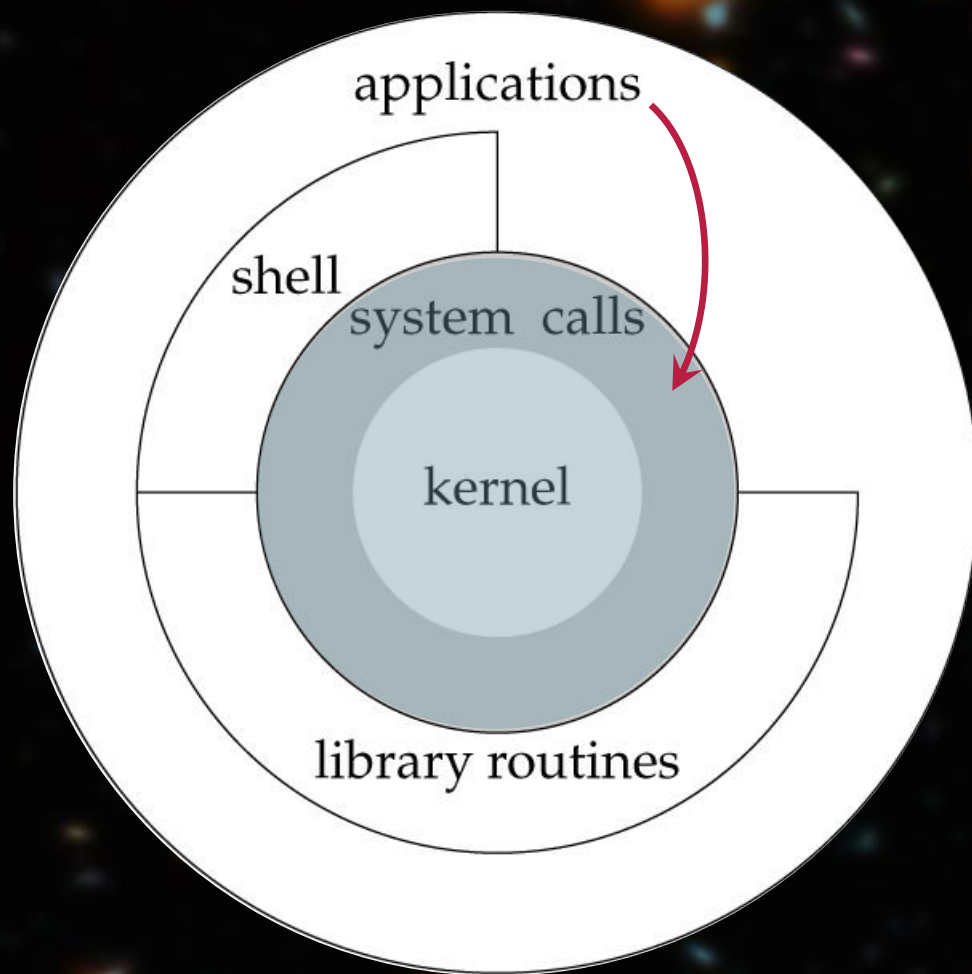
Kernel: Network Manager

Kernel: Process Manager

Bus

Processor





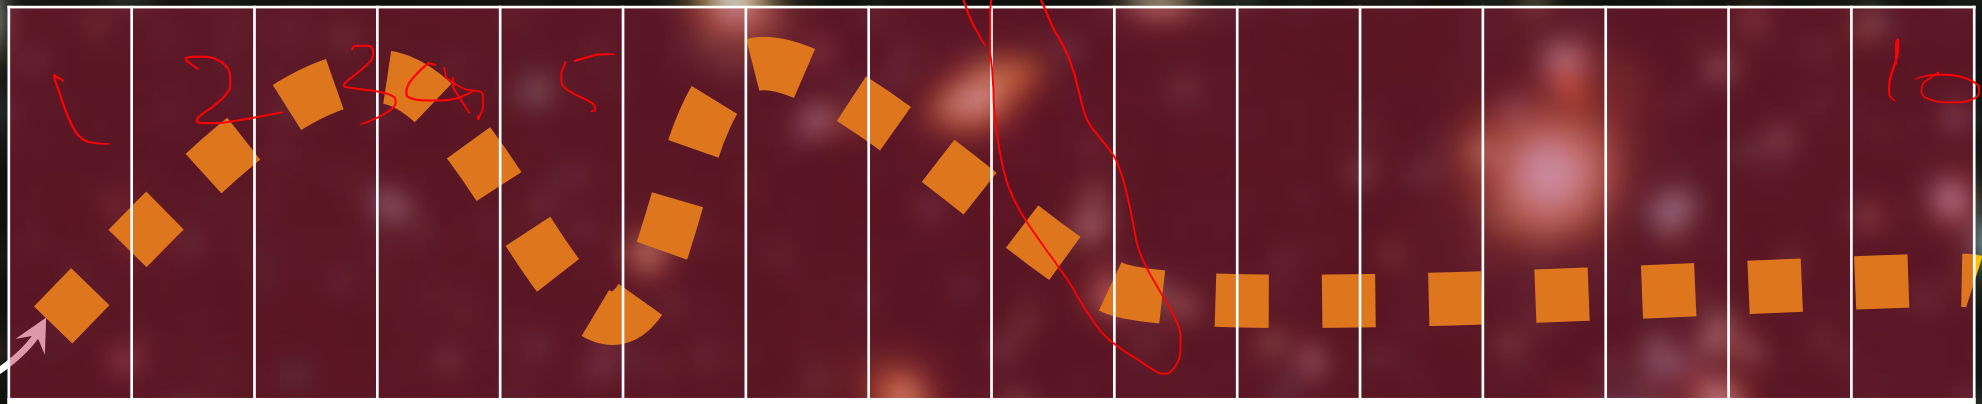
Header	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Description
< aio.h>	•	•	•	•	asynchronous I/O
< cpio.h>	•	•	•	•	cpio archive values
< dirent.h>	•	•	•	•	directory entries (Section 4.22)
< dlfcn.h>	•	•	•	•	dynamic linking
< fcntl.h>	•	•	•	•	file control (Section 3.14)
< fnmatch.h>	•	•	•	•	filename-matching types
< glob.h>	•	•	•	•	pathname pattern-matching and generation
< grp.h>	•	•	•	•	group file (Section 6.4)
< iconv.h>	•	•	•	•	codeset conversion utility
< langinfo.h>	•	•	•	•	language information constants
< monetary.h>	•	•	•	•	monetary types and functions
< netdb.h>	•	•	•	•	network database operations
< nl_types.h>	•	•	•	•	message catalogs
< poll.h>	•	•	•	•	poll function (Section 14.4.2)
< pthread.h>	•	•	•	•	threads (Chapters 11 and 12)
< pwd.h>	•	•	•	•	password file (Section 6.2)
< regex.h>	•	•	•	•	regular expressions
< sched.h>	•	•	•	•	execution scheduling
< semaphore.h>	•	•	•	•	semaphores
< strings.h>	•	•	•	•	string operations
< tar.h>	•	•	•	•	tar archive values
< termios.h>	•	•	•	•	terminal I/O (Chapter 18)
< unistd.h>	•	•	•	•	symbolic constants
< wordexp.h>	•	•	•	•	word-expansion definitions
< arpa/inet.h>	•	•	•	•	Internet definitions (Chapter 16)
< net/if.h>	•	•	•	•	socket local interfaces (Chapter 16)
< netinet/in.h>	•	•	•	•	Internet address family (Section 16.3)
< netinet/tcp.h>	•	•	•	•	Transmission Control Protocol definitions
< sys/mman.h>	•	•	•	•	memory management declarations
< sys/select.h>	•	•	•	•	select function (Section 14.4.1)
< sys/socket.h>	•	•	•	•	sockets interface (Chapter 16)
< sys/stat.h>	•	•	•	•	file status (Chapter 4)
< sys/statvfs.h>	•	•	•	•	file system information
< sys/times.h>	•	•	•	•	process times (Section 8.17)
< sys/types.h>	•	•	•	•	primitive system data types (Section 2.8)
< sys/un.h>	•	•	•	•	UNIX domain socket definitions (Section 17.2)
< sys/utsname.h>	•	•	•	•	system name (Section 6.9)
< sys/wait.h>	•	•	•	•	process control (Section 8.6)

lseek

A Better Use Case: Binary Search in Sorted Elements

$\mathcal{O}(n) = 1M$
 $\mathcal{O}(n) \log$
 $\mathcal{O}(\log n)$

fd



dup

```
#include <unistd.h>  
int dup(int fd);
```

the lowest-numbered available file descriptor to the same file if OK, -1 on error

dup vs. multiple open

open () : each returned fd has its own properties and current offset

fd1 = open ("test.txt", O_RDONLY)

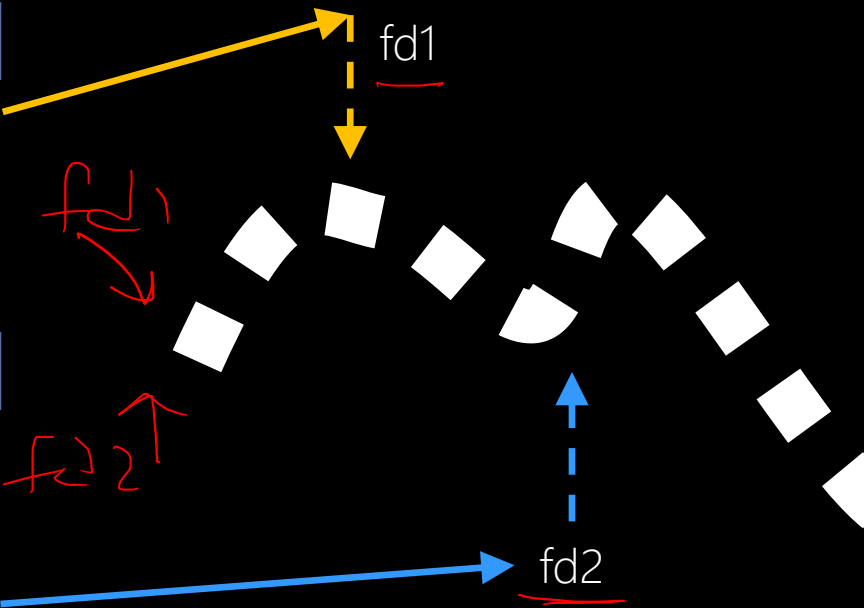
fd2 = open ("test.txt", O_WRONLY)

Kernel
File System

Process1

File Descriptors	File Pointer
fd1	
fd2	
...	

File Pointer
current file offset
other flags RD
pointer to first byte
File Pointer
current file offset
other flags WR
pointer to first byte



dup vs. multiple open

dup () : each returned fd has the same properties and current offset
fd2 = dup (fd1)

Kernel
File System

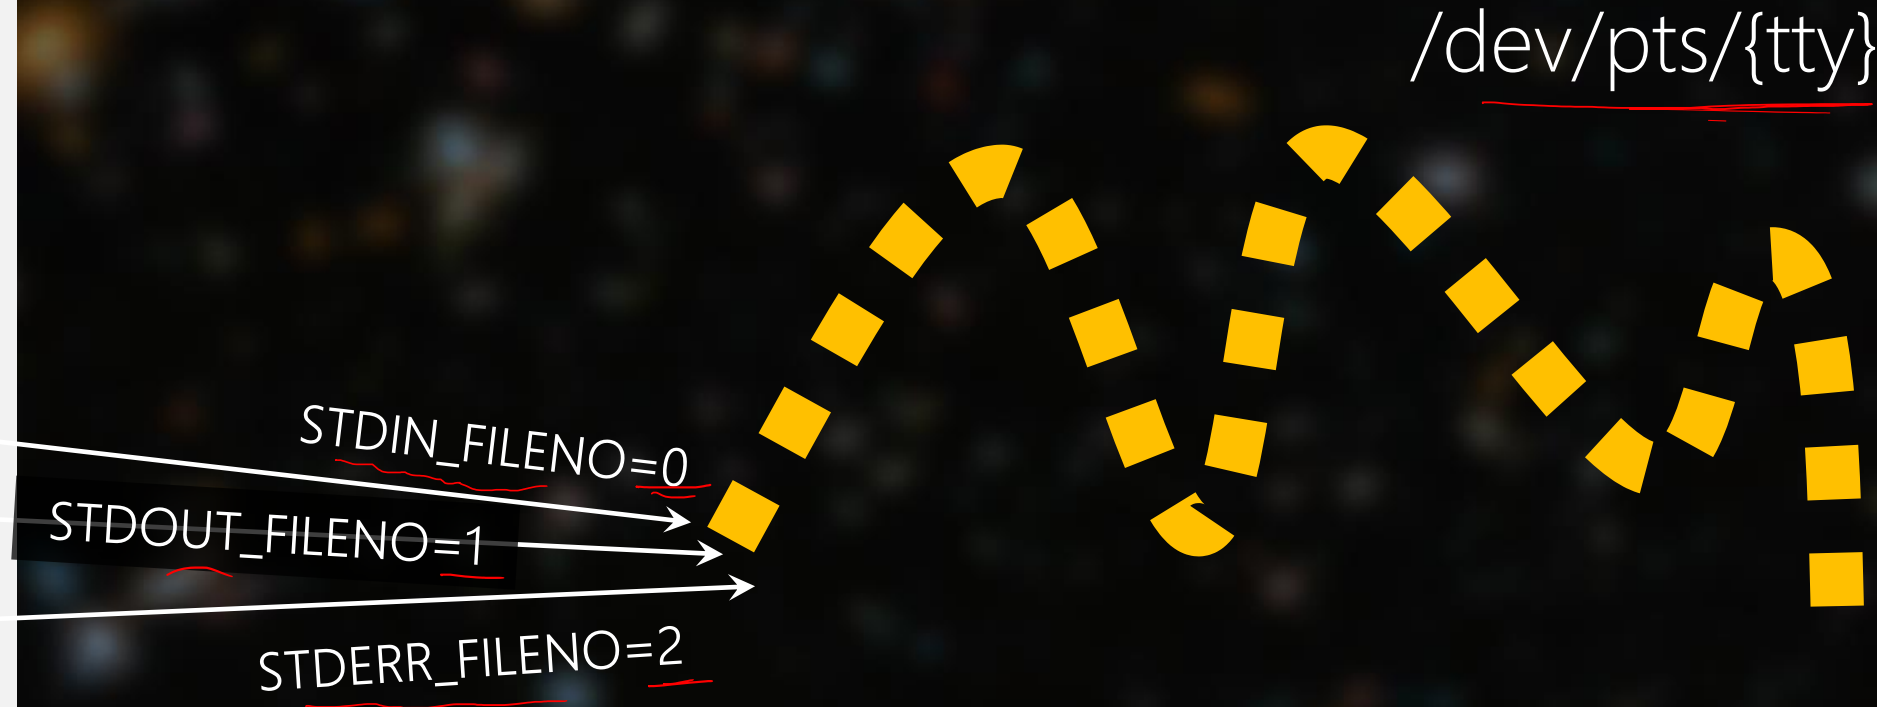
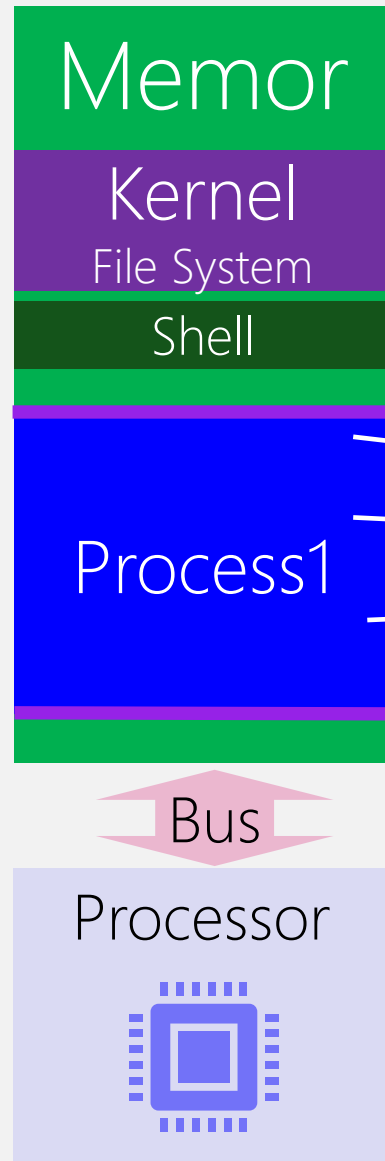
Process1

File Descriptors		File Pointer
fd1		
fd2		
...		

File Pointer
current file offset
other flags
pointer to first byte

fd1 fd2

Computer



When Shell bootstraps a program, it automatically opens three fds for the program (process):

STDIN_FILENO = 0 : O_RDONLY
STDOUT_FILENO = 1 : O_WRONLY
STDERR_FILENO = 2 : O_WRONLY

```
#include <fcntl.h>
#include <unistd.h>
void main(void){

char buf_rd[20];

int fd_rd = open("/dev/fd/0" O_RDONLY);
int res = read(fd_rd, buf_rd, 10);
int fd_wr = open("/dev/fd/1", O_WRONLY);
write(fd_wr, buf_rd, 10);
```

```
#include <fcntl.h>
#include <unistd.h>
void main(void){

char buf_rd[20];

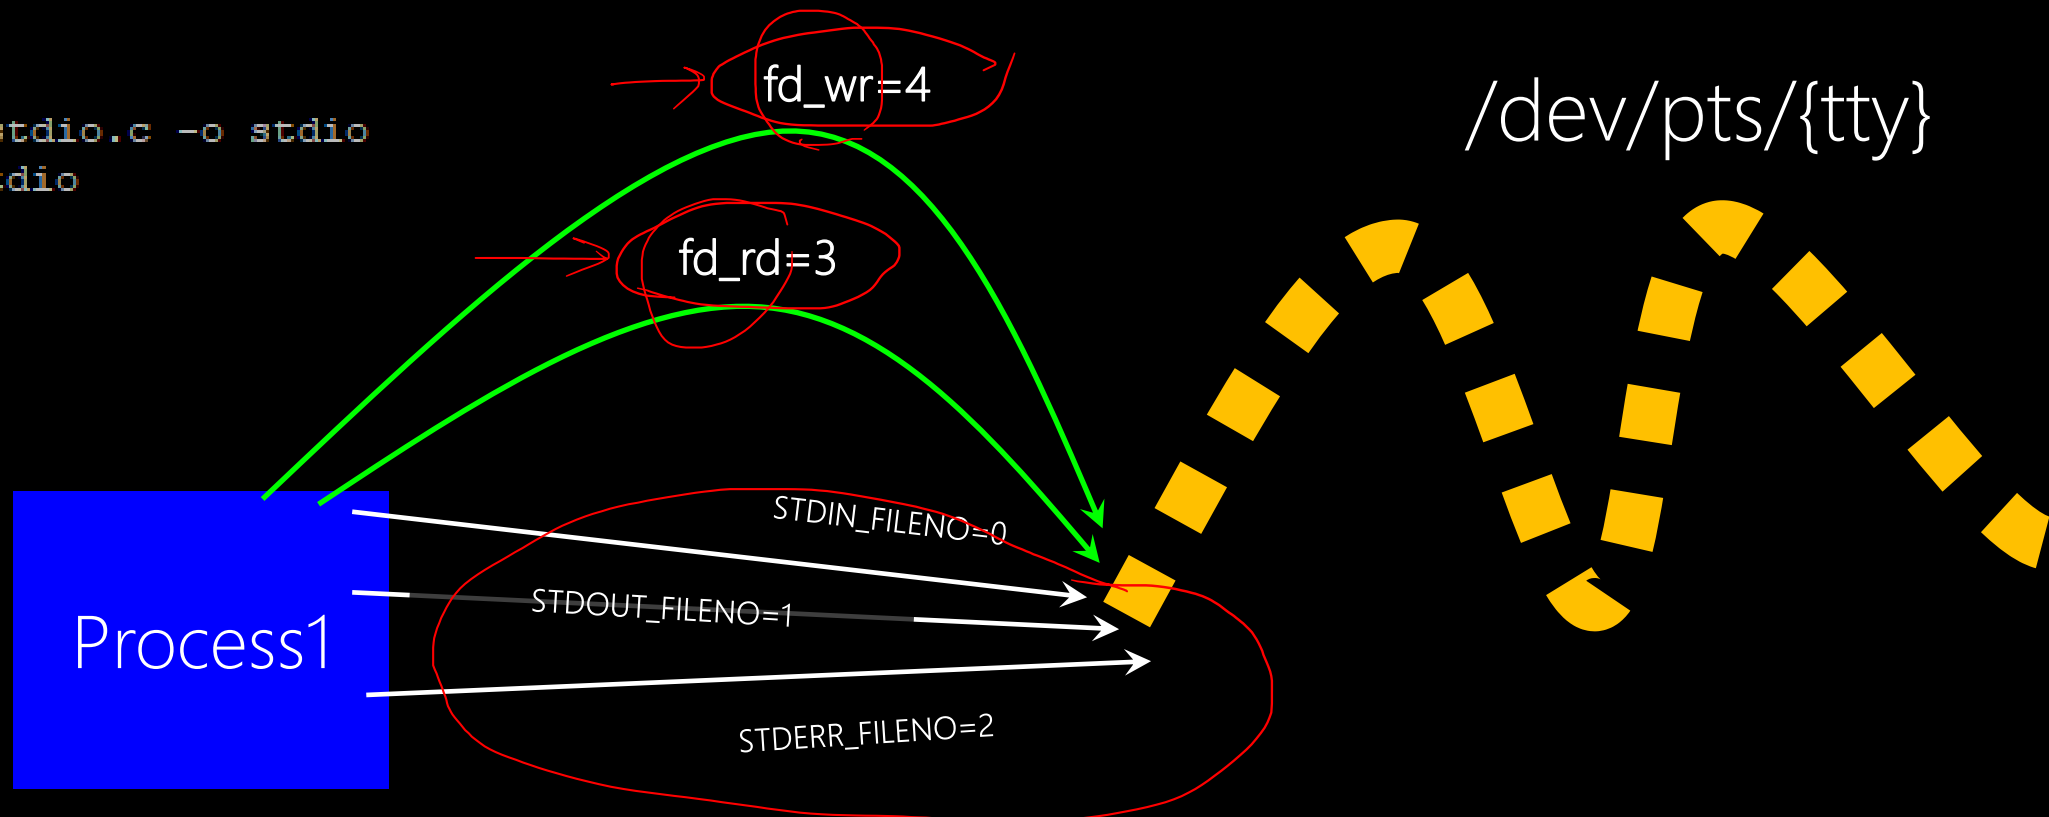
read(STDIN_FILENO, buf_rd, 10);
write(STDOUT_FILENO, buf_rd, 10);
```

```
#include <fcntl.h>
#include <unistd.h>
void main(void){

char buf_rd[20];

read(0, buf_rd, 10);
write(1, buf_rd, 10);
```

```
hfani@charlie:~$ cc stdio.c -o stdio
hfani@charlie:~$ ./stdio
comp2560
comp2560
hfani@charlie:~$
```



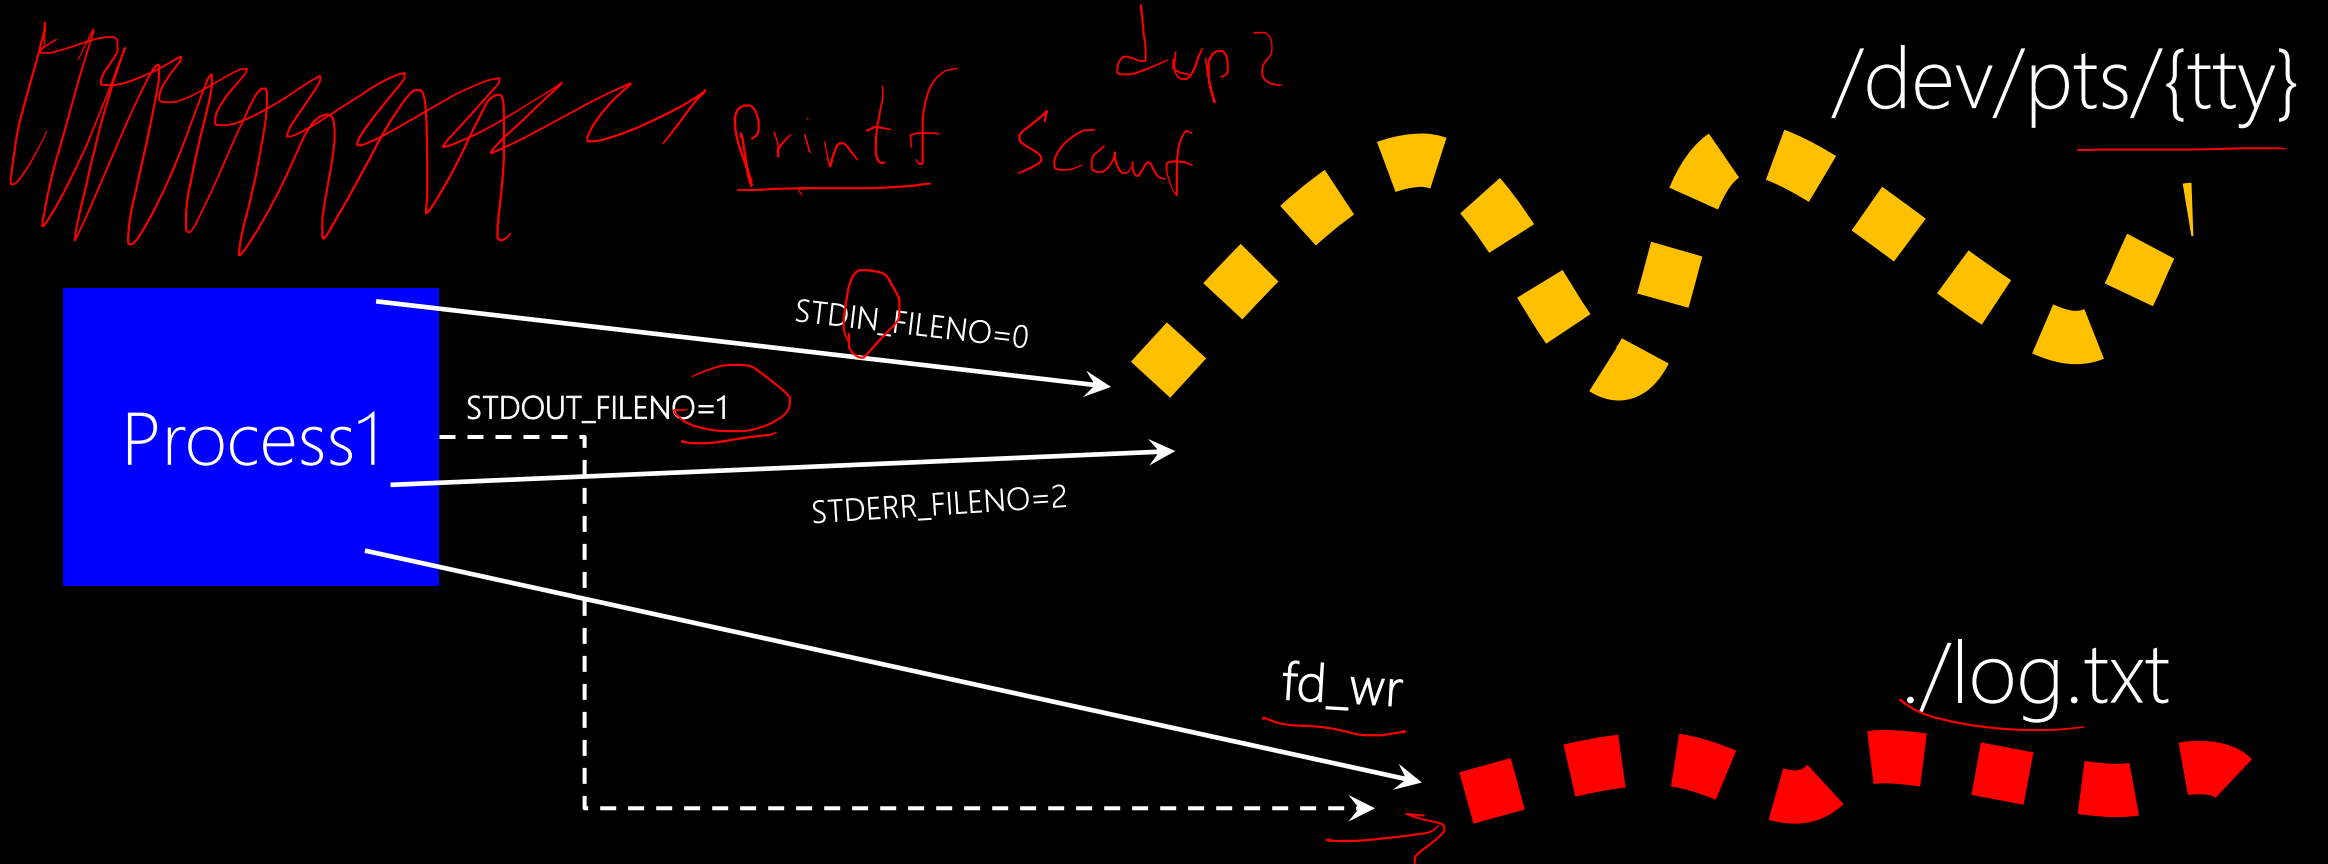
```

#include <fcntl.h>
#include <unistd.h>
void main(void) {
    char buf_rd[20];

    int fd_wr = open("../log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    //now the fd with value 1 is free
    //let's get it for our log file
    int new_fd = dup(fd_wr);
    //now the value of new_fd is 1
    //both fd_wr and new_fd are pointing to log.txt

```

← the lowest-numbered available file descriptor




```

#include <fcntl.h>
#include <unistd.h>
void main(void){
    char buf_rd[20];

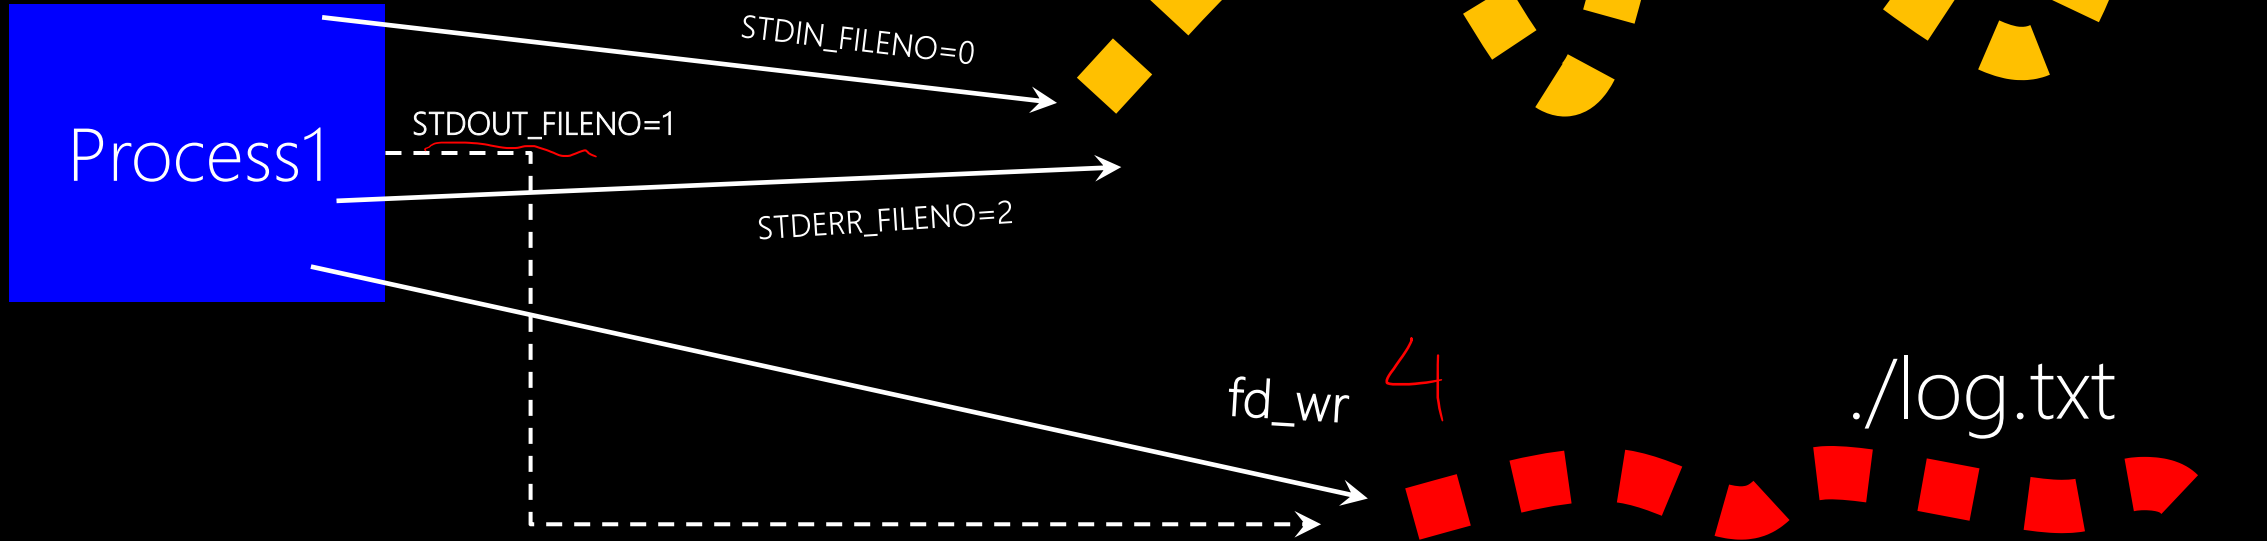
    int fd_wr = open("./log.txt", O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR);
    close(1);
    //now the fd with value 1 is free
    //let's get it for our log file
    int new_fd = dup(fd_wr);
    //now the value of new_fd is 5
    //both fd_wr and new_fd are pointing to log.txt

```

```

read(STDIN_FILENO, buf_rd, 10);
write(STDOUT_FILENO, buf_rd, 10);

```



Shell

\$./program > log.txt

pr < log.txt
sound & moun

We can ask the shell to do this redirection for us
{program file} > {new destination for STDOUT_FILENO}

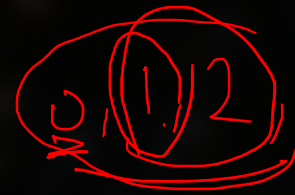


Story

How about STDIN_FILENO? STDERR_FILENO?

1 → out

use(1)

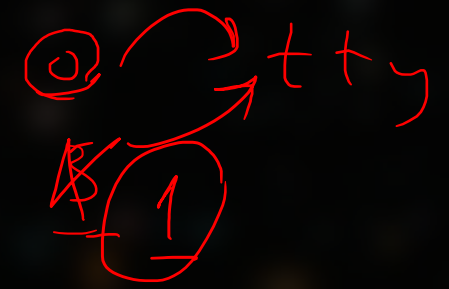
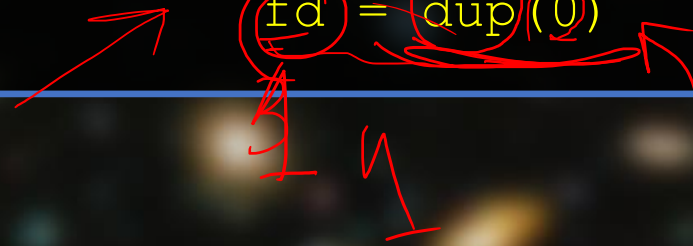


1-out

Story

What does this mean and what is the benefit?

`fd = dup(0)`



dup2(fdl, fdl)

At Home

I/O Efficiency buffered vs. unbuffered

System Call

`unistd.h` vs. `stdio.h`

Library Routine

At Home.



File Sharing

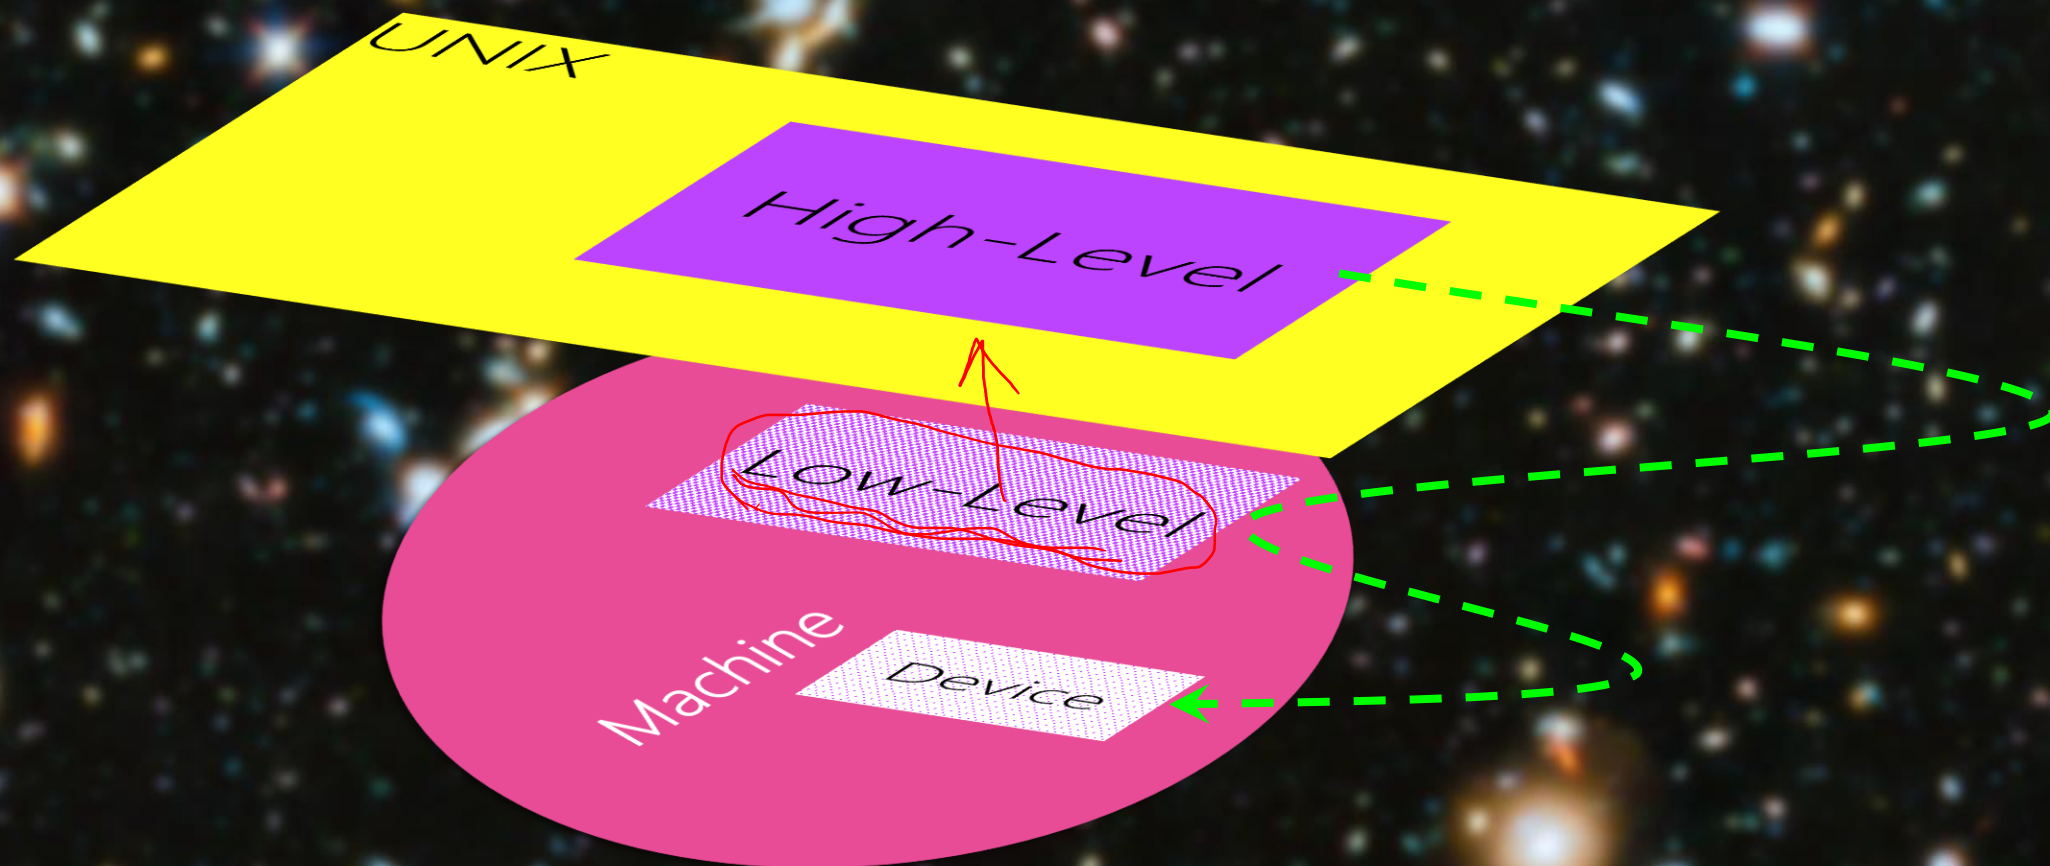
Advanced! We won't cover it.



Atomic Operation

Advanced! We won't cover it.

- 
- Review File System `lseek`, `dup`, I/O Redirection
 - File System for Storage Devices: `i-node`, Directories



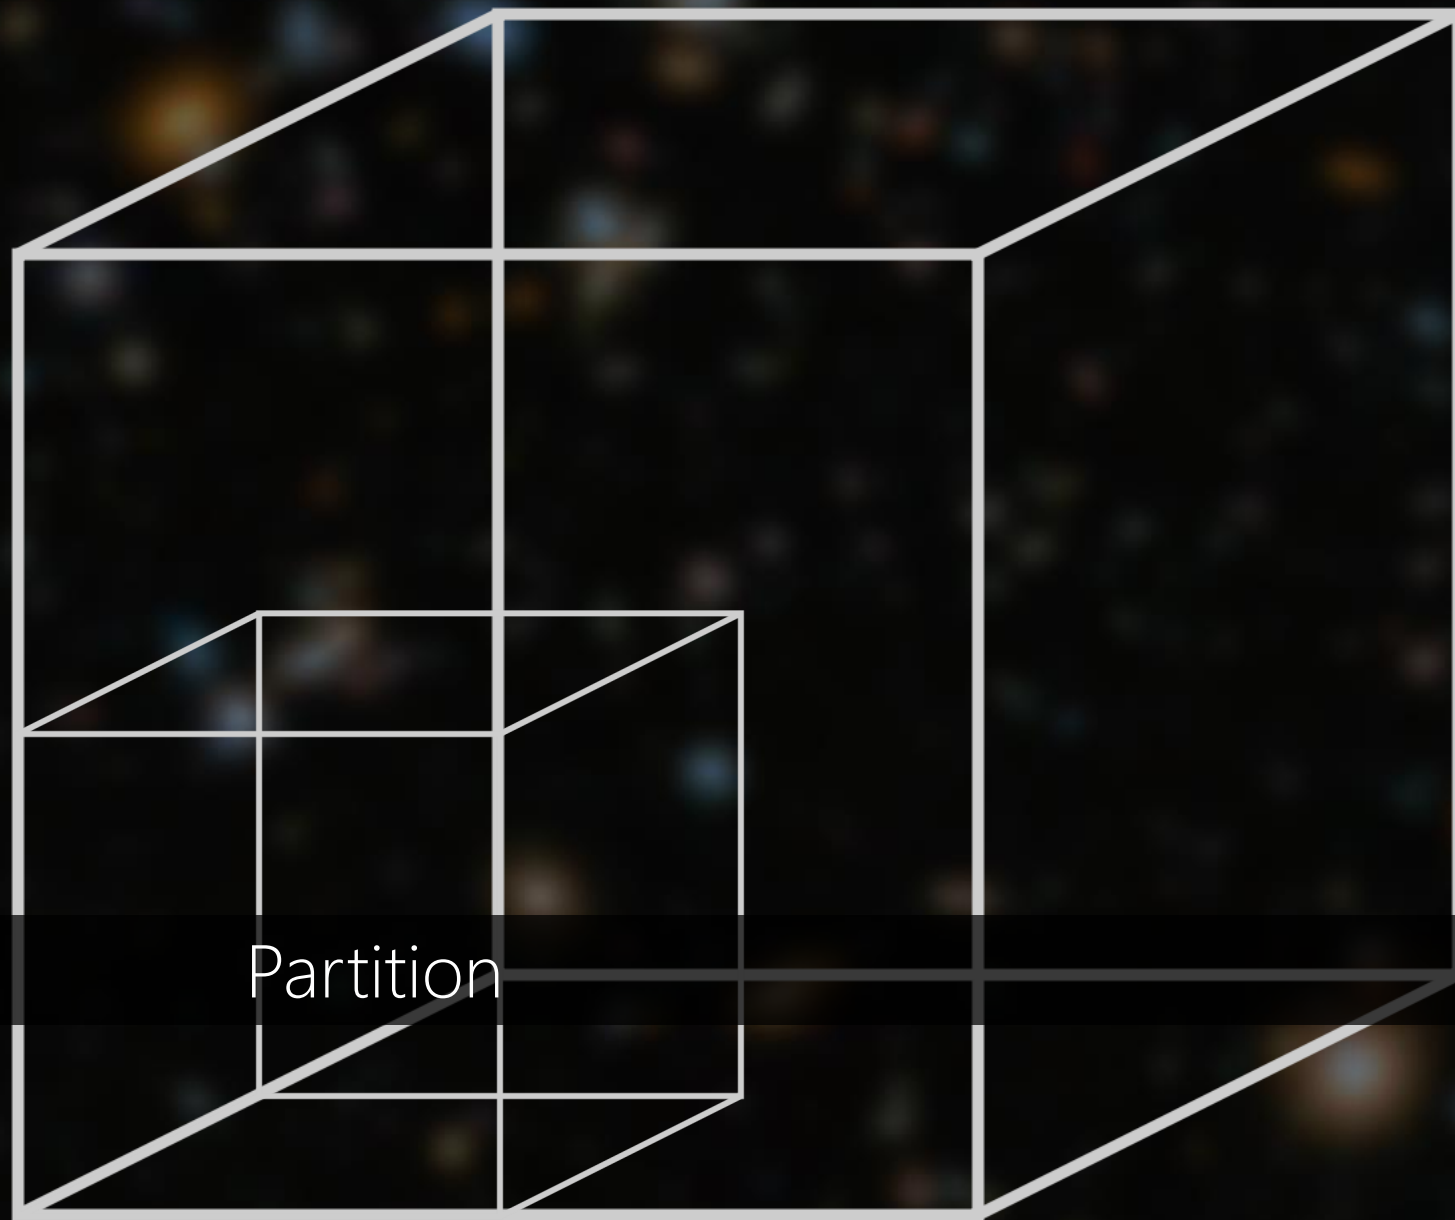
Storage File System

Disk File System

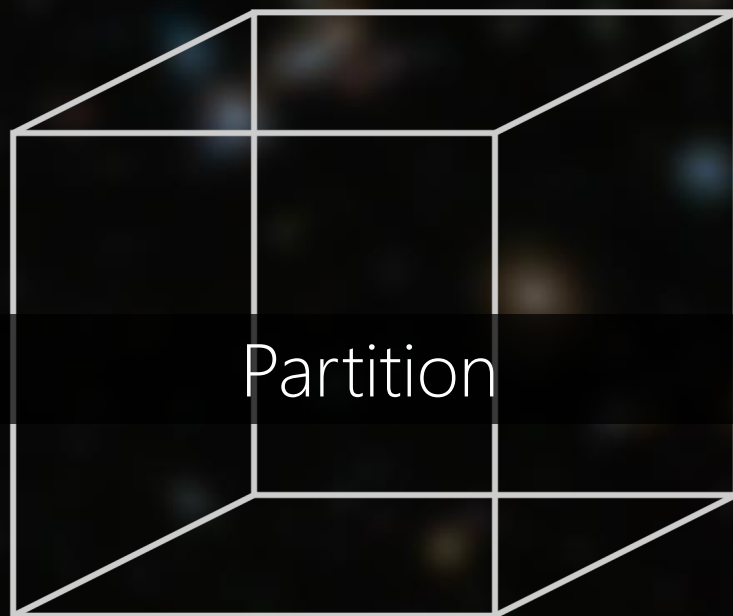
File System: Low Level

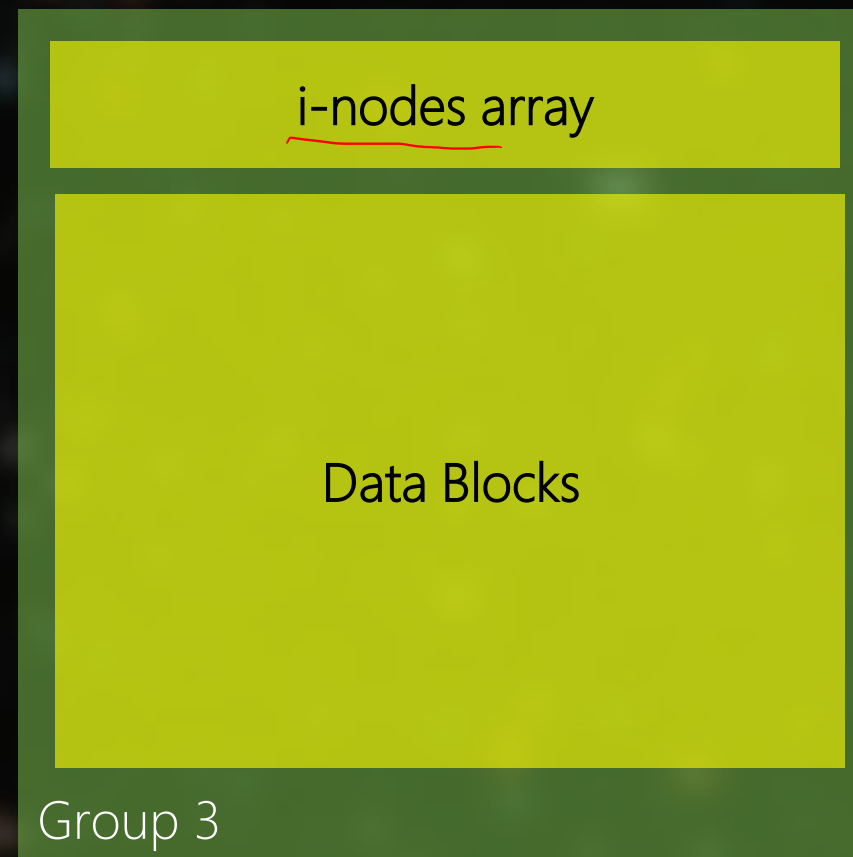
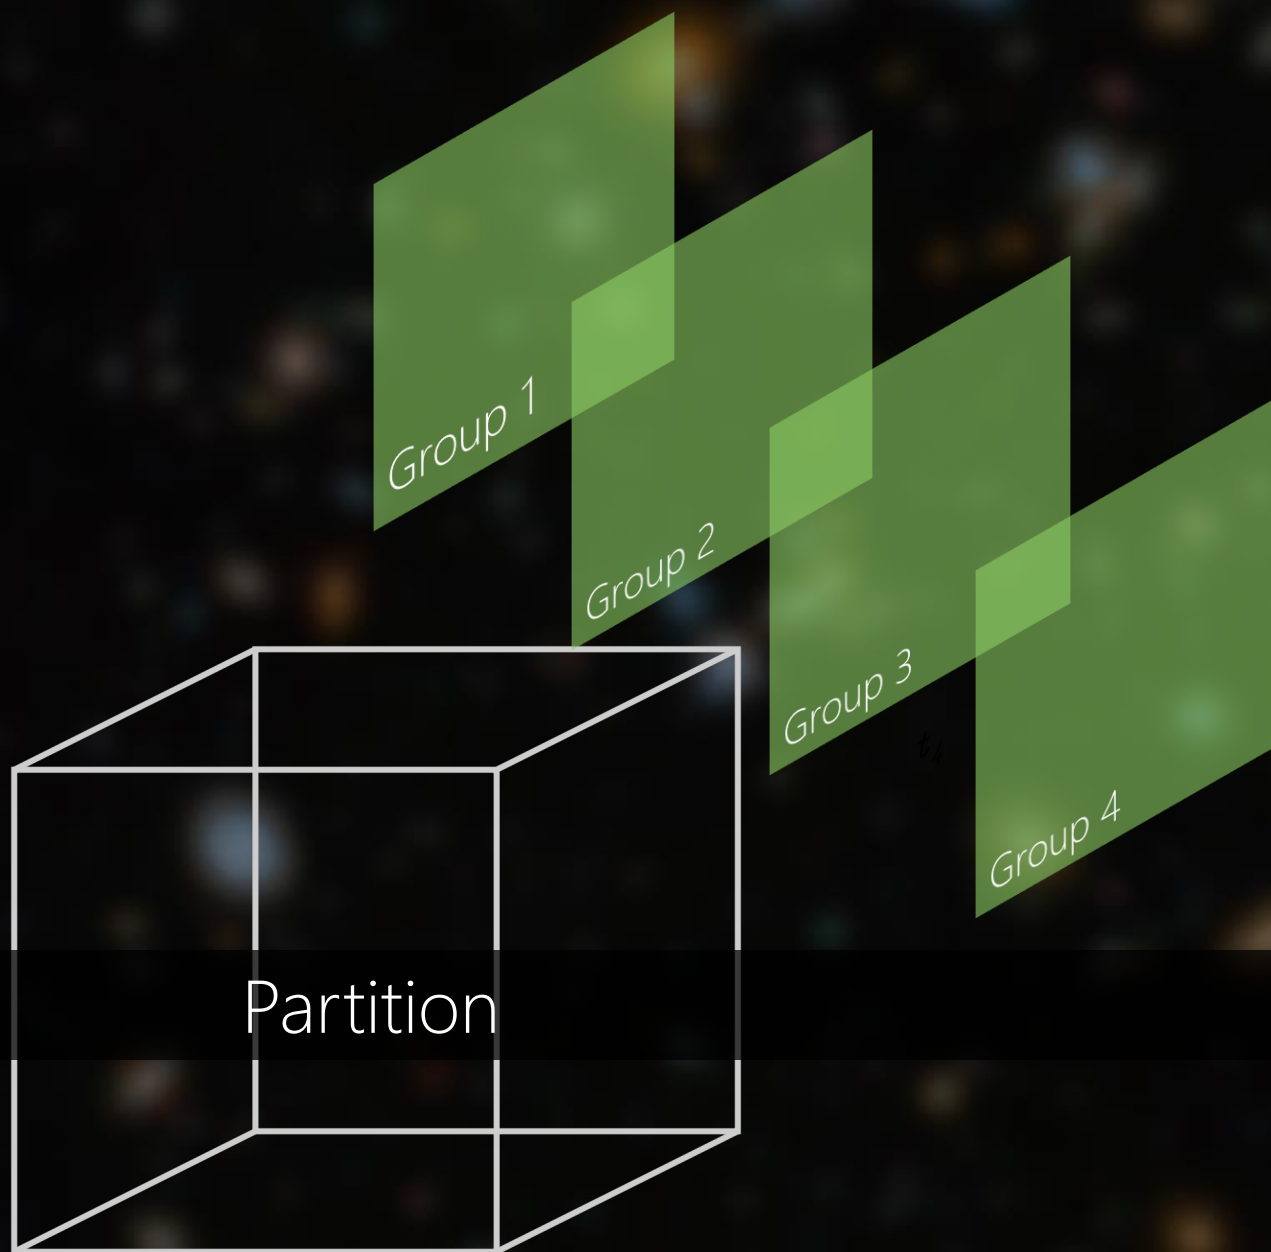


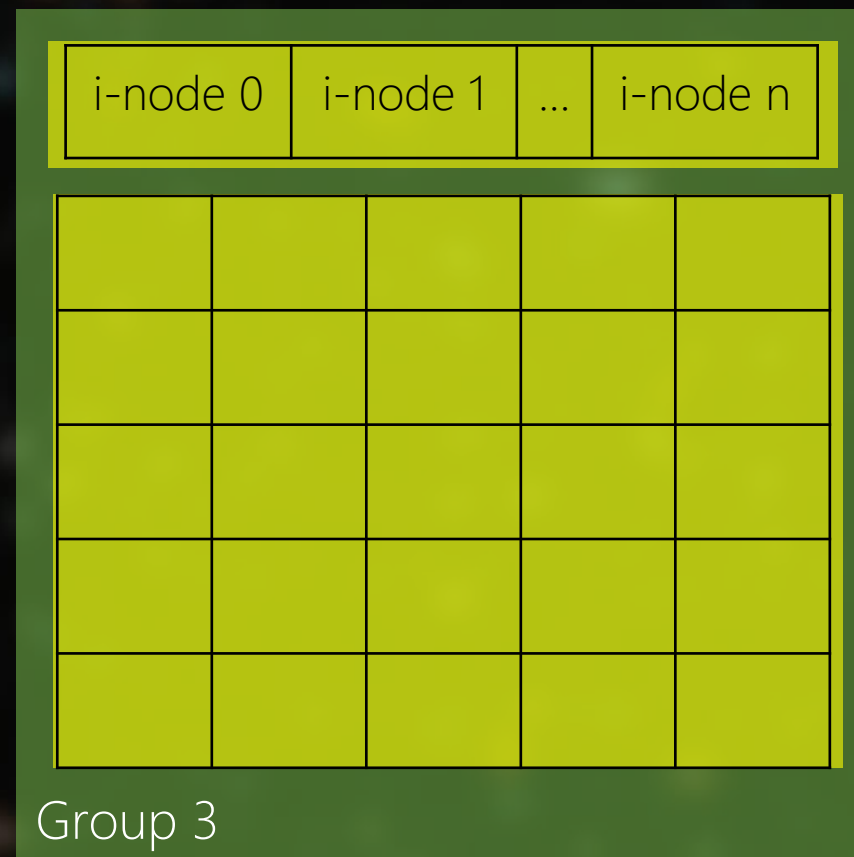
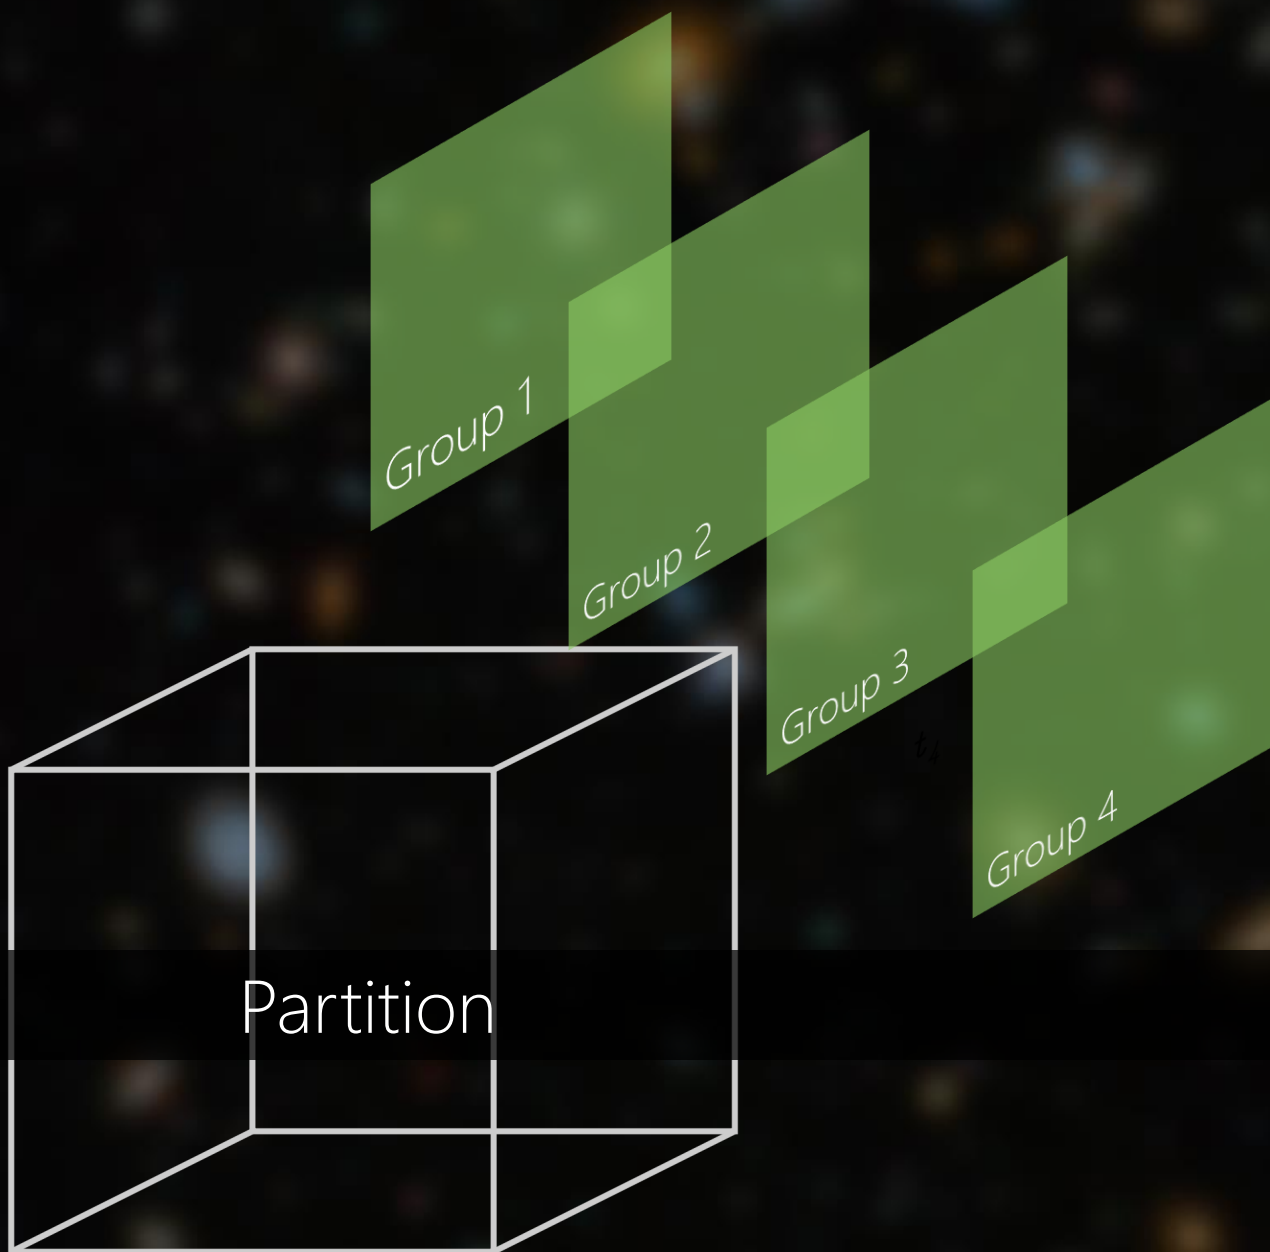
Storage

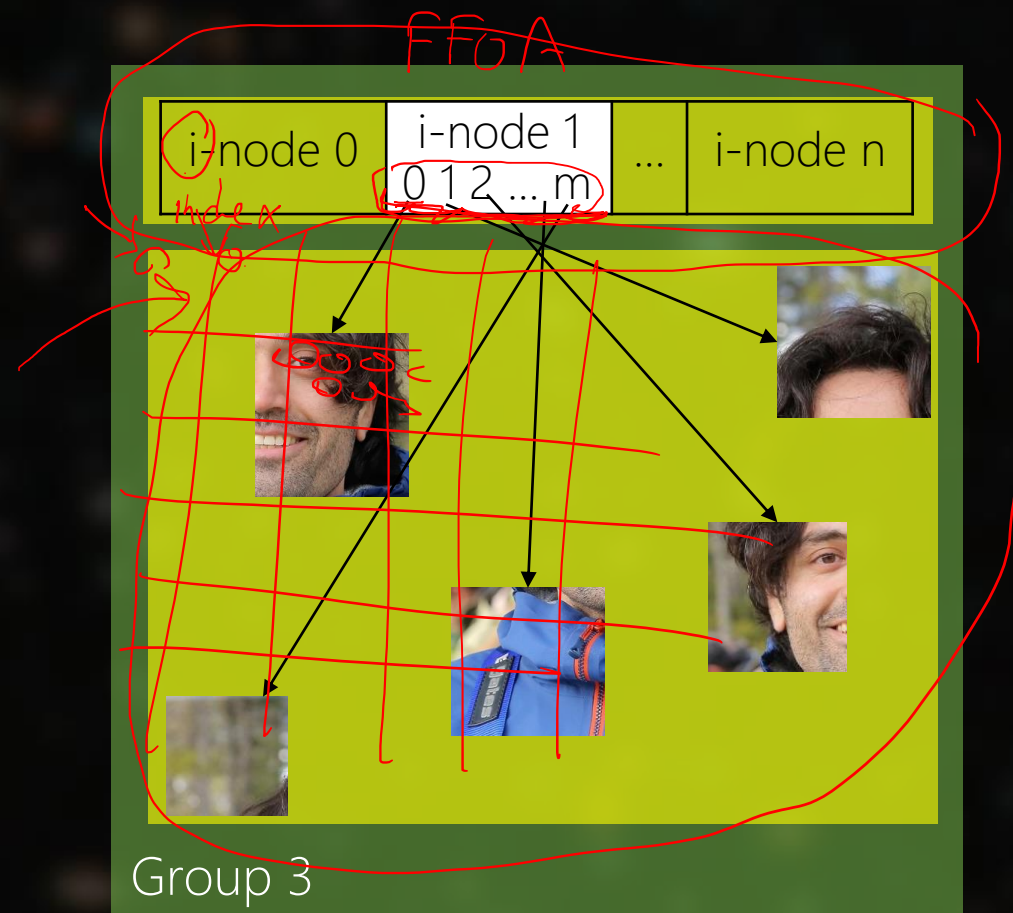
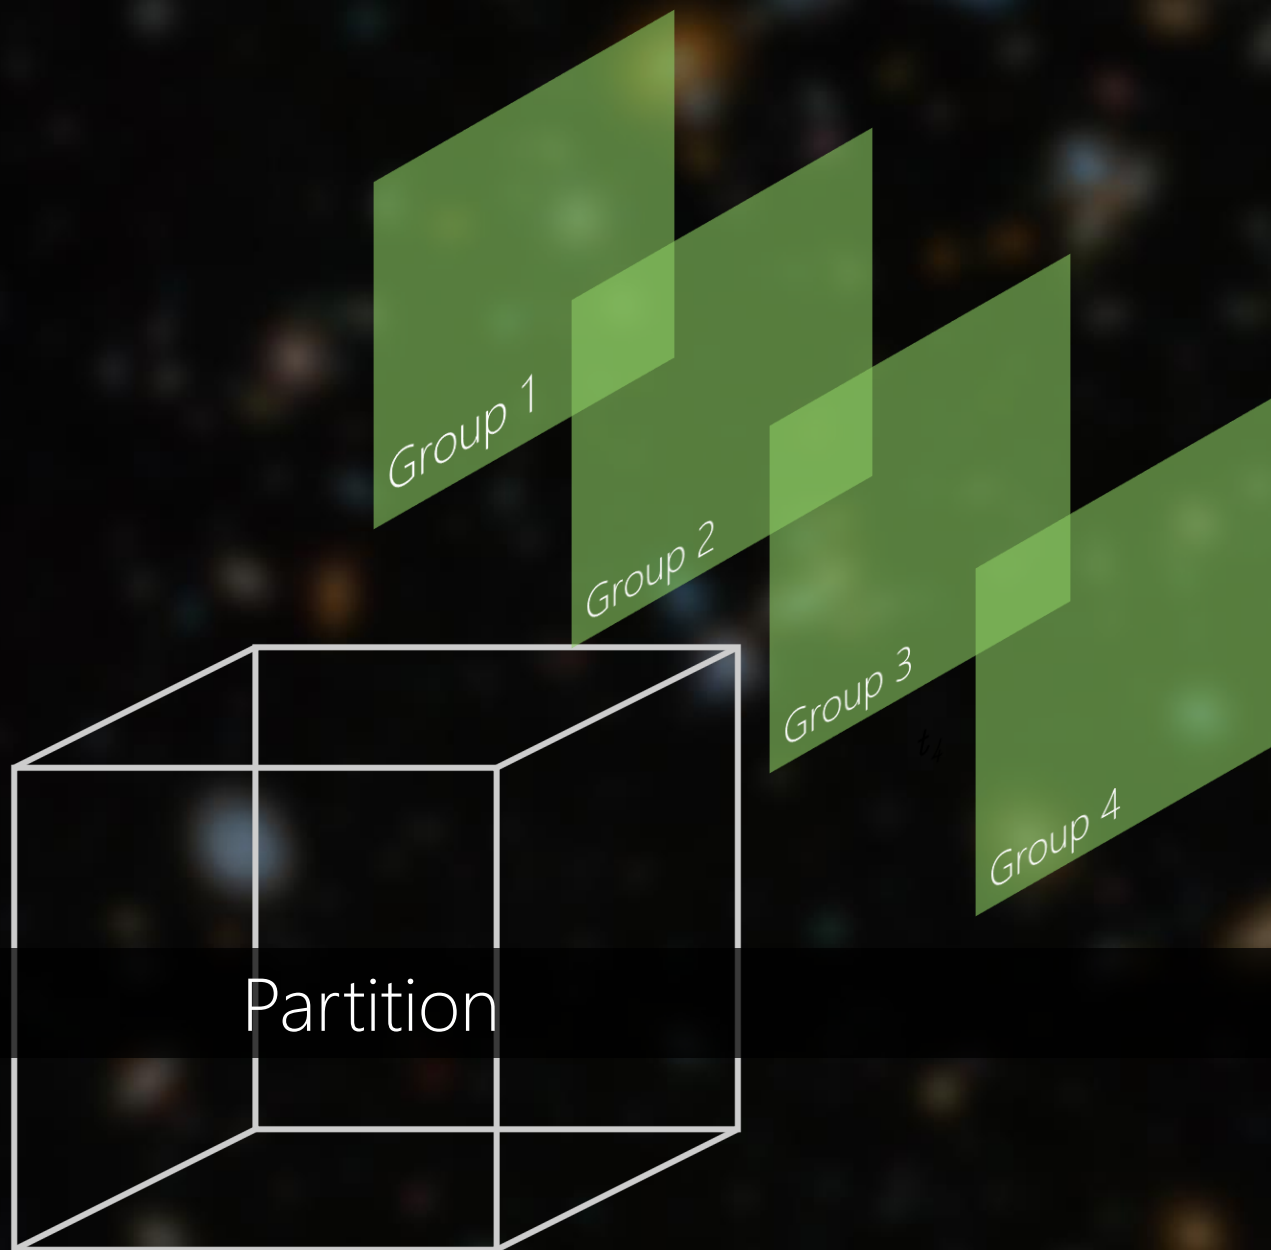


Partition









Each File has one i-node

```
$ ls -i {filename}
```

```
hfani@bravo:~$ ls -i hfani.jpeg  
99306609 hfani.jpeg
```



Each File has many data blocks

`$ stat {filename}`

```
hfani@bravo:~$ stat hfani.jpeg
  File: hfani.jpeg
  Size: 1552      Blocks: 5      IO Block: 1048576 regular file
Device: 2ch/44d Inode: 99306609   Links: 1
Access: (0644/-rw-r--r--)  Uid: (239080/   hfani)   Gid: (   400/   temp)
Access: 2022-09-14 11:17:45.909419913 -0400
Modify: 2022-09-14 11:17:45.927970184 -0400
Change: 2022-10-19 00:16:57.563217031 -0400
 Birth: -
```


Handwritten annotations:

- A red box around "Size: 1552" with an arrow pointing to it.
- A yellow box around "Blocks: 5" with a red circle around the "5".
- A yellow box around "IO Block: 1048576" with a red circle around "1048576".
- A red "X" above "Blocks: 5".
- Red handwritten text "Byte" above "IO Block: 1048576".

Can we vi an i-node?

NO

/Path/Filename → i-node → Data Blocks

"/home/hfani/hfani.jpeg" → 99306609 → 

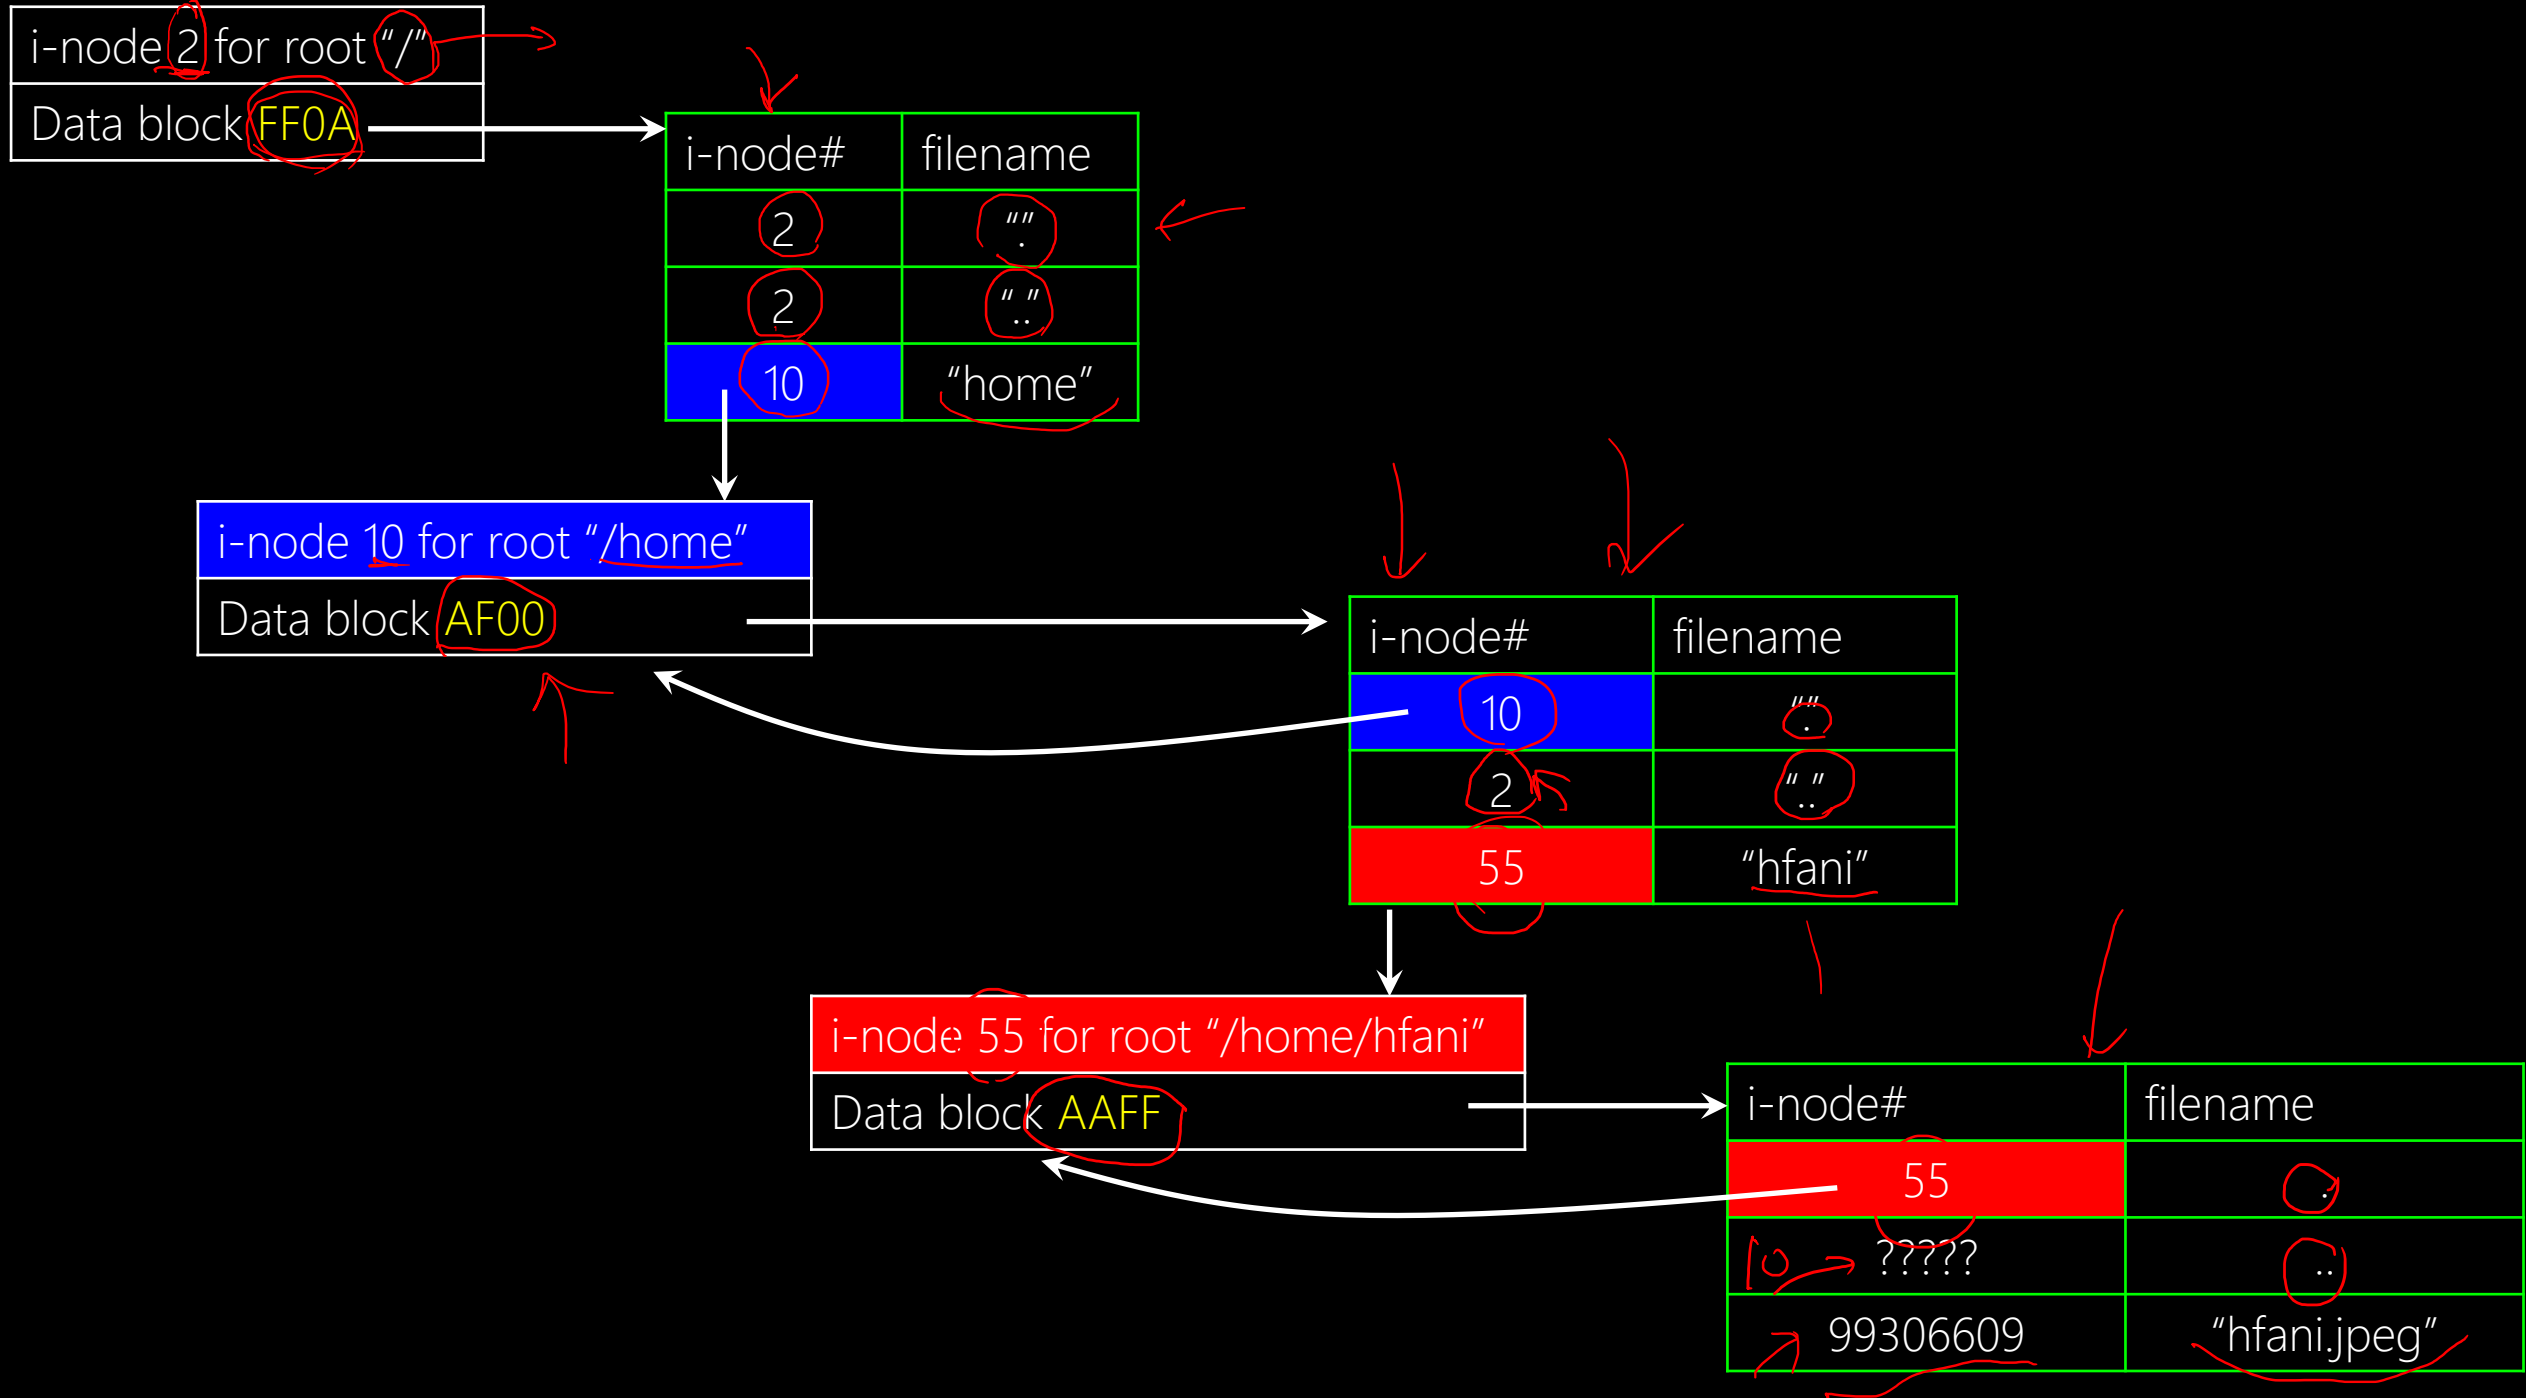
Red annotations: A red circle around the path, a red circle around the i-node number, and red arrows pointing from the path and i-node to the image thumbnails.

creat() or open()

Directories

They are files. So, each of them has one i-node.

The content is not an image, audio, ... but mapping between filenames and their i-nodes



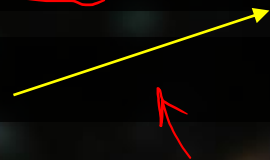
Directories

Is it possible to have multiple links to the same file? Yes.
How?

"/home/hfani/hfani.jpeg" → 1407374883666816 →



"/nosseinfani.png"



599

i-node 2 for root /
Data block **FF0A**

i-node#	filename
2	.
2	..
99306609	hosseinfani.png
10	

No Duplication!
Still one single file.
Two or more links.
hardlink, shortcut,

i-node 10 for root '/home'
Data block **AF00**

i-node#	filename
10	.
2	..
55	hfani

i-node 55 for root '/home/hfani'
Data block **AAFF**

i-node#	filename
55	.
?????	..
99306609	hfani.jpeg

hf

↓ /home

← 993-hf

hfani



What happens when you delete a file?



What happens when you move a file?



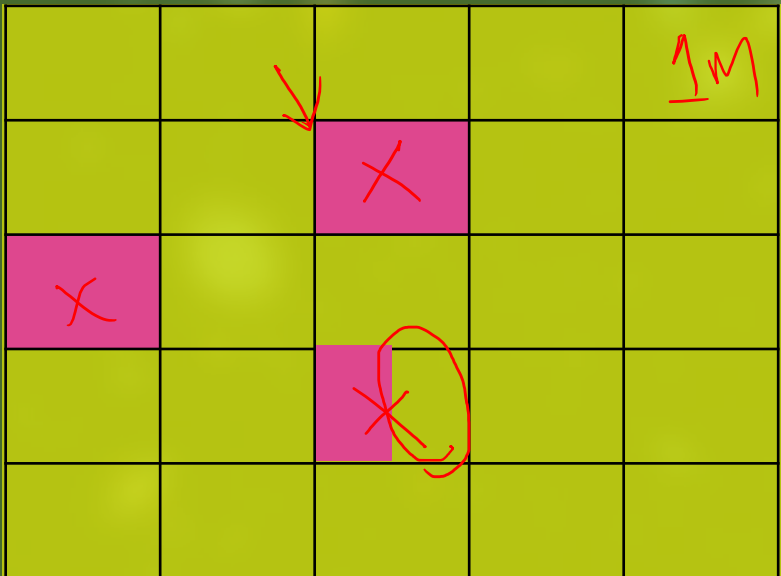
What happens when you copy a file?

Who determines the #partitions, #groups, data block size?

Partitioning (Formatting)

Large Data Blocks vs Small Data Blocks

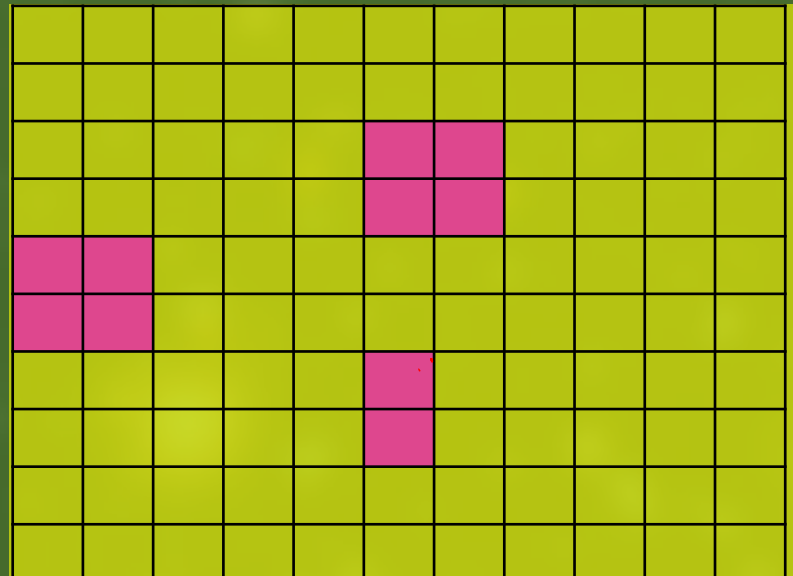
i-node 0	i-node 1	...	i-node n
----------	----------	-----	----------



3 × 1M Blocks!

~~2.5M~~
3M

i-node 0	i-node 1	...	i-node n
----------	----------	-----	----------



10 × 256K Blocks.

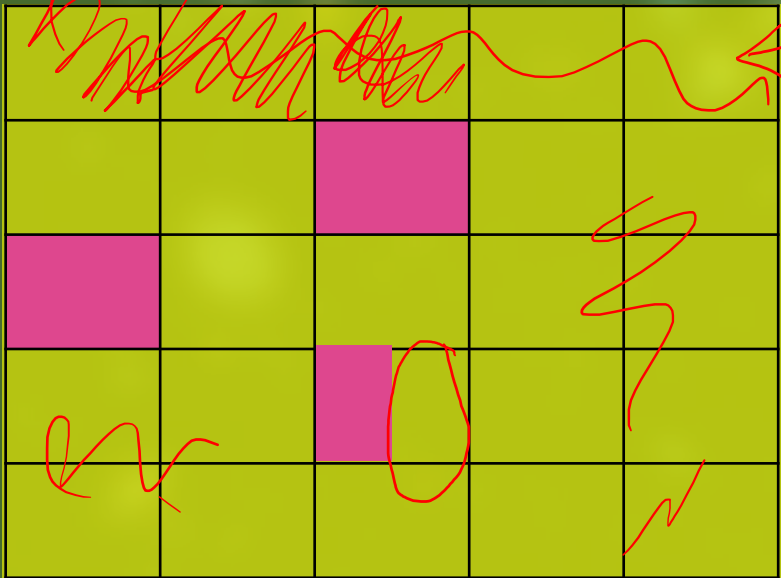
2.5M

Large Data Blocks vs Small Data Blocks

Internal vs. External Fragmentation

Small vs. Large i-node

i-node 0	i-node 1	...	i-node n
----------	----------	-----	----------



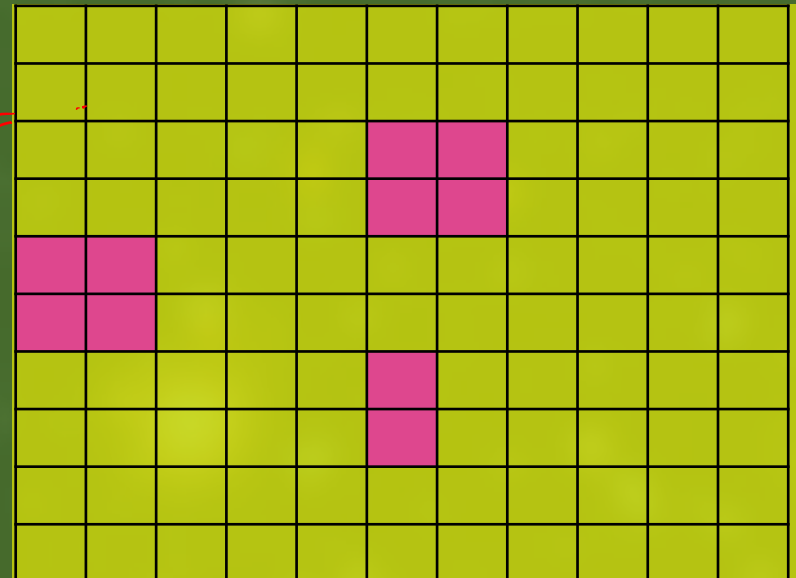
3 × 1M Blocks!

Movies?
Documents?

2-3 GB

256K

i-node 0	i-node 1	...	i-node n
----------	----------	-----	----------



10 × 256K Blocks.

File Types (High-Level to Low-Level)

- Regular Files → text or binary
- Directory Files → (i-node, filename) pairs
- Special Files (Devices)
 - Block → HDD, SSD, CD, ... ←
 - Character (Stream) → TTY, Keyboard, Mouse, Printer, ...
- Socket (Networking)
- FIFO (Pipes)
- Symbolic Link

RD/WR on a Regular File: `open("./myfile.txt")`
vs.

RD/WR on a Storage (Lab06): `open("/dev/sda1")`

-
- Regular Files → text or binary
 - Directory Files → (i-node, filename) pairs
 - Special Files (Devices)
 - Block → HDD, SSD, CD, ...
 - Character (Stream) → TTY, Keyboard, Mouse, Printer, ...
 - Socket (Networking)
 - FIFO (Pipes)
 - Symbolic Link

Ca + sda1