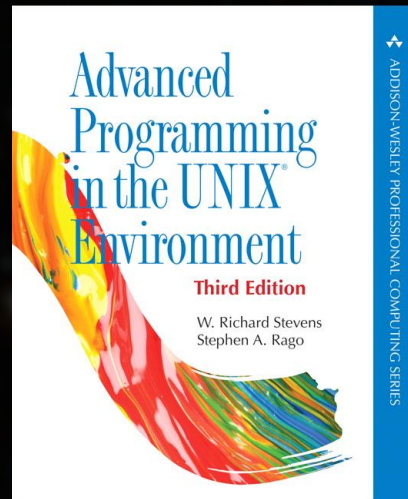


# Into the Wild



Into the Wild (2007) - Sean Penn  
Eddie Vedder - Hard Sun

<https://www.youtube.com/watch?v=Ez8b2VHjVB0>



# Chapter 07: Process Environment

## Chapter 08: Process Control

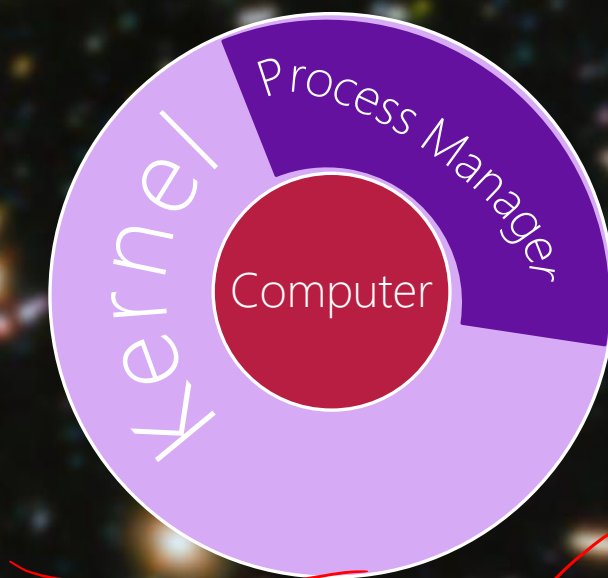
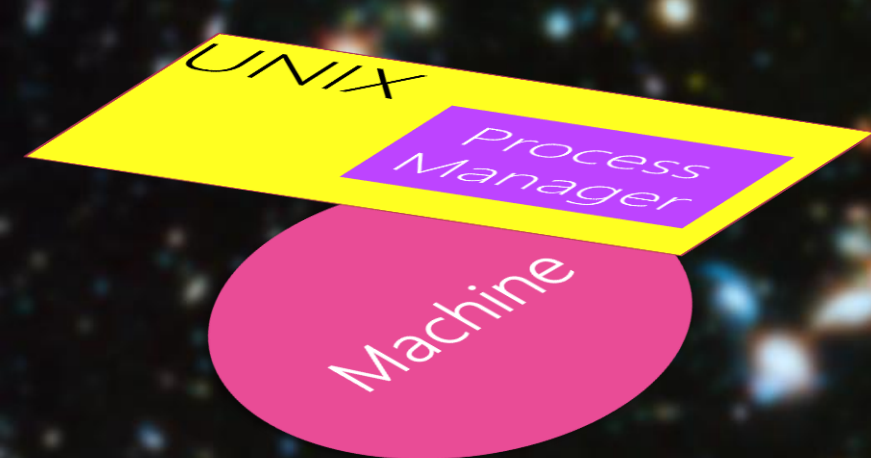
---

# Process Manager

aka. Process Control

---





# Computer

Memory

Kernel: Device Manager

Kernel: Memory Manager

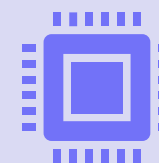
Kernel: File Manager

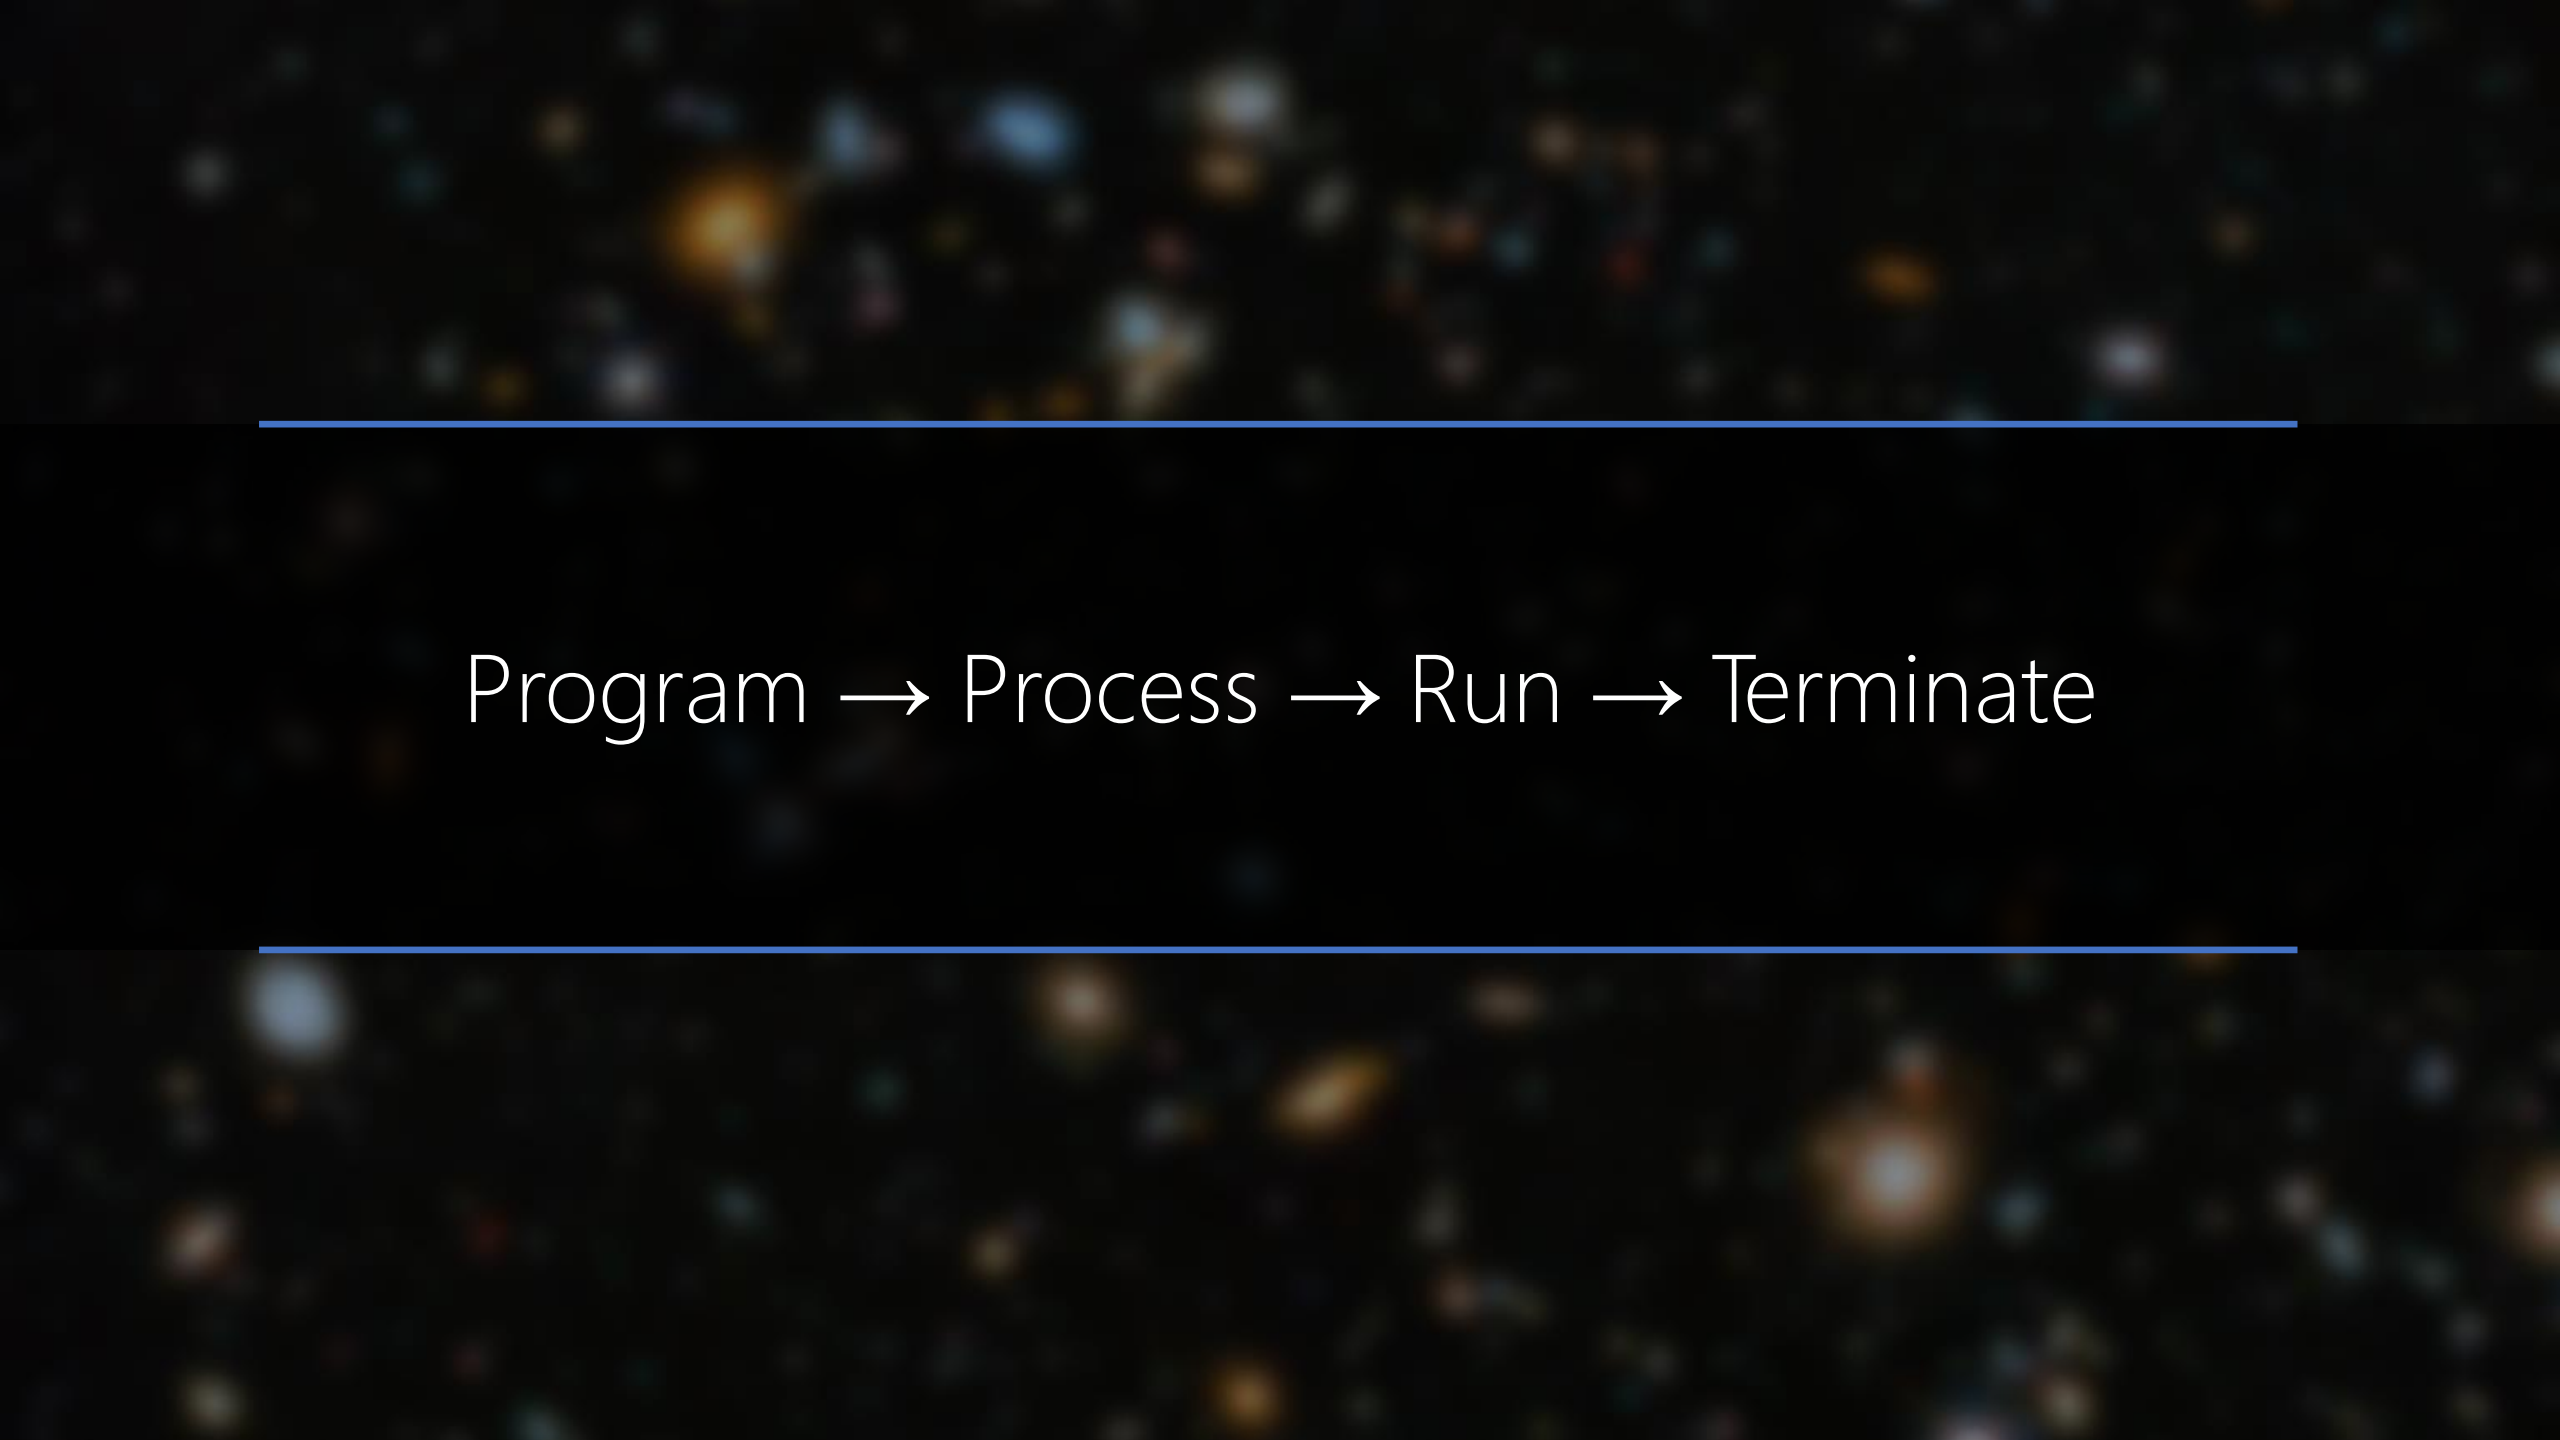
Kernel: Network Manager

Kernel: Process Manager

Bus

Processor

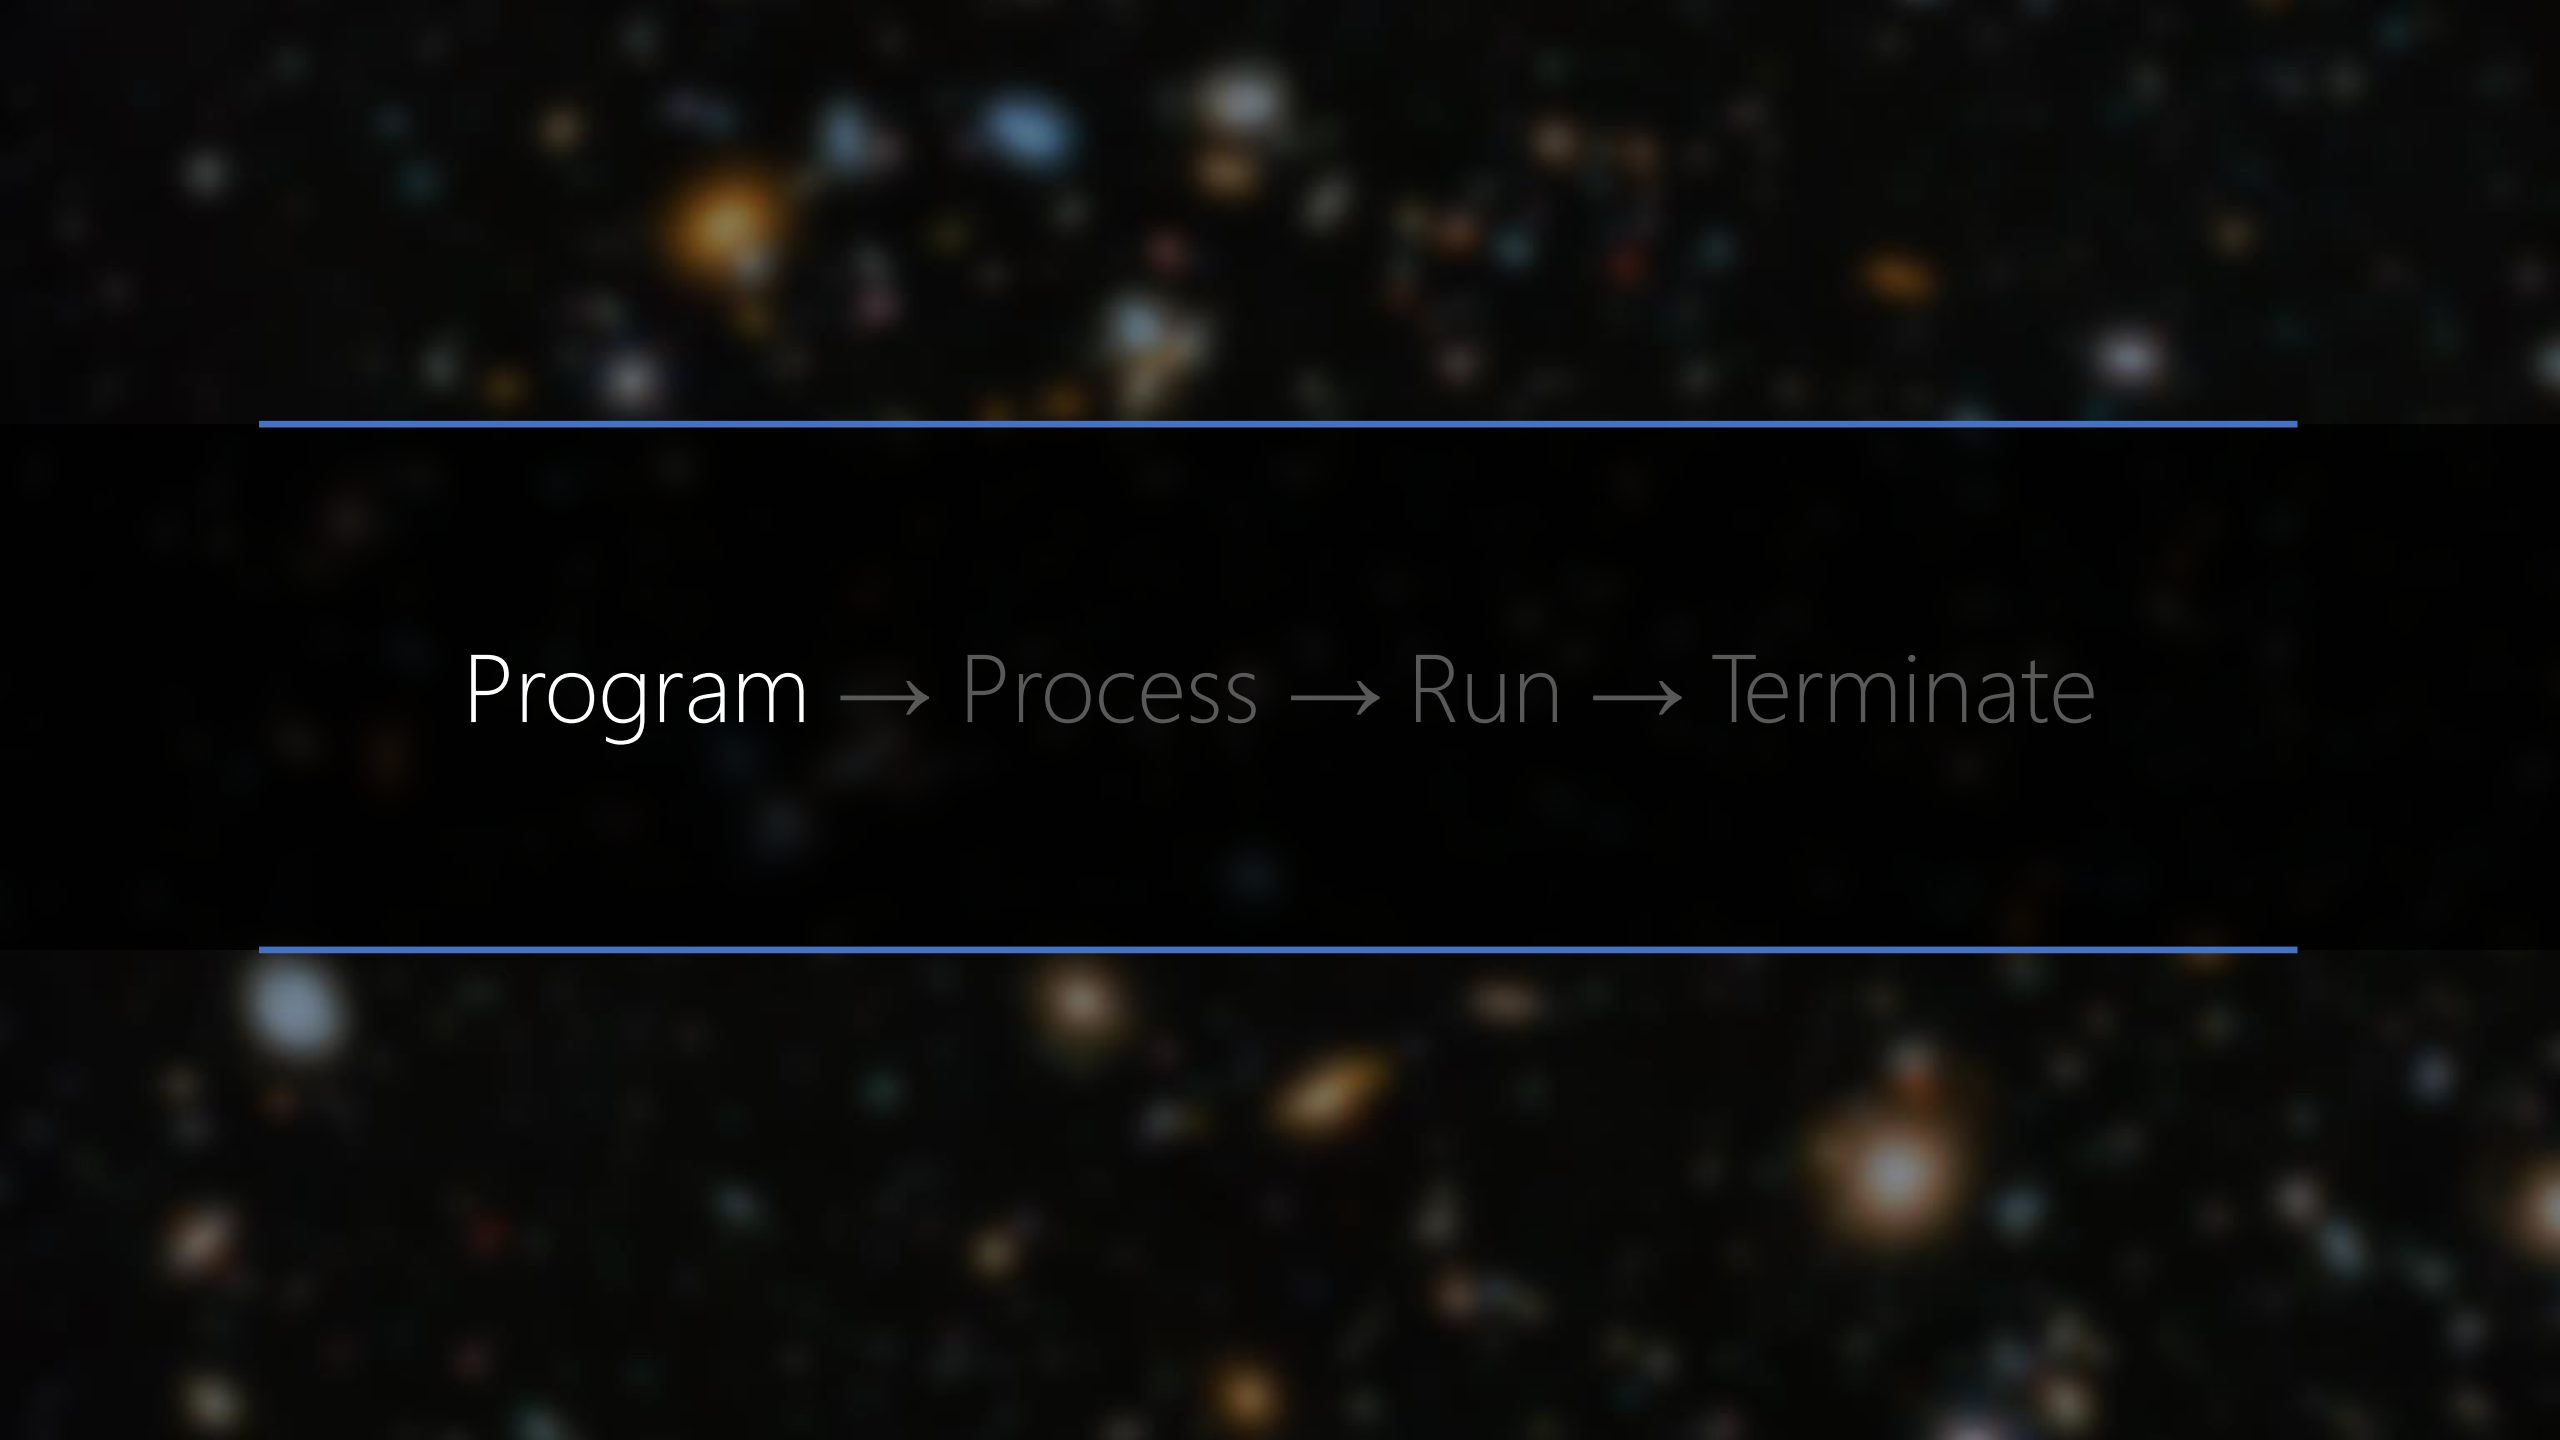




---

Program → Process → Run → Terminate

---



---

Program → Process → Run → Terminate

---

---

Any Program MUST have an entry point

---

What part of the code has the first opcode?

---

```
void main(void)
```

---

```
shell$ ./program
```



---

scanf( )

void main(int argc, char \*argv[])  
int main(int argc, char \*argv[])

---

shell\$ ./program arg1 arg2 arg3 ....

```
hfani@charlie:~$ vi main_args.c
#include <stdio.h>
int main(int argc, char *argv[]){
    printf("there are %d arguments in the shell:\n", argc);
    for(int i=0; i < argc; ++i){
        printf("arg%d: %s\n", i, argv[i]);
    }
    return 0;
}
```

Name of the program file is the first argument!

```
hfani@charlie:~$ cc main_args.c -o main_args
hfani@charlie:~$ ./main_args
there are 1 arguments in the shell:
arg0: ./main_args
hfani@charlie:~$ ./main_args param1 param2
there are 3 arguments in the shell:
arg0: ./main_args
arg1: param1
arg2: param2
hfani@charlie:~$
```

```
hfani@charlie:~$ vi main_add.c
```

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

Arguments are string of chars!

ASCII to Integer

int atoi (const char \*str);

```
hfani@charlie:~$ cc main_add.c -o main_add
```

```
hfani@charlie:~$ ./main_add 2 2
```

```
2 + 2 = 4
```

```
hfani@charlie:~$ ./main_add 2 4
```

```
2 + 4 = 6
```

```
hfani@charlie:~$
```



Into the Wild (2007) - Sean Penn





Program → Process → Run → Terminate





---

# Memory Layout of a C Program

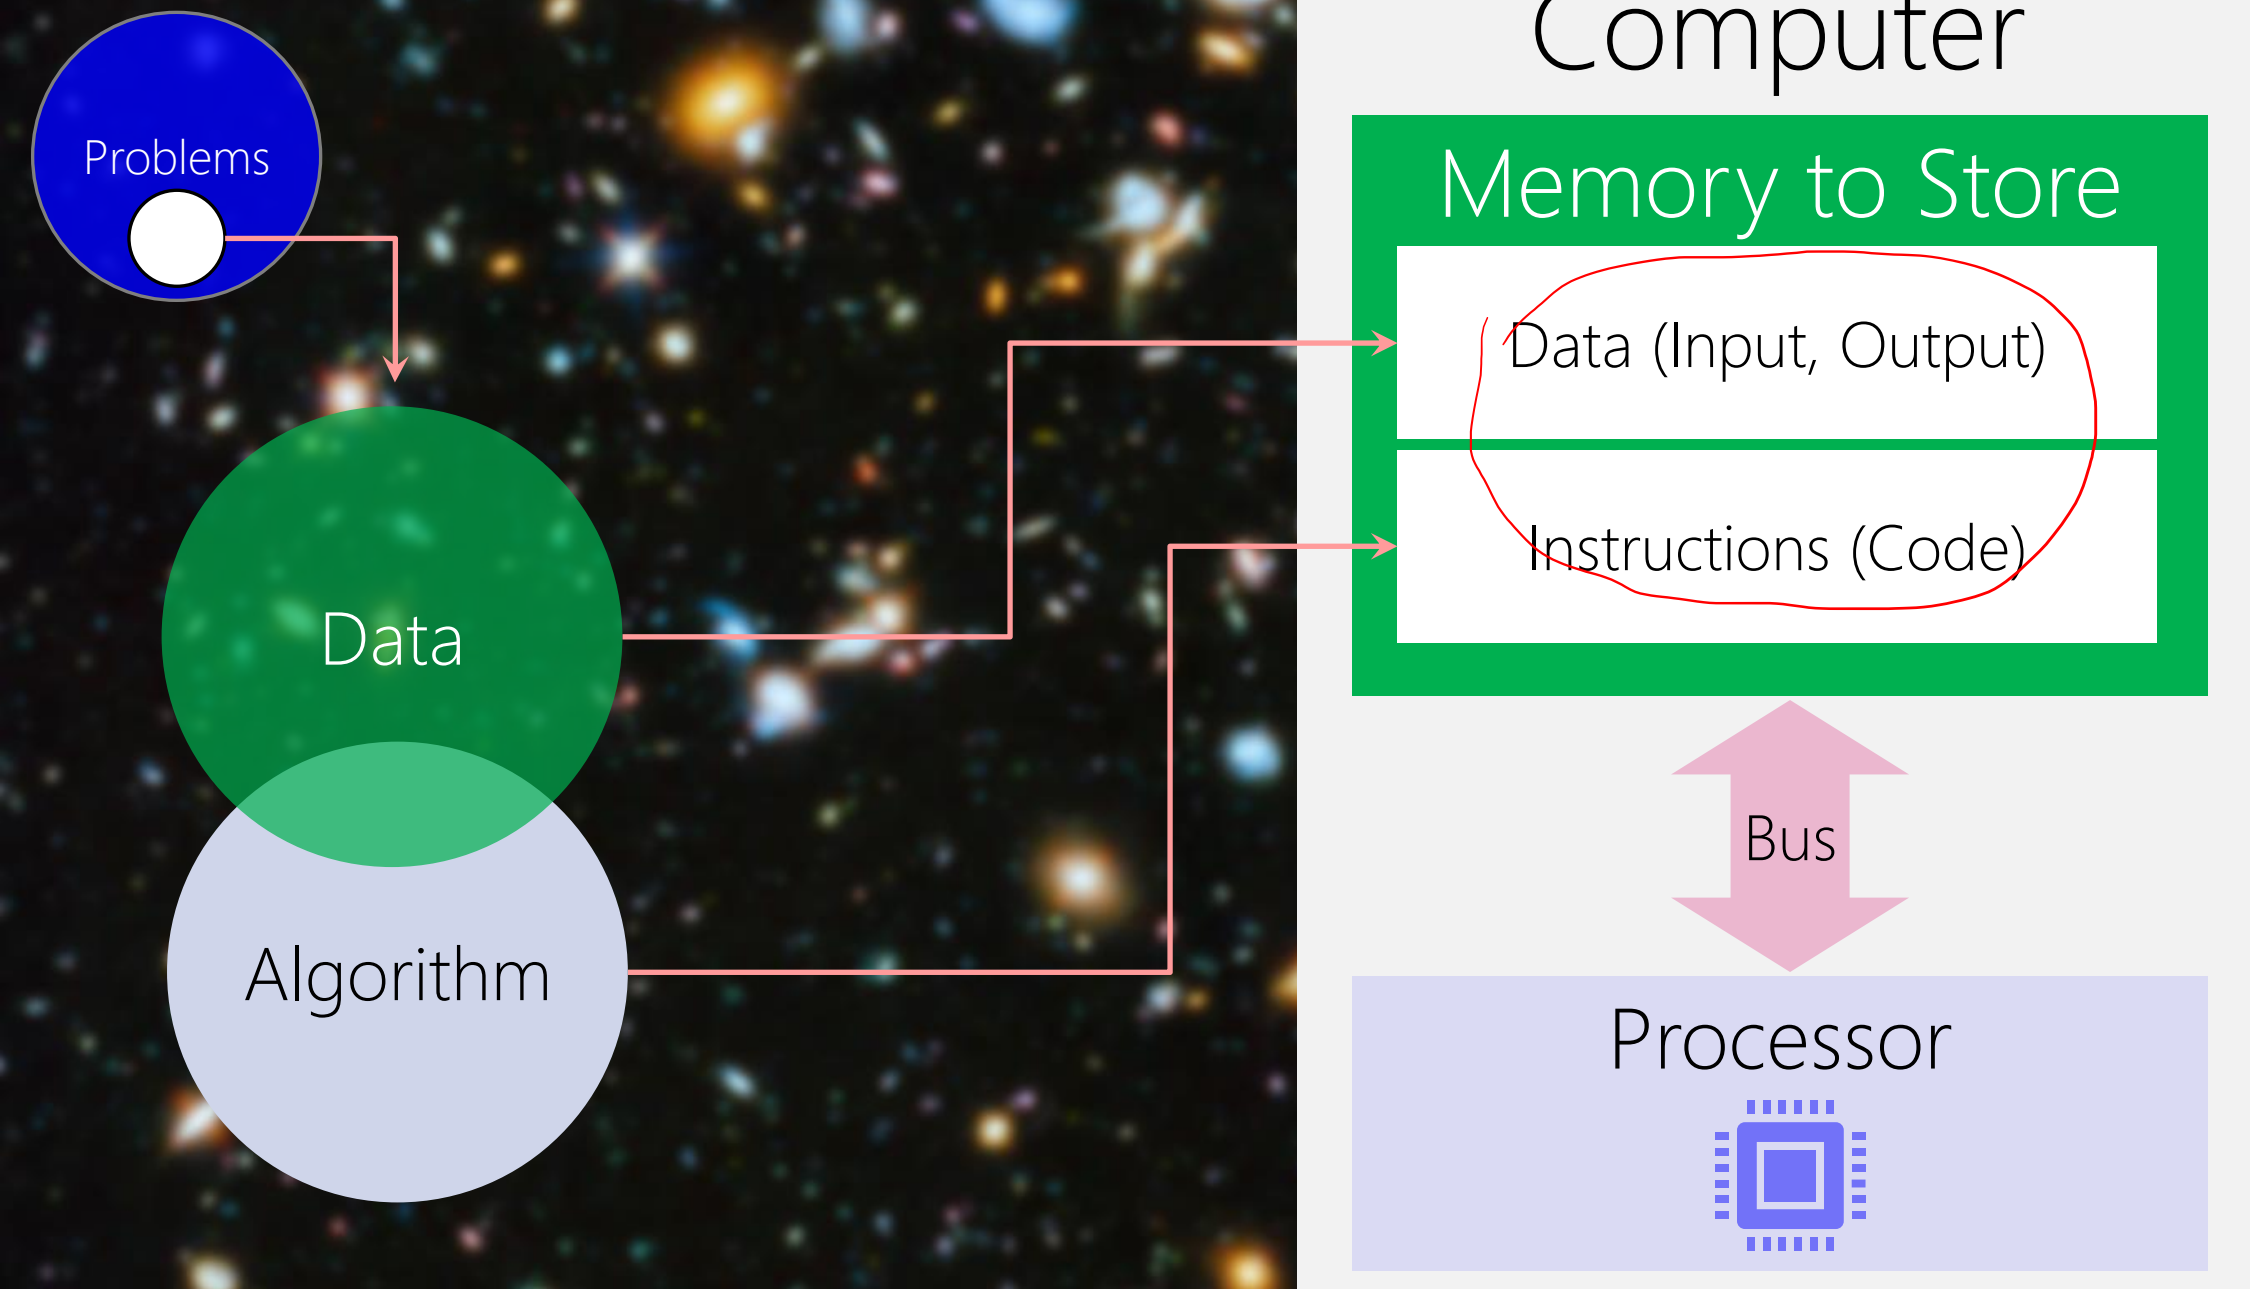
---















142

FAIRBANKS CITY TRANSIT SYSTEM



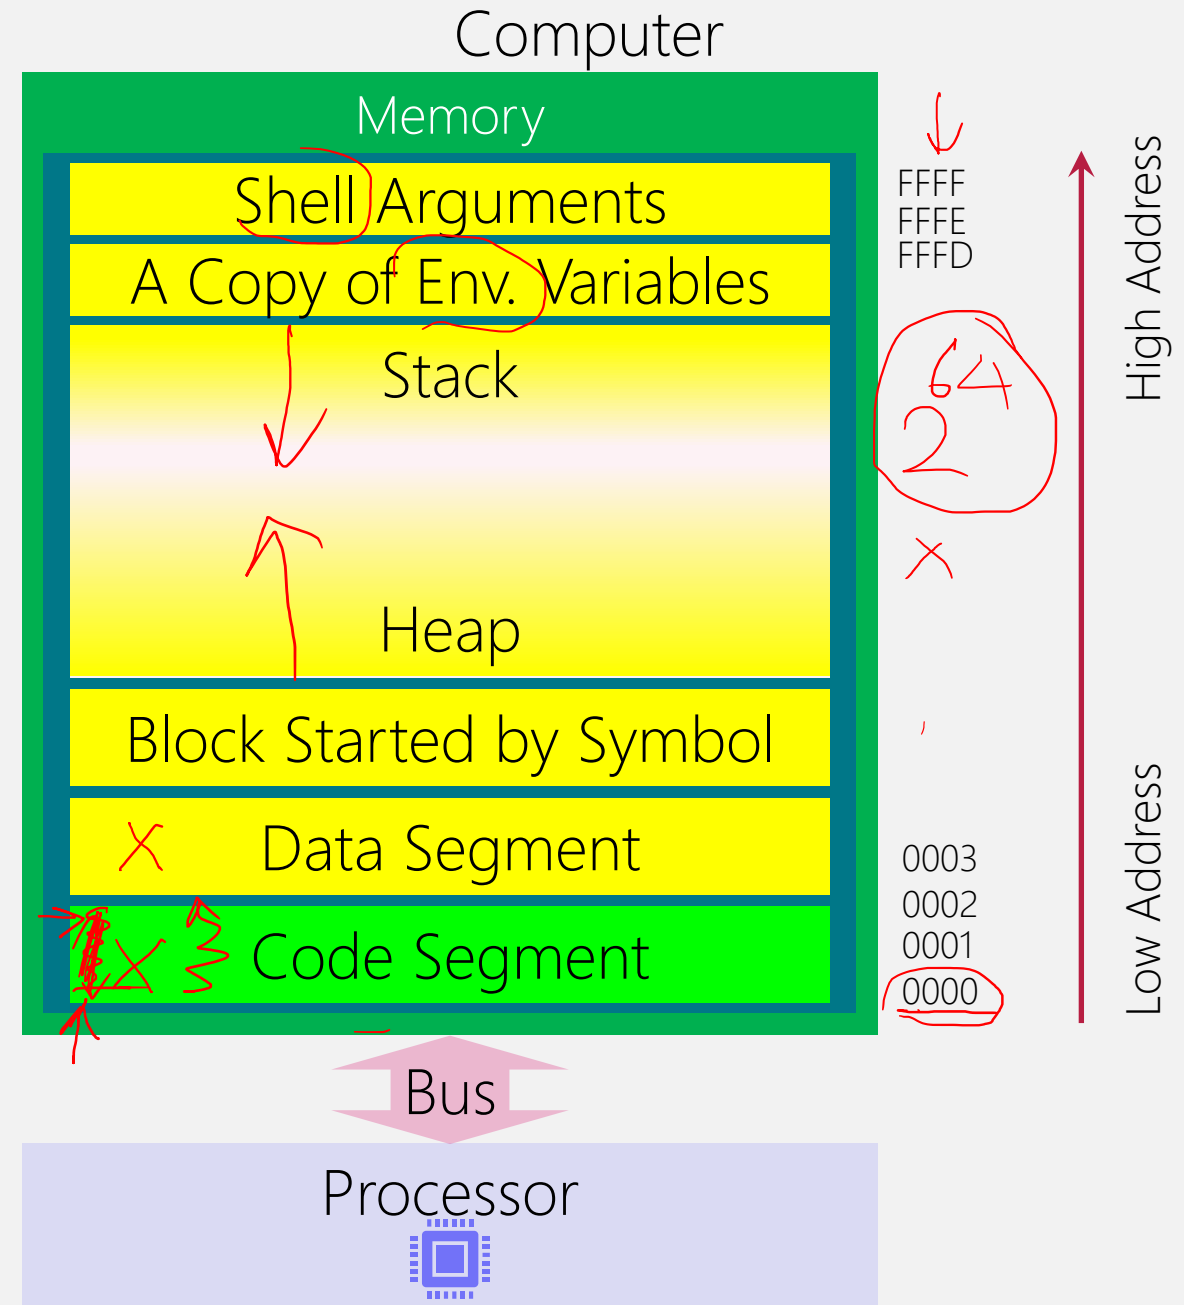
```

#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}

```





# Code Segment (CS)

aka. text

```
#include <stdio.h>
#include <stdlib.h>
int main(int argc, char *argv[]) {
    int a = atoi(argv[1]);
    int b = atoi(argv[2]);
    printf("%d + %d = %d\n", a, b, a + b);
    return 0;
}
```

0002e10	1000	0000	0000	0000	000d	0000
0002e20	1214	0000	0000	0000	0019	0000
0002e30	3de8	0000	0000	0000	001b	0000
0002e40	0008	0000	0000	0000	001a	0000
0002e50	3df0	0000	0000	0000	001c	0000
0002e60	0008	0000	0000	0000	fef5	6fff
0002e70	0308	0000	0000	0000	0005	0000
0002e80	03f0	0000	0000	0000	0006	0000
0002e90	0330	0000	0000	0000	000a	0000

# Data Segment (DS)

global variables, and variables inside main, that are initialized to a default value! (compile time)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

Memory

Data Segment

0F13

00000000

0F12

00000000

0F11

00000000

0F10

00000000

# Block Started by a Symbol (BSS)

For uninitialized variables! (compile time)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

*main()*

Memory

BSS Segment

0F15

0F14

→ ????????

→ ????????

<https://en.wikipedia.org/wiki/.bss>

# Shell Argument + Environment Variables

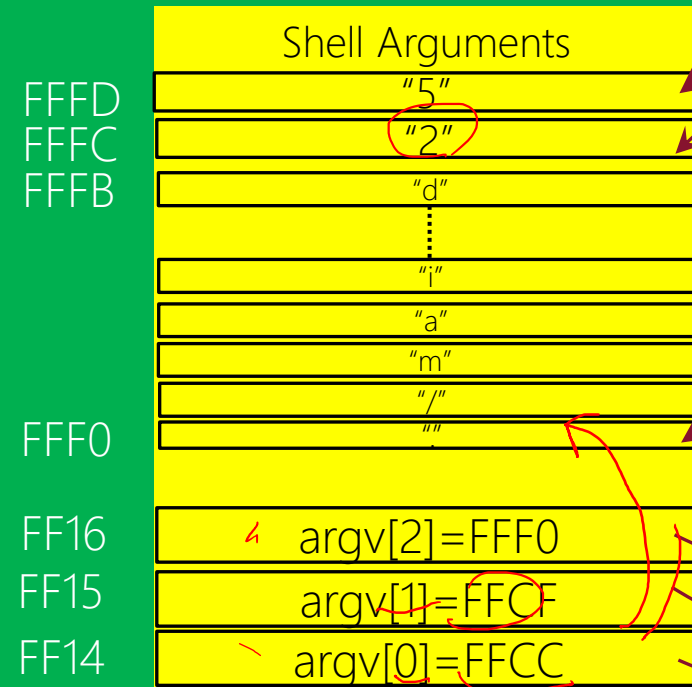
Provided by the Shell (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

## Memory



# Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>

int result;

int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef atoi
/* Convert a string to an int. */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_TYPE **en
                 int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base, group, _
}
libc_hidden_def (INTERNAL (strtol))
```





Into the Wild (2007) - Sean Penn

# Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>

int result;

int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

My Return Address 1

# Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    → return (int) strtol (nptr, (char **) NULL, 10);
    lib_hidden_def (atoi)
}
```

```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_T
                  int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base,
    lib_hidden_def (INTERNAL (strtol))
}
```

My Return Address 2

My Return Address 1



# Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

Diagram illustrating the stack frame for the `atoi` function. The function signature is `int atoi (const char *nptr)`. The return type is `int`. The arguments are `const char *nptr`. The function body shows a call to `strtol` with arguments `nptr`, `(char **) NULL`, and `10`. The return value of `strtol` is cast to `int` and returned. The function is defined in `libc_hidden_def (atoi)`.

Diagram illustrating the stack frame for the `strtol` function. The function signature is `INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_T int base, int group)`. The return type is `INTERNAL`. The arguments are `const STRING_TYPE *nptr`, `STRING_T int base`, and `int group`. The function body shows a call to `INTERNAL (_strtol_l)` with arguments `nptr`, `endptr`, `base`, and `group`. The return value of `INTERNAL (_strtol_l)` is returned. The function is defined in `libc_hidden_def (INTERNAL (strtol))`.

Where should I come back?

My Return Address 2

My Return Address 1

# Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

```
INT
INTERNAL (strtoul) (const STRING_TYPE *nptr, STRING_T
                  int base, int group)
{
    return INTERNAL (__strtoul_1) (nptr, endptr, base,
    libc_hidden_def (INTERNAL (strtoul))
```

My Return Address 2

My Return Address 1



# Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_T
                  int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base,
    libc_hidden_def (INTERNAL (strtol))
```

Where should I  
come back?

My Return Address 1

# Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);
    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef atoi
/* Convert a string to an int. */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

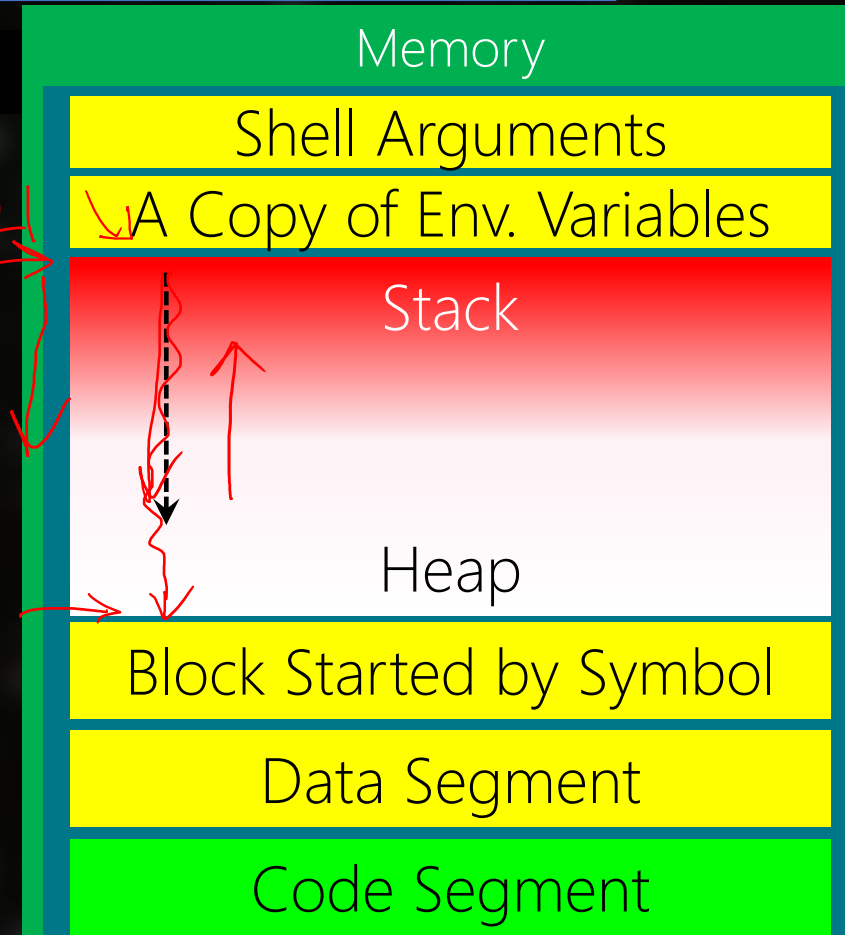
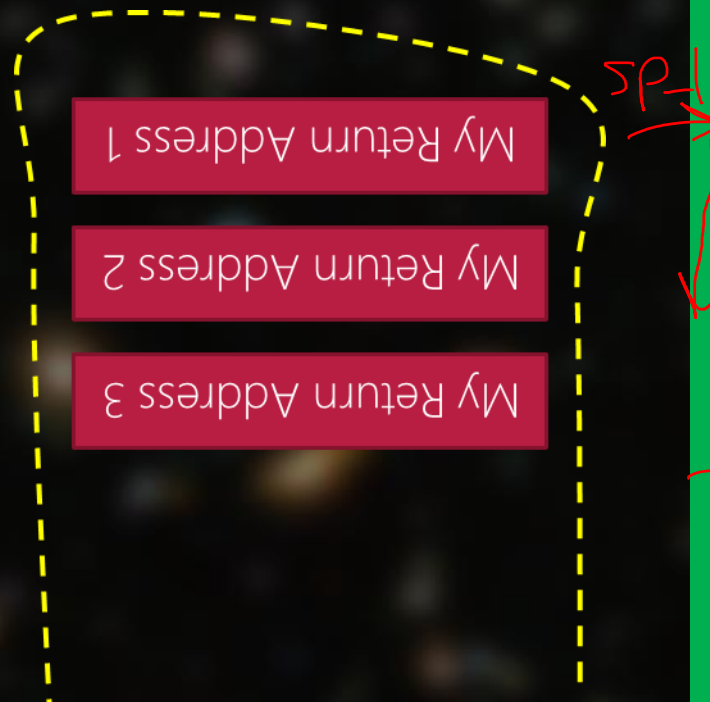
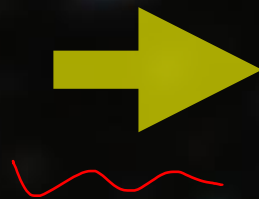
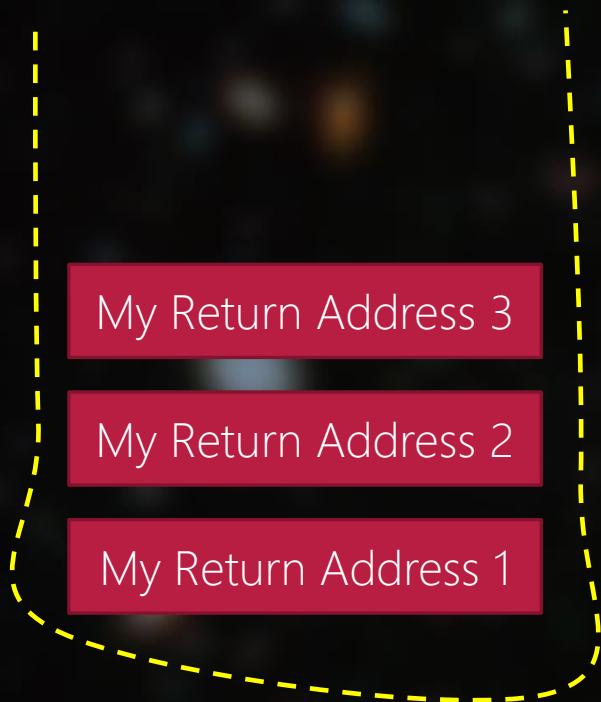
```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_T
                  int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base,
}
libc_hidden_def (INTERNAL (strtol))
```

My Return Address 1

# Stack

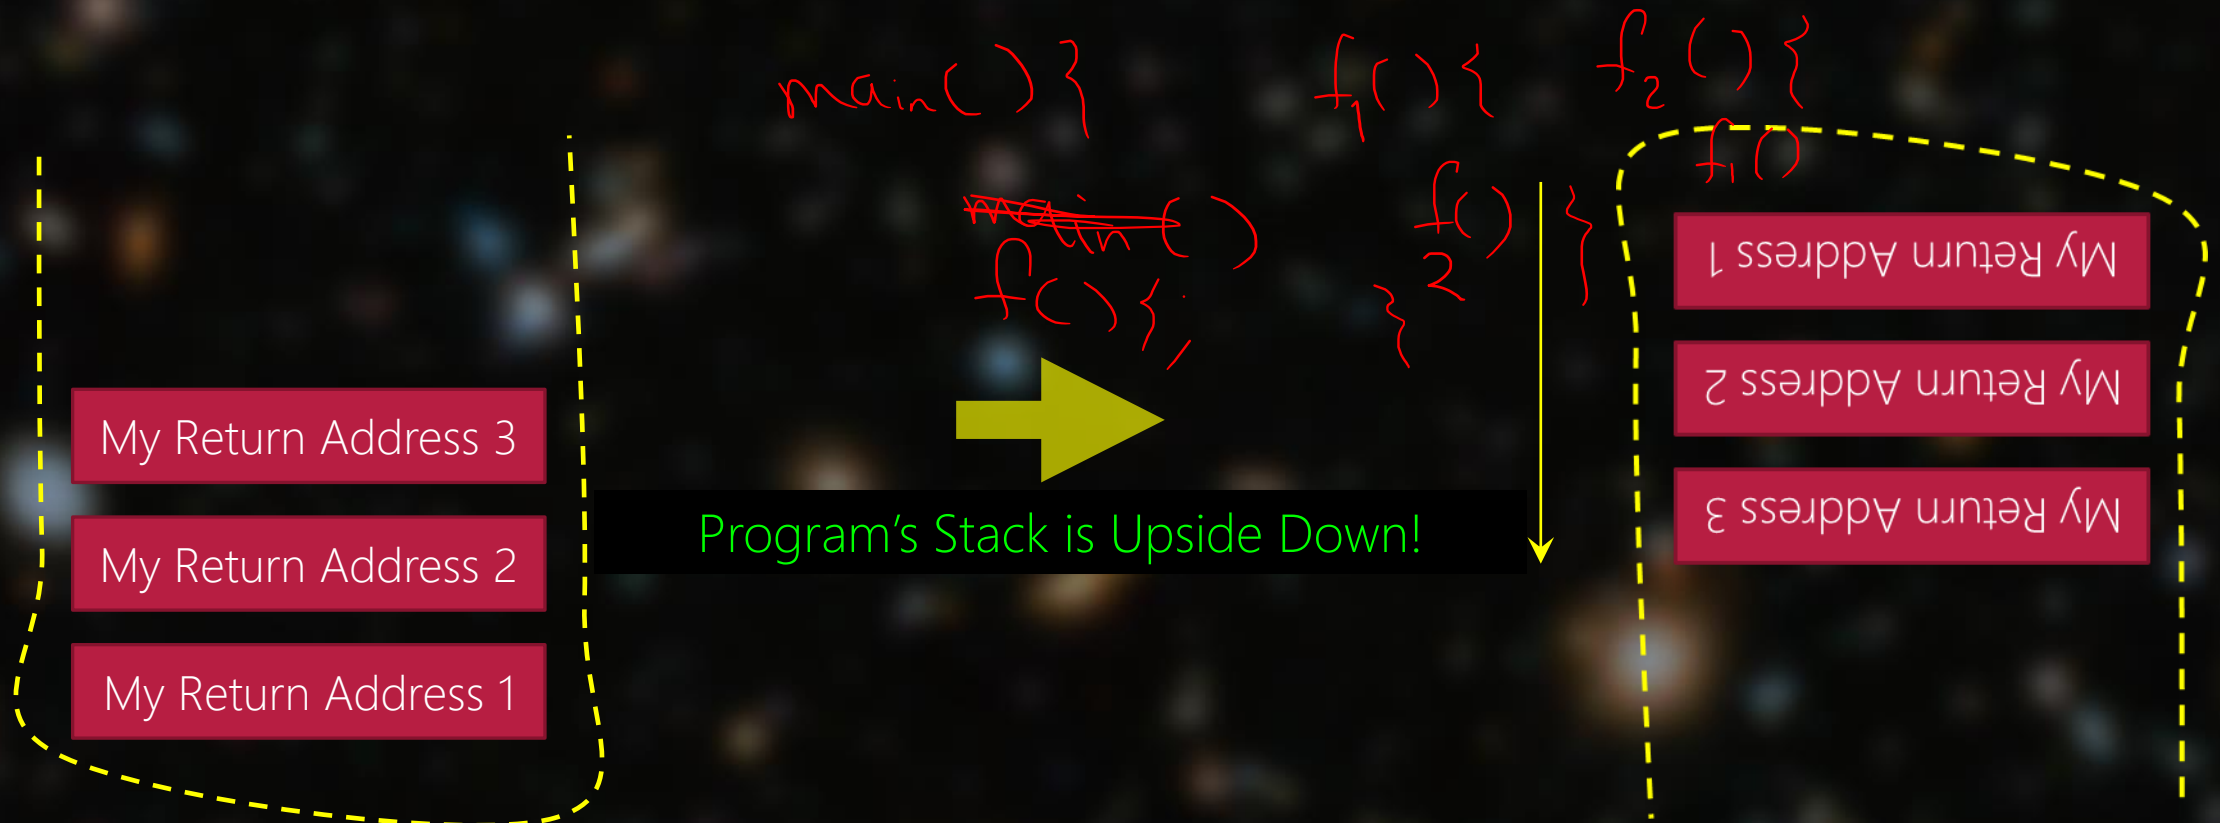
Functions Arguments, Local Variables, Return Address (runtime)

Program's Stack is Upside Down!



# Stack Overflow?

Functions Arguments, Local Variables, Return Address (runtime)





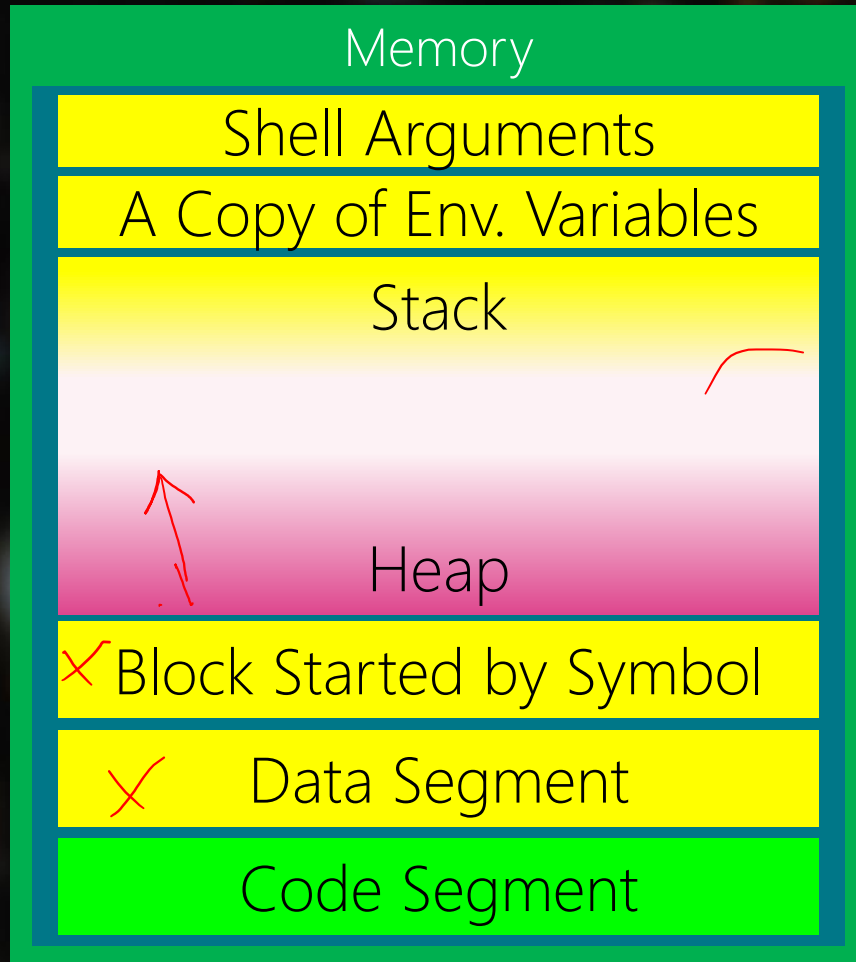


Into the Wild (2007) - Sean Penn



# Heap

Dynamic memory allocation (runtime)



## Memory Allocators by Library Routines

```
#include <stdlib.h>
void *malloc(size_t size)
void *realloc(void *ptr, size_t newsize)
```

Size is fixed during compile time  
Value is dynamic during runtime

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

*Handwritten notes:*  
-  $\text{printf}(\&a)$   
-  $\text{int}^* a$   
-  $\underline{a} = \text{malloc}(\text{int})$

```
hfani@charlie:~$ ./main_add 2 2
2 + 2 = 4
hfani@charlie:~$ ./main_add 2 4
2 + 4 = 6
hfani@charlie:~$
```

Size is dynamic during runtime  
Value is dynamic during runtime

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]){
    int size_a = 0;
    int size_b = 0;
    size_a = atoi(argv[1]);
    size_b = atoi(argv[2]);

    int *a = malloc(size_a * sizeof(int));
    printf("enter the first number with %d digits:\n", size_a);
    for(int i = 0; i < size_a; ++i){
        scanf("%d", a + i);
    }

    int *b = malloc(size_b * sizeof(int));
    printf("enter the first number with %d digits:\n", size_b);
    for(int i = 0; i < size_b; ++i){
        scanf("%d", b + i);
    }
```

Handwritten notes:

- $\text{int} \Rightarrow 4 \text{ Bytes}$
- $\text{int}^* a$ ;  $a = \text{FF0A}$
- $(*a) = 2$
- 3 (circled, next to `size_a`)
- 4 (next to `size_b`)
- 32 (next to `sizeof(int)`)
- 2 (next to `scanf`)
- 31 (circled, next to `scanf`)
- 2 -1 (circled, next to `scanf`)

```
hfani@charlie:~$ ./main_malloc 3 4
enter the first number with 3 digits:
1
3
9
enter the first number with 4 digits:
6
5
7
2
139 + 6572
```

Handwritten notes:

- $\text{int}^*$

---

Size is **dynamic** during **runtime**  
Value is **dynamic** during **runtime**

---

```
hfani@charlie:~$ ./main_malloc 1000000000000000 10000000000000000
```

What happens?

Size is **dynamic** during **runtime**  
Value is **dynamic** during **runtime**

```
hfani@charlie:~$ ./main_malloc 1000000000000000 10000000000000000
```





---

# Heap

Dynamic memory allocation (runtime)

---

## Memory Allocators by Library Routines

```
#include <stdlib.h>
void *malloc(size_t size)
void *realloc(void *ptr, size_t newsize)
```

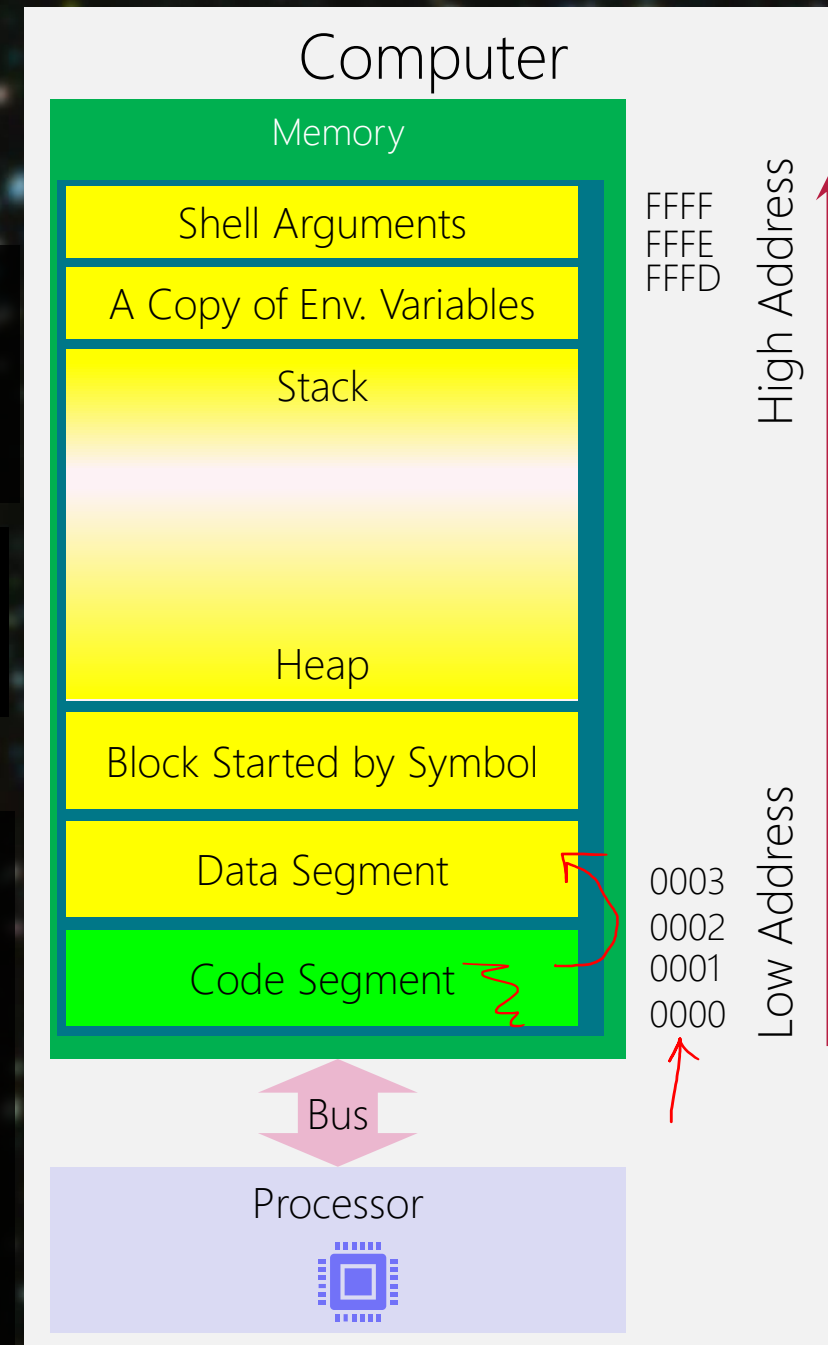
## Memory Allocators by System Calls?

# Shell's `size` command

```
hfani@charlie:~$ size ./main_malloc
text      data      bss      dec      hex filename
2239      616         8     2863     b2f ./main_malloc
```

Is this info for:

- { Compile time?
- { Runtime?



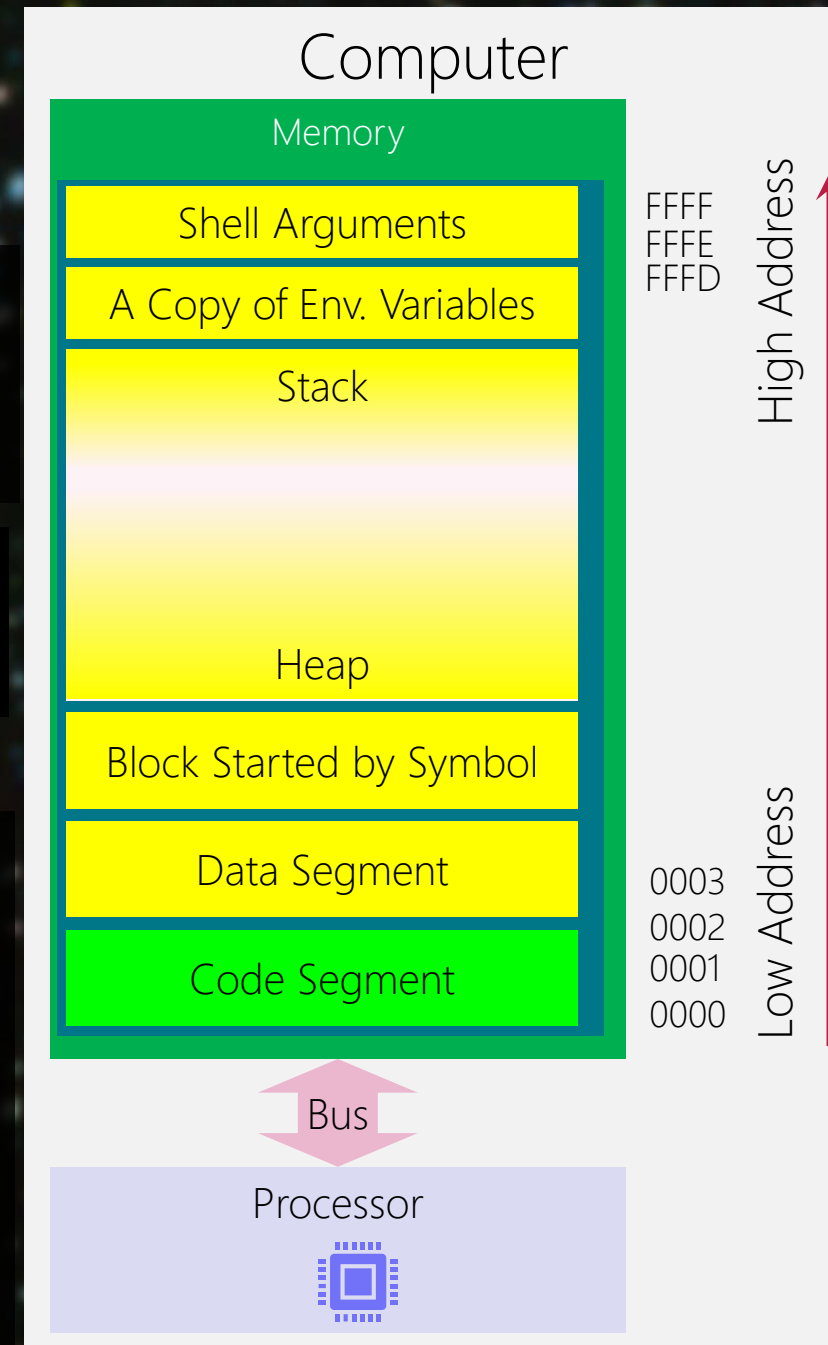
# Shell's `size` command

```
hfani@charlie:~$ size ./main_malloc
text    data     bss      dec       hex filename
2239     616         8    2863     b2f ./main_malloc
```

Why is not any info for:

- Stack?
- Heap?

Profiler





---

## Process ~~Identifier~~ (pid)

---

Non-negative

Unique among **processes** (live programs)

Not an identifier! It can be reused (delay reuse)

# Process Identifier by System Call `getpid()`

```
#include <unistd.h>
pid_t getpid(void);
```

Return process ID of calling process

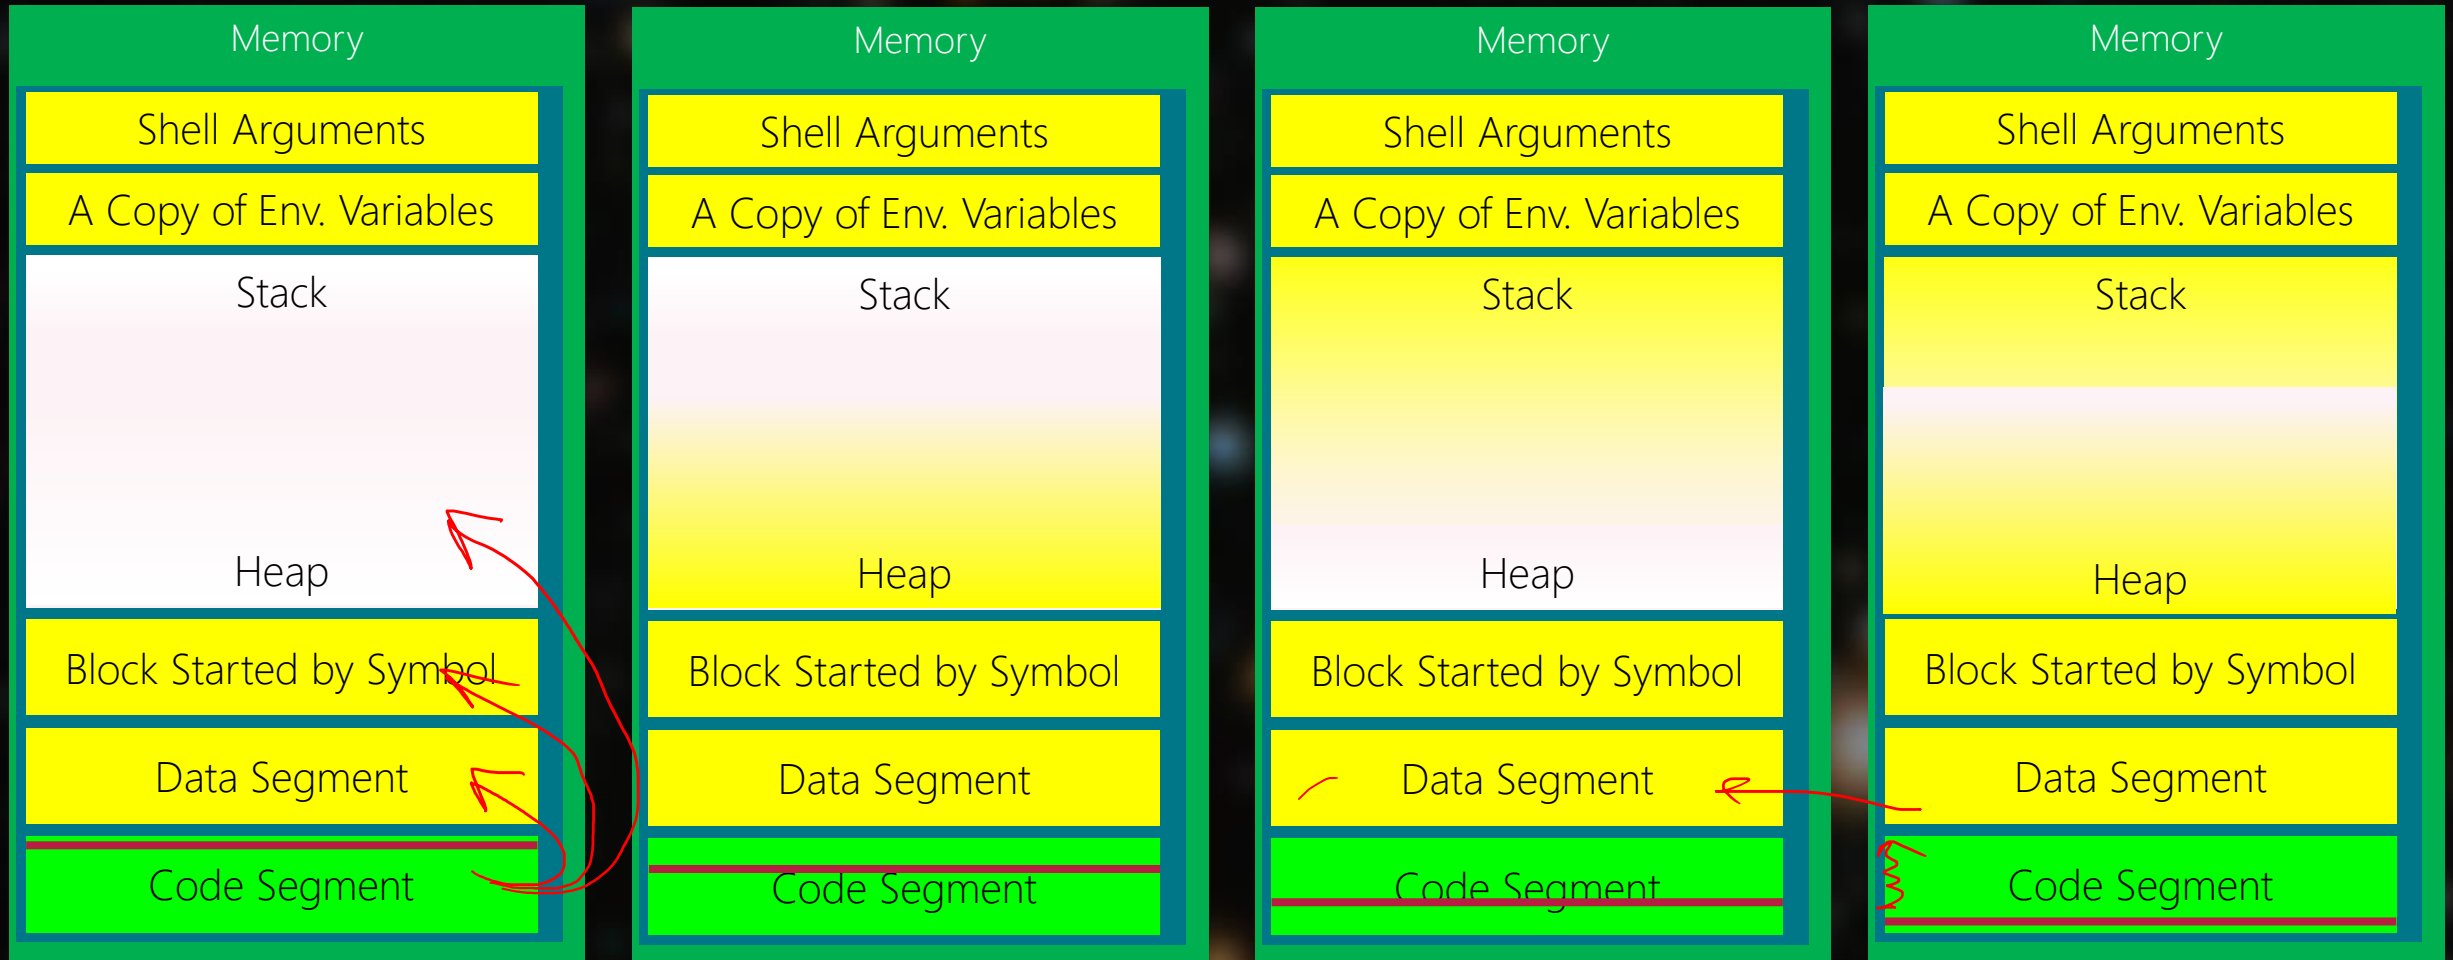
```
#include <unistd.h>
#include <stdio.h>
int main(void){
    printf("%d\n", getpid());
    return 0;
}
```

```
hfani@alpha:~$ ./getpid
871198
hfani@alpha:~$ ./getpid
871217
```

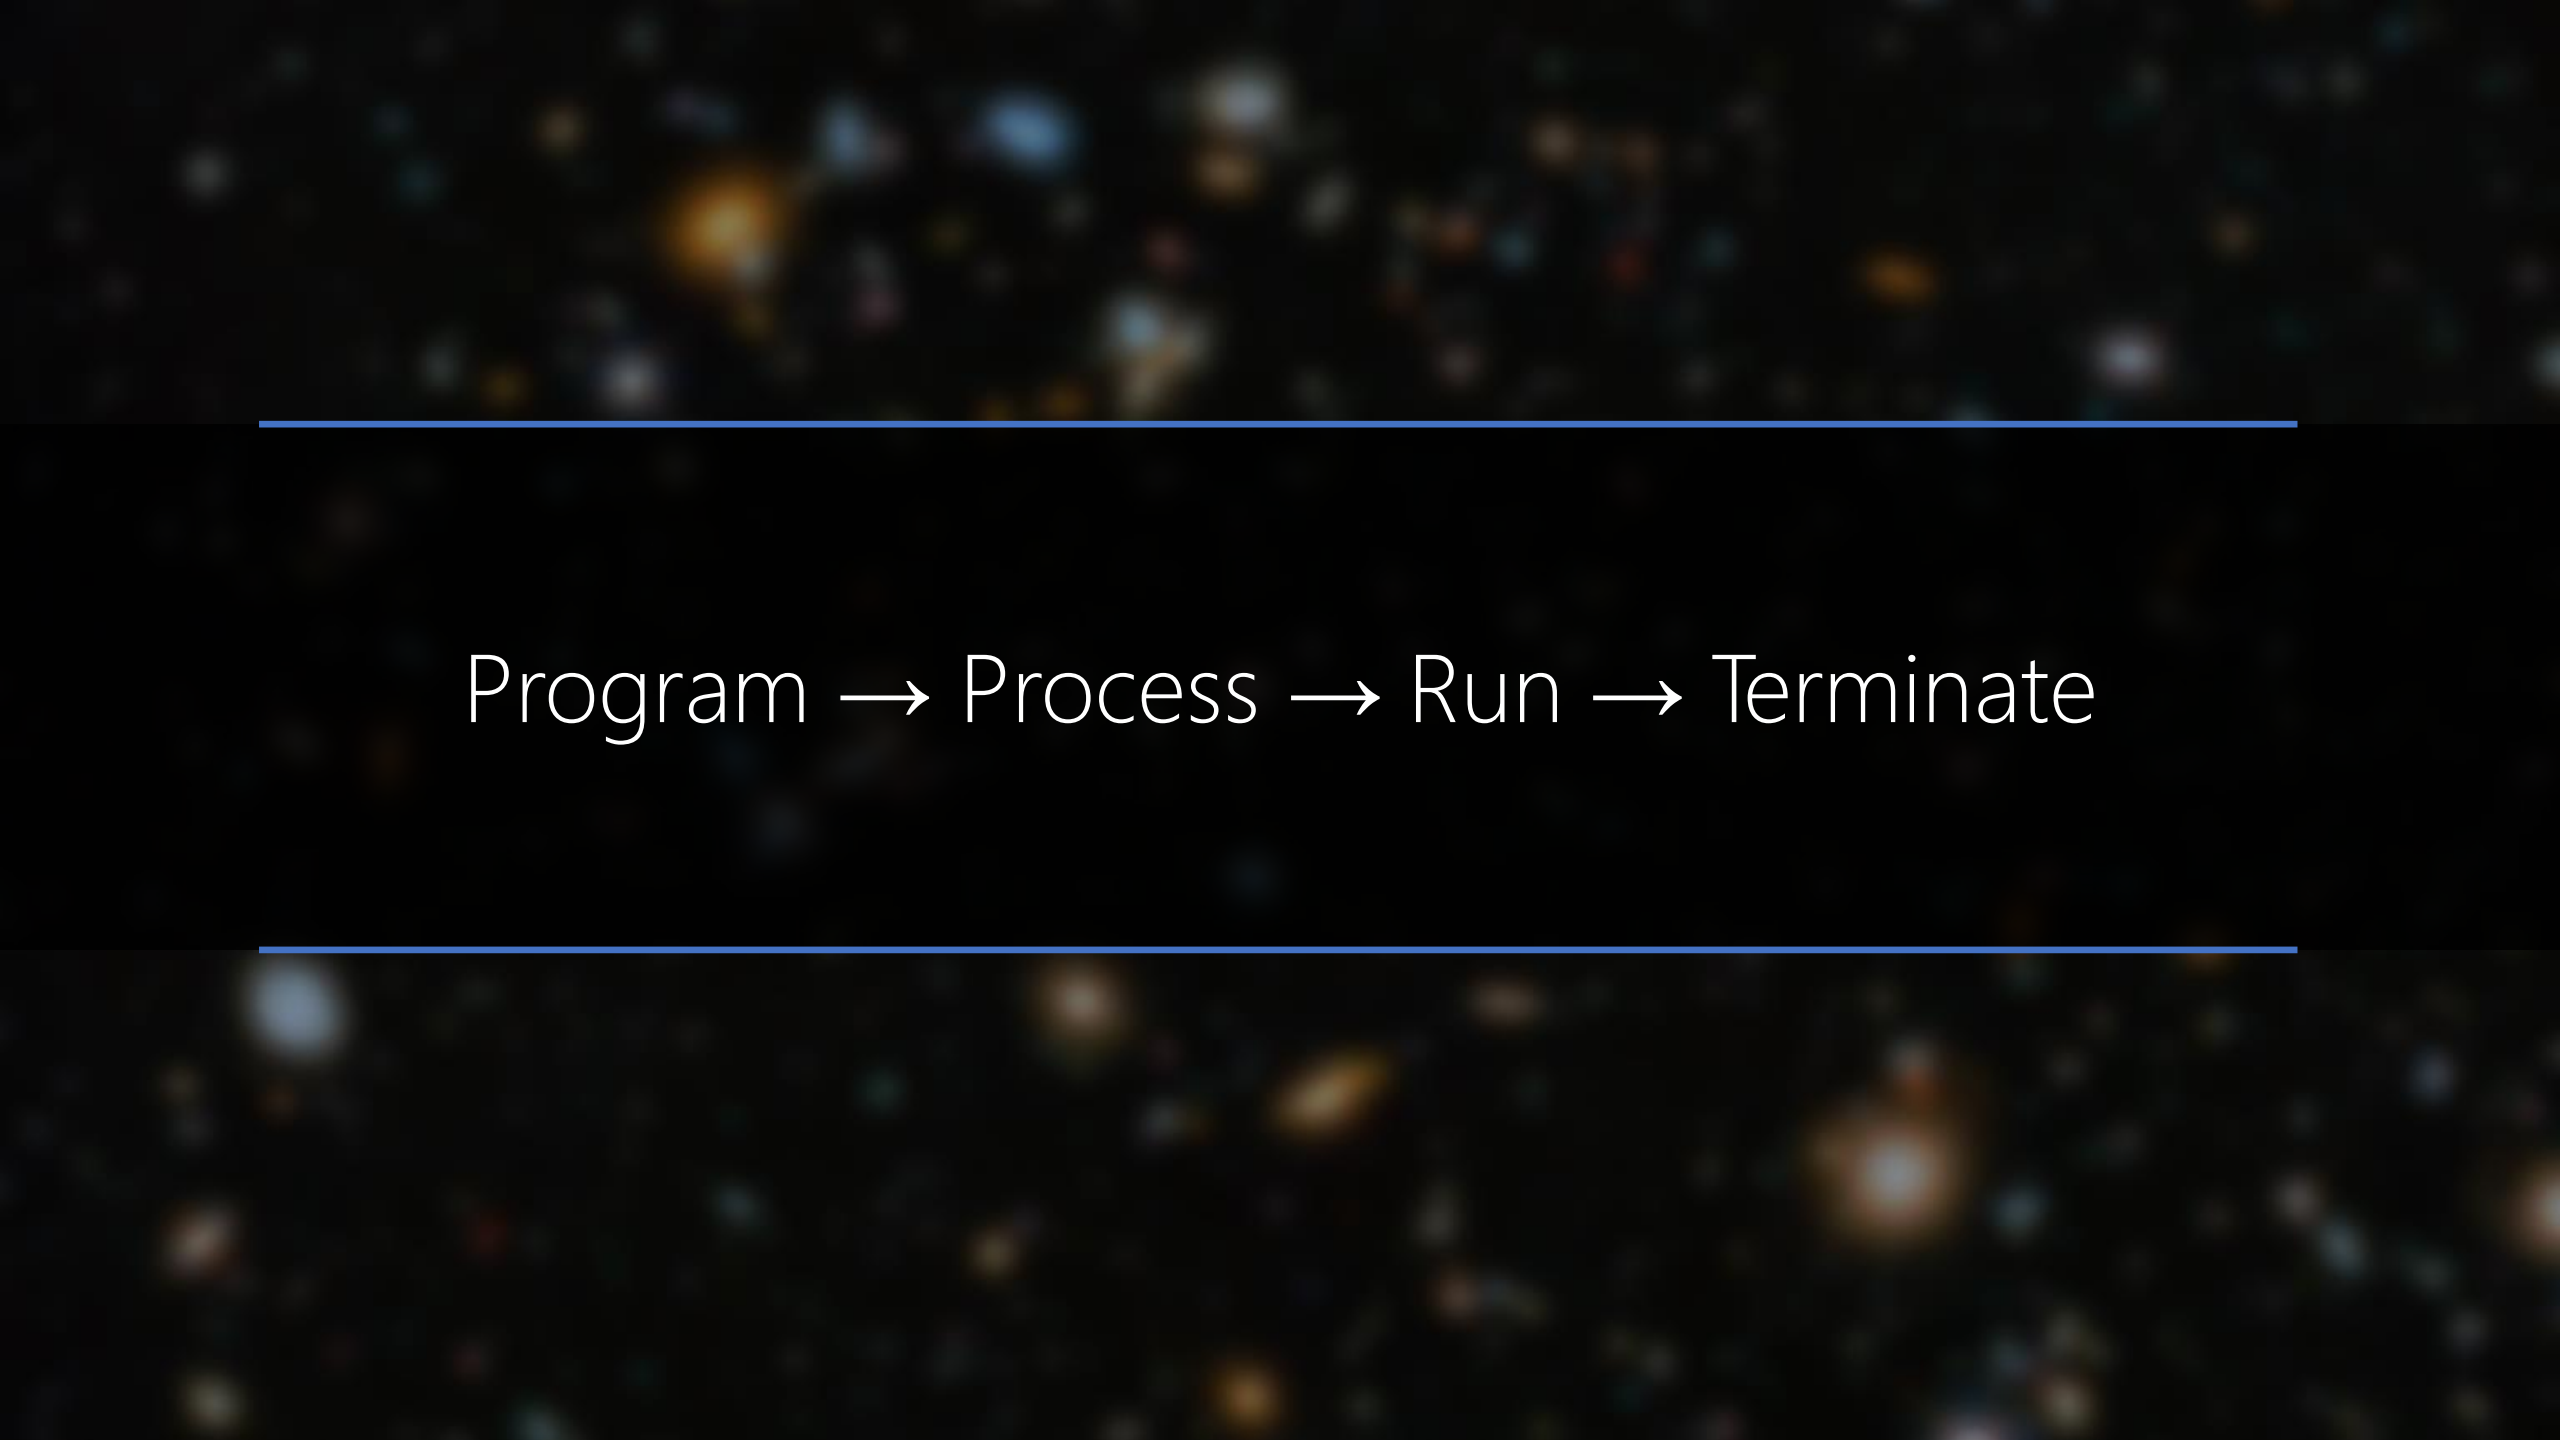


Into the Wild (2007) - Sean Penn

Program → Process → Run → Terminate







---

Program → Process → Run → Terminate

---

---

C does not have error/exception handling!

---



---

C has exit status (code)

Normal vs. Abnormal Exits

---

# C has exit status (code)

Normal

```
void main(void){  
    ///lines of codes  
}
```

```
void main(void){  
    ///lines of codes  
    return;  
}
```

```
int main(void){  
    ///lines of codes  
    return 0;  
}
```

```
#include <unistd.h>  
int main(void){  
    ///lines of codes  
    _exit(0);  
}
```

```
#include <stdlib.h>  
int main(void){  
    ///lines of codes  
    exit(0);  
}
```

```
#include <stdlib.h>  
int main(void){  
    ///lines of codes  
    exit(EXIT_SUCCESS);  
}
```



---

C has exit status (code)

---

Normal

Clean up procedure

- Flushes unwritten buffered data.
- Closes all open file descriptors.
- Returns an integer exit status to the operating system.



---

C has exit status (code)

Abnormal

- 
- Any non-zero number less than 256
  - Receiving a SIGNAL  
e.g., SIGABRT raised by `abort()`

# C has exit status (code)

## Abnormal

```
hfani@charlie:~$ kill -l
```

1) <u>SIGHUP</u>	2) SIGINT	3) SIGQUIT	4) SIGILL	5) SIGTRAP
6) <u>SIGABRT</u>	7) <u>SIGBUS</u>	8) SIGFPE	9) <u>SIGKILL</u>	10) SIGUSR1
11) <u>SIGSEGV</u>	12) SIGUSR2	13) SIGPIPE	14) <u>SIGALRM</u>	15) SIGTERM
16) SIGSTKFLT	17) SIGCHLD	18) SIGCONT	19) SIGSTOP	20) SIGTSTP
21) SIGTTIN	22) SIGTTOU	23) SIGURG	24) SIGXCPU	25) SIGXFSZ
26) SIGVTALRM	27) SIGPROF	28) SIGWINCH	29) SIGIO	30) SIGPWR
31) SIGSYS	34) SIGRTMIN	35) SIGRTMIN+1	36) SIGRTMIN+2	37) SIGRTMIN+3
38) SIGRTMIN+4	39) SIGRTMIN+5	40) SIGRTMIN+6	41) SIGRTMIN+7	42) SIGRTMIN+8
43) SIGRTMIN+9	44) SIGRTMIN+10	45) SIGRTMIN+11	46) SIGRTMIN+12	47) SIGRTMIN+13
48) SIGRTMIN+14	49) SIGRTMIN+15	50) SIGRTMAX-14	51) SIGRTMAX-13	52) SIGRTMAX-12
53) SIGRTMAX-11	54) SIGRTMAX-10	55) SIGRTMAX-9	56) SIGRTMAX-8	57) SIGRTMAX-7
58) SIGRTMAX-6	59) SIGRTMAX-5	60) SIGRTMAX-4	61) SIGRTMAX-3	62) SIGRTMAX-2
63) SIGRTMAX-1	64) <u>SIGRTMAX</u>			

## Shell's Variable for Exit Status

echo \$?

```
hfani@charlie:~$ ./main_exit_normal_2
hfani@charlie:~$ echo $?
0
```

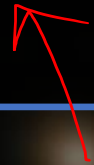
```
hfani@charlie:~$ ./main_malloc (2) (3)
enter the first number with 2 digits:
^C
hfani@charlie:~$ echo $?
130
```



---

Suicide

`_exit(1), hang up!`



`echo $?`

---



---

Suicide

`_exit(1), hang up!`

---

```
#include <fcntl.h>
#include <unistd.h>
```

Library Routines!

```
#include <stdio.h>
#include <errno.h>
#include <string.h>
```

```
void main(void){
```

```
    int fd;
```

```
    do{
```

```
        fd = open("./not_exist/test.txt", O_WRONLY | O_CREAT, S_IWUSR);
```

```
        if (fd == -1){
```

```
            printf("kernel rejected to open the file!\n");
```

```
            printf("the error is %d\n and the error string is %s\n", errno, strerror(errno));
```

```
            if(errno == 2){
```

```
                _exit(1);
```

```
            }
```

```
        }
```

```
    }while(fd > -1);
```

```
hfani@alpha:~/comp2560_f2022$ ./suicide
```

```
kernel rejected to open the file!
```

```
the error is 2
```

```
and the error string is 'No such file or directory'
```

```
hfani@alpha:~/comp2560_f2022$ echo $?
```

```
1
```

## Thread: Error-Reporting in Unix (strerror(errno))



Refresh



Alignments



Search

Select: [All](#) [None](#)

Thread 12 of 13



Message Actions ▼

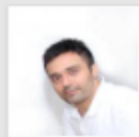
Expand All

Collapse All

2 Post(s) in this Thread

0 Unread

0 Unread Replies to Me

**Hossein Fani** 🟢

2 days ago

**Error-Reporting in Unix (strerror(errno))**

Overall Rating: ★★★★★

Dear Students,

In class, there was a question about how to know more about the reason(s) when kernel rejects our system calls like `creat()` or `open()` a file or device.

Part of the library routine in `<stdlib.h>`, the C standard library, which is part of posix, we have a function `strerror()` that accepts a variable called `errno`. The `errno` is a number that shows the reason. In order to see the explanation, in `<errno.h>` `errno` is mapped to string value.

You can find an example of this error-reporting mechanism here:

**C library function - `strerror()`**

C library function - `strerror()`, The C library function `char *strerror(int errnum)` searches an internal array for the error number `errnum` and returns a pointer to an error message string. The e

Next round of questions are:

1) `errno` is a variable inside kernel or inside `stdlib.h` header file?





Into the Wild (2007) - Sean Penn