

Into the Wild

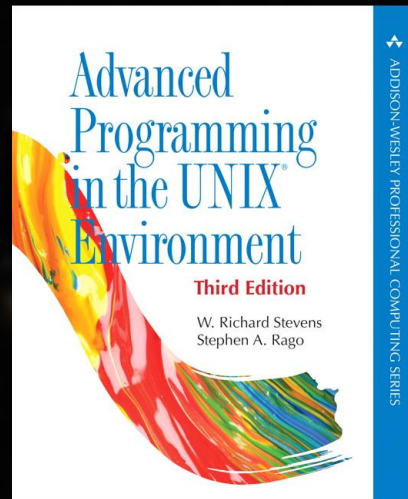


Into the Wild (2007) - Sean Penn
Eddie Vedder - Hard Sun

<https://www.youtube.com/watch?v=Ez8b2VHjVB0>



Lab08, Lec08





Chapter 07: Process Environment Chapter 08: Process Control



Process Manager

aka. Process Control



Computer

Memory

Kernel: Device Manager

Kernel: Memory Manager

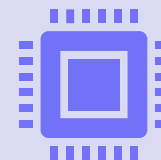
Kernel: File Manager

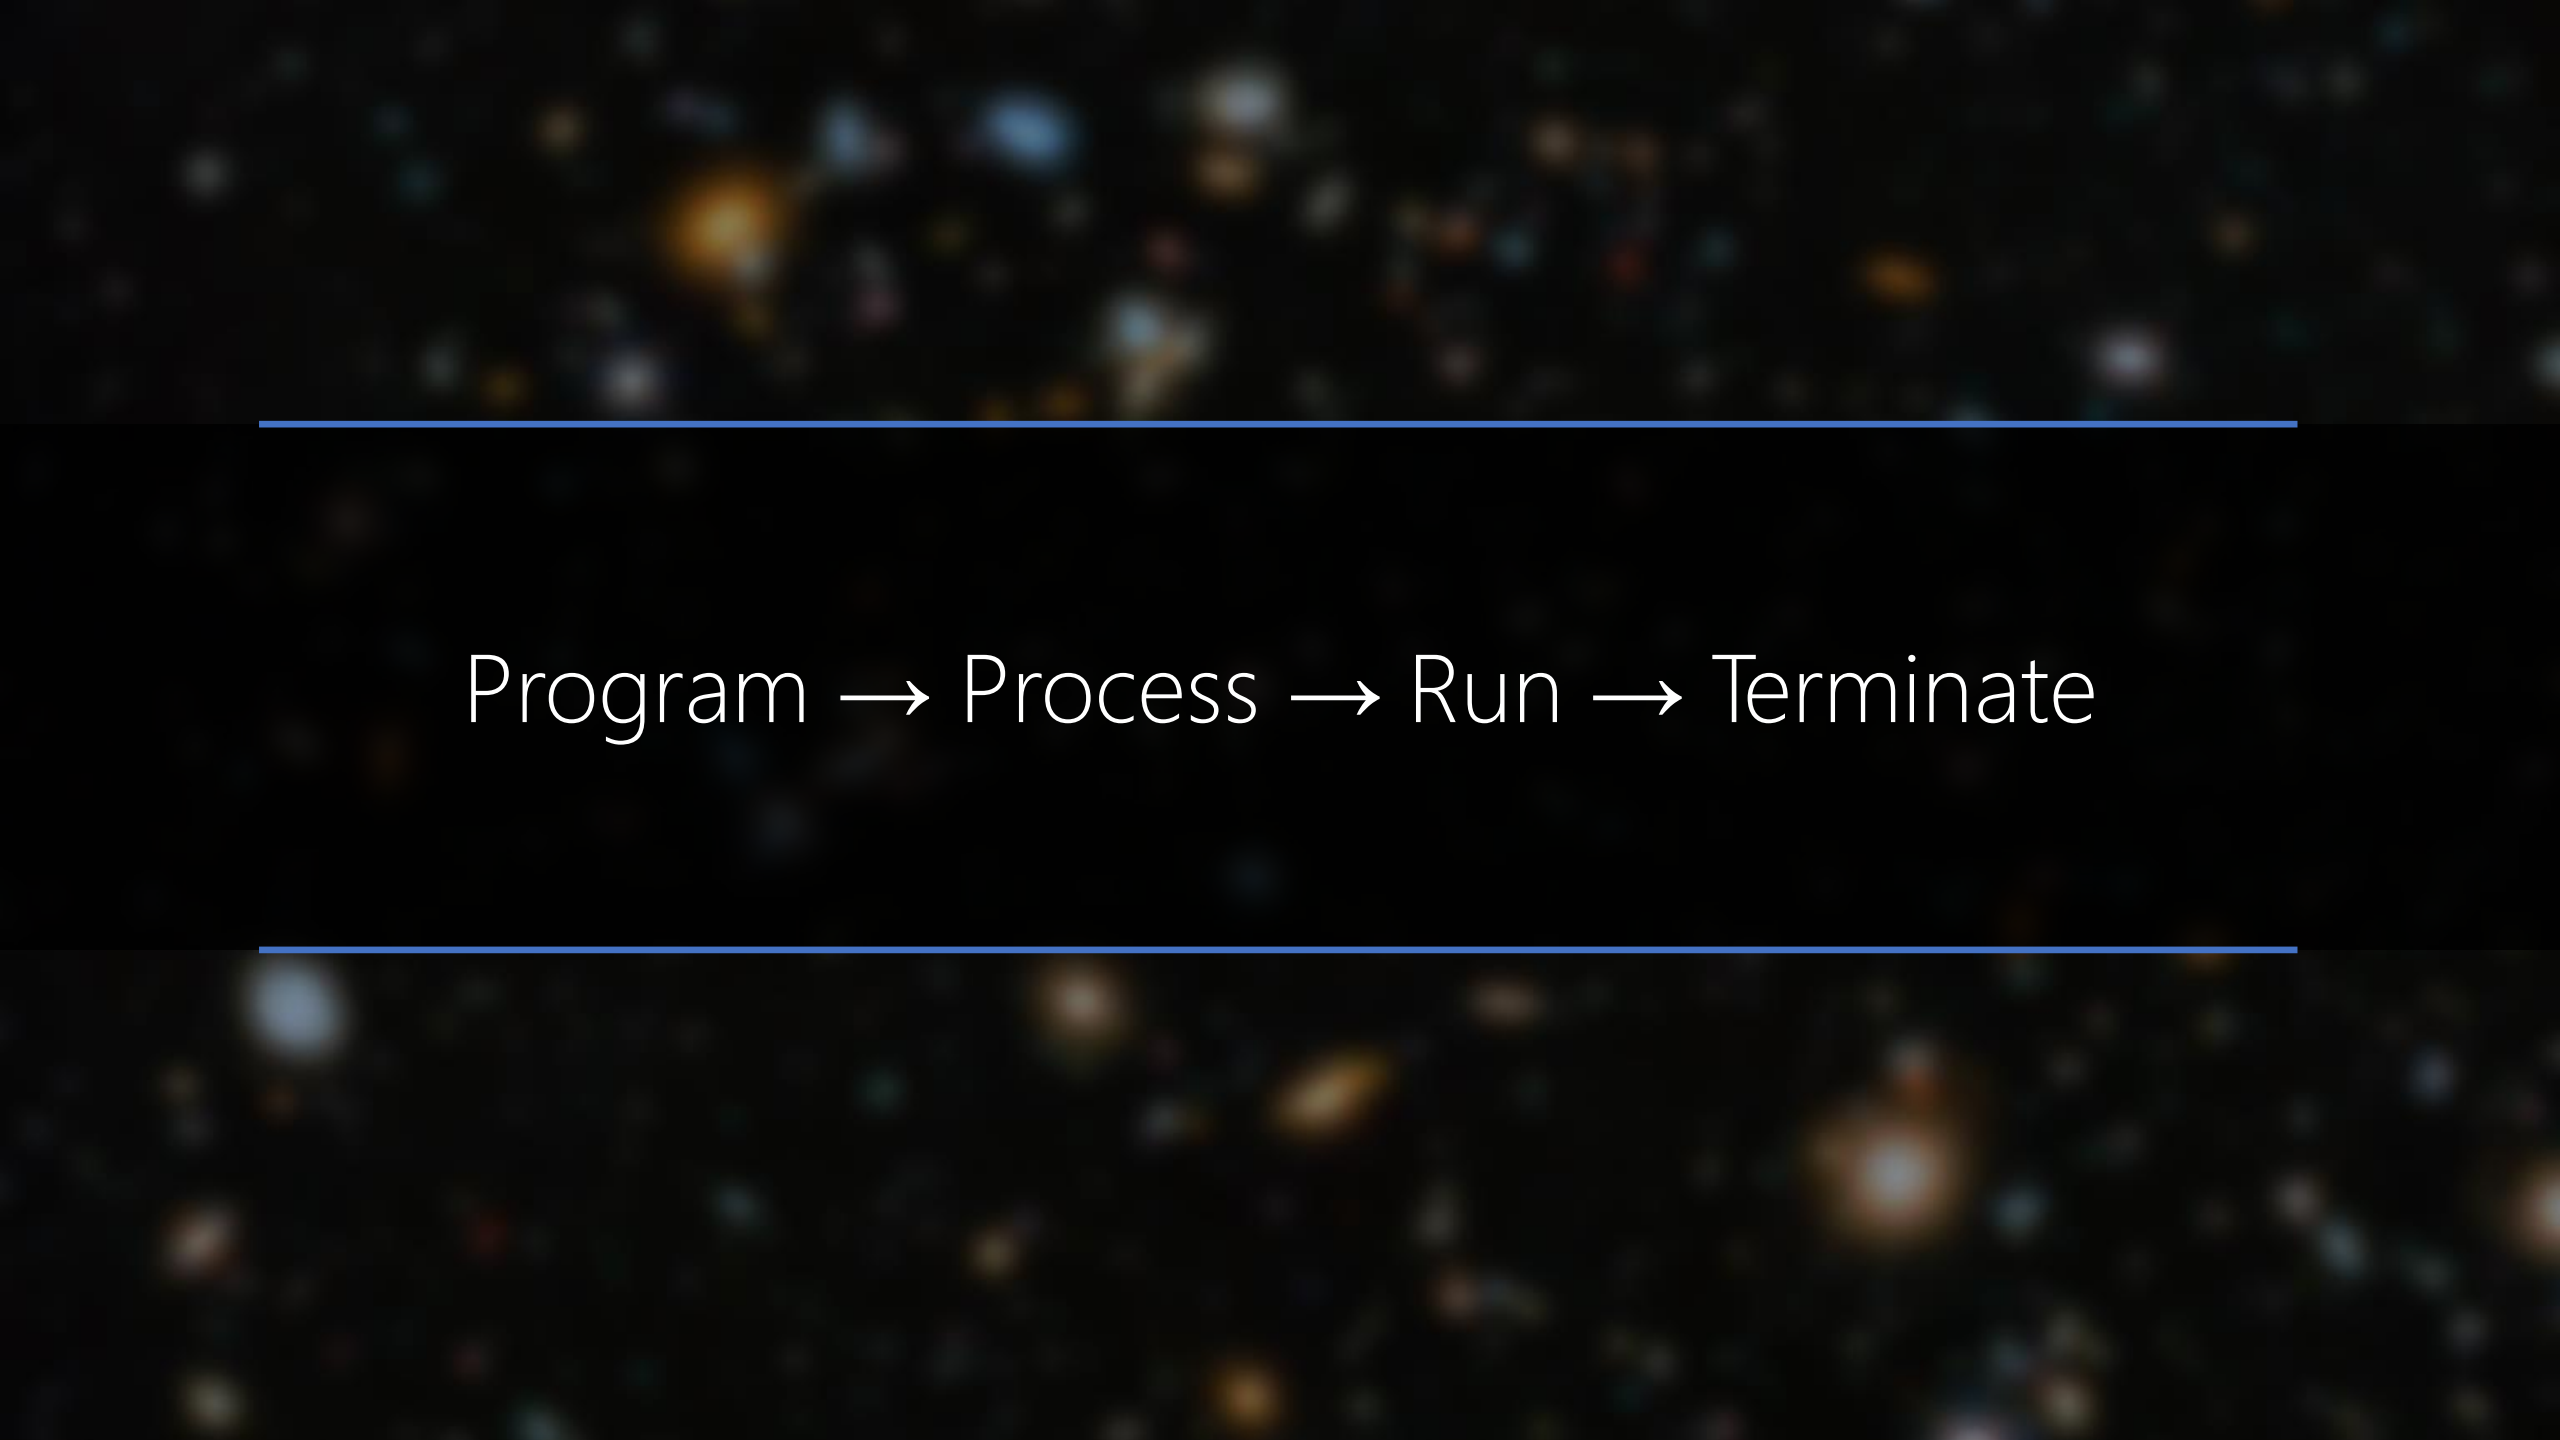
Kernel: Network Manager

Kernel: Process Manager

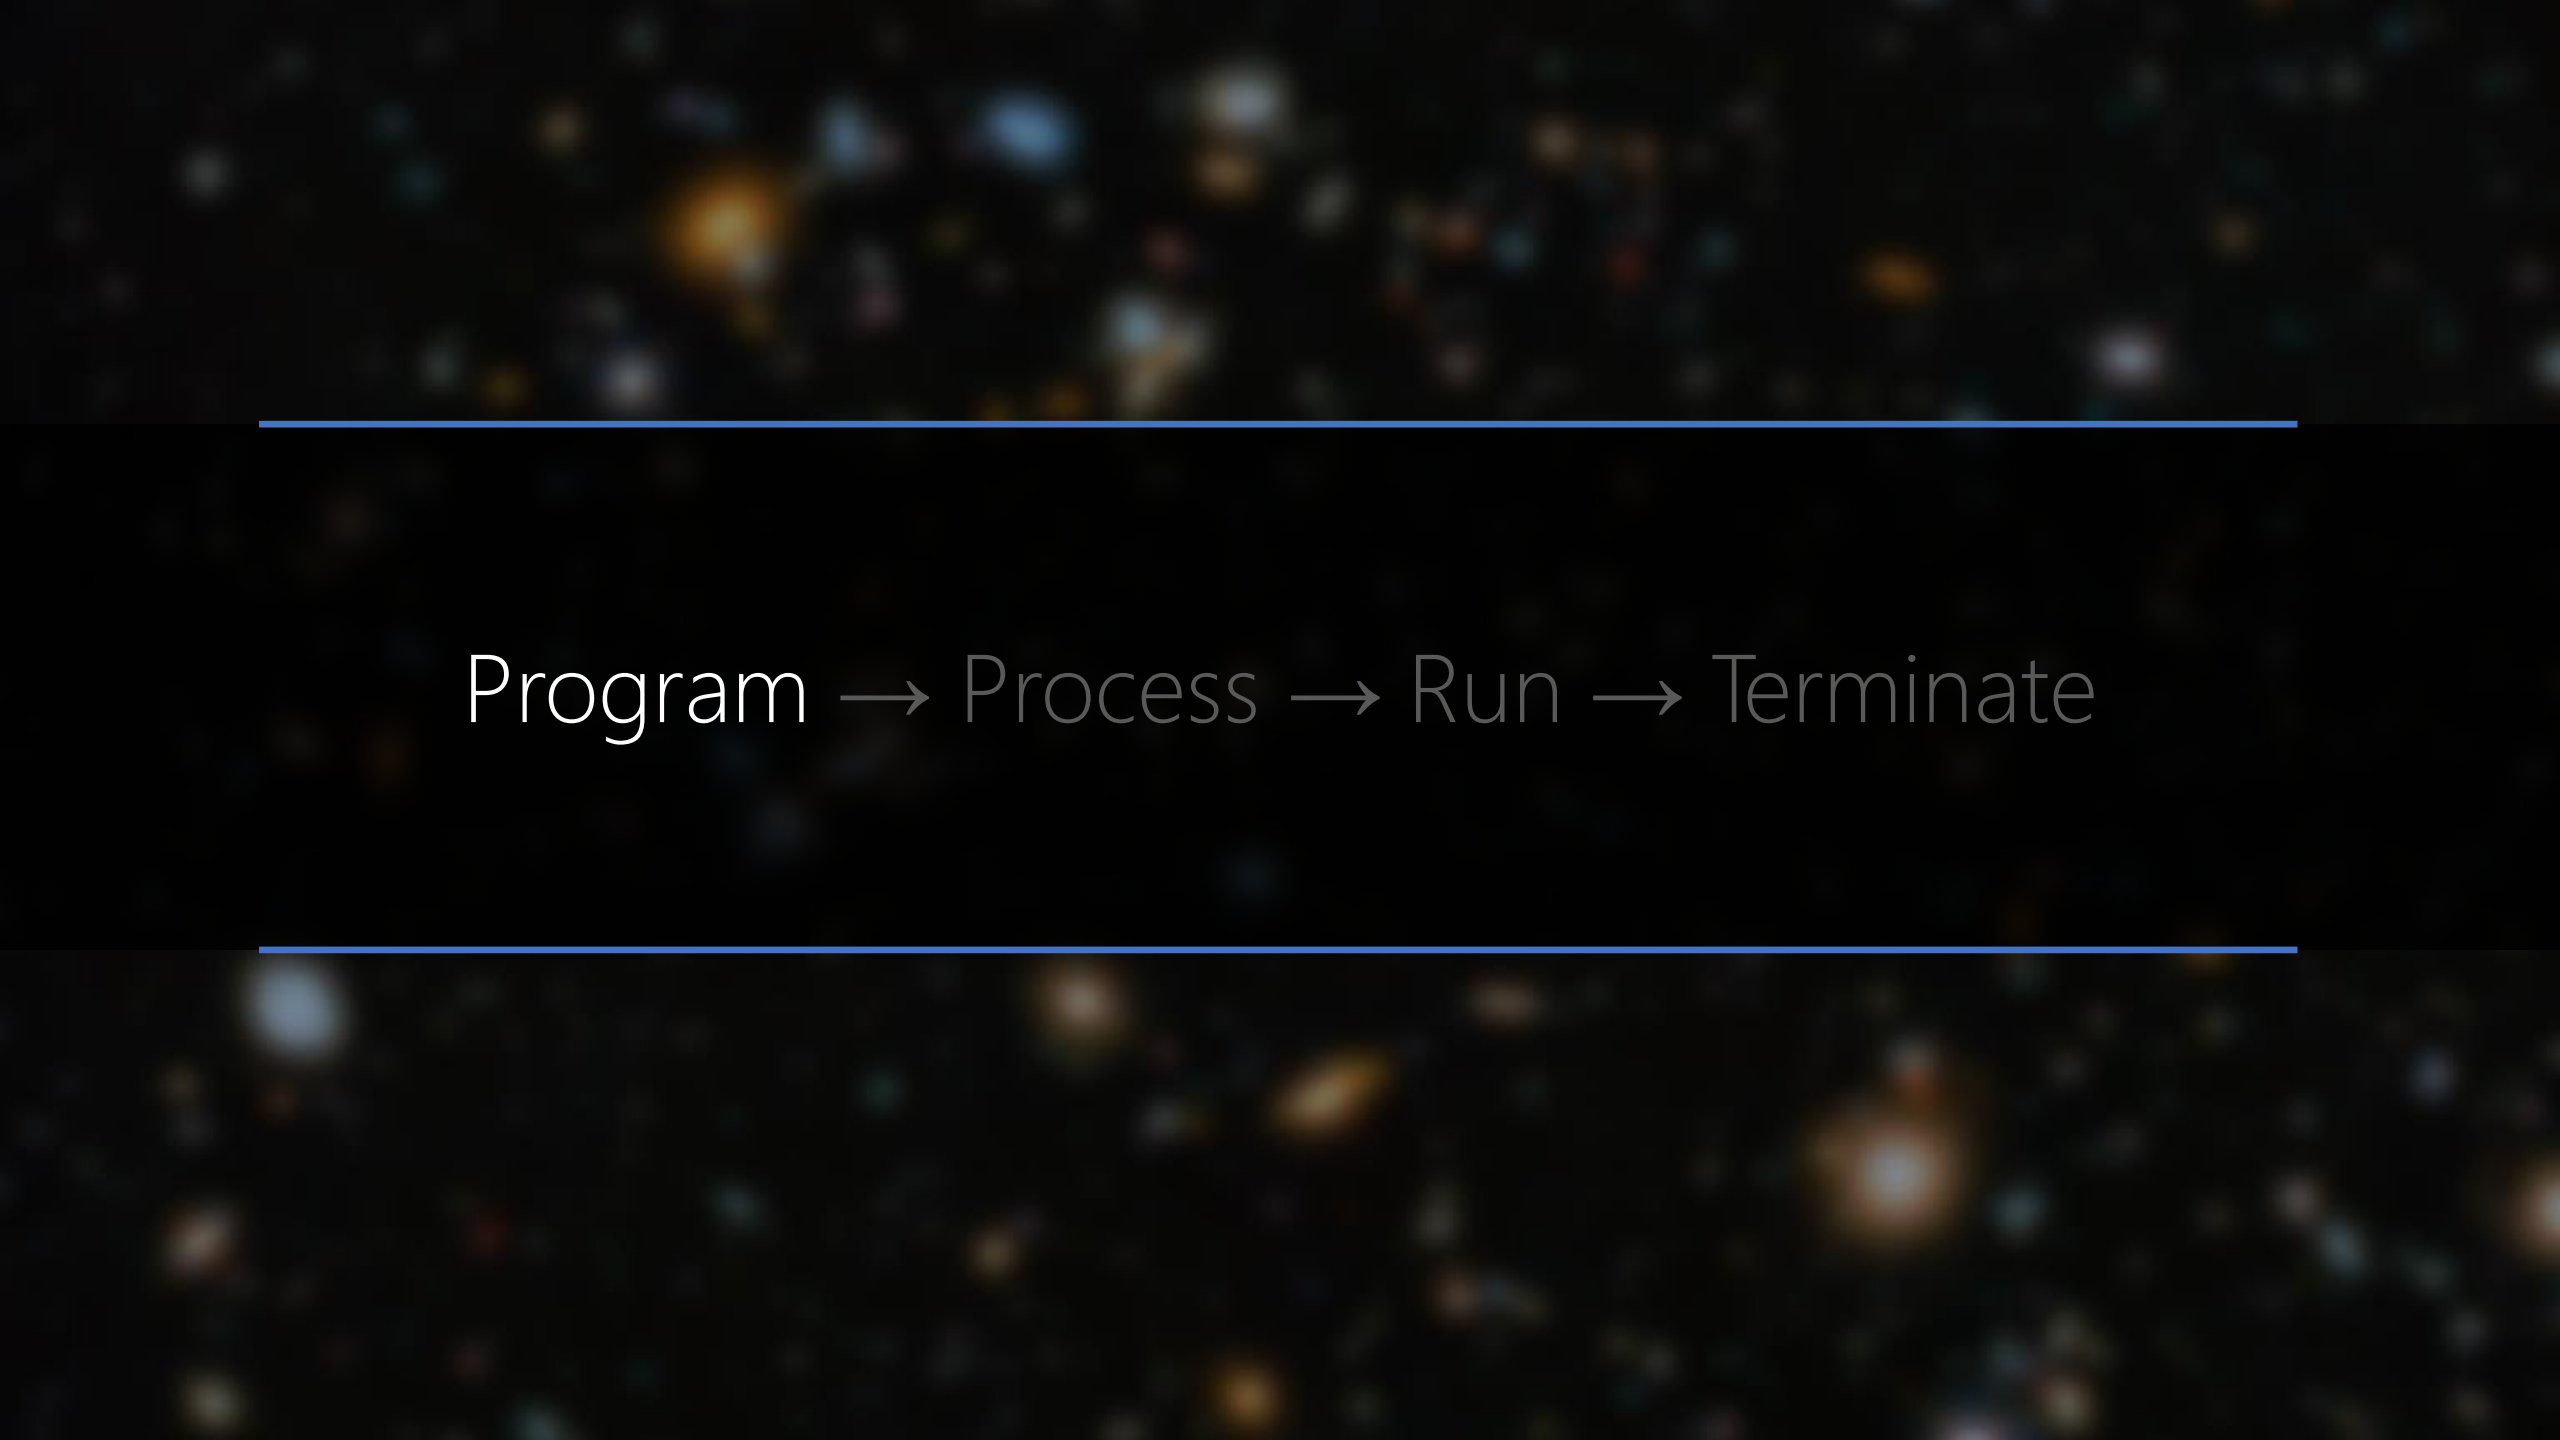
Bus

Processor





Program → Process → Run → Terminate



Program → Process → Run → Terminate



Any Program MUST have an entry point

What part of the code has the first opcode?

```
void main(void)
```

```
shell$ ./program
```

```
void main(int argc, char *argv[])  
int  main(int argc, char *argv[])
```

```
shell$ ./program arg1 arg2 arg3 ....
```

```
hfani@charlie:~$ vi main_args.c
#include <stdio.h>
int main(int argc, char *argv[]){
    printf("there are %d arguments in the shell:\n", argc);
    for(int i=0; i < argc; ++i){
        printf("arg%d: %s\n", i, argv[i]);
    }
    return 0;
}
```

Name of the program file is the first argument!

```
hfani@charlie:~$ cc main_args.c -o main_args
hfani@charlie:~$ ./main_args
there are 1 arguments in the shell:
arg0: ./main_args
hfani@charlie:~$ ./main_args param1 param2
there are 3 arguments in the shell:
arg0: ./main_args
arg1: param1
arg2: param2
hfani@charlie:~$
```

```
hfani@charlie:~$ vi main_add.c
```

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

Arguments are string of chars!

ASCII to Integer

int atoi (const char *str);

```
hfani@charlie:~$ cc main_add.c -o main_add
```

```
hfani@charlie:~$ ./main_add 2 2
```

```
2 + 2 = 4
```

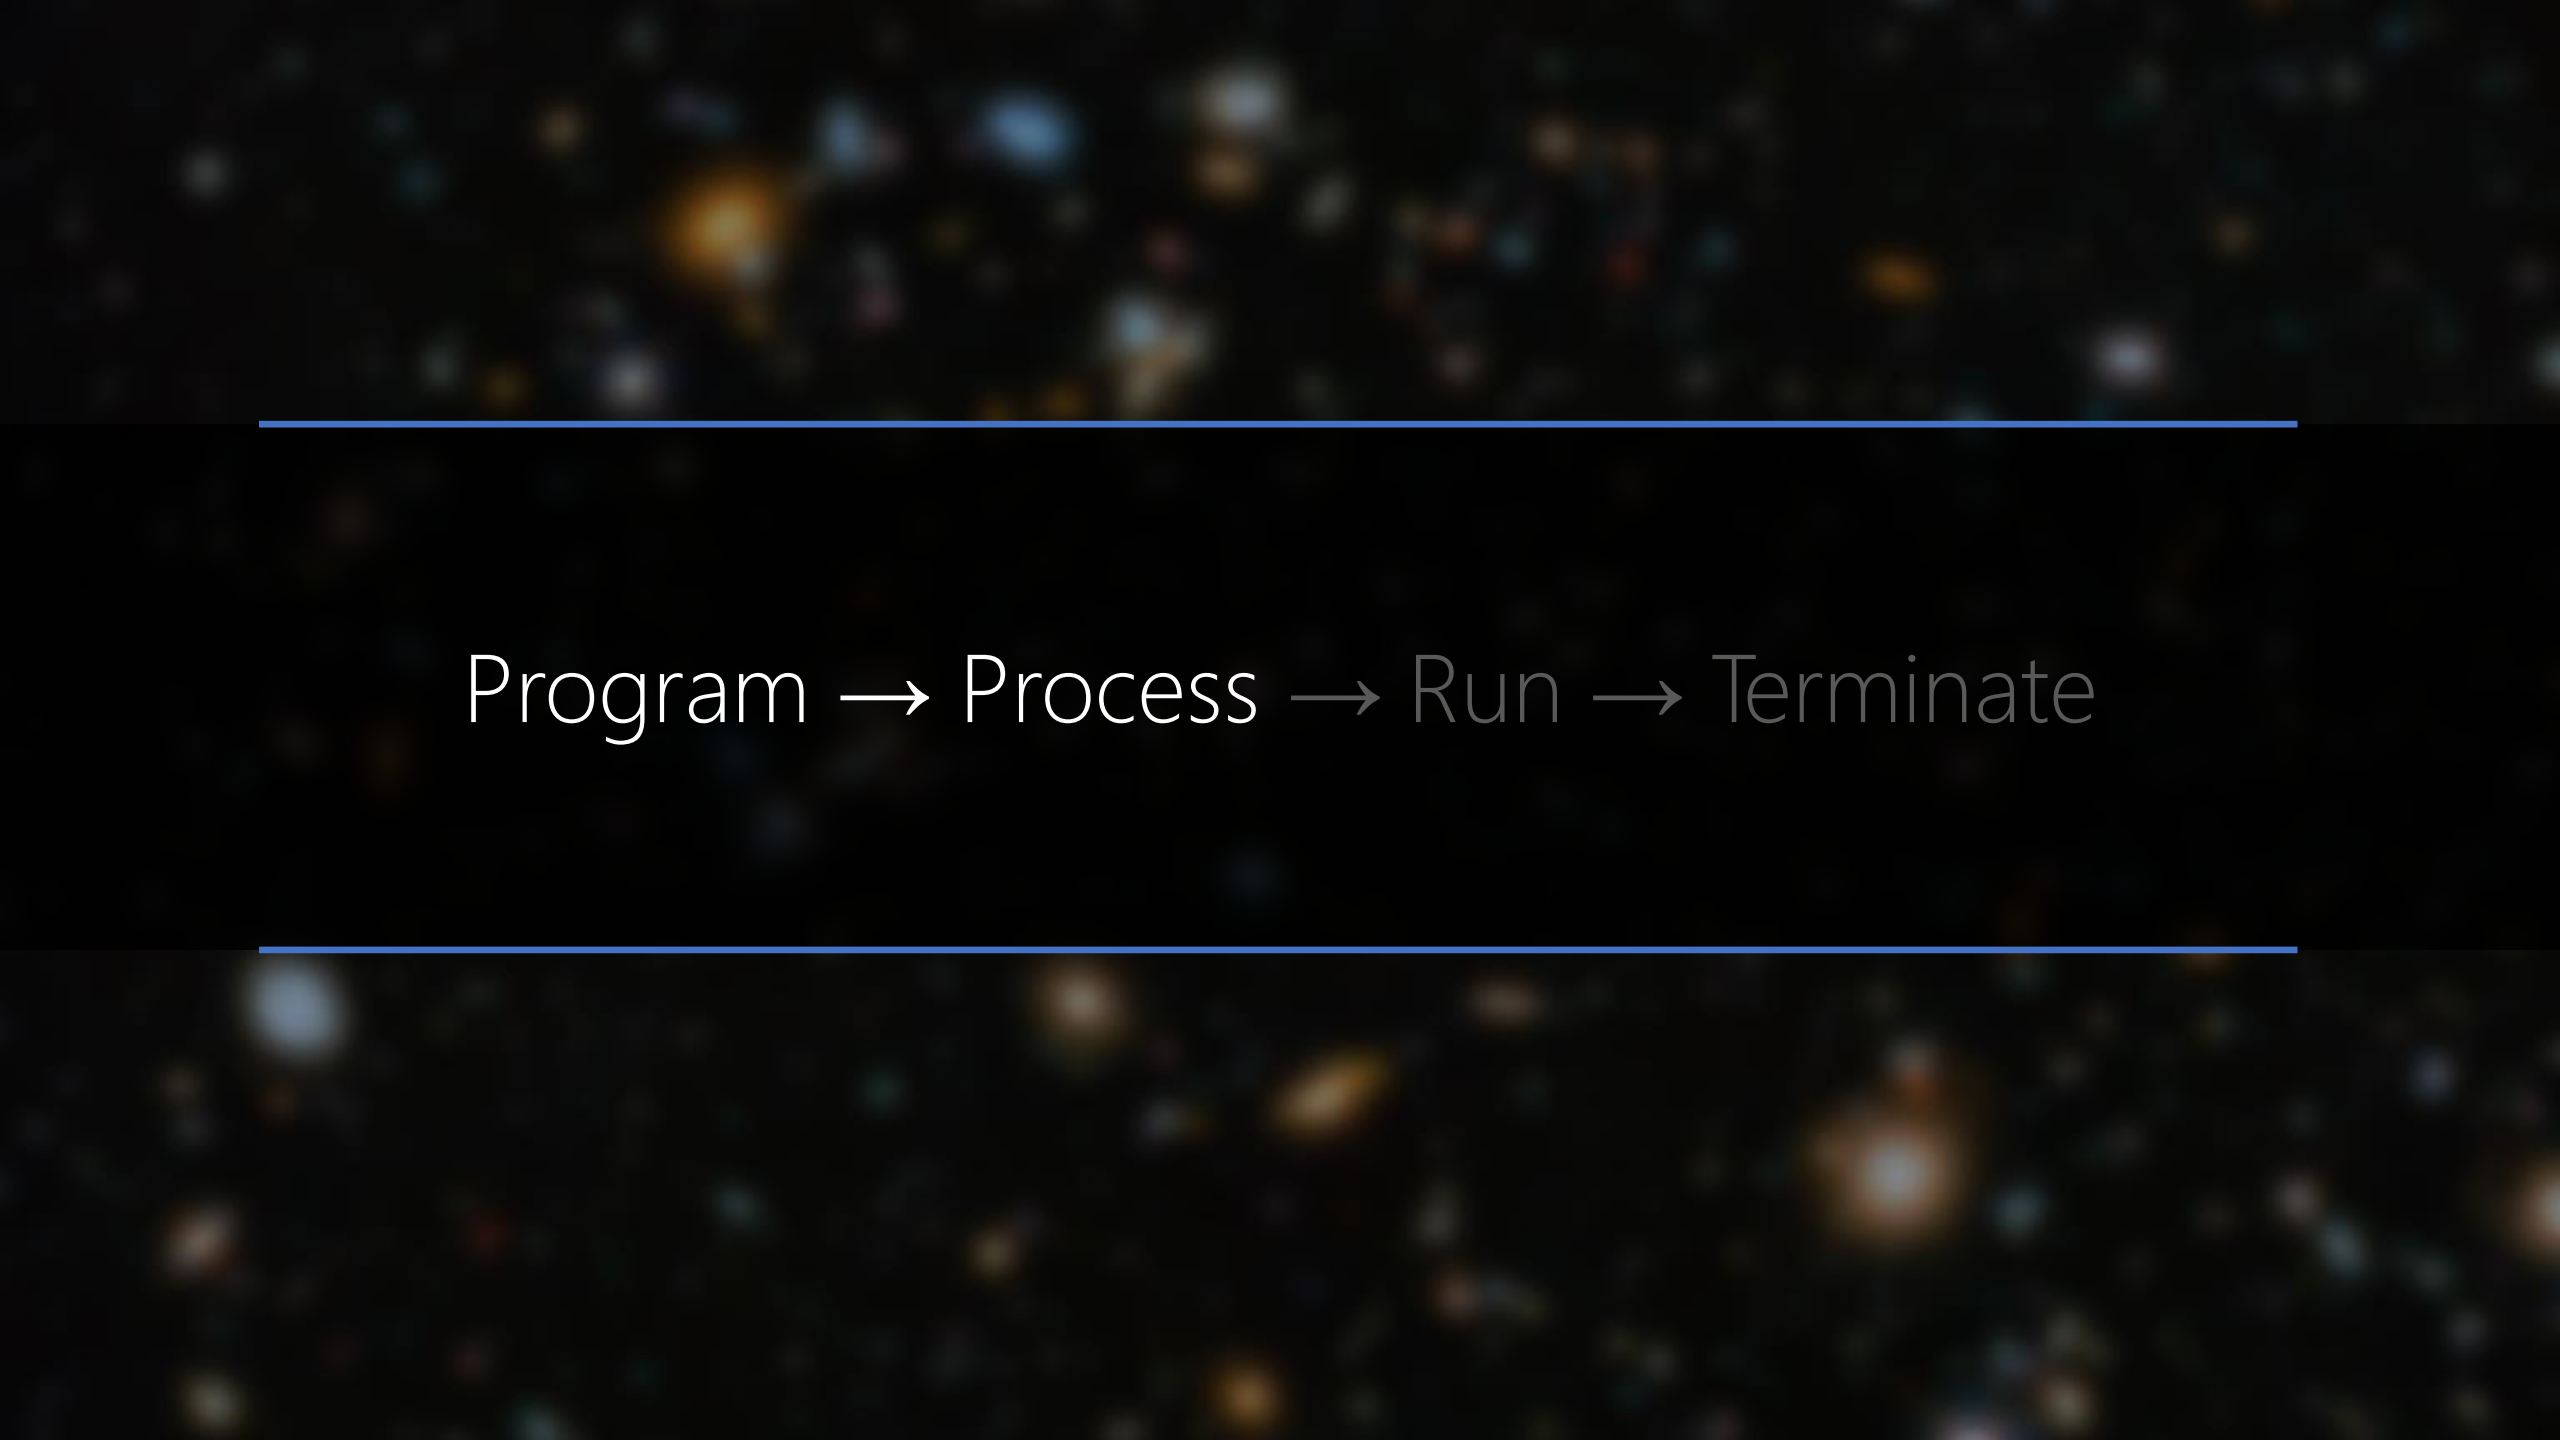
```
hfani@charlie:~$ ./main_add 2 4
```

```
2 + 4 = 6
```

```
hfani@charlie:~$
```




Into the Wild (2007) - Sean Penn

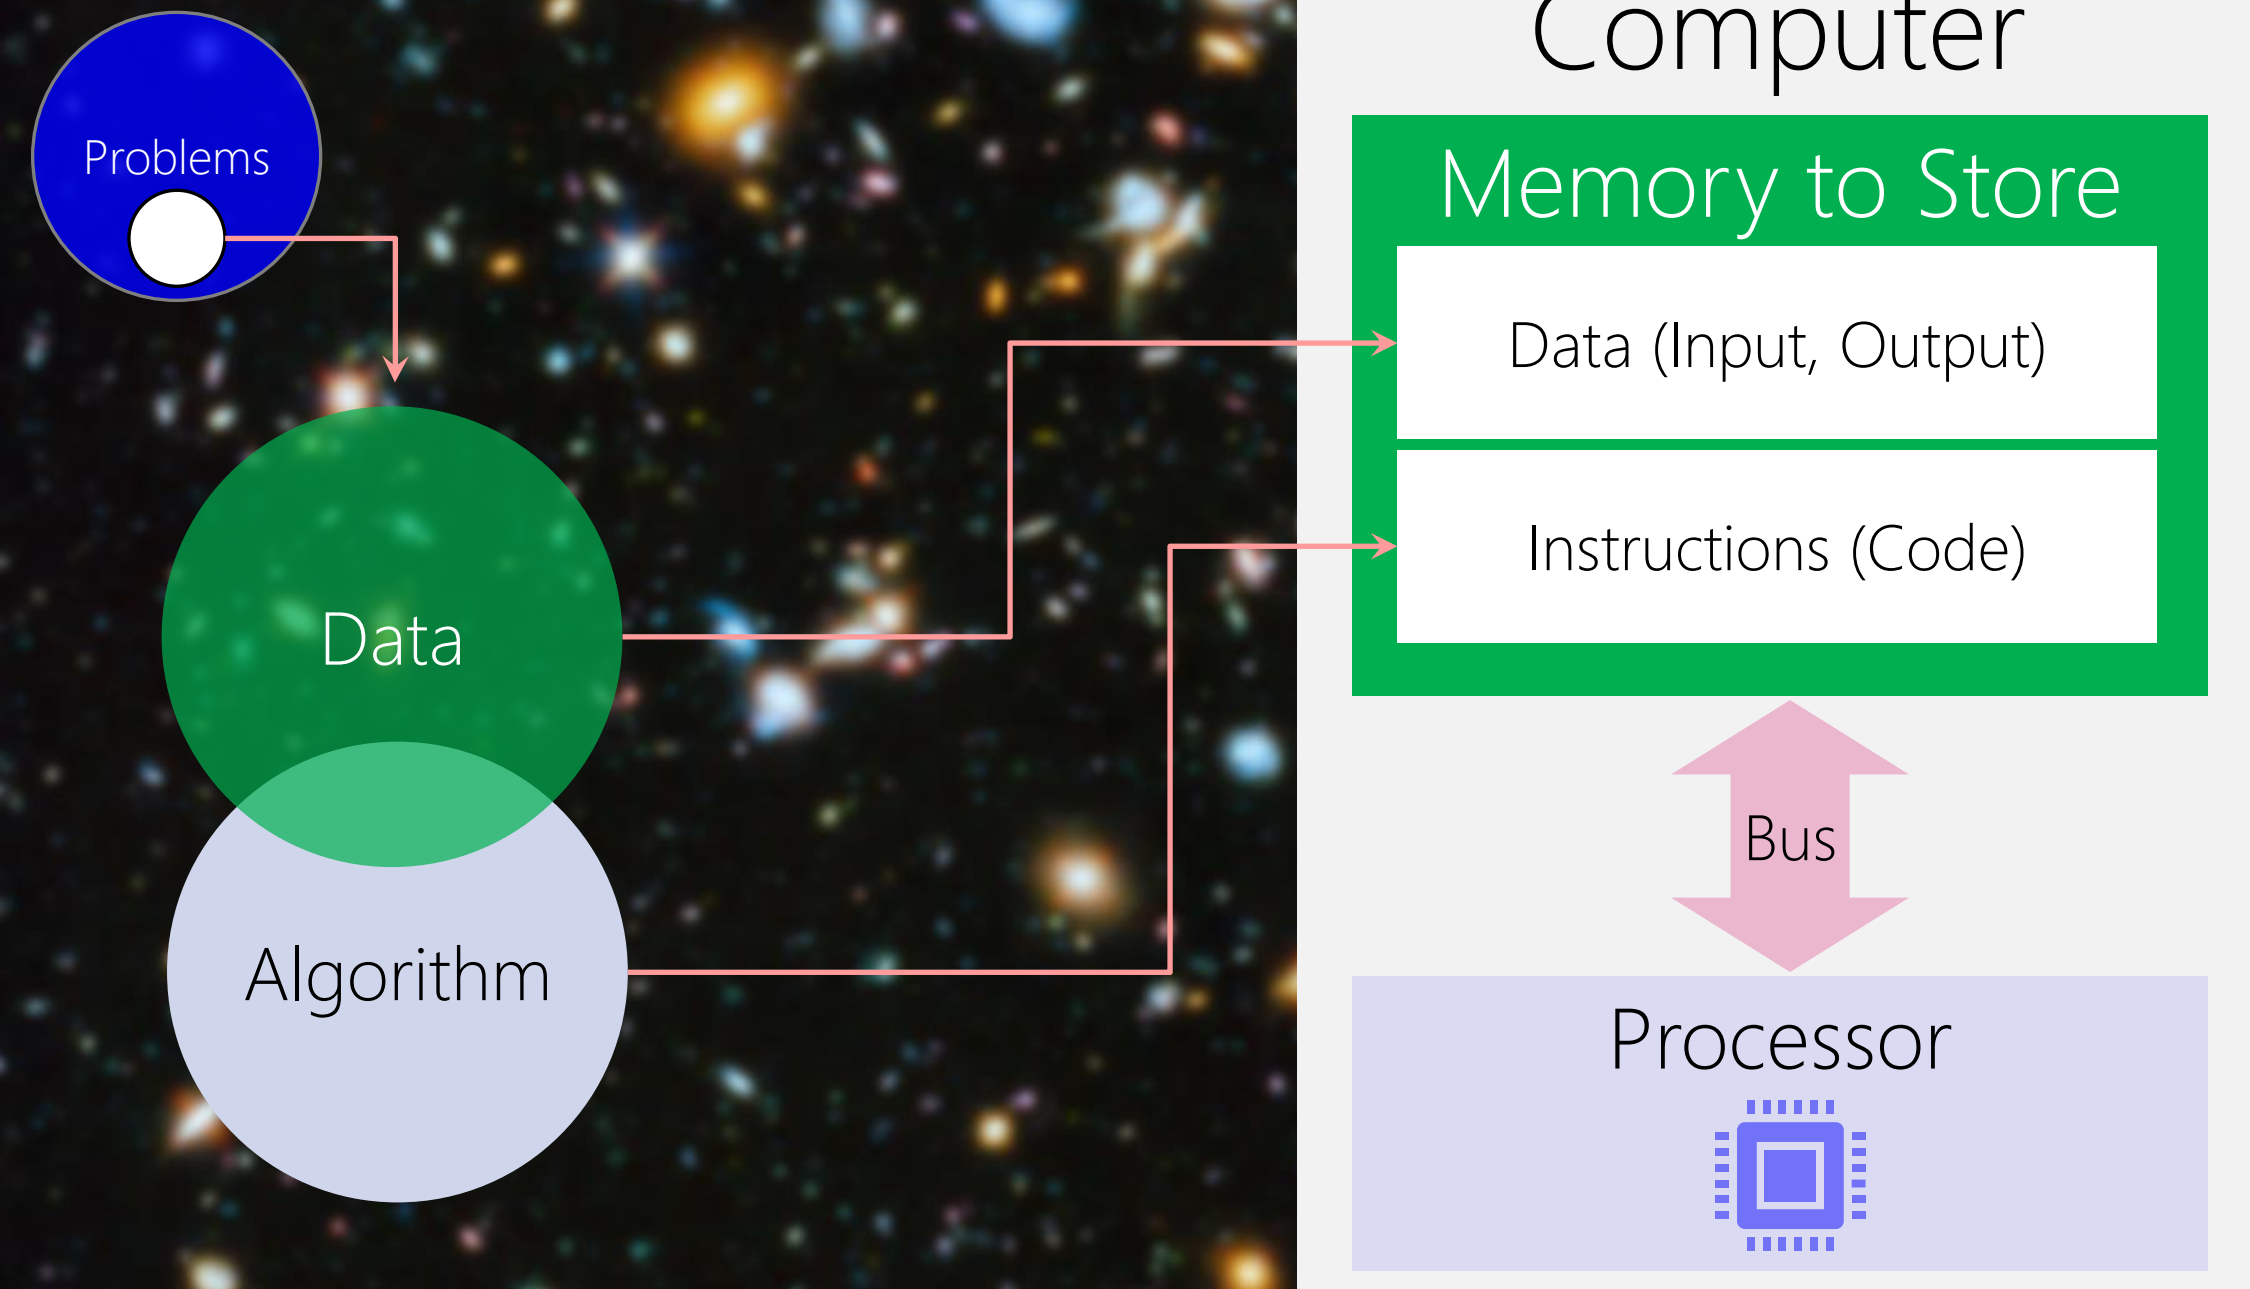


Program → Process → Run → Terminate

The background of the slide is a deep space image showing numerous distant galaxies in various colors (blue, orange, white) against a black background. Two horizontal blue lines are positioned above and below the title text.

Memory Layout of a C Program





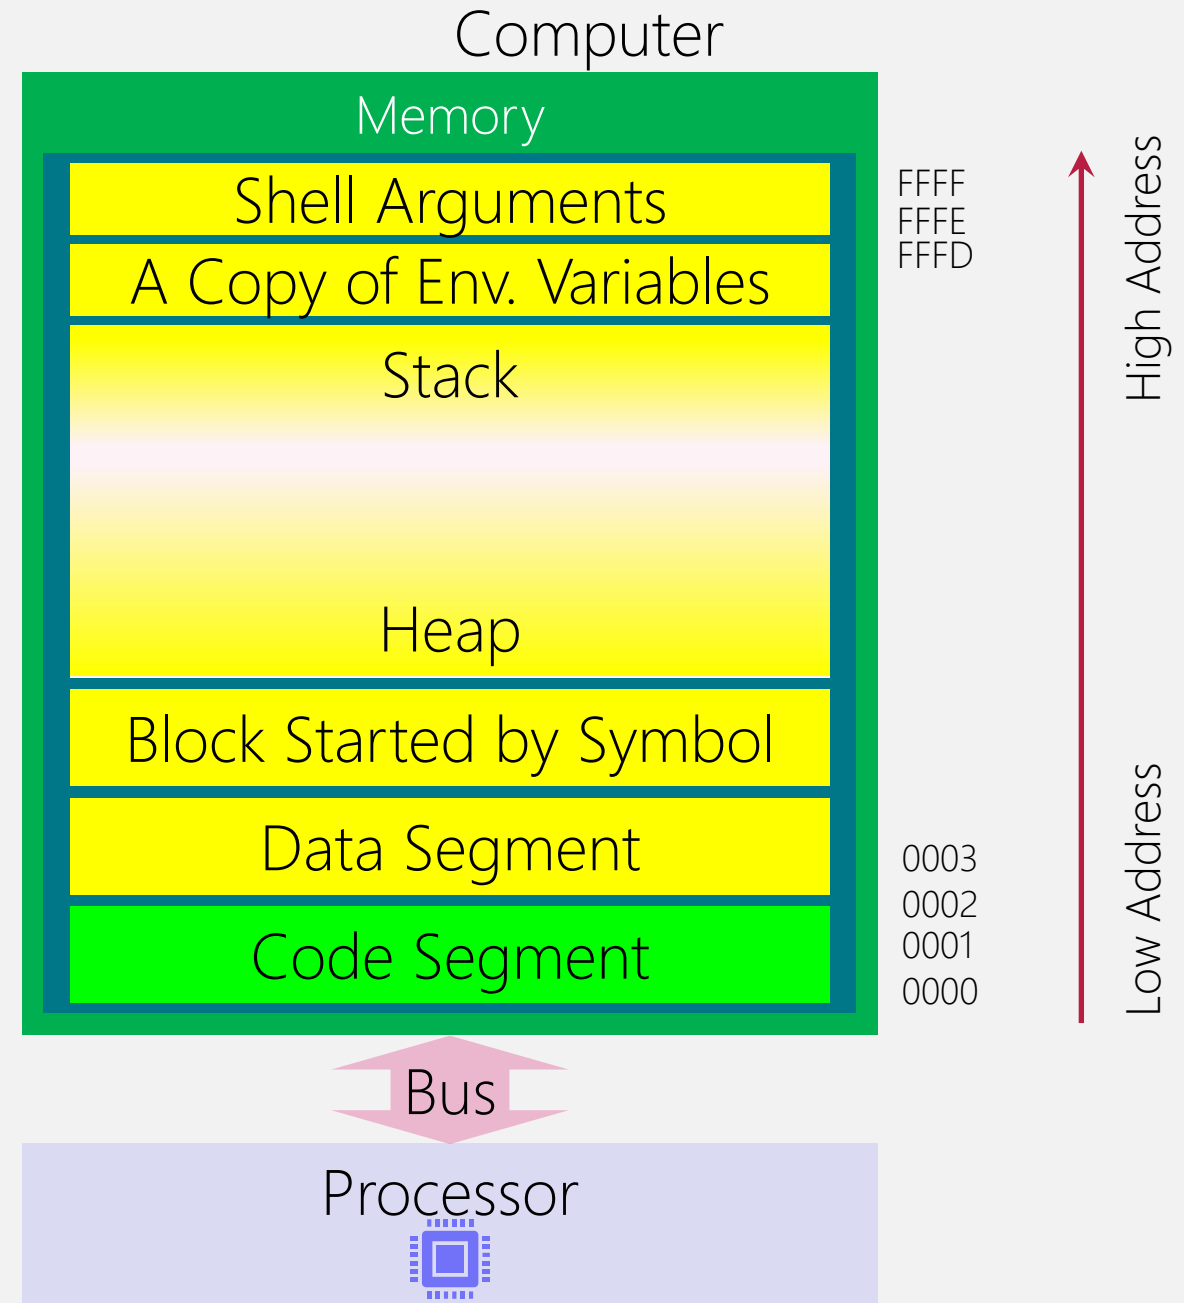
```

#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}

```



Code Segment (CS)

aka. text

```
#include <stdio.h>
#include <stdlib.h>
int main(int argc, char *argv[]) {
    int a = atoi(argv[1]);
    int b = atoi(argv[2]);
    printf("%d + %d = %d\n", a, b, a + b);
    return 0;
}
```

0002e10	1000	0000	0000	0000	000d	0000
0002e20	1214	0000	0000	0000	0019	0000
0002e30	3de8	0000	0000	0000	001b	0000
0002e40	0008	0000	0000	0000	001a	0000
0002e50	3df0	0000	0000	0000	001c	0000
0002e60	0008	0000	0000	0000	fef5	6fff
0002e70	0308	0000	0000	0000	0005	0000
0002e80	03f0	0000	0000	0000	0006	0000
0002e90	0330	0000	0000	0000	000a	0000

Data Segment (DS)

global variables, and variables inside `main`, that are initialized to a default value! (compile time)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

Memory

Data Segment

0F13
0F12
0F11
0F10

00000000
00000000
00000000
00000000

Block Started by a Symbol (BSS)

For uninitialized variables! (compile time)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

Memory

BSS Segment

0F15
0F14

????????

????????

<https://en.wikipedia.org/wiki/.bss>

Shell Argument + Environment Variables

Provided by the Shell (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

Memory

Shell Arguments	
FFFD	"5"
FFFC	"2"
FFFB	"d"
	⋮
	"i"
	"a"
	"m"
	"/"
FFF0	""
FF16	argv[2]=FFF0
FF15	argv[1]=FFCF
FF14	argv[0]=FFCC

Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_TYPE **en
                  int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base, group, _
}
libc_hidden_def (INTERNAL (strtol))
```



Into the Wild (2007) - Sean Penn

Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);
    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

My Return Address 1

Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_T
                  int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base,
    libc_hidden_def (INTERNAL (strtol))
```

My Return Address 2

My Return Address 1

Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int. */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_T
                  int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base,
    libc_hidden_def (INTERNAL (strtol))
}
```

Where should I
come back?

My Return Address 2

My Return Address 1

Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_T
                  int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base,
}
libc_hidden_def (INTERNAL (strtol))
```

My Return Address 2

My Return Address 1

Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_T
                  int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base,
    libc_hidden_def (INTERNAL (strtol))
```

Where should I
come back?

My Return Address 1

Stack

Functions Arguments, Local Variables, Return Address (runtime)

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);
    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}
```

```
#include <stdlib.h>
#undef      atoi
/* Convert a string to an int.  */
int
atoi (const char *nptr)
{
    return (int) strtol (nptr, (char **) NULL, 10);
}
libc_hidden_def (atoi)
```

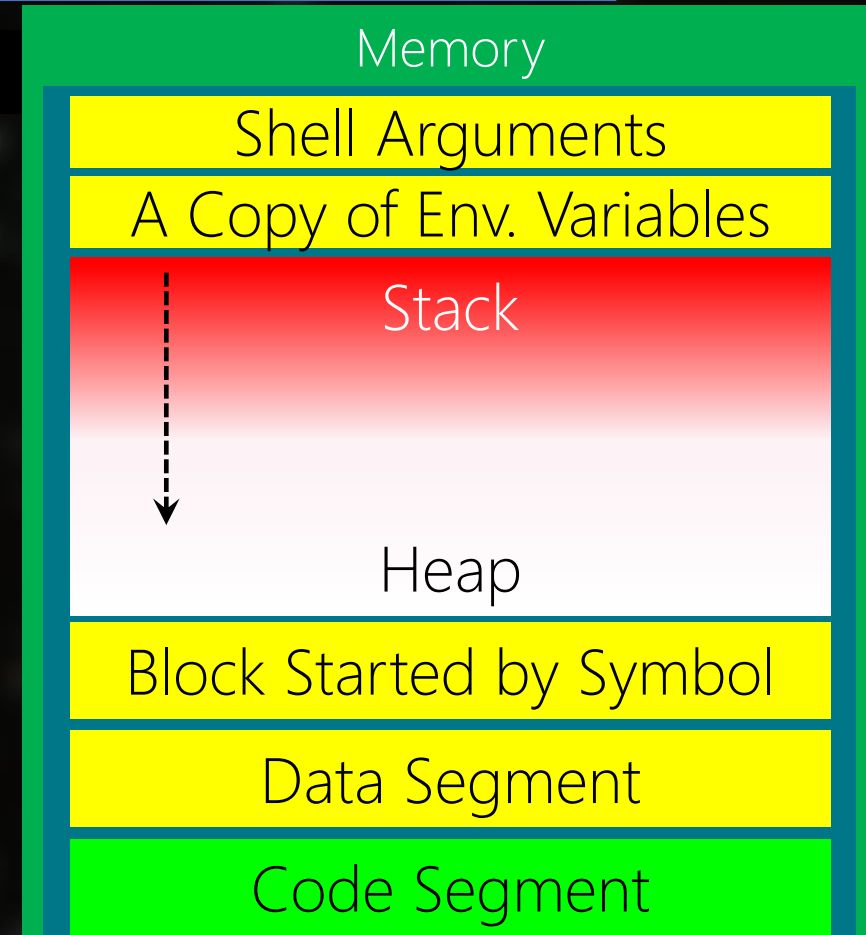
```
INT
INTERNAL (strtol) (const STRING_TYPE *nptr, STRING_T
                  int base, int group)
{
    return INTERNAL (__strtol_1) (nptr, endptr, base,
}
libc_hidden_def (INTERNAL (strtol))
```

My Return Address 1

Stack


Functions Arguments, Local Variables, Return Address (runtime)

Program's Stack is Upside Down!



Stack Overflow?

Functions Arguments, Local Variables, Return Address (runtime)

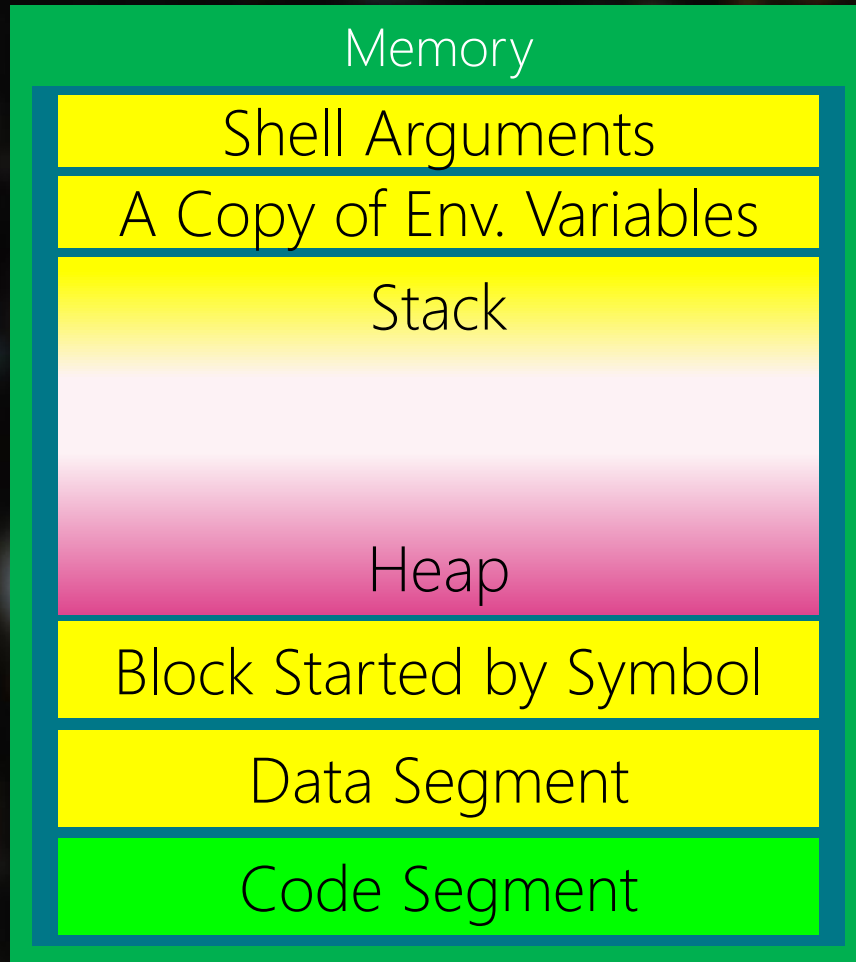




Into the Wild (2007) - Sean Penn

Heap

Dynamic memory allocation (runtime)



Memory Allocators by Library Routines

```
#include <stdlib.h>
void *malloc(size_t size)
void *realloc(void *ptr, size_t newsize)
```


Size is **fixed** during **compile** time
Value is **dynamic** during **runtime**

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]) {
    int a = 0;
    int b = 0;
    a = atoi(argv[1]);
    b = atoi(argv[2]);

    result = a + b;

    printf("%d + %d = %d\n", a, b, result);
    return 0;
}

hfani@charlie:~$ ./main_add 2 2
2 + 2 = 4
hfani@charlie:~$ ./main_add 2 4
2 + 4 = 6
hfani@charlie:~$
```

Size is dynamic during runtime
Value is dynamic during runtime

```
#include <stdio.h>
#include <stdlib.h>
int result;
int main(int argc, char *argv[]){
    int size_a = 0;
    int size_b = 0;
    size_a = atoi(argv[1]);
    size_b = atoi(argv[2]);

    int *a = malloc(size_a * sizeof(int));
    printf("enter the first number with %d digits:\n", size_a);
    for(int i = 0; i < size_a; ++i){
        scanf("%d", a + i);
    }

    int *b = malloc(size_b * sizeof(int));
    printf("enter the first number with %d digits:\n", size_b);
    for(int i = 0; i < size_b; ++i){
        scanf("%d", b + i);
    }
}
```

```
hfani@charlie:~$ ./main_malloc 3 4
enter the first number with 3 digits:
1
3
9
enter the first number with 4 digits:
6
5
7
2
139 + 6572
```


Size is **dynamic** during **runtime**
Value is **dynamic** during **runtime**

```
hfani@charlie:~$ ./main_malloc 1000000000000000 10000000000000000
```

What happens?

Size is **dynamic** during **runtime**
Value is **dynamic** during **runtime**

```
hfani@charlie:~$ ./main_malloc 1000000000000000 10000000000000000
```



Heap

Dynamic memory allocation (runtime)

Memory Allocators by Library Routines

```
#include <stdlib.h>
void *malloc(size_t size)
void *realloc(void *ptr, size_t newsize)
```


Memory Allocators by System Calls?

Shell's `size` command

```
hfani@charlie:~$ size ./main_malloc
text    data    bss     dec     hex filename
2239    616      8    2863    b2f ./main_malloc
```

Is this info for:

- Compile time?
- Runtime?




Shell's `size` command

```
hfani@charlie:~$ size ./main_malloc
text    data    bss     dec      hex filename
2239     616      8     2863     b2f ./main_malloc
```

Why is not any info for:

- Stack?
- Heap?



Process Identifier (pid)

Non-negative

Unique among **processes** (live programs)

Not an identifier! It can be reused (delay reuse)

Process Identifier by System Call `getpid()`

```
#include <unistd.h>
pid_t getpid(void);
```

Return process ID of calling process

```
#include <unistd.h>
#include <stdio.h>
int main(void){
    printf("%d\n", getpid());
    return 0;
}
```

```
hfani@alpha:~$ ./getpid
871198
hfani@alpha:~$ ./getpid
871217
```



Into the Wild (2007) - Sean Penn