



Lab09: Shared File | Pipe between Children



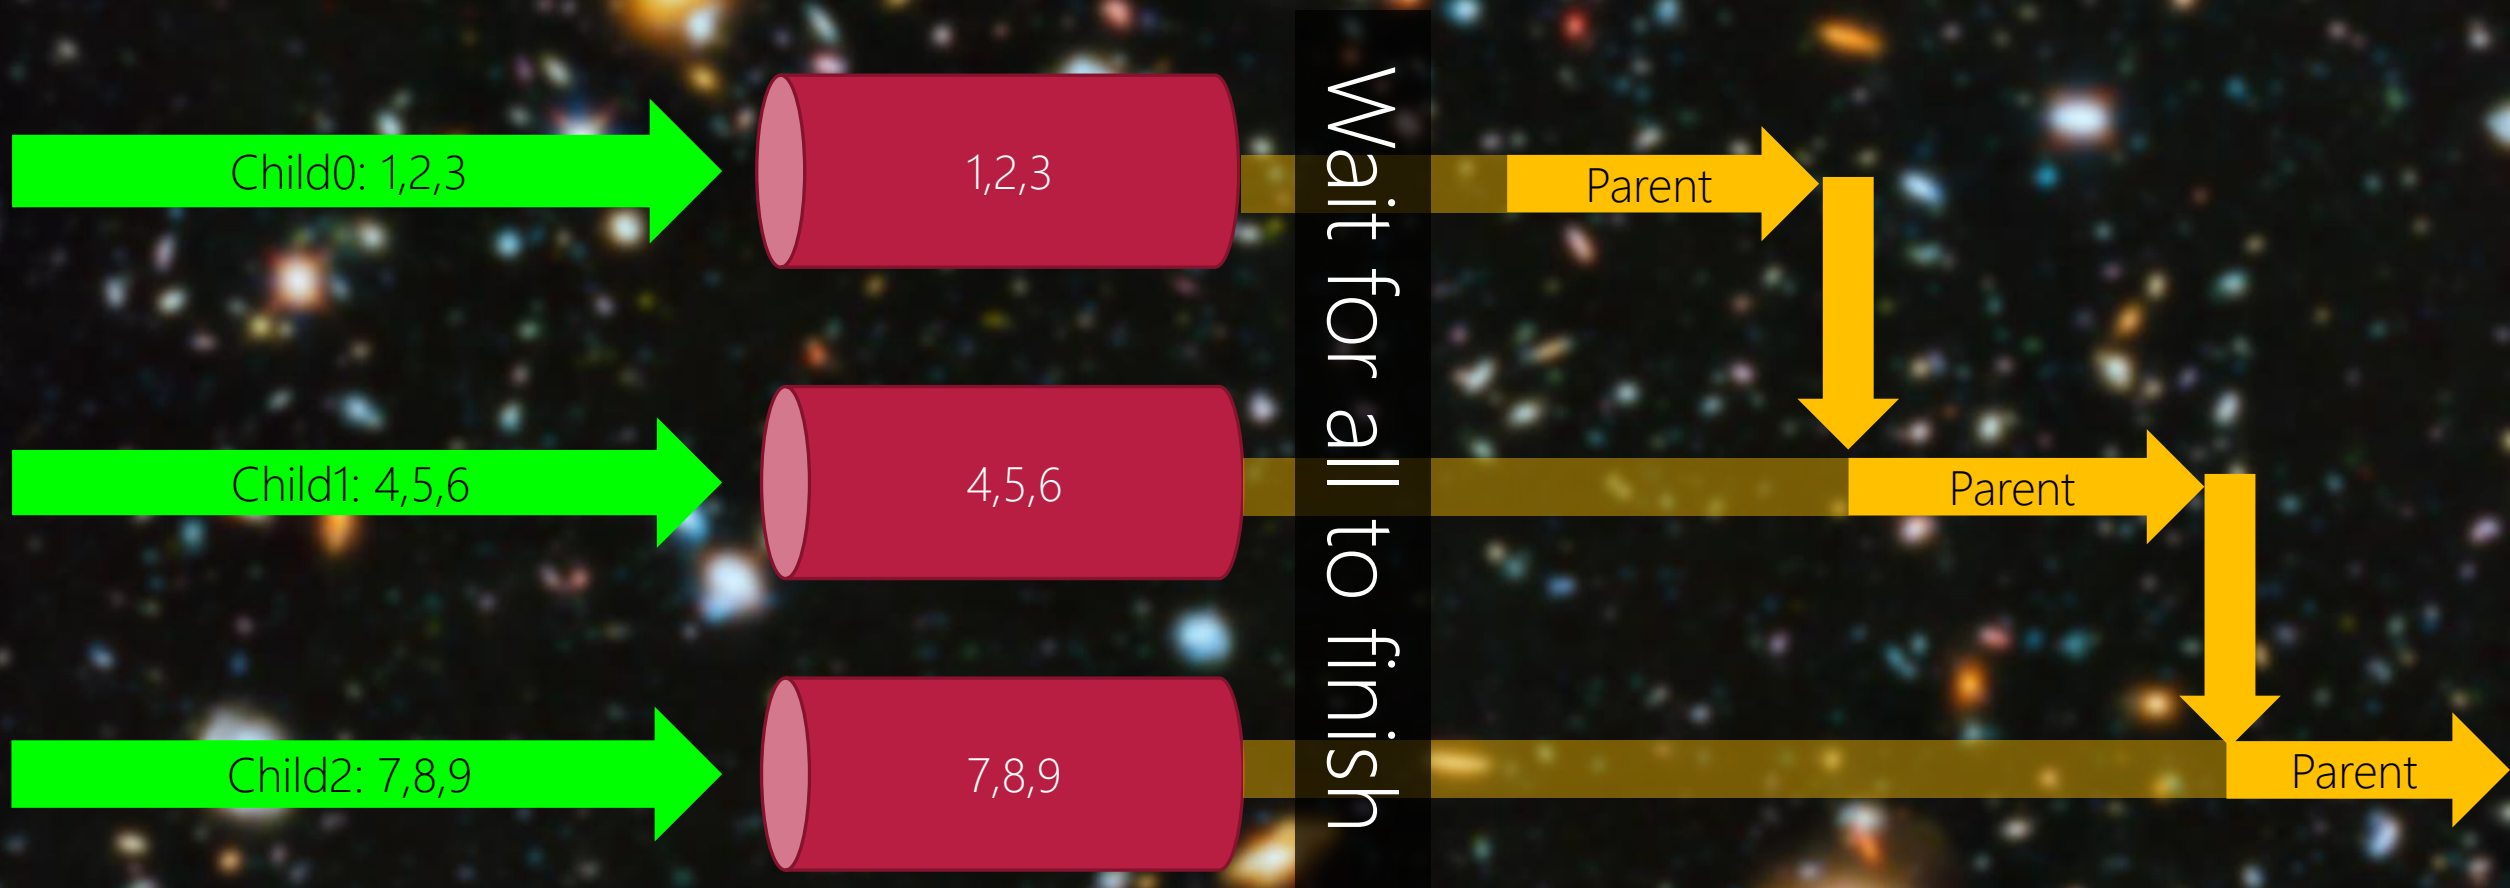
No Assumption about
Processor Sharing
by the Kernel



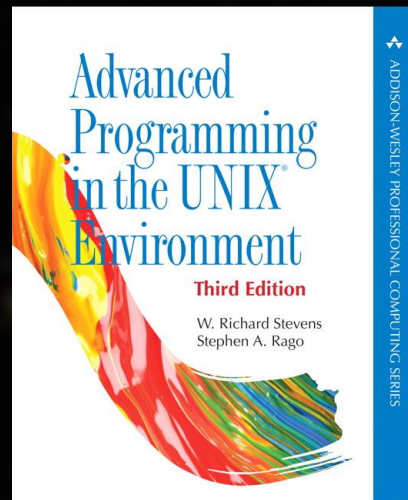
Sleep



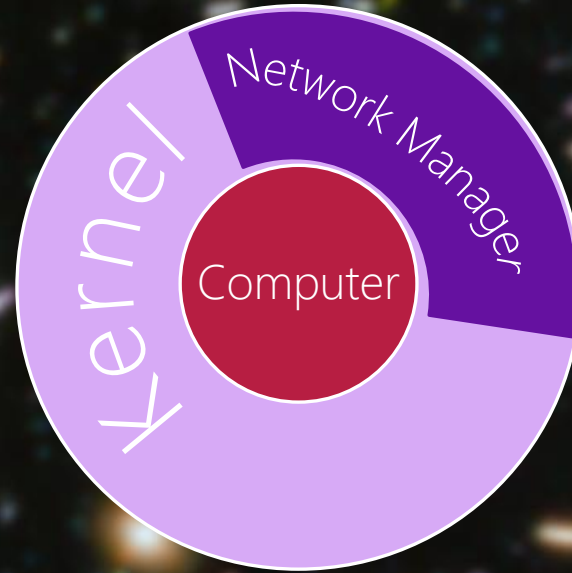
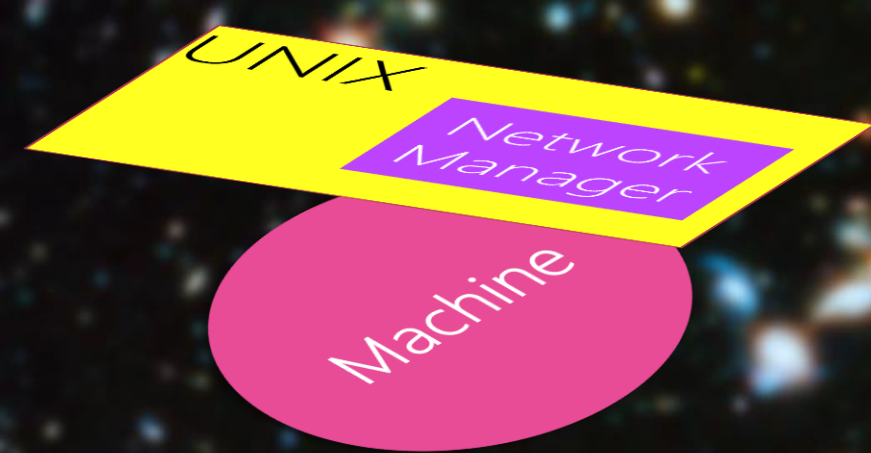
Still, There is a Problem.
Why?



Always correct regardless of processor scheduling.



Chapter 16: Network IPC (Sockets)



Computer

Memory

Kernel: Device Manager

Kernel: Memory Manager

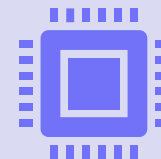
~~Kernel: File Manager~~

Kernel: Network Manager

~~Kernel: Process Manager~~

Bus

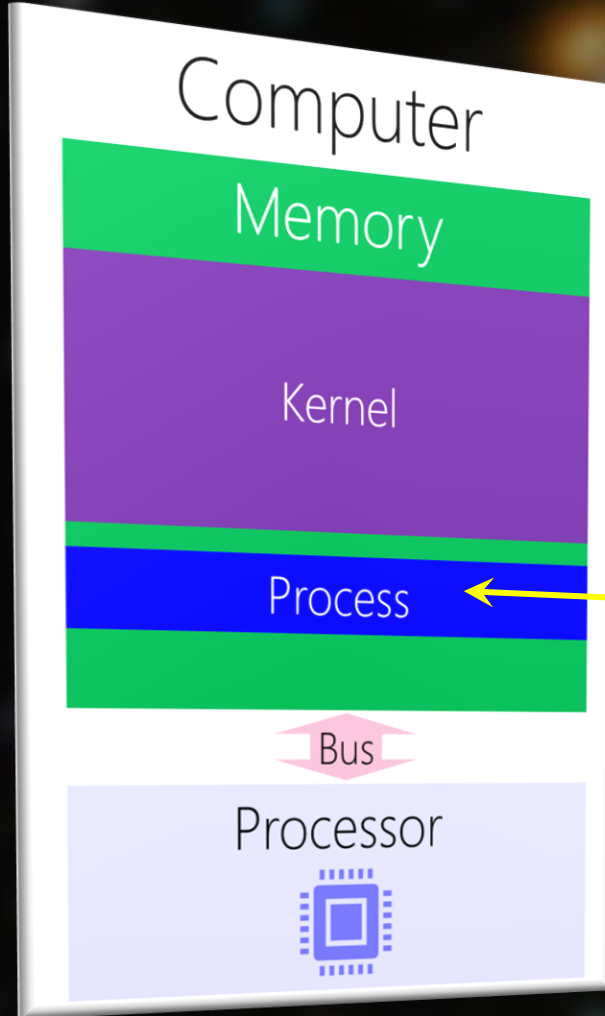
Processor



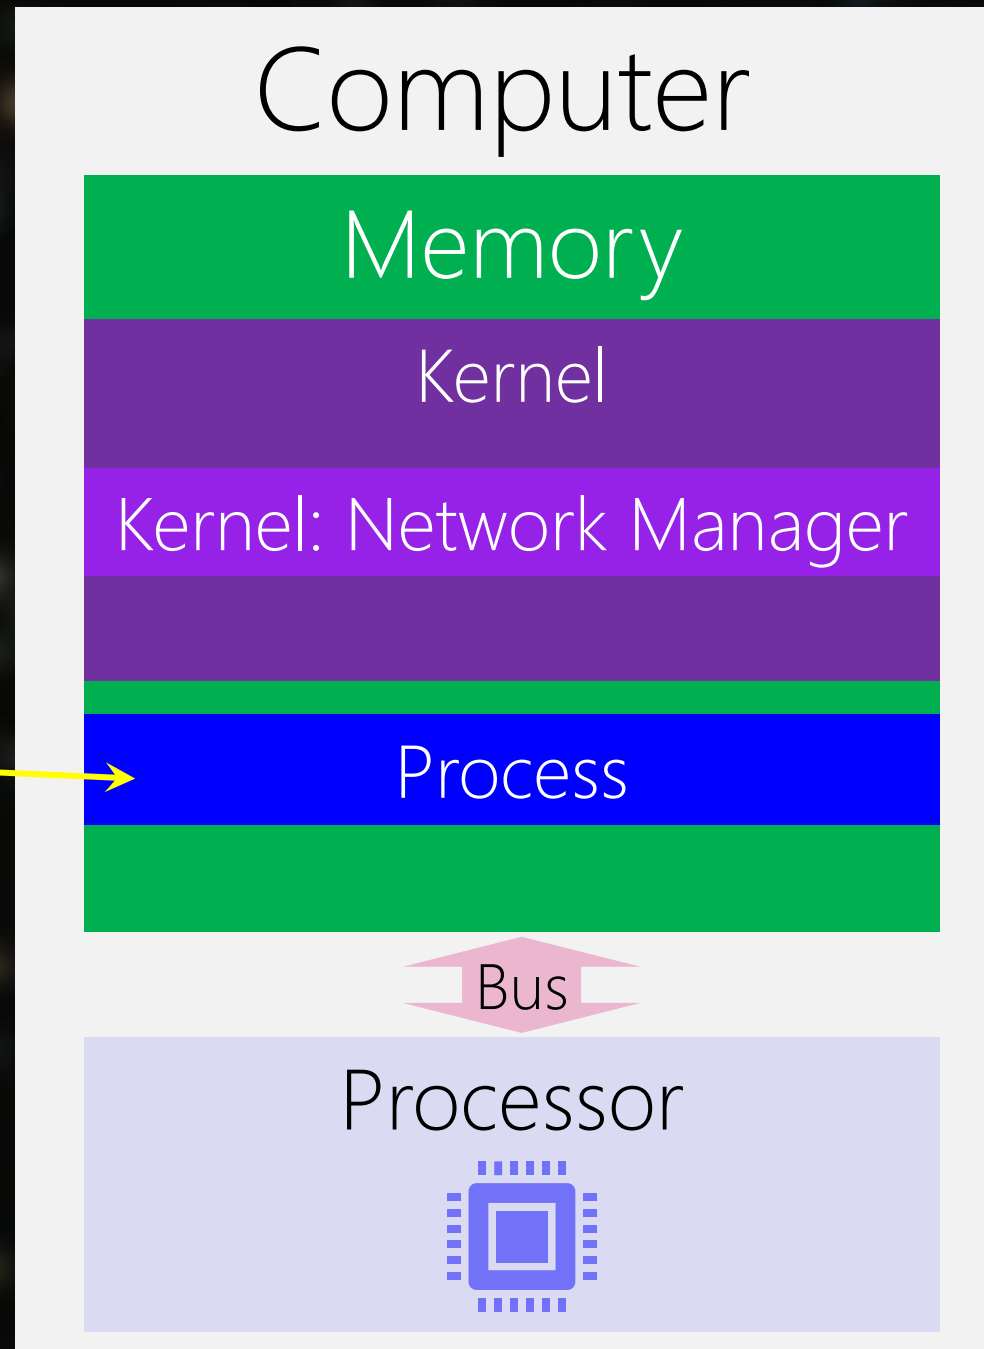
Multiprocessing Computers

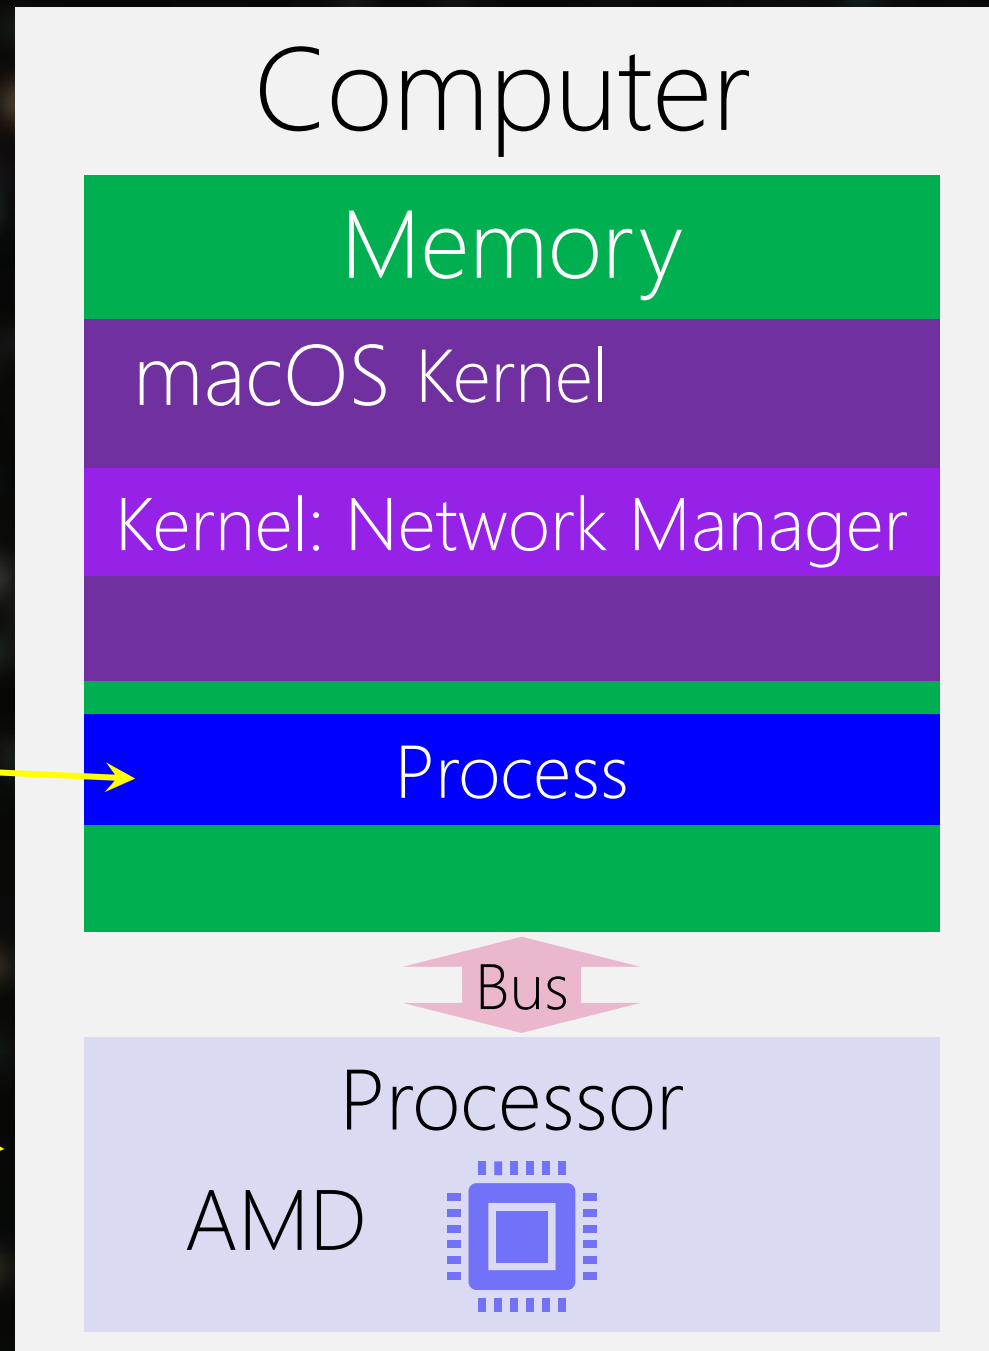

aka Computer Network

Multiple Single Processor Multiprocessor




Network IPC





Network IPC


Physical Connection
Wired/Wireless



COMP3670: Computer Networks

Network Topology


The way computer systems are connected.
By software control, we can convert one to another.




Ring



Line



Star



Bus

Network Protocol

Conversation Protocol between Computers

Language, Order (Who Talks, Who Listens), Addressing (Finding Each Other), ...

1975: 2-network communications between Stanford and University College London

1977: 3-network between sites in the US, the UK, and Norway

TCP/IP

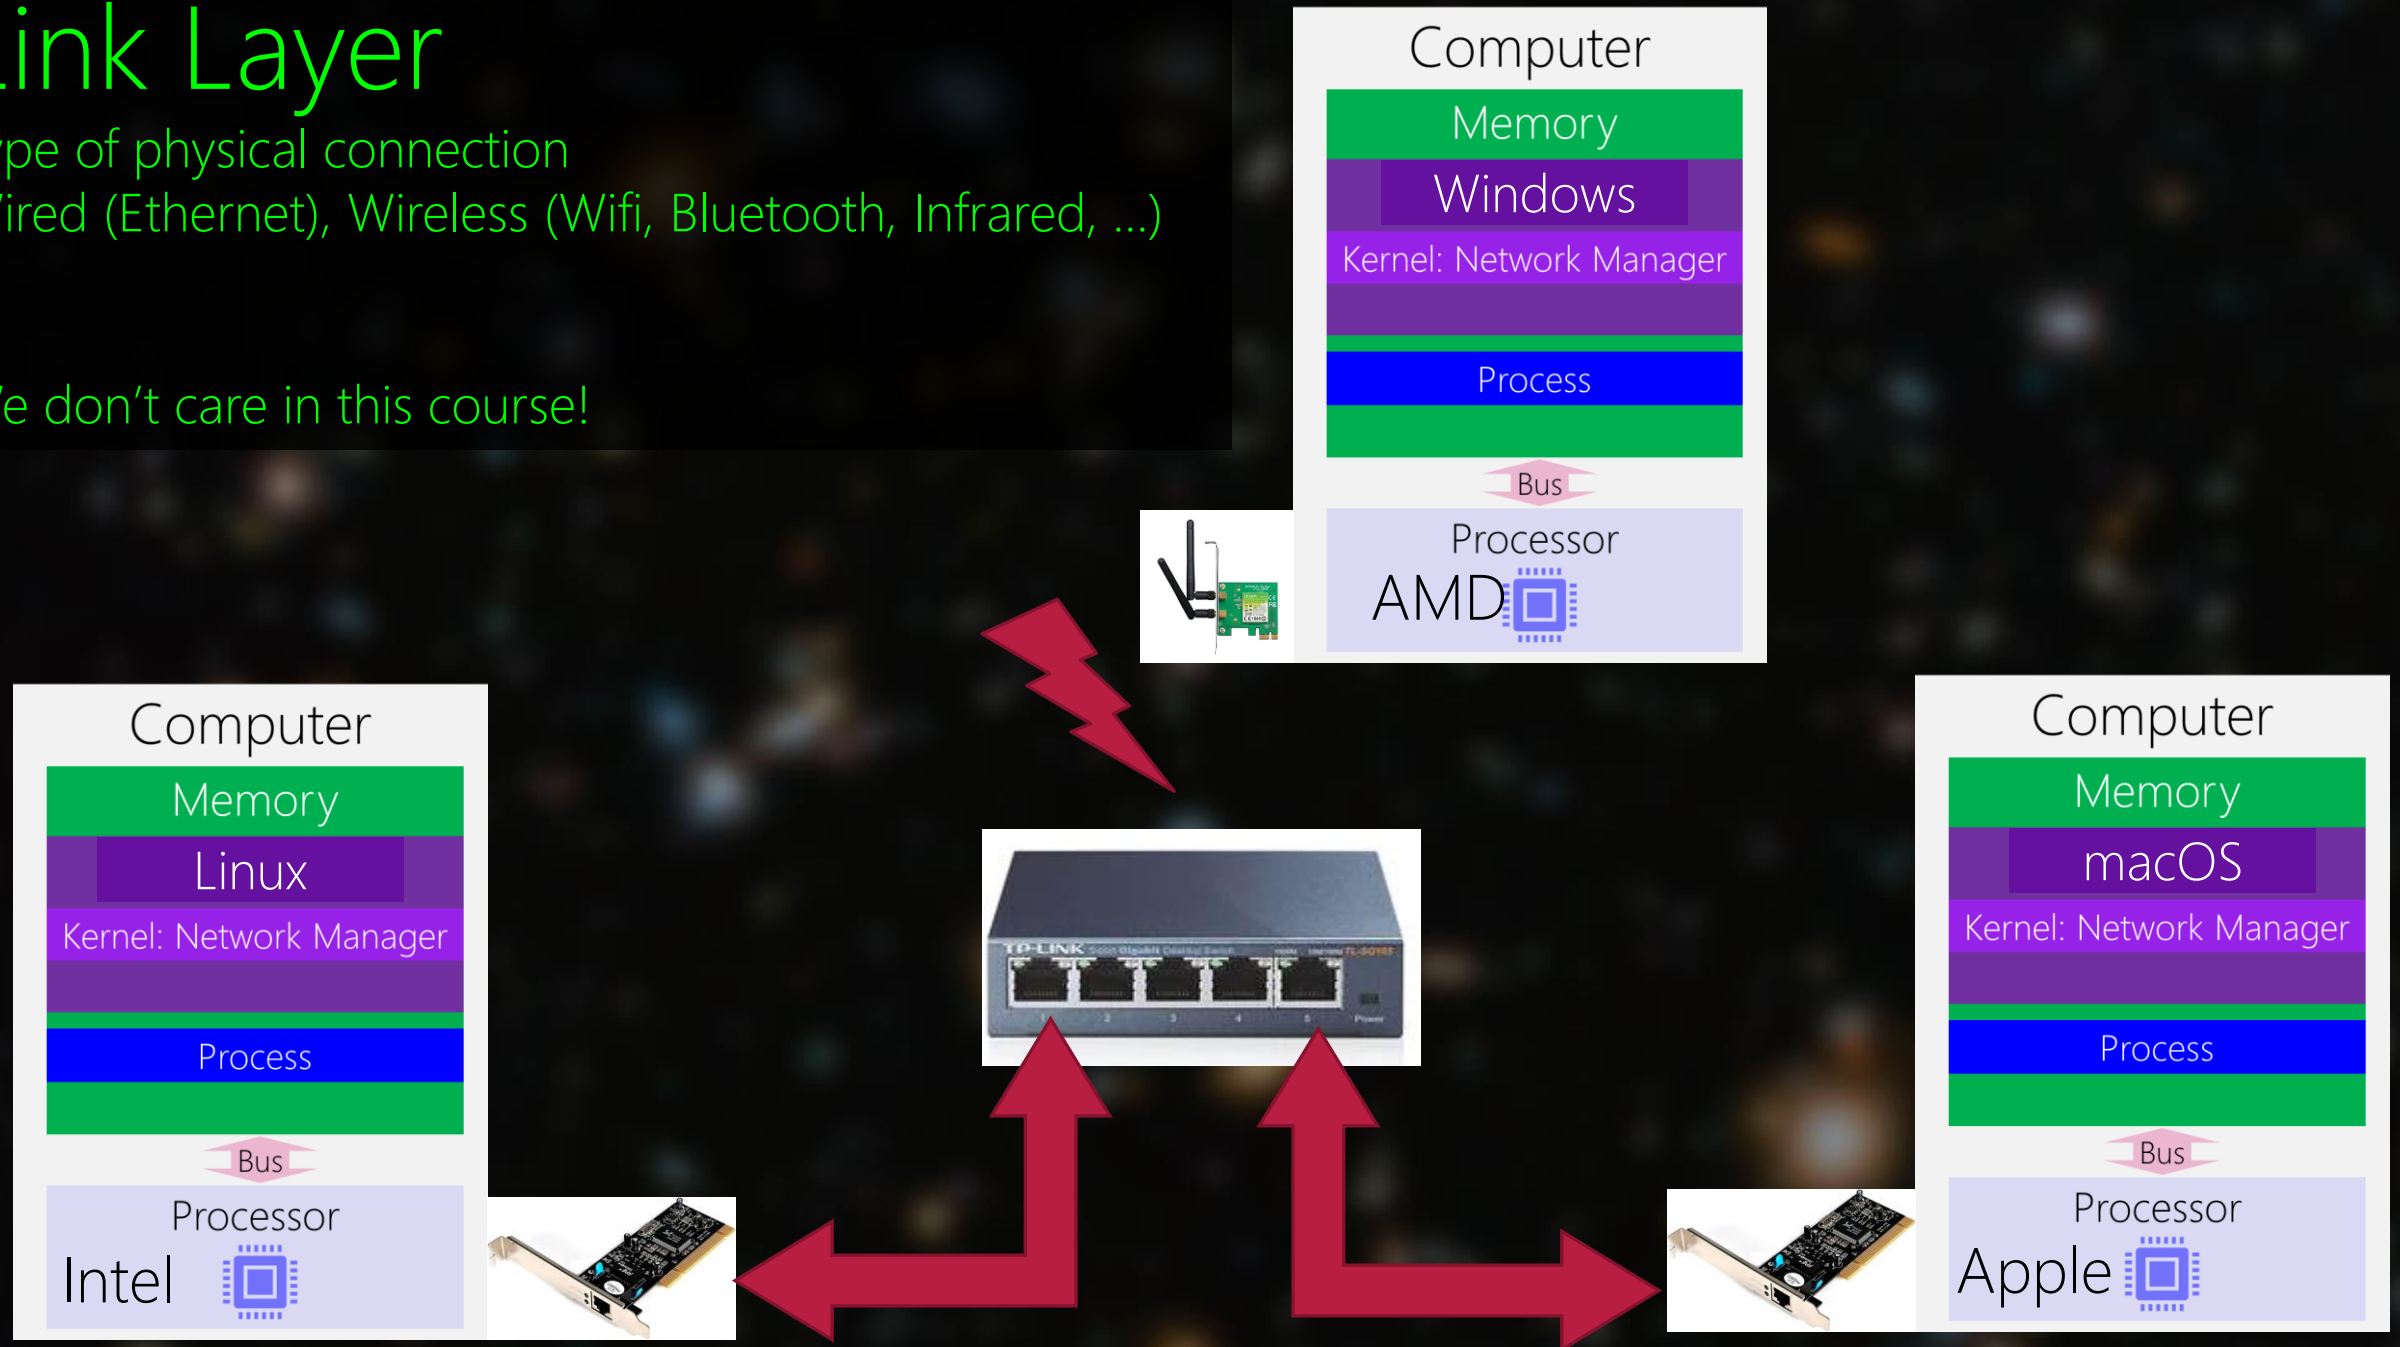
There are other network protocols!
TCP/IP is just a name. It does not represent all this protocol offers!

Link Layer

Type of physical connection

Wired (Ethernet), Wireless (Wifi, Bluetooth, Infrared, ...)

We don't care in this course!



Inter-Network → Internet (Network) Layer → Internet Protocol (IP)

Computers' Address, Names,


We use the addresses in this course.

We don't care about the rest.

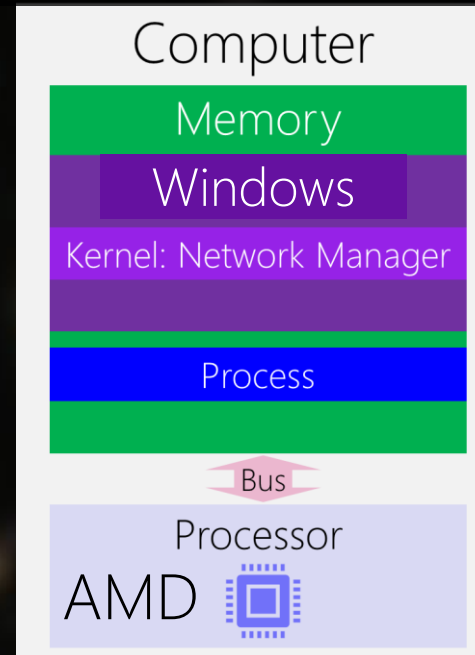
Why the format is like this?

Who assigns the addresses?

...

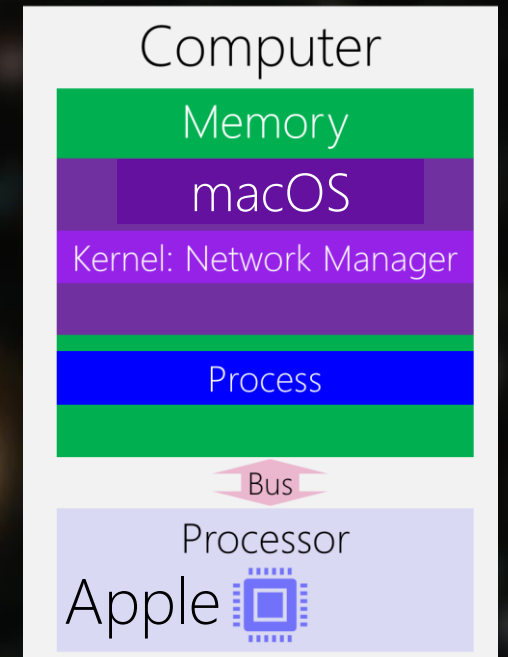


137.207.82.52



4.2.2.2

137.207.140.134



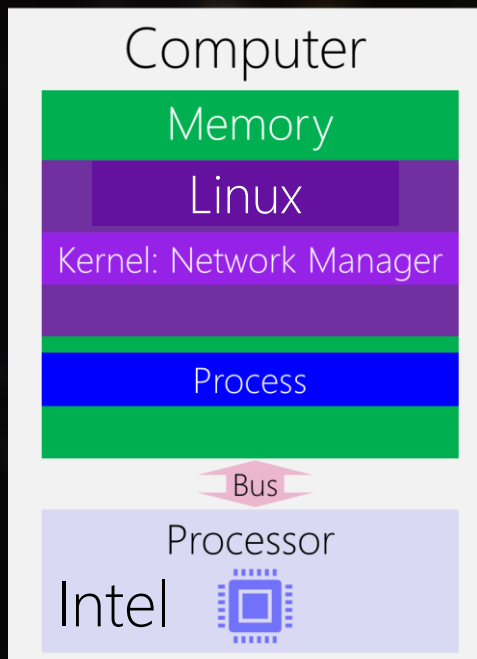
Transport Layer

Agreement on communication protocol

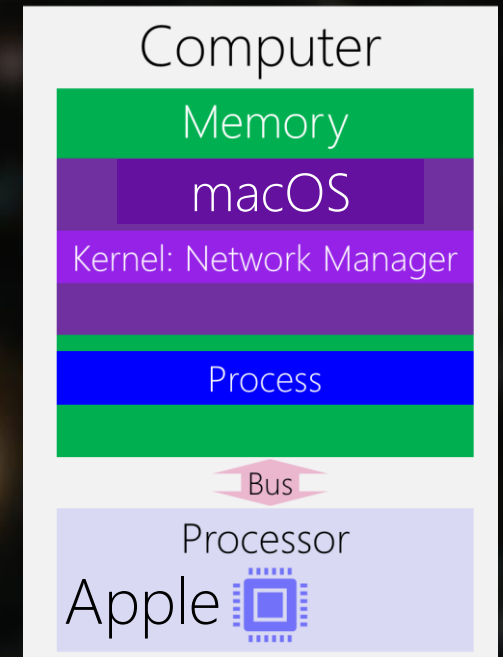
1) Connectionless == Sending a mail

- No order (a mail may be sent sooner, but received later)
- No reliability (non-tracking mail.) Cannot see whether it is received or lost
- Each message is self-contained (Does not depend on previous or next mails)
- Simple and light (no overhead for sender, like PR card by government of Canada)

User Datagram Protocol (UDP)



137.207.82.52



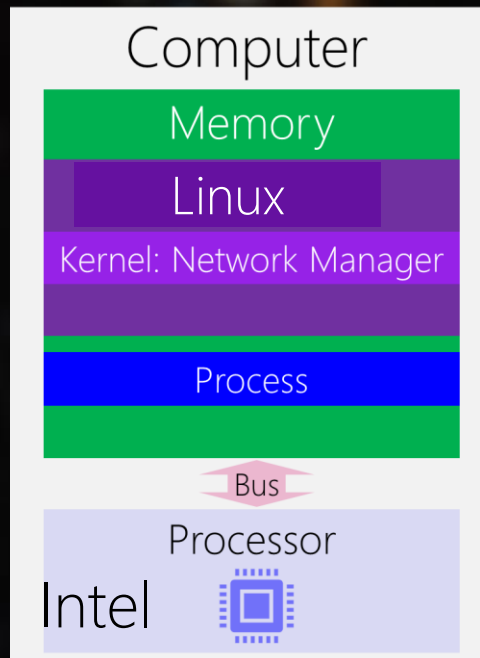
Transport Layer

Agreement on communication protocol

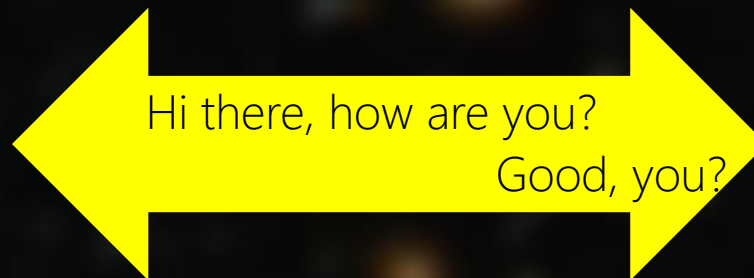
2) Connection-Oriented == Phone Call

- Foremost setup a connection to make sure there is a receiver ready
- Ordered (when you talk on the phone, the words are transferred in order)
- Reliability (there is an active listener)
- Each *packet* depends on previous or next packets
- Connection overhead for sender

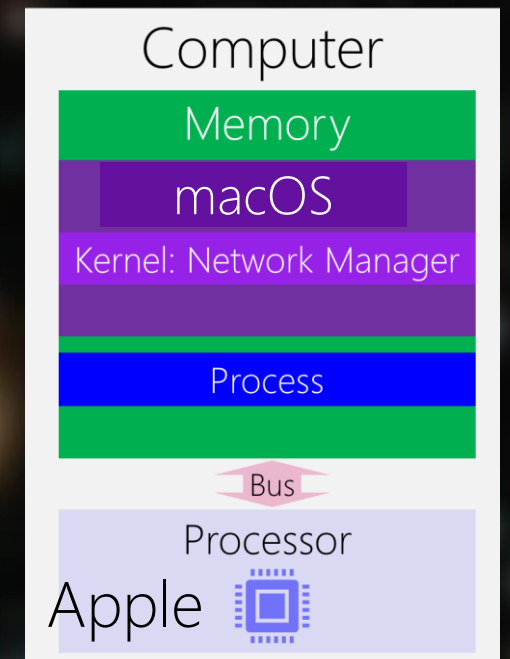
Transmission Control Protocol (TCP)



137.207.82.52



137.207.140.134




Application Layer

Any process that wants to communicate via the network



TCP/IP

Just a name for [Link | Internet | Transport | Application] network protocol



TCP/IP

Just a name for [Link | Internet | Transport | Application] network protocol



TCP/IP

Just a name for [Link | Internet | Transport | Application] network protocol





Socket Programming


TCP/IP: UDP

TCP/IP: UDP at Sender

Connectionless Communication

Sending a mail

- 1) Creating Socket
- 2) Binding to an Address (Optional)
- 3) Find the Receiver's Address
- 4) Send the Mail to the Receiver




TCP/IP: UDP at Sender

Connectionless Communication

Sending a mail

1) Creating Socket

```
#include <sys/socket.h>
int socket(int domain, int type, int protocol);
Returns socket descriptor if OK, -1 on error
```




```
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#include <stdio.h>
#include <string.h>
int main(void) {
    int domain = AF_INET;//Network Protocol: TCP/IP
    int type = SOCK_DGRAM;//Connectionless
    int protocol = 0;//Default transport: UDP for Internet connectionless
```

Set up the type of network communication

Domain	Description
AF_INET	IPv4 Internet domain
AF_INET6	IPv6 Internet domain (optional in POSIX.1)
AF_UNIX	UNIX domain
AF_UNSPEC	unspecified

Figure 16.1 Socket communication domains

Type	Description
SOCK_DGRAM	fixed-length, connectionless, unreliable messages
SOCK_RAW	datagram interface to IP (optional in POSIX.1)
SOCK_SEQPACKET	fixed-length, sequenced, reliable, connection-oriented messages
SOCK_STREAM	sequenced, reliable, bidirectional, connection-oriented byte streams

Figure 16.2 Socket types

Protocol	Description
IPPROTO_IP	IPv4 Internet Protocol
IPPROTO_IPV6	IPv6 Internet Protocol (optional in POSIX.1)
IPPROTO_ICMP	Internet Control Message Protocol
IPPROTO_RAW	Raw IP packets protocol (optional in POSIX.1)
IPPROTO_TCP	Transmission Control Protocol
IPPROTO_UDP	User Datagram Protocol

Figure 16.3 Protocols defined for Internet domain sockets

```
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#include <stdio.h>
#include <string.h>
int main(void) {
    int domain = AF_INET; //Network Protocol: TCP/IP
    int type = SOCK_DGRAM; //Connectionless
    int protocol = 0; //Default transport: UDP for Internet connectionless

    int sender_sd; //socket descriptor ~= file descriptor
    sender_sd = socket(domain, type, protocol); ←
    if (sender_sd == -1) {
        printf("error in creating socket!\n");
        exit(1);
    }
    else
        printf("socket has created for sender with sd:%d\n", sender_sd);
}
```

Open a socket and receive a socket descriptor

Very similar to `open()` a file and file descriptor


Indeed, behind the scene, there are implemented very similar!

TCP/IP: UDP at Sender

Connectionless Communication

Sending a mail

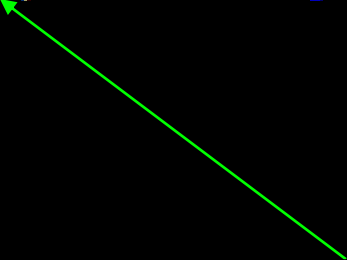
- 1) Creating Socket
- 2) Binding to an Address (Optional)



```
#include <sys/socket.h>
int bind(int sockfd, const struct sockaddr *addr, socklen_t len);
Returns 0 if OK, -1 on error
```



```
//Binding to an address is optional for sender!  
struct in_addr sender_sin_address;  
sender_sin_address.s_addr = inet_addr("137.207.140.134"); //nslookup `hostname`
```



IP Address of the computer system
[[0-255][0-255][0-255][0-255]]
4 bytes

There are many ways to know the IP of a computer, e.g.,

```
hfani@bravo:~$ nslookup `hostname`
```

```
//Binding to an address is optional for sender!  
struct in_addr sender_sin_address;  
sender_sin_address.s_addr = inet_addr("137.207.140.134 "); //nslookup `hostname`  
int sender_sin_port = htons(2000); //larger than 1024
```

htons(0); → random port by the kernel



In a computer system, there are multiple processes
To distinguish which process is the sender → PORT

End Point (Full Address) → IP:PORT

It is unique all around the world! Why?

```
struct sockaddr_in sender_sin;  
sender_sin.sin_family = domain;  
sender_sin.sin_addr = sender_sin_address;  
sender_sin.sin_port = sender_sin_port;  
int result = bind(sender_sd, (struct sockaddr *) &sender_sin, sizeof(sender_sin));  
if (result == -1) {  
    printf("error in binding sender to the address:port = %d:%d\n", sender_sin.sin_addr.s_addr, sender_sin.sin_port);  
    exit(1);  
}  
else  
    printf("sender bound to the address:port = %d:%d\n", sender_sin.sin_addr.s_addr, sender_sin.sin_port);
```


Socket ↔ IP:PORT

TCP/IP: UDP at Sender

Connectionless Communication

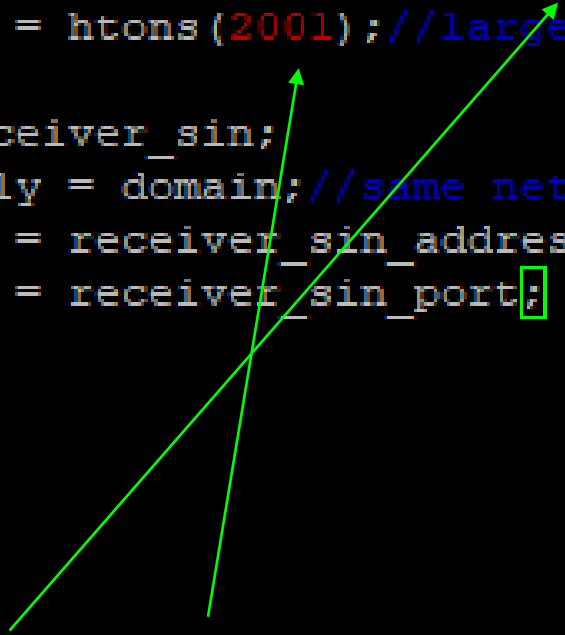
Sending a mail

- 1) Creating Socket
- 2) Binding to an Address (Optional)
- 3) Find the Receiver's Address




```
//Sending a message to the receiver at 137.207.82.52:2001
struct in_addr receiver_sin_address;
receiver_sin_address.s_addr = inet_addr("137.207.82.52");//nslookup `hostname` at the ta
int receiver_sin_port = htons(2001);//larger than 1024

struct sockaddr_in receiver_sin;
receiver_sin.sin_family = domain;//same network protocol
receiver_sin.sin_addr = receiver_sin_address;
receiver_sin.sin_port = receiver_sin_port;
```



IP:PORT of the Receiver


It is unique all around the world! Why?

TCP/IP: UDP at Sender

Connectionless Communication

Sending a mail

- 1) Creating Socket
- 2) Binding to an Address (Optional)
- 3) Find the Receiver's Address
- 4) Send the Mail to the Receiver



```
#include <sys/socket.h>
ssize_t sendto(int sockfd, const void *buf, size_t nbytes, int flags, const struct sockaddr *destaddr, socklen_t destlen);
```

Returns number of bytes sent if OK, -1 on error

```
char *mail = "a 10 percent promotion for candian tire!";
```

← Message

```
result = sendto(sender_sd, mail, strlen(mail), 0, (struct sockaddr *)&receiver_sin, sizeof(receiver_sin));  
if (result == -1) {  
    printf("error in sending message to the receiver!\n");  
    exit(1);  
}  
else  
{  
    printf("a mail has sent to the receiver at address:port = %d:%d\n", receiver_sin.sin_addr.s_addr, receiver_sin.sin_port);  
    printf("the content of the mail is <%s>\n", mail);  
}
```

Very similar to `write()` to a file

Receiver's IP:PORT

```
hfani@bravo:~$ ./sender
socket has created for sender with sd:3
sender bound to the address:port = -2037592183 :53255
a mail has sent to the receiver at address:port = 877842313:53511
the content of the mail is <a 10 percent promotion for candian tire!>
```

Why IP:PORT does not look like familiar?

Sender's IP:PORT = 137.207.140.134:2000

Receiver's IP:PORT = 137.207.82.52:2001


```
hfani@bravo:~$ ./sender
socket has created for sender with sd:3
sender bound to the address:port = -2037592183 :53255
a mail has sent to the receiver at address:port = 877842313:53511
the content of the mail is <a 10 percent promotion for candian tire!>
```

But there no receiver!
What happen to the mail?!

TCP/IP

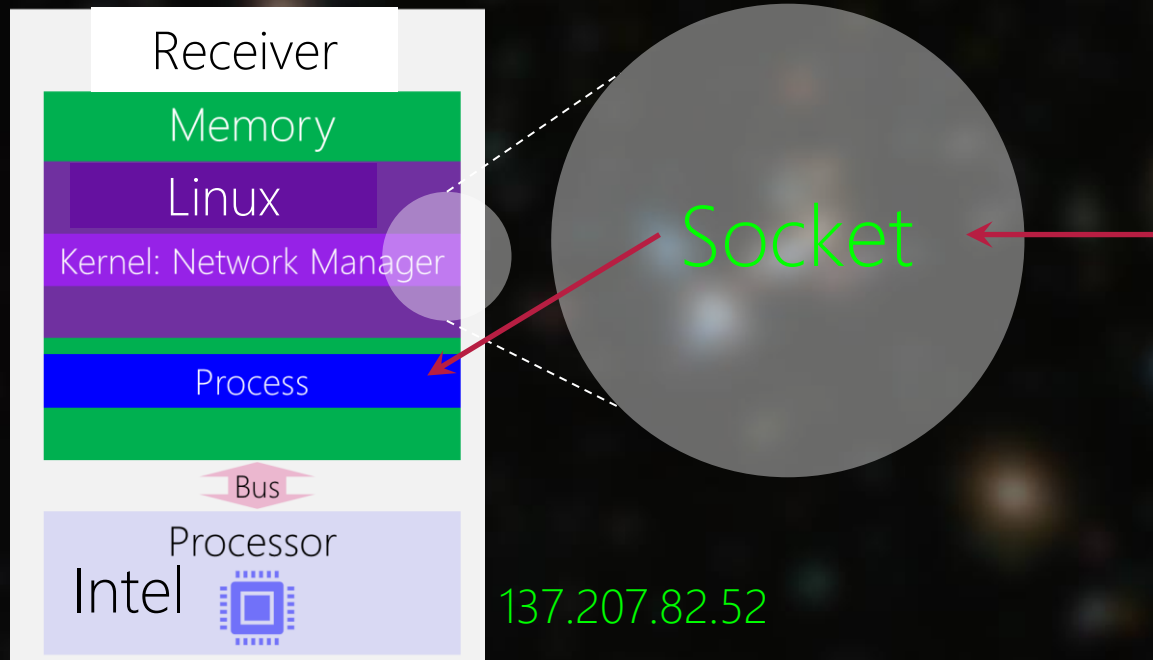
Just a name for [Link | Internet | Transport | Application] network protocol



TCP/IP: UDP at Receiver

Connectionless Communication

Sending a mail



- 1) Creating Socket
- 2) Binding to an Address (MUST)
- 3) Wait to receive a mail
- 4) Find the Sender's Address (Optional)
- 5) Read the Mail from the Sender